Amazon Braket

Developer Guide

Amazon Braket: Developer Guide

Table of Contents

What is Amazon Braket?

Amazon Braket is a fully managed AWS service that helps researchers, scientists, and developers get started with quantum computing. Quantum computing has the potential to solve computational problems that are beyond the reach of classical computers because it harnesses the laws of quantum mechanics to process information in new ways.

Gaining access to quantum computing hardware can be expensive and inconvenient. Limited access makes it difficult to run algorithms, optimize designs, evaluate the current state of the technology, and plan for when to invest your resources for maximum benefit. Braket helps you overcome these challenges.

Braket offers a single point of access to a variety of quantum computing technologies. With Braket, you can:

- Explore and design quantum and hybrid algorithms.
- Test algorithms on different quantum circuit simulators.
- Run algorithms on different types of quantum computers.
- Create proof of concept applications.

Defining quantum problems and programming quantum computers to solve them requires a new set of skills. To help you gain these skills, Braket offers different environments to simulate and run your quantum algorithms. You can find the approach that best suits your requirements and get started quickly with a set of example environments called *notebooks*.

Braket development has three stages — build, test, and run:

Build - Braket provides fully managed Jupyter notebook environments that make it easy to get started. Braket notebooks are pre-installed with sample algorithms, resources, and developer tools, including the Amazon Braket SDK. With the Amazon Braket SDK, you can build quantum algorithms and then test and run them on different quantum computers and simulators by changing a single line of code.

Test - Braket provides access to fully managed, high-performance quantum circuit simulators. You can test and validate your circuits. Braket handles all the underlying software components and Amazon Elastic Compute Cloud (Amazon EC2) clusters to take away the burden of simulating quantum circuits on classical high performance computing (HPC) infrastructure.

Run - Braket provides secure, on-demand access to different types of quantum computers. You have access to gate-based quantum computers from IonQ, OQC, Rigetti, and Xanadu as well as an analog Hamiltonian simulator from QuEra. You also have no upfront commitment, and no need to procure access through individual providers.

About quantum computing and Braket

Quantum computing is in its early developmental stage. It's important to understand that no universal, fault-tolerant quantum computer exists at present. Therefore, certain types of quantum hardware are better suited for each use case and it's crucial to have access to a variety of computing hardware. Braket offers a variety of hardware through third-party providers.

Existing quantum hardware is limited due to noise, which introduces errors. The industry is in the Noisy Intermediate Scale Quantum (NISQ) era. In the NISQ era, quantum computing devices are too noisy to sustain pure quantum algorithms, such as *Shor's algorithm* or *Grover's algorithm*. Until better quantum error correction is available, the most practical quantum computing requires the combination of classical (traditional) computing resources with quantum computers to create hybrid algorithms. Braket helps you work with *hybrid quantum algorithms.*

In hybrid quantum algorithms, quantum processing units (QPUs) are used as co-processors for CPUs, thus speeding up specific calculations in a classical algorithm. These algorithms utilize iterative processing, in which computation moves between classical and quantum computers. For example, current applications of quantum computing in chemistry, optimization, and machine learning are based on *variational quantum algorithms*, which are a type of *hybrid quantum algorithm*. In variational quantum algorithms, classical optimization routines adjust the parameters of a parameterized quantum circuit iteratively, much in the same way the weights of a neural network are adjusted iteratively based on the error in a machine learning training set. Braket offers access to the PennyLane open source software library, which assists you with *variational quantum algorithms.*

Quantum computing is gaining traction for computations in four main areas:

- **Number theory** — including factoring and cryptography (for example, *Shor's algorithm* is a primary quantum method for number theory computations)
- **Optimization** — including constraint satisfaction, solving linear systems, and machine learning
- **Oracular computing** — including search, hidden subgroups, and order finding (for example, *Grover's algorithm* is a primary quantum method for oracular computations)
- **Simulation** — including direct simulation, knot invariants, and quantum approximate optimization algorithm (QAOA) applications

Applications for these categories of computations can be found in financial services, biotechnology, manufacturing, and pharmaceuticals, to name a few. Braket offers capabilities and example notebooks that can already be applied to many proof of concept problems in addition to certain practical problems.

Amazon Braket terms and concepts

The following terms and concepts are used in Braket:

Analog Hamiltonian simulation

Analog Hamiltonian simulation (AHS) is a distinct quantum computing paradigm for direct simulation of time-dependent quantum dynamics of many-body systems. In AHS, users directly specify a time-dependent Hamiltonian and the quantum computer is tuned in such a way that it directly emulates the continuous time evolution under this Hamiltonian. AHS devices are typically special-purpose devices and not universal quantum computers like gate-based devices. They are limited to a class of Hamiltonians they can simulate. However, since these Hamiltonians are naturally implemented on the device, AHS does not suffer from the overhead required to formulate algorithms as circuits and implement gate operations.

Braket

We named the Braket service after the bra-ket notation, a standard notation in quantum mechanics. It was introduced by Paul Dirac in 1939 to describe the state of quantum systems, and it is also known as the Dirac notation.

Braket job

Amazon Braket has a feature called Amazon Braket Hybrid Jobs (or Braket Jobs for short) that provides fully managed executions of hybrid algorithms. A Braket job consists of three components:

1. The definition of your algorithm, which can be provided as a script, Python module, or Docker container.
2. The *job instance*, based on Amazon EC2, on which to run your algorithm. The default is an ml.m5.xlarge instance.
3. The *quantum device* on which to run the *quantum tasks* that are part of your algorithm. A single job typically contains a collection of many tasks.

Device

In Amazon Braket, a device is a backend that can run *quantum tasks*. A device can be a *QPU* or a *quantum circuit simulator*. To learn more, see Amazon Braket supported devices (p. 8).

Gate-based quantum computing

In gate-based quantum computing (QC), also called circuit-based QC, computations are broken down into elementary operations (gates). Certain sets of gates are universal, meaning that every computation can be expressed as a finite sequence of those gates. Gates are the building blocks of *quantum circuits* and are analogous to the logic gates of classical digital circuits.

Hamiltonian

The quantum dynamics of a physical system are determined by its Hamiltonian, which encodes all information about the interactions between constituents of the system and the effects of exogenous driving forces. The Hamiltonian of an N-qubit system is commonly represented as a 2^N by 2^N matrix of complex numbers on classical machines. By running an analog Hamiltonian simulation on a quantum device, you can avoid these exponential resource requirements.

Pulse

A pulse is a transient physical signal transmitted to the qubits. It is described by a waveform played in a frame that serves as a support for the carrier signal and is bound to the hardware channel or port. Customers can design their own pulses by providing the analog envelope that modulates the high-frequency sinusoidal carrier signal. The frame is uniquely described by a frequency and a phase that are often chosen to be on resonance with the energy separation between the energy levels for |0⟩ and |1⟩ of the qubit. Gates are thus enacted as pulses with a predetermined shape and calibrated parameters such as its amplitude, frequency and duration. Use cases that are not covered by template waveforms will be enabled via custom waveforms which will be specified at the single sample resolution by providing a list of values separated by a fixed, physical cycle-time.

Quantum circuit

A quantum circuit is the instruction set that defines a computation on a gate-based quantum computer. A quantum circuit is a sequence of quantum gates, which are reversible transformations on a qubit register, together with measurement instructions.

Quantum circuit simulator

A quantum circuit simulator is a computer program that runs on classical computers and calculates the measurement outcomes of a *quantum circuit*. For general circuits, the resource requirements of a quantum simulation grow exponentially with the number of qubits to simulate. Braket provides access to both managed (accessed through the Braket API) and local (part of the Amazon Braket SDK) quantum circuit simulators.

Quantum computer

A quantum computer is a physical device that uses quantum-mechanical phenomena, such as superposition and entanglement, to perform computations. There are different paradigms to quantum computing (QC), such as *gate-based* QC.

Quantum processing unit (QPU)

A QPU is a physical quantum computing device that can run on a quantum task. QPUs can be based on different QC paradigms, such as gate-based QC. To learn more, see Amazon Braket supported devices (p. 8).

QPU native gates

QPU native gates can be directly mapped to control pulses by the QPU control system. Native gates can be run on the QPU device without further compilation. Subset of *QPU supported gates*. You can find the native gates of a device on the **Devices** page in the Amazon Braket console and through the Braket SDK.

QPU supported gates

QPU supported gates are the gates accepted by the QPU device. These gates might not be able to directly run on the QPU, meaning that they might need to be decomposed into native gates. You can find the supported gates of a device on the **Devices** page in the Amazon Braket console and through the Amazon Braket SDK.

Quantum task

In Braket, a quantum task is the atomic request to a *device*. For *gate-based QC* devices, this includes the quantum circuit (including the measurement instructions and number of shots) and other request metadata. You can create quantum tasks through the Amazon Braket SDK or by using the CreateQuantumTask API operation directly. After you create a task, it will be queued until the requested device becomes available. You can view your quantum tasks on the **Tasks** page of the Amazon Braket console or by using the GetQuantumTask or SearchQuantumTasks API operations.

Qubit

The basic unit of information in a quantum computer is called a qubit (quantum bit), much like a bit in classical computing. A qubit is a two-level quantum system that can be realized by different physical implementations, such as superconducting circuits or individual ions and atoms. Other qubit types are based on photons, electronic or nuclear spins, or more exotic quantum systems.

Shots

Since quantum computing is inherently probabilistic, any circuit needs to be evaluated multiple times to get an accurate result. A single circuit execution and measurement is called a shot. The number of shots (repeated executions) for a circuit is chosen based on the desired accuracy for the result.

AWS terminology and tips for Amazon Braket

IAM users

An IAM user is an identity that you create in AWS. It represents the person or application that interacts with AWS services and resources. It consists of a name and credentials. By default, when you create a new IAM user in AWS, it has no permissions associated with it. To allow the IAM user to perform specific actions in AWS, such as launching an Amazon EC2 instance or creating an Amazon S3 bucket, you must grant the IAM user the necessary permissions.

- **Best practice:** We recommend that you create an individual IAM user for each person who needs access to AWS. Even if you have multiple employees who require the same level of access, create individual IAM users for each of them. This approach provides additional security by allowing each IAM user to have a unique set of security credentials.

IAM policies

An IAM policy is a document that allows or denies permissions to AWS services and resources. IAM policies enable you to customize users' levels of access to resources. For example, you can allow users access to all of the Amazon S3 buckets within your AWS account, or only a specific bucket.

- **Best practice:** Follow the security principle of *least privilege* when granting permissions. By following this principle, you help to prevent users or roles from having more permissions than needed to perform their tasks. For example, if an employee needs access to only a specific bucket, specify the bucket in the IAM policy instead of granting the employee access to all of the buckets in your AWS account.

IAM roles

An IAM role is an identity that you can assume to gain temporary access to permissions. Before an IAM user, application, or service can assume an IAM role, they must be granted permissions to switch to the role. When someone assumes an IAM role, they abandon all previous permissions that they had under a previous role and assume the permissions of the new role.

- **Best practice:** IAM roles are ideal for situations in which access to services or resources needs to be granted temporarily, instead of long-term.

Amazon S3 bucket

Amazon Simple Storage Service (Amazon S3) is an AWS service that lets you store data as *objects* in *buckets*. Amazon S3 buckets offer unlimited storage space. The maximum size for an object in an Amazon S3 bucket is 5 TB. You can upload any type of file data to an Amazon S3 bucket, such as images, videos, text files, backup files, media files for a website, archived documents, and your Braket task results.

- **Best practice:** You can set permissions to control access to your S3 bucket. For more information, see Bucket policies and user policies in the Amazon S3 documentation.

How Amazon Braket works

Amazon Braket provides on-demand access to quantum computing devices, including managed circuit simulators and different types of QPUs. In Amazon Braket, the atomic request to a device is a task. For gate-based QC devices, this request includes the quantum circuit (including the measurement instructions and number of shots) and other request metadata. For analog Hamiltonian simulators, the task contains the physical layout of the quantum register and the time- and space-dependence of the manipulating fields.

In this section, we are going to learn about the high-level flow of running tasks on Amazon Braket.

Amazon Braket task flow

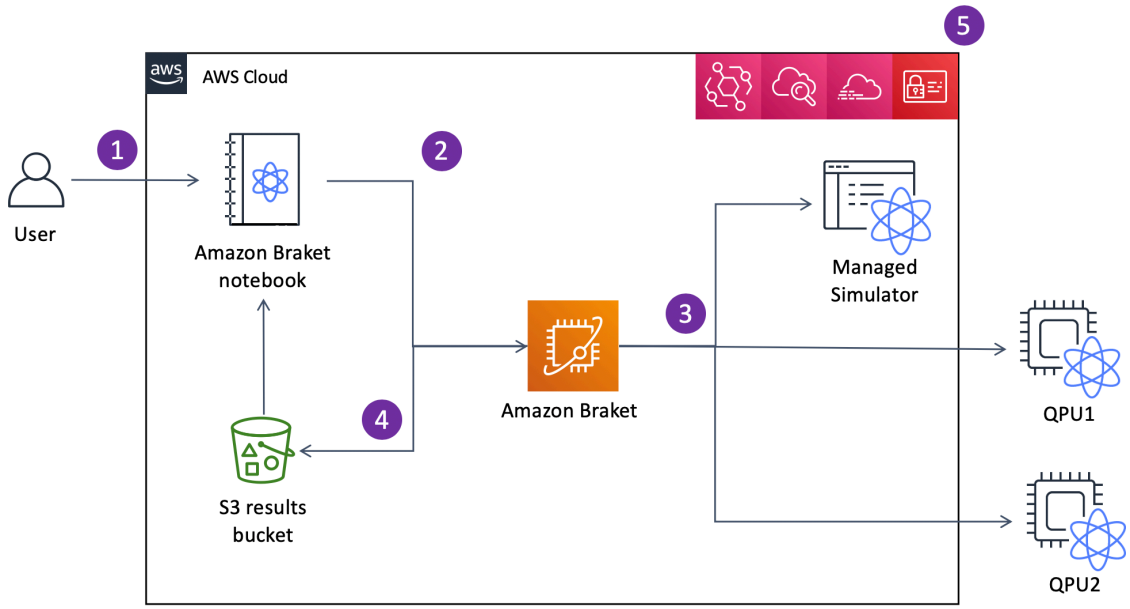

With Jupyter notebooks, you can conveniently define, submit, and monitor your tasks from the Amazon Braket Console or using the Amazon Braket SDK. You can build your quantum circuits directly in the SDK. However, for analog Hamiltonian simulators, you define the register layout and the controlling fields. After your task is defined, you can choose a device to run it on and submit it to the Amazon Braket API (2). Depending on the device you chose, the task is queued until the device becomes available and the task is sent to the QPU or simulator for implementation (3). Amazon Braket gives you access to five different types of QPUs (IonQ, Oxford Quantum Circuits (OQC), QuEra, Rigetti, Xanadu) and three on-demand simulators (SV1, DM1, TN1). To learn more, see Amazon Braket supported devices (p. 8).

After processing your task, Amazon Braket returns the results to an Amazon S3 bucket, where the data is stored in your AWS account (4). At the same time, the SDK polls for the results in the background and loads them into the Jupyter notebook at task completion. You can also view and manage your tasks on the **Tasks** page in the Amazon Braket console or by using the `GetQuantumTask` operation of the Amazon Braket API.

Amazon Braket is integrated with AWS Identity and Access Management (IAM), Amazon CloudWatch, AWS CloudTrail and Amazon EventBridge for user access management, monitoring and logging as well as for event based processing (5).

Third-party data processing

Tasks that are submitted to a QPU device are processed on quantum computers located in facilities operated by third party providers. To learn more about security and third-party processing in Amazon Braket, see Security of Amazon Braket Hardware Providers (p. 166).

Core repositories and plugins for Braket

Core repositories

The following displays a list of core repositories that contain key packages that are used for Braket:

- Braket Python SDK - Use the Braket Python SDK to set up your code on Jupyter notebooks in the Python programming language. After your Jupyter notebooks are set up, you can run your code on Braket devices and simulators
- Braket Schemas - The contract between the Braket SDK and the Braket service.
- Braket Default Simulator - All our local quantum simulators for Braket (state vector and density matrix).

Plugins

Then there are the various plugins that are used along with various devices and programming tools. These include Braket supported plugins as well as plugins that are supported by third parties as shown below.

Amazon Braket supported:

- Amazon Braket algorithm library - A catalog of pre-built quantum algorithms written in Python. Run them as they are or use them as a starting point to build more complex algorithms.
- Braket-PennyLane plugin - Use PennyLane as the QML framework on Braket.
- Braket-Strawberry Fields plugin - Interact with Xanadu's Borealis device.

Third-party (Braket team monitors and contributes):

- Qiskit-Braket provider - Use the Qiskit SDK to access Braket resources.
- Braket-Julia SDK - (EXPERIMENTAL) A Julia native version of the Braket SDK

Amazon Braket supported devices

In Amazon Braket, a device represents a QPU or simulator that you can call to run quantum tasks. Amazon Braket provides access to five QPU devices — from IonQ, Oxford Quantum Circuits, QuEra, Rigetti, and Xanadu, three on-demand simulators, and embedded simulators. For all devices, you can find further device properties, such as device topology, calibration data, and native gate sets, on the **Devices** tab of the Amazon Braket console or by means of the `GetDevice` API. When constructing a circuit with the simulators, Amazon Braket currently requires that you use contiguous qubits or indices. If you are working with the Amazon Braket SDK, you have access to device properties as shown in the following code example.

```
from braket.aws import AwsDevice
from braket.devices import LocalSimulator

device = AwsDevice('arn:aws:braket:::device/quantum-simulator/amazon/sv1')
 #SV1
# device = LocalSimulator()
 #Local State Vector Simulator
# device = LocalSimulator("default")
 #Local State Vector Simulator
# device = LocalSimulator(backend="default")
 #Local State Vector Simulator
# device = LocalSimulator(backend="braket_sv")
 #Local State Vector Simulator
# device = LocalSimulator(backend="braket_dm")
 #Local Density Matrix Simulator
# device = AwsDevice('arn:aws:braket:::device/quantum-simulator/amazon/tn1')
 #TN1
# device = AwsDevice('arn:aws:braket:::device/quantum-simulator/amazon/dm1')
 #DM1
# device = AwsDevice('arn:aws:braket:::device/qpu/ionq/ionQdevice')
 #IonQ
# device = AwsDevice('arn:aws:braket:::device/qpu/rigetti/Aspen-11')
 #Rigetti Aspen-11
# device = AwsDevice('arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2')
 #Rigetti Aspen M-2
# device = AwsDevice('arn:aws:braket:eu-west-2::device/qpu/oqc/Lucy')
 #OQC Lucy
# device = AwsDevice('arn:aws:braket:us-east-1::device/qpu/xanadu/Borealis')
 #Xanadu Borealis
# device = AwsDevice('arn:aws:braket:us-east-1::device/qpu/quera/Aquila')
 #QuEra Aquila

# get device properties
device.properties
```

Supported QPUs

- IonQ (p. 10)
- OQC Lucy (p. 11)
- QuEra Aquila (p. 11)
- Rigetti Aspen-11 (p. 10)
- Rigetti Aspen-M-2 (p. 10)

- Xanadu Borealis (p. 12)

Supported simulators

- Local state vector simulator (braket_sv) ('Default Simulator') (p. 12)
- Local density matrix simulator (`braket_dm`) (p. 12)
- State vector simulator (SV1) (p. 13)
- Density matrix simulator (DM1) (p. 13)
- Tensor network simulator (TN1) (p. 14)
- PennyLane's Lightning Simulators (p. 15)

Choose the best simulator for your task

- Compare simulators (p. 15)

> **Note**
> To view the available AWS Regions for each device, scroll right across the following table.

Amazon Braket devices

Provider	Device Name	Paradigm	Type	Device ARN	Region
IonQ	ionQdevice	gate-based	QPU	arn:aws:braket:::device/qpu/ionq/ionQdevice	us-east-1
Oxford Quantum Circuits	Lucy	gate-based	QPU	arn:aws:braket:eu-west-2::device/qpu/oqc/Lucy	eu-west-2
QuEra	Aquila	analog Hamiltonian simulation	QPU	arn:aws:braket:us-east-1::device/qpu/quera/Aquila	us-east-1
Rigetti	Aspen-11	gate-based	QPU	arn:aws:braket:::device/qpu/rigetti/Aspen-11	us-west-1
Rigetti	Aspen M-2	gate-based	QPU	arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2	us-west-1
Xanadu	Borealis	continuous-variable	QPU	arn:aws:braket:us-east-1::device/qpu/xanadu/Borealis	us-east-1
AWS	braket_sv	gate-based	Simulator	N/A (local simulator in Braket SDK)	N/A
AWS	braket_dm	gate-based	Simulator	N/A (local simulator in Braket SDK)	N/A
AWS	SV1	gate-based	Simulator	arn:aws:braket:::device/quantum-simulator/amazon/sv1	All Regions where Amazon Braket is available.

Provider	Device Name	Paradigm	Type	Device ARN	Region
AWS	DM1	gate-based	Simulator	arn:aws:braket:::device/quantum-simulator/amazon/dm1	All Regions where Amazon Braket is available.
AWS	TN1	gate-based	Simulator	arn:aws:braket:::device/quantum-simulator/amazon/tn1	us-west-2, us-east-1, and eu-west-2

To view additional details about the QPUs you can use with Amazon Braket, see Amazon Braket Hardware Providers.

IonQ

IonQ offers a gate-based QPU based on ion trap technology. IonQ's trapped ion QPUs are built on a chain of trapped 171Yb+ ions that are spatially confined by means of a microfabricated surface electrode trap within a vacuum chamber.

The IonQ device supports the following quantum gates.

```
'x', 'y', 'z', 'rx', 'ry', 'rz', 'h', 'cnot', 's', 'si', 't', 'ti', 'v', 'vi', 'xx', 'yy',
 'zz', 'swap', 'i'
```

With verbatim compilation, the IonQ device supports the following native gates.

```
'gpi', 'gpi2', 'ms'
```

These native gates can only be used with verbatim compilation. To learn more about verbatim compilation, see Verbatim Compilation.

Rigetti

Rigetti quantum processors are universal, gate-model machines based on all-tunable superconducting qubits. The Rigetti Aspen-11 system is based on scalable 40-qubit node technology. The Rigetti Aspen-M-2 system leverages their proprietary multi-chip technology and is assembled from two 40-qubit processors.

The Rigetti devices support the following quantum gates.

```
'cz', 'xy', 'ccnot', 'cnot', 'cphaseshift', 'cphaseshift00', 'cphaseshift01',
 'cphaseshift10', 'cswap', 'h', 'i', 'iswap', 'phaseshift', 'pswap', 'rx', 'ry', 'rz', 's',
 'si', 'swap', 't', 'ti', 'x', 'y', 'z'
```

With verbatim compilation, the Rigetti devices support the following native gates.

```
'rx', 'rz', 'cz', 'cphaseshift', 'xy'
```

Rigetti superconducting quantum processors can run the 'rx' gate with only the angles of ±π/2 or ±π.

Pulse-level control is available on the Rigetti devices, which support a set of predefined frames that are of the following types:

```
'rf', 'rf_f12', 'ro_rx', 'ro_rx', 'cz', 'cphase', 'xy'
```

See Roles of frames and ports (p. 110) for more information about these frames.

Oxford Quantum Circuits (OQC)

OQC quantum processors are universal, gate-model machines, built using scalable Coaxmon technology. The OQC Lucy system is an 8-qubit device with the topology of a ring in which each qubit is connected to its two nearest neighbors.

The Lucy device supports the following quantum gates.

```
'ccnot', 'cnot', 'cphaseshift', 'cswap', 'cy', 'cz', 'h', 'i', 'phaseshift', 'rx', 'ry',
 'rz', 's', 'si', 'swap', 't', 'ti', 'v', 'vi', 'x', 'y', 'z', 'ecr'
```

With verbatim compilation, the OQC device supports the following native gates.

```
'i', 'rz', 'v', 'x', 'ecr'
```

Pulse-level control is available on the OQC devices. The OQC devices support a set of predefined frames that are of the following types:

```
'drive', 'second_state', 'measure', 'acquire', 'cross_resonance',
 'cross_resonance_cancellation'
```

OQC devices support dynamic declaration of frames provided that you supply a valid port identifier. See Roles of frames and ports (p. 110) for more information about these frames and ports.

> **Note**
> When using pulse control with OQC, the length of your programs cannot exceed a maximum of 90 microseconds. It is important to remember that the duration's upper bound of single-qubit and two-qubit gates is approximately 50 nanoseconds and 1 microsecond, respectively. These numbers may vary depending on the qubits used, the device's current calibration and the circuit compilation.

QuEra

QuEra offers neutral-atom based devices that can run analog Hamiltonian simulation (AHS) tasks. These special-purpose devices faithfully reproduce the time-dependent quantum dynamics of hundreds of simultaneously interacting qubits.

One can program these devices in the paradigm of analog Hamiltonian simulation by prescribing the layout of the qubit register and the temporal and spatial dependence of the manipulating fields. Amazon Braket provides utilities to construct such programs via the AHS module of the python SDK, `braket.ahs`.

For more information, see the Analog Hamiltonian Simulation example notebooks or the Submit an analog program using QuEra's Aquila (p. 86) page.

Xanadu

Xanadu builds photonic quantum computers that use continuous variables for quantum computing known as qumodes instead of traditional two-level systems or qubits. In the case of the Borealis QPU, each qumode represents the quantized electromagnetic field of a laser pulse traveling through the device. Gates between two qumodes are implemented by interfering two temporally separated qumodes via programmable beam splitters and delay lines. Instead of the usual one- and two-qubit gates, such as Hadamard or CNOT, continuous variable quantum computing uses gates such as rotation, displacement, beam-splitting, and squeezing.

Xanadu's photonic quantum computer Borealis is not a universal machine, capable of arbitrary quantum computation, but instead implements a specific protocol known as Gaussian boson sampling (GBS). GBS is a model of photonic quantum computation first introduced by Hamilton et al. that consists of multi-mode linear optical operations followed by photon-counting measurements. The Borealis device implements GBS with 216 temporally-spaced qumodes. Amazon Braket offers access to Borealis via the open-source Strawberry Fields library for photonic quantum computation.

For more information, see the Borealis example notebook

Local state vector simulator (`braket_sv`)

The local state vector simulator (`braket_sv`) is part of the Amazon Braket SDK that runs locally in your environment. It is well-suited for rapid prototyping on small circuits (up to 25 qubits) depending on the hardware specifications of your Braket notebook instance or your local environment.

The simulator supports all gates in the Amazon Braket SDK, but QPU devices support a smaller subset. You can find the supported gates of a device in the device properties.

> **Note**
> The local simulator supports advanced OpenQASM features which may not be supported on QPU devices or other simulators. For more information on supported features, see the examples provided in the OpenQASM Local Simulator notebook.

For more information about how to work with simulators, see the Amazon Braket examples.

Local density matrix simulator (`braket_dm`)

The local density matrix simulator (`braket_dm`) is part of the Amazon Braket SDK that runs locally in your environment. It is well-suited for rapid prototyping on small circuits with noise (up to 12 qubits) depending on the hardware specifications of your Braket notebook instance or your local environment.

You can build common noisy circuits from the ground up using gate noise operations such as bit-flip and depolarizing error. You can also apply noise operations to specific qubits and gates of existing circuits that are intended to run both with and without noise.

The `braket_dm` local simulator can provide the following results, given the specified number of shots:

- Reduced density matrix: Shots = 0

> **Note**
> The local simulator supports advanced OpenQASM features, which may not be supported on QPU devices or other simulators. For more information about supported features, see the examples provided in the OpenQASM Local Simulator notebook.

To learn more about the local density matrix simulator, see the Braket introductory noise simulator example.

State vector simulator (SV1)

SV1 is a fully managed, high-performance, universal state vector simulator. It can simulate circuits of up to 34 qubits. You can expect a 34-qubit, dense, and square circuit (circuit depth = 34) to take approximately 1–2 hours to complete, depending on the type of gates used and other factors. Circuits with all-to-all gates are well suited for SV1. It returns results in forms such as a full state vector or an array of amplitudes.

SV1 has a maximum runtime of 6 hours. It has a default of 35 concurrent tasks, and a maximum of 100 (50 in us-west-1 and eu-west-2) concurrent tasks.

SV1 results

SV1 can provide the following results, given the specified number of shots:

- Sample: Shots > 0
- Expectation: Shots >= 0
- Variance: Shots >= 0
- Probability: Shots > 0
- Amplitude: Shots = 0

For more about results, see Result types.

SV1 is always available, it runs your circuits on demand, and it can run multiple circuits in parallel. The runtime scales linearly with the number of operations and exponentially with the number of qubits. The number of shots has a small impact on the runtime. To learn more, visit Compare simulators.

Simulators support all gates in the Braket SDK, but QPU devices support a smaller subset. You can find the supported gates of a device in the device properties.

Density matrix simulator (DM1)

DM1 is a fully managed, high-performance, density matrix simulator. It can simulate circuits of up to 17 qubits.

DM1 has a maximum runtime of 6 hours, a default of 35 concurrent tasks, and a maximum of 50 concurrent tasks.

DM1 results

DM1 can provide the following results, given the specified number of shots:

- Sample: Shots > 0
- Expectation: Shots >= 0
- Variance: Shots >= 0
- Probability: Shots > 0
- Reduced density matrix: Shots = 0, up to max 8 qubits

For more information about results, see Result types.

DM1 is always available, it runs your circuits on demand, and it can run multiple circuits in parallel. The runtime scales linearly with the number of operations and exponentially with the number of qubits. The number of shots has a small impact on the runtime. To learn more, see Compare simulators.

Noise gates and limitations

```
AmplitudeDamping
    Probability has to be within [0,1]
BitFlip
    Probability has to be within [0,0.5]
Depolarizing
    Probability has to be within [0,0.75]
GeneralizedAmplitudeDamping
    Probability has to be within [0,1]
PauliChannel
    The sum of the probabilities has to be within [0,1]
Kraus
    At most 2 qubits
    At most 4 (16) Kraus matrices for 1 (2) qubit
PhaseDamping
    Probability has to be within [0,1]
PhaseFlip
    Probability has to be within [0,0.5]
TwoQubitDephasing
    Probability has to be within [0,0.75]
TwoQubitDepolarizing
    Probability has to be within [0,0.9375]
```

Tensor network simulator (TN1)

TN1 is a fully managed, high-performance, tensor network simulator. TN1 can simulate certain circuit types with up to 50 qubits and a circuit depth of 1,000 or smaller. TN1 is particularly powerful for sparse circuits, circuits with local gates, and other circuits with special structure, such as quantum Fourier transform (QFT) circuits. TN1 operates in two phases. First, the *rehearsal phase* attempts to identify an efficient computational path for your circuit, so TN1 can estimate the runtime of the next stage, which is called the *contraction phase*. If the estimated contraction time exceeds the TN1 simulation runtime limit, TN1 does not attempt contraction.

TN1 has a runtime limit of 6 hours. It is limited to a maximum of 10 (5 in eu-west-2) concurrent tasks.

TN1 results

The contraction phase consists of a series of matrix multiplications. The series of multiplications continues until a result is reached or until it is determined that a result cannot be reached.

Note: Shots must be > 0 for the TN1 simulator.

Result types include:

- Sample
- Expectation
- Variance

For more about results, see Result types.

TN1 is always available, it runs your circuits on demand, and it can run multiple circuits in parallel. To learn more, see Compare simulators.

Simulators support all gates in the Braket SDK, but QPU devices support a smaller subset. You can find the supported gates of a device in the device properties.

Visit the Amazon Braket GitHub repository for a TN1 example notebook to help you get started with TN1.

Best practices for working with the TN1 simulator

- Avoid all-to-all circuits.
- Test a new circuit or class of circuits with a small number of shots, to learn the circuit's "hardness" for TN1.
- Split large shot simulations over multiple tasks.

PennyLane's lightning simulators

In addition to the Braket simulators, we also support PennyLane's lightning simulators. With PennyLane's lightning simulators, you can leverage advanced gradient computation methods, such as adjoint differentiation, to evaluate gradients faster. The lightning.qubit simulator is available as a device via Braket NBIs and as an embedded simulator, whereas the lightning.gpu simulator needs to be run as an embedded simulator with a GPU instance. See the Embedded simulators in Braket Jobs notebook for an example of using lightning.gpu. For more information about embedded simulators, see the Run a job with Amazon Braket Hybrid Jobs.

Compare simulators

This section helps you select the Amazon Braket simulator that's best suited for your task, by describing some concepts, limitations, and use cases.

Choosing between local simulators and managed simulators (SV1, TN1, DM1)

The performance of *local simulators* depends on the hardware that hosts the local environment, such as a Braket notebook instance, used to run your simulator. *Managed simulators* run in the AWS cloud and are designed to scale beyond typical local environments. Managed simulators are optimized for larger circuits, but add some latency overhead per task or batch of tasks. This can imply a trade-off if many tasks are involved. Given these general performance characteristics, the following guidance can help you choose how to run simulations, including ones with noise.

For **simulations**:

- When employing fewer than 18 qubits, use a local simulator.
- When employing 18–24 qubits, choose a simulator based on the workload.
- When employing more than 24 qubits, use a managed simulator.

For **noise simulations**:

- When employing fewer than 9 qubits, use a local simulator.
- When employing 9–12 qubits, choose a simulator based on the workload.
- When employing more than 12 qubits, use the DM1 simulator.

What is a state vector simulator?

The Amazon Braket state vector simulator (SV1) is a universal state vector simulator. It stores the full wave function of the quantum state and sequentially applies gate operations to the state. It stores all

possibilities, even the extremely unlikely ones. The SV1 simulator's run time for a task increases linearly with the number of gates in the circuit.

What is a density matrix simulator?

The Amazon Braket density matrix simulator (DM1) simulates quantum circuits with noise. It stores the full density matrix of the system and sequentially applies the gates and noise operations of the circuit. The final density matrix contains complete information about the quantum state after the circuit runs. The runtime generally scales linearly with the number of operations and exponentially with the number of qubits.

What is a tensor network simulator?

The Amazon Braket tensor network simulator (TN1) encodes quantum circuits into a structured graph.

- The nodes of the graph consist of quantum gates, or qubits.
- The edges of the graph represent connections between gates.

As a result of this structure, TN1 can find simulated solutions for relatively large and complex quantum circuits.

The TN1 simulator requires two phases

Typically, TN1 operates in a two-phase approach to simulating quantum computation.

- **The rehearsal phase:** In this phase, TN1 comes up with a way to traverse the graph in an efficient manner, which involves visiting every node so that you can obtain the measurement you desire. As a customer, you do not see this phase because TN1 performs both phases together for you. It completes the first phase and determines whether to perform the second phase on its own based on practical constraints. You have no input into that decision after the simulation has begun.
- **The contraction phase:** This phase is analogous to the execution phase of a computation in a classical computer. The phase consists of a series of matrix multiplications. The order of these multiplications has a great effect on the difficulty of the computation. Therefore, the rehearsal phase is accomplished first in order to find the most effective computation paths across the graph. After it finds the contraction path during the rehearsal phase, TN1 contracts together the gates of your circuit to produce the results of the simulation.

TN1 graphs are analogous to a map

Metaphorically, you can compare the underlying TN1 graph to the streets of a city. In a city with a planned grid, it is easy to find a route to your destination using a map. In a city with unplanned streets, duplicate street names, and so forth, it can be difficult to find a route to your destination by looking at a map.

If TN1 did not perform the rehearsal phase, it would be like walking around the streets of the city to find your destination, instead of looking at a map first. It can really pay off in terms of walking time to spend more time looking at the map. Similarly, the rehearsal phase provides valuable information.

You might say that the TN1 has a certain "awareness" of the structure of the underlying circuit that it traverses. It gains this awareness during the rehearsal phase.

Types of problems best suited for each of these types of simulators

SV1 is well-suited for any class of problems that rely primarily on having a certain number of qubits and gates. Generally, the time required grows linearly with the number of gates, while it does not depend on the number of shots. SV1 is generally faster than TN1 for circuits under 28 qubits.

The SV1 simulator can be slower for higher qubit numbers because it actually simulates all possibilities, even the extremely unlikely ones. It has no way to determine which outcomes are likely. Thus, for a 30-

qubit evaluation, SV1 must calculate 2^30 configurations. The limit of 34 qubits for the Amazon Braket SV1 simulator is a practical constraint due to memory and storage limitations. You can think of it like this: Each time you add a qubit to the SV1 simulator, the problem becomes twice as hard.

For many classes of problems, the TN1 simulator can evaluate much larger circuits in realistic time than the SV1 simulator because TN1 takes advantage of the structure of the graph. It essentially tracks the evolution of solutions from its starting place and it retains only the configurations that contribute to an efficient traversal. Put another way, it saves the configurations to create an ordering of matrix multiplication that results in a simpler evaluation process.

For TN1, the number of qubits and gates matters, but the structure of the graph matters a lot more. For example, TN1 is very good at evaluating circuits (graphs) in which the gates are short-range (that is, each qubit is connected by gates only to its nearest neighbour qubits), and circuits (graphs) in which the connections (or gates) have similar range. A typical range for TN1 is having each qubit talk only to other qubits that are 5 qubits away. If most of the structure can be decomposed into simpler relationships such as these, which can be represented in *more*, *smaller*, or *more uniform* matrices, TN1 performs the evaluation easily.

Limitations of the TN1 simulator

The TN1 simulator can be slower than the SV1 simulator depending on the graph's structural complexity. For certain graphs, TN1 terminates the simulation after the rehearsal stage, and shows a status of `FAILED`, for either of these two reasons:

- **Cannot find a path** — If the graph is too complex, it is too difficult to find a good traversal path and the simulator gives up on the computation. TN1 cannot perform the contraction. You may see an error message similar to this one: `No viable contraction path found.`
- **Contraction stage is too difficult** — In some graphs, TN1 can find a traversal path, but it is very long and extremely time-consuming to evaluate. In this case, the contraction is so expensive that the cost would be prohibitive and instead, TN1 exits after the rehearsal phase. You may see an error message similar to this one: `Predicted runtime based on best contraction path found exceeds TN1 limit.`

> **Note**
> You are billed for the rehearsal stage of TN1 even if contraction is not performed and you see a `FAILED` status.

The predicted runtime also depends on the shot count. In worst-case scenarios, TN1 contraction time depends linearly on the shot count. The circuit may be contractable with fewer shots. For example, you might submit a task with 100 shots, which TN1 decides is uncontractable, but if you resubmit with only 10, the contraction proceeds. In this situation, to attain 100 samples, you could submit 10 tasks of 10 shots for the same circuit and combine the results in the end.

As a best practice, we recommend that you always test your circuit or circuit class with a few shots (for example, 10) to find out how hard your circuit is for TN1, before you proceed with a higher number of shots.

> **Note**
> The series of multiplications that forms the contraction phase begins with small, NxN matrices. For example, a 2-qubit gate requires a 4x4 matrix. The intermediate matrices required during a contraction that is adjudged to be too difficult are gigantic. Such a computation would require days to complete. That's why Amazon Braket does not attempt extremely complex contractions.

Concurrency

All Braket simulators give you the ability to run multiple circuits concurrently. Concurrency limits vary by simulator and region. For more information on concurrency limits, see the Quotas page.

Example notebooks

Amazon Braket provides a variety of example notebooks showing the types of circuits that can either work well for, or challenge, the TN1 and SV1 simulators, such as the quantum Fourier transformation (QFT).

Amazon Braket Regions and endpoints

Amazon Braket is available in the following AWS Regions:

Region availability of Amazon Braket

Region Name	Region	Braket Endpoint	QPU
US East (N. Virginia)	us-east-1	braket.us-east-1.amazonaws.com	IonQ
US East (N. Virginia)	us-east-1	braket.us-east-1.amazonaws.com	Xanadu
US East (N. Virginia)	us-east-1	braket.us-east-1.amazonaws.com	QuEra
US West (N. California)	us-west-1	braket.us-west-1.amazonaws.com	Rigetti
EU West 2 (London)	eu-west-2	braket.eu-west-2.amazonaws.com	OQC

You can run Amazon Braket from any Region in which it is available, but each QPU is available only in a single Region. Tasks that run on a QPU device can be viewed in the Amazon Braket console in the Region of that device. If you are using the Amazon Braket SDK, you can submit tasks to any QPU device, regardless of the Region in which you are working. The SDK automatically creates a session to the Region for the QPU specified.

For general information about how AWS works with Regions and endpoints, see AWS service endpoints in the *AWS General Reference*.

Amazon Braket pricing

With Amazon Braket, you have access to quantum computing resources on demand without upfront commitment. You pay only for what you use. To learn more about pricing, please visit our pricing page.

Near real-time cost tracking

The Braket SDK offers you the option to add a near real-time cost tracking to your quantum workloads. Each of our example notebooks includes cost tracking code to provide you with a maximum cost estimate on Braket's quantum processing units (QPUs) and managed simulators. Maximum cost estimates will be shown in USD and are not inclusive of any credits or discounts.

> **Note**
> Charges shown are estimates based on your Amazon Braket simulator and quantum processing unit (QPU) task usage. Estimated charges shown may differ from your actual charges. Estimated charges do not factor in any discounts or credits and you may experience additional charges based on your use of other services such as Amazon Elastic Compute Cloud (Amazon EC2).

Cost tracking for the SV1 Simulator

In order to demonstrate how the cost tracking function can be used, we will be constructing a Bell State circuit and running it on our SV1 simulator. Begin by importing the Braket SDK modules, defining a Bell State and adding the `Tracker()` function to our circuit:

```
#import any required modules
from braket.aws import AwsDevice
from braket.circuits import Circuit
from braket.tracking import Tracker

#create our bell circuit
circ = Circuit().h(0).cnot(0,1)
device = AwsDevice("arn:aws:braket:::device/quantum-simulator/amazon/sv1")
with Tracker() as tracker:
task = device.run(circ, shots=1000).result()

#Your results
print(task.measurement_counts)
```

When you run your Notebook, you can expect the following output for your Bell State simulation. The tracker function will show you the number of shots sent, tasks completed, the execution duration, the billed execution duration, and your maximum cost in USD. Your execution time may vary for each simulation.

```
tracker.quantum_tasks_statistics()
 {'arn:aws:braket:::device/quantum-simulator/amazon/sv1':
  {'shots': 1000,
  'tasks': {'COMPLETED': 1},
  'execution_duration': datetime.timedelta(microseconds=4000),
  'billed_execution_duration': datetime.timedelta(seconds=3)}}

tracker.simulator_tasks_cost()
$0.00375
```

Using the cost tracker to set maximum costs

The cost tracker can also be utilized to set maximum costs on a program. You may have a maximum threshold for how much you want to spend on a given program. The cost tracker can be used to build out cost control logic in your execution code. The following example takes the same circuit on a Rigetti QPU and limits the cost to 1 USD. The cost to run one iteration of the circuit in our code is 0.37 USD. We have set the logic to repeat the iterations until the total cost exceeds 1 USD; hence, the code snippet will run 3 times until the next iteration exceeds 1 USD. Generally, a program would continue to iterate until it reaches your desired maximum cost, in this case - three iterations.

```
device = AwsDevice("arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2")
with Tracker() as tracker:
while tracker.qpu_tasks_cost() < 1:
result = device.run(circ, shots=200).result()
print(tracker.quantum_tasks_statistics())
print(tracker.qpu_tasks_cost(), "USD")
```

```
{'arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2': {'shots': 600, 'tasks':
 {'COMPLETED': 3}}}
1.11 USD
```

> **Note**
> The cost tracker will not track duration for failed TN1 simulator tasks. During a TN1 simulation, if your rehearsal completes, but the contraction step fails, your rehearsal charge will not be shown in the cost tracker.

Best practices for cost savings

Consider the following best practices for using Amazon Braket. Save time, minimize costs, and avoid common errors.

Verify with simulators

- Verify your circuits using a simulator before you run it on a QPU, so you can fine-tune your circuit without incurring charges for QPU usage.
- Although the results from running the circuit on a simulator may not be identical to the results from running the circuit on a QPU, you can identify coding errors or configuration issues using a simulator.

Restrict user access to certain devices

- You can set up restrictions that keep unauthorized users from submitting tasks on certain devices. The recommended method for restricting access is with AWS IAM. For more information about how to do that, see Restrict access.
- We recommend that you do **not** use your **admin** account as a way to give or restrict user access to Amazon Braket devices.

Set billing alarms

- You can set a billing alarm to notify you when your bill reaches a preset limit. The recommended way to set up an alarm is through AWS Budgets. You can set custom budgets and receive alerts when your costs or usage may exceed your budgeted amount. Information is available at AWS Budgets.

Test TN1 simulator tasks with low shot counts

- Simulators cost less than QHPs, but certain simulators can be expensive if tasks are run with high shot counts. We recommend that you test your TN1 simulator tasks with a low shot count. Shot count does not affect the cost for SV1 and local simulator tasks.

Check all Regions for tasks

- The console displays tasks only for your current AWS Region. When looking for billable tasks that have been submitted, be sure to check all Regions.
- You can view a list of devices and their associated Regions on the Supported Devices (p. 8) documentation page.

Amazon Braket Quotas

The following table lists the service quotas for Amazon Braket. Service quotas, also referred to as limits, are the maximum number of service resources or operations for your AWS account.

Some quotas can be increased. For more information, see AWS service quotas.

- Burst rate quotas cannot be increased.
- The maximum rate increase for adjustable quotas (except burst rate, which cannot be adjusted) is 2X the specified default rate limit. For example, a default quota of 60 can be adjusted to a maximum of 120.
- The adjustable quota for concurrent SV1 (DM1) tasks allows a maximum of 60 per AWS Region.

Resource	Description	Limit	Adjustable
Rate of API requests	The maximum number of requests per second that you can send in this account in the current Region.	140	Yes
Burst rate of API requests	The maximum number of additional requests per second (RPS) that you can send in one burst in this account in the current Region.	600	No
Rate of `CreateQuantumTask` requests	The maximum number of `CreateQuantumTask` requests you can send per second in this account per Region.	20	Yes
Burst rate of `CreateQuantumTask` requests	The maximum number of additional `CreateQuantumTask` requests per second (RPS) that you can send in one burst in this account in the current Region.	40	No
Rate of `SearchQuantumTasks` requests	The maximum number of `SearchQuantumTasks` requests you can send per second in this account per Region.	5	Yes
Burst rate of `SearchQuantumTasks` requests	The maximum number of additional `SearchQuantumTasks`	50	No

Resource	Description	Limit	Adjustable
	requests per second (RPS) that you can send in one burst in this account in the current Region.		
Rate of `GetQuantumTask` requests	The maximum number of `GetQuantumTask` requests you can send per second in this account per Region.	100	Yes
Burst rate of `GetQuantumTask` requests	The maximum number of additional `GetQuantumTask` requests per second (RPS) that you can send in one burst in this account in the current Region.	500	No
Rate of `CancelQuantumTask` requests	The maximum number of `CancelQuantumTask` requests you can send per second in this account per Region.	2	Yes
Burst rate of `CancelQuantumTask` requests	The maximum number of additional `CancelQuantumTask` requests per second (RPS) that you can send in one burst in this account in the current Region.	20	No
Rate of `GetDevice` requests	The maximum number of `GetDevice` requests you can send per second in this account per Region.	5	Yes
Burst rate of `GetDevice` requests	The maximum number of additional `GetDevice` requests per second (RPS) that you can send in one burst in this account in the current Region.	50	No
Rate of `SearchDevices` requests	The maximum number of `SearchDevices` requests you can send per second in this account per Region.	5	Yes

Resource	Description	Limit	Adjustable
Burst rate of `SearchDevices` requests	The maximum number of additional `SearchDevices` requests per second (RPS) that you can send in one burst in this account in the current Region.	50	No
Rate of `CreateJob` requests	The maximum number of `CreateJob` requests you can send per second in this account per Region.	1	Yes
Burst rate of `CreateJob` requests	The maximum number of additional `CreateJob` requests per second (RPS) that you can send in one burst in this account in the current Region.	5	No
Rate of `SearchJob` requests	The maximum number of `SearchJob` requests you can send per second in this account per Region.	5	Yes
Burst rate of `SearchJob` requests	The maximum number of additional `SearchJob` requests per second (RPS) that you can send in one burst in this account in the current Region.	50	No
Rate of `GetJob` requests	The maximum number of `GetJob` requests you can send per second in this account per Region.	5	Yes
Burst rate of `GetJob` requests	The maximum number of additional `GetJob` requests per second (RPS) that you can send in one burst in this account in the current Region.	25	No
Rate of `CancelJob` requests	The maximum number of `CancelJob` requests you can send per second in this account per Region.	2	Yes

Resource	Description	Limit	Adjustable
Burst rate of `CancelJob` requests	The maximum number of additional `CancelJob` requests per second (RPS) that you can send in one burst in this account in the current Region.	5	No
Number of concurrent **SV1** tasks	The maximum number of concurrent tasks running on the state vector simulator (SV1) in the current Region.	100 (50 in us-west-1 and eu-west-2)	No
Number of concurrent **DM1** tasks	The maximum number of concurrent tasks running on the density matrix simulator (DM1) in the current Region.	100 (50 in us-west-1 and eu-west-2)	No
Number of concurrent **TN1** tasks	The maximum number of concurrent tasks running on the tensor network simulator (TN1) in the current Region.	10 (5 in eu-west-2)	Yes
Number of concurrent jobs	The maximum number of concurrent jobs in the current Region.	5	Yes

The following are the default classical compute instance quotas for Hybrid Jobs. If you wish to raise these quotas, please contact AWS Support.

Resource	Description	Limits	Adjustable
Maximum number of instances of ml.c4.xlarge for jobs	The maximum number of instances of type ml.c4.xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes
Maximum number of instances of ml.c4.2xlarge for jobs	The maximum number of instances of type ml.c4.2xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes
Maximum number of instances of ml.c4.4xlarge for jobs	The maximum number of instances of type ml.c4.4xlarge allowed for all Amazon Braket	5	Yes

Resource	Description	Limits	Adjustable
	Hybrid Jobs in this account and region.		
Maximum number of instances of ml.c4.8xlarge for jobs	The maximum number of instances of type ml.c4.8xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes
Maximum number of instances of ml.c5.xlarge for jobs	The maximum number of instances of type ml.c5.xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes
Maximum number of instances of ml.c5.2xlarge for jobs	The maximum number of instances of type ml.c5.2xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes
Maximum number of instances of ml.c5.4xlarge for jobs	The maximum number of instances of type ml.c5.4xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	1	Yes
Maximum number of instances of ml.c5.9xlarge for jobs	The maximum number of instances of type ml.c5.9xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	1	Yes
Maximum number of instances of ml.c5.18xlarge for jobs	The maximum number of instances of type ml.c5.18xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.c5n.xlarge for jobs	The maximum number of instances of type ml.c5n.xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes

Resource	Description	Limits	Adjustable
Maximum number of instances of ml.c5n.2xlarge for jobs	The maximum number of instances of type ml.c5n.2xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.c5n.4xlarge for jobs	The maximum number of instances of type ml.c5n.4xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.c5n.9xlarge for jobs	The maximum number of instances of type ml.c5n.9xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.c5n.18xlarge for jobs	The maximum number of instances of type ml.c5n.18xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.g4dn.xlarge for jobs	The maximum number of instances of type ml.g4dn.xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.g4dn.2xlarge for jobs	The maximum number of instances of type ml.g4dn.2xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.g4dn.4xlarge for jobs	The maximum number of instances of type ml.g4dn.4xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.g4dn.8xlarge for jobs	The maximum number of instances of type ml.g4dn.8xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes

Resource	Description	Limits	Adjustable
Maximum number of instances of ml.g4dn.12xlarge for jobs	The maximum number of instances of type ml.g4dn.12xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.g4dn.16xlarge for jobs	The maximum number of instances of type ml.g4dn.16xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.m4.xlarge for jobs	The maximum number of instances of type ml.m4.xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes
Maximum number of instances of ml.m4.2xlarge for jobs	The maximum number of instances of type ml.m4.2xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes
Maximum number of instances of ml.m4.4xlarge for jobs	The maximum number of instances of type ml.m4.4xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	2	Yes
Maximum number of instances of ml.m4.10xlarge for jobs	The maximum number of instances of type ml.m4.10xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.m4.16xlarge for jobs	The maximum number of instances of type ml.m4.16xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.m5.large for jobs	The maximum number of instances of type ml.m5.large allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes

Resource	Description	Limits	Adjustable
Maximum number of instances of ml.m5.xlarge for jobs	The maximum number of instances of type ml.m5.xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes
Maximum number of instances of ml.m5.2xlarge for jobs	The maximum number of instances of type ml.m5.2xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes
Maximum number of instances of ml.m5.4xlarge for jobs	The maximum number of instances of type ml.m5.4xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	5	Yes
Maximum number of instances of ml.m5.12xlarge for jobs	The maximum number of instances of type ml.m5.12xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.m5.24xlarge for jobs	The maximum number of instances of type ml.m5.24xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.p2.xlarge for jobs	The maximum number of instances of type ml.p2.xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.p2.8xlarge for jobs	The maximum number of instances of type ml.p2.8xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.p2.16xlarge for jobs	The maximum number of instances of type ml.p2.16xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes

Resource	Description	Limits	Adjustable
Maximum number of instances of ml.p3.2xlarge for jobs	The maximum number of instances of type ml.p3.2xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.p3.8xlarge for jobs	The maximum number of instances of type ml.p3.8xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum number of instances of ml.p3.16xlarge for jobs	The maximum number of instances of type ml.p3.16xlarge allowed for all Amazon Braket Hybrid Jobs in this account and region.	0	Yes
Maximum allowed compute instances for a job	The maximum allowed number of compute instances for a job.	5	Yes

Requesting limit updates

If you receive a ServiceQuotaExceeded exception for an instance type and do not have sufficient instances available for it, you may request a limit increase from the Service Quotas page in the AWS console.

> **Note**
> p3 instances are not available in us-west-1. If your job is unable to provision requested ML compute capacity, use another region. In addition, if you do not see an instance in the table, it is not available for Hybrid Jobs.

Additional quotas and limits

- The Amazon Braket quantum task action is limited to 3MB in size.
- The maximum number of shots per task allowed for the SV1 managed simulator, DM1 managed simulator and Rigetti device is 100,000.
- The maximum number of shots per task allowed for the TN1 managed simulator is 1000.
- For IonQ and OQC devices, the maximum is 10,000 shots per task.
- For QuEra, the maximum allowed shots per task is 1000.
- For TN1 and the QPU devices, shots per task must be > 0.

When will my task run?

When you submit a circuit, Amazon Braket sends it to the device you specify. QPU and simulator tasks are queued and processed in the order they are received. The time required to process your task after you submit it varies depending on the number and complexity of tasks submitted by other Amazon Braket customers and on the availability of the QPU you selected.

Status change notifications in email or SMS

Amazon Braket sends events to Amazon EventBridge when the availability of a QPU changes or when your task's state changes. Follow these steps to receive device and task status change notifications by email or SMS message:

1. Create an Amazon SNS topic and a subscription to email or SMS. Availability of email or SMS depends on your Region. For more information, see Getting started with Amazon SNS and Sending SMS messages.
2. Create a rule in EventBridge that triggers the notifications to your SNS topic. For more information, see Monitoring Amazon Braket with Amazon EventBridge (p. 183).

Task completion alerts

You can set up notifications through the Amazon Simple Notification Service (SNS) so that you receive an alert when your Amazon Braket task is complete. Active notifications are useful if you expect a long wait time — for example, when you submit a large task or when you submit a task outside of a device's availability window. If you do not want to wait for the task to complete, you can set up an SNS notification.

An Amazon Braket notebook walks you through the setup steps. For more information, see the Amazon Braket example notebook for setting up notifications.

QPU availability windows and status

QPU availability varies among the Quantum Hardware Providers (QHPs).

In the **Devices** page of the Amazon Braket console, you can see the current and upcoming availability windows for each device. You also can view the status of each device.

A device is considered *offline* if is not available to customers, regardless of availability window. For example, it could be offline due to scheduled maintenance, upgrades, or operational issues.

The **Announcements** page in the Amazon Braket console shows scheduled downtime in advance to all Amazon Braket customers.

Enable Amazon Braket

You can enable Amazon Braket in your account through the AWS console.

Prerequisites

To enable and run Amazon Braket, you must have an IAM user or role with permission to initiate Amazon Braket actions. These permissions are included in the **AmazonBraketFullAccess** IAM policy (ARN:arn:aws:iam::aws:policy/AmazonBraketFullAccess).

> **Note**
> *If you are an administrator:*
> To enable Amazon Braket for other IAM users, for IAM roles, or for an IAM group in an account, you must grant permissions to each user, role, or group. You can grant these permissions by attaching the **AmazonBraketFullAccess** policy or by attaching a custom policy that you create. To learn more about the permissions necessary to use Amazon Braket, see Managing access to Amazon Braket.

Steps to enable Amazon Braket

1. Sign in to the Amazon Braket console with your AWS account.
2. Open the **Amazon Braket** console.
3. When you open the **Amazon Braket** console, you land by default on the **Devices** page. The message boxed in red at the top of the page informs you that you need to create a service-linked role to secure the permissions required to create resources with Amazon Braket.

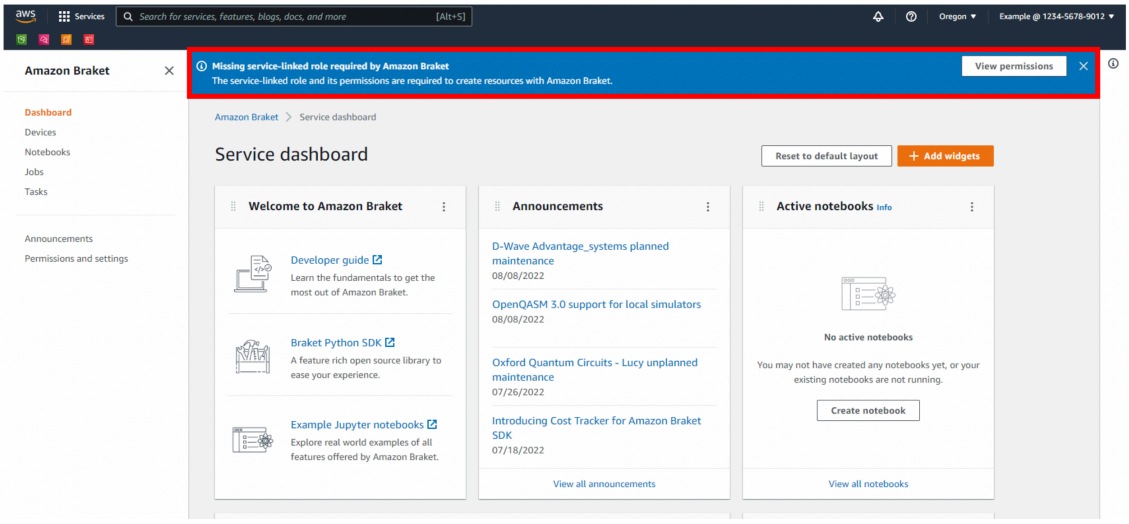

4. To create a service-linked role for the AWS account, select the **View Permissions** button on the right of the message at the top of the page or choose the **Permissions and settings** tab on the left side of the **Amazon Braket** page and then the **Execution roles** tab on the **Permissions and settings** page. To create this service-linked role, select the **Create service-linked role** button.

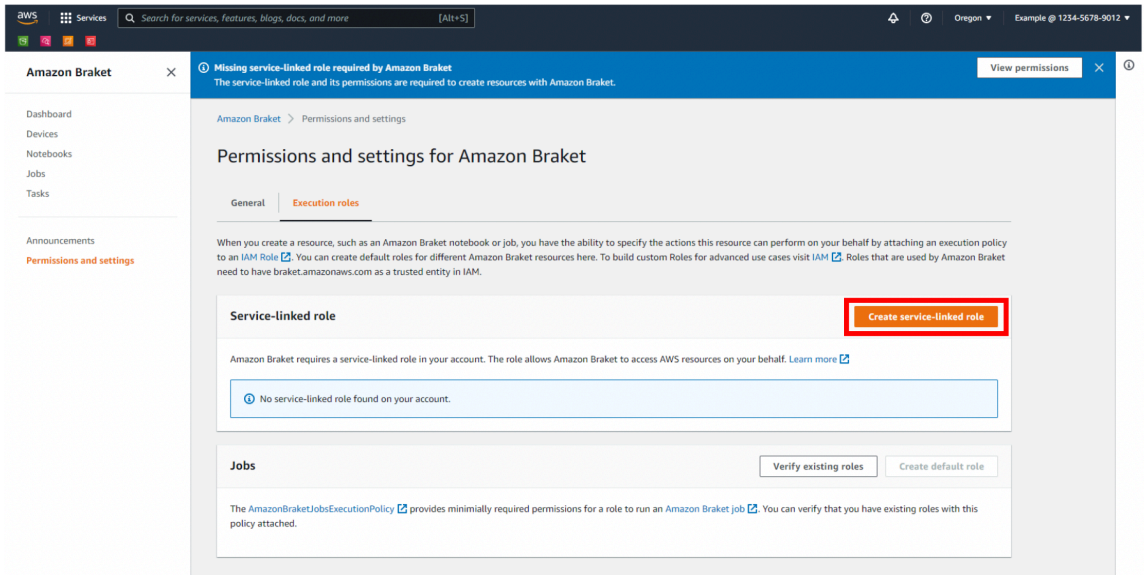

5. In order to use third-party devices, you have to agree to certain conditions regarding data transfer to third parties. For more information, see Enable third-party devices (p. 35).

Note
Systems that do not involve third-party providers, such as Amazon's simulators, can be used without agreeing to the **Enable third-party devices** agreement. But to use third party devices, such as QPU's, this agreement must be signed.

Enable third-party devices

Amazon Braket allows you to run quantum tasks or jobs on any Amazon owned device as long as you have a service-linked role (SLR) in your account. For more information on how to acquire an SLR, see Steps to enable Amazon Braket (p. 33) and see the Service-linked role (p. 164) section to understand what an SLR actually is. To access and enable third-party devices such as QPUs with Amazon Braket, you must agree to certain terms and conditions. The terms and conditions of this agreement are provided on the **General** tab of the **Permissions and settings** page as shown in the following figure. Select the **Accept terms and enable** button to agree to them in order to access third-party devices.

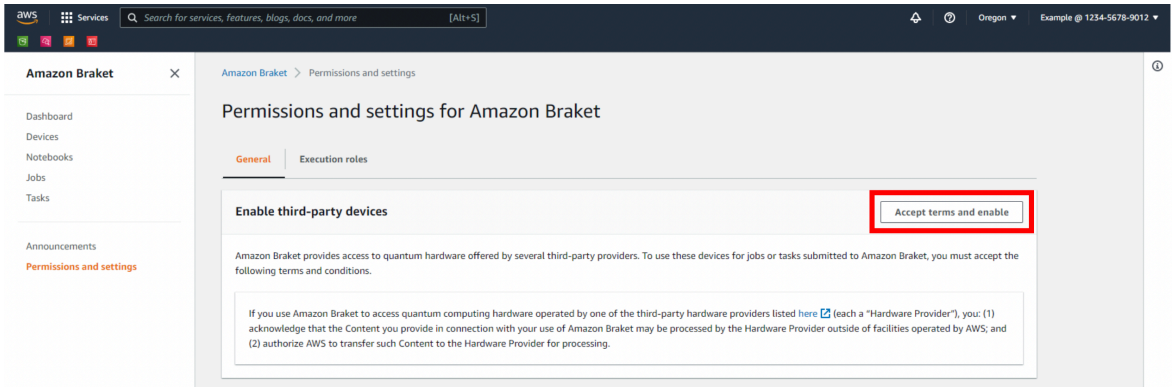

Note
Accepting these terms to enable use of third-party devices only needs to be done **once per account** and only needs to be done if you are accessing third-party hardware. You do not need this to access Amazon Braket's local or managed simulators.

Get started with Amazon Braket

After you have followed the instructions in Enable Amazon Braket (p. 33), you can get started with Amazon Braket.

The steps to get started include:

- Create an Amazon Braket notebook instance (p. 36)
- Run your first circuit using the Amazon Braket Python SDK (p. 38)
- Explore the algorithm library (p. 40)
- Hello AHS: Run your first Analog Hamiltonian Simulation (p. 41)

Create an Amazon Braket notebook instance

Amazon Braket provides fully-managed Jupyter notebooks to get you started. The Amazon Braket notebook instances are based on Amazon SageMaker notebook instances. You can learn more about notebook instances. To get started with Braket, follow these steps to create an Amazon Braket notebook instance.

1. Open the Amazon Braket console.
2. Choose **Notebooks** in the left pane, then choose **Create notebook**.
3. In **Notebook instance settings**, enter a **Notebook instance name** using only alphanumeric and hyphen characters.
4. Select the **Notebook instance type**. Choose the smallest type you need. To get started, choose a cost-effective instance type, such as **ml.t3.medium**. The instance types are Amazon SageMaker notebook instances. To learn more, see Amazon SageMaker pricing.
5. In **Permissions and encryption**, select **Create a new role** (a new role with a name that begins with `AmazonBraketServiceSageMakerNotebook` is created).
6. If you want to associate a public Github repository with your notebook instance, click on the **Git repository** dropdown and select **Clone a public git repository from url** from the **Repository** dropdown menu. Enter the URL of the repo in the **Git repository URL** text bar.

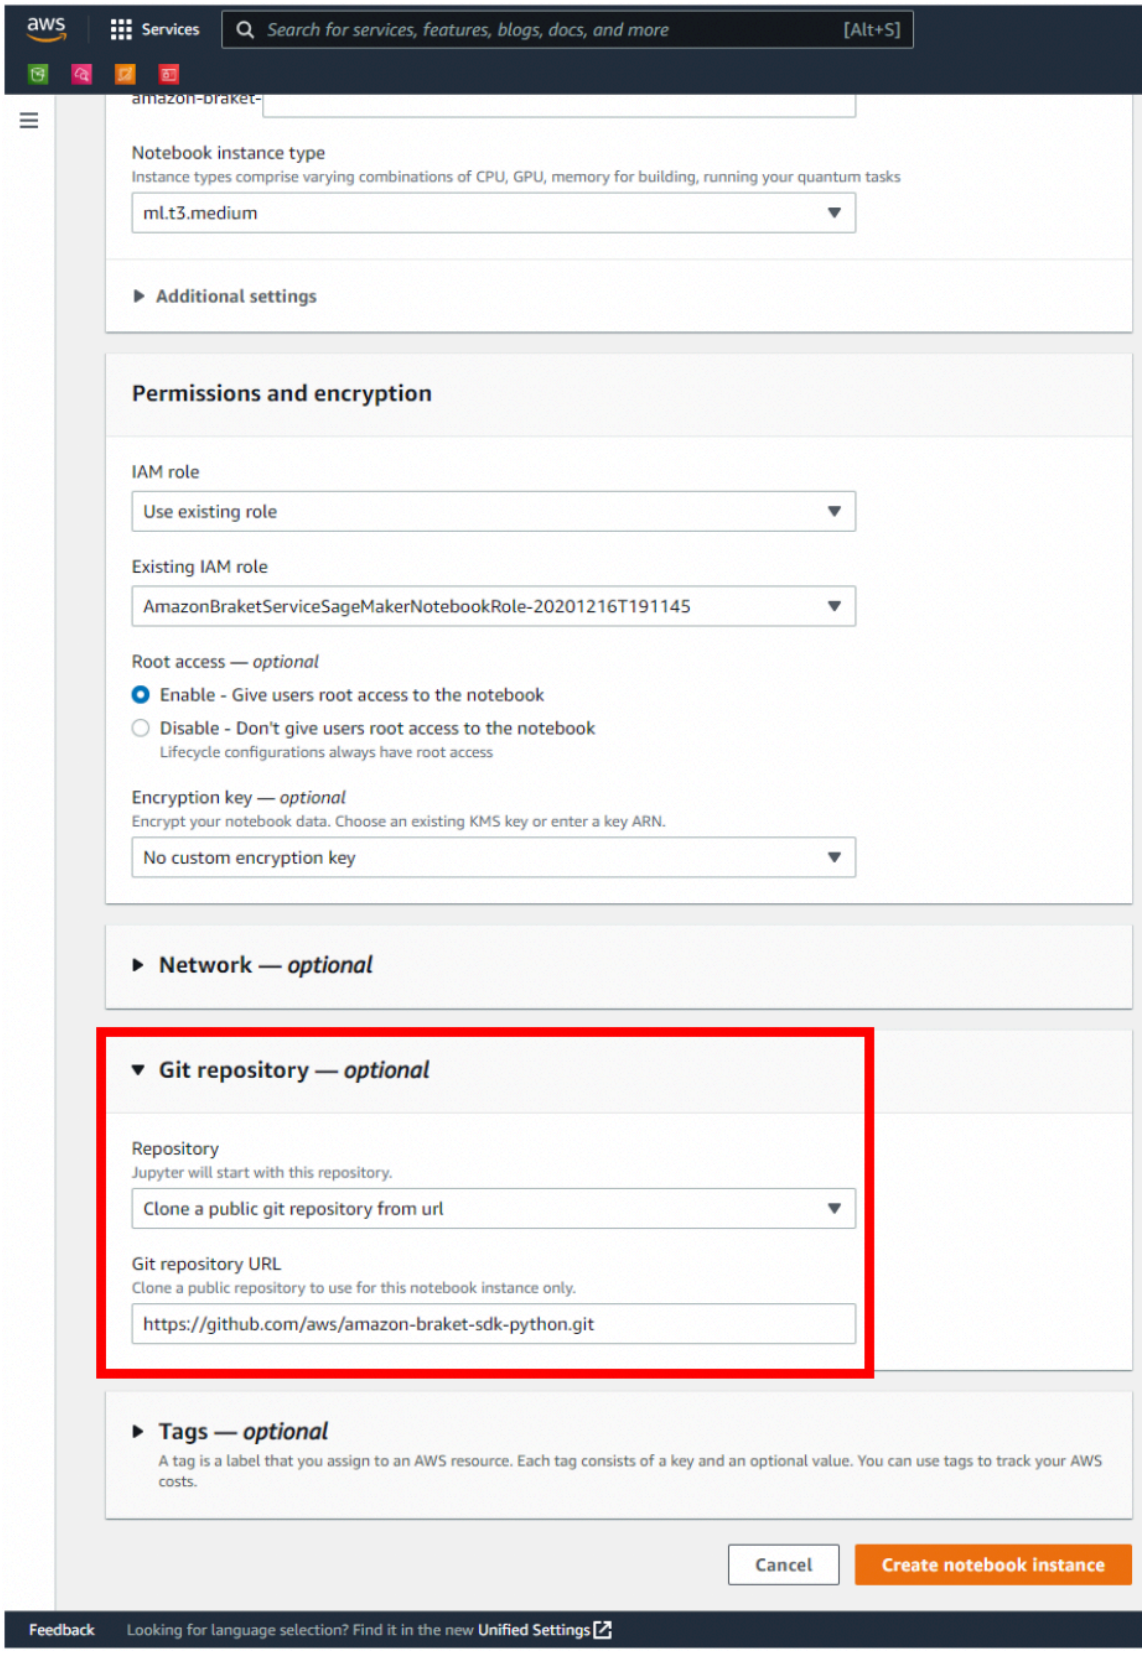
aws
Services
Search for services, features, blogs, docs, and more
[Alt+S]
amazon-braket-
Notebook instance type
Instance types comprise varying combinations of CPU, GPU, memory for building, running your quantum tasks
ml.t3.medium
Additional settings
Permissions and encryption
IAM role
Use existing role
Existing IAM role
AmazonBraketServiceSageMakerNotebookRole-20201216T191145
Root access — optional
Enable - Give users root access to the notebook
Disable - Don't give users root access to the notebook
Lifecycle configurations always have root access
Encryption key — optional
Encrypt your notebook data. Choose an existing KMS key or enter a key ARN.
No custom encryption key
Network — optional
Git repository — optional
Repository
Jupyter will start with this repository.
Clone a public git repository from url
Git repository URL
Clone a public repository to use for this notebook instance only.
https://github.com/aws/amazon-braket-sdk-python.git
Tags — optional
A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to track your AWS costs.
Cancel
Create notebook instance
Feedback
Looking for language selection? Find it in the new Unified Settings

7. Choose **Create notebook instance**.

   It takes several minutes to create the notebook. The notebook is displayed on the **Notebooks** page with a status of **Pending**. When the notebook instance is ready to use, the status changes to **InService**. You may need to refresh the page to display the updated status for the notebook.

   Note
   You can view and manage your Amazon Braket notebook instances in the Amazon Braket and Amazon SageMaker consoles. Additional Amazon Braket notebook settings are available through the SageMaker console.

If you're working in the Amazon Braket console within AWS, as given previously, the Amazon Braket SDK and plugins are preloaded in the notebooks you just created. If you want to run on your own machine, you can install the SDK and plugins when you run the command `pip install amazon-braket-sdk` or when you run the command `pip install amazon-braket-pennylane-plugin` (the latter for use with PennyLane plugins).

Run your first circuit using the Amazon Braket Python SDK

After your notebook instance has launched, open the instance with a standard Jupyter interface by choosing the notebook you just created.

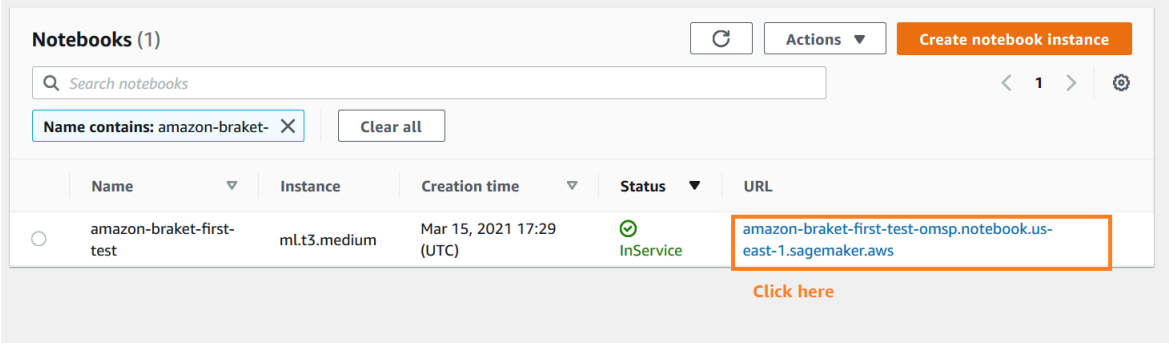

Amazon Braket notebook instances are pre-installed with the Amazon Braket SDK and all its dependencies. Start by creating a new notebook with `conda_braket` kernel.

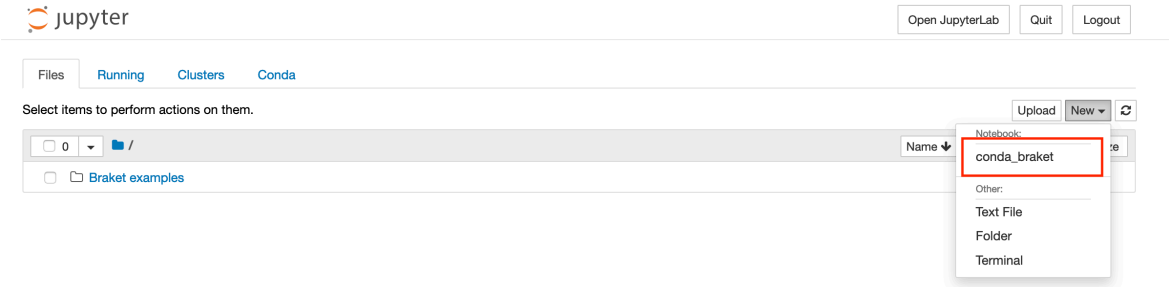

You can start with a simple "Hello, world!" example. First, construct a circuit that prepares a Bell state, and then run that circuit on different devices to obtain the results.

Begin by importing the Amazon Braket SDK modules and defining a simple Bell State circuit.

```
import boto3
from braket.aws import AwsDevice
from braket.devices import LocalSimulator
from braket.circuits import Circuit

# create the circuit
bell = Circuit().h(0).cnot(0, 1)
```

You can visualize the circuit with this command:

```
print(bell)
```

Run your circuit on the local simulator

Next, choose the quantum device on which to execute the circuit. The Amazon Braket SDK comes with a local simulator for rapid prototyping and testing. We recommend using the local simulator for smaller circuits, which can be up to 25 qubits (depending on your local hardware).

Here's how to instantiate the local simulator:

```
# instantiate the local simulator
local_sim = LocalSimulator()
```

and run the circuit:

```
# run the circuit
result = local_sim.run(bell, shots=1000).result()
counts = result.measurement_counts
print(counts)
```

You should see a result something like this:

```
Counter({'11': 503, '00': 497})
```

The specific Bell state you have prepared is an equal superposition of |00> and |11>, and you'll find a roughly equal (up to shot noise) distribution of 00 and 11 as measurement outcomes, as expected.

Run your circuit on a managed simulator

Amazon Braket also provides access to a fully-managed, high-performance simulator, SV1, for running larger circuits. SV1 is a state-vector simulator that allows for simulation of quantum circuits of up to 34 qubits. You can find more information on SV1 in the Supported Devices (p. 8) section and in the AWS console. When running tasks on SV1 (and on TN1 or any QPU), the results of your task are stored in an S3 bucket in your account. If you do not specify a bucket, the Braket SDK creates a default bucket `amazon-braket-{region}-{accountID}` for you. To learn more, see Managing access to Amazon Braket (p. 155).

> **Note**
> Fill in your actual, existing bucket name where the following example shows `example-bucket` as your bucket name. Bucket names for Amazon Braket always begin with `amazon-braket-` followed by other identifying characters you add.

```
# get the account ID
aws_account_id = boto3.client("sts").get_caller_identity()["Account"]
# the name of the bucket
```

```
my_bucket = "example-bucket"
# the name of the folder in the bucket
my_prefix = "simulation-output"
s3_folder = (my_bucket, my_prefix)
```

To run a circuit on SV1, you must provide the location of the S3 bucket you previously selected as a positional argument in the `.run()` call.

```
# choose the cloud-based managed simulator to run your circuit
device = AwsDevice("arn:aws:braket:::device/quantum-simulator/amazon/sv1")

# execute the circuit
task = device.run(bell, s3_folder, shots=100)
# display the results
print(task.result().measurement_counts)
```

The Amazon Braket console provides further information about your task. Navigate to the **Tasks** tab in the console and your task should be on the top of the list. Alternatively, you can search for your task using the unique task ID or other criteria.

> **Note**
> After 90 days, Amazon Braket automatically removes all task IDs and other metadata associated with your tasks. For more information, see Data retention.

Running on a QPU

With Amazon Braket, you can run the previous quantum circuit example on a physical quantum computer by just changing a single line of code. Amazon Braket provides access to QPU devices from IonQ, Oxford Quantum Circuits, QuEra, Rigetti, and Xanadu. You can find information about the different devices and availability windows in the Supported Devices (p. 8) section, and in the AWS console under the **Devices** tab. The following example shows how to instantiate a Rigetti device.

```
# choose the Rigetti hardware to run your circuit
device = AwsDevice("arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2")
```

Choose an IonQ device with this code:

```
# choose the Ionq device to run your circuit
device = AwsDevice("arn:aws:braket:::device/qpu/ionq/ionQdevice")
```

When you execute your task, the Amazon Braket SDK polls for a result (with a default timeout of 5 days). You can change this default by modifying the `poll_timeout_seconds` parameter in the the `.run()` command as shown in the example that follows. Keep in mind that if your polling timeout is too short, results may not be returned within the polling time, such as when a QPU is unavailable and a local timeout error is returned. You can restart the polling by calling the `task.result()` function.

```
# define task with 1 day polling timeout
task = device.run(bell, s3_folder, poll_timeout_seconds=24*60*60)
print(task.result().measurement_counts)
```

Explore the algorithm library

The Amazon Braket algorithm library is a catalog of pre-built quantum algorithms written in Python. You can run these algorithms as they are or use them as a starting point to build more complex algorithms. You can access the algorithm library from the Braket console.

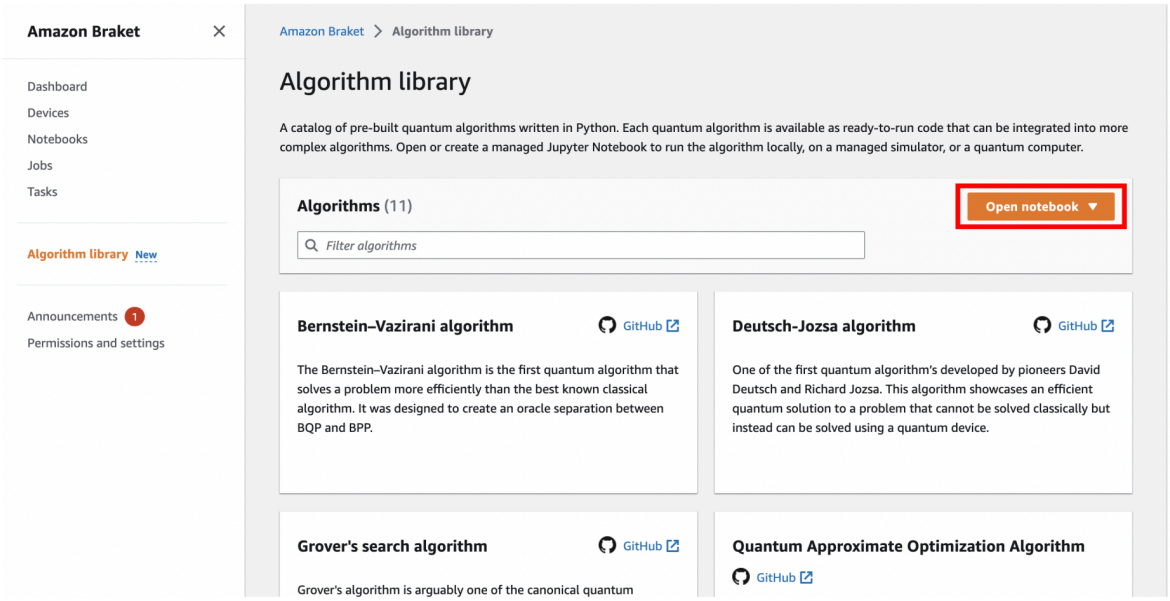

The Braket console provides a description of each available algorithm in the algorithm library. Choose a GitHub link to see the details of each algorithm, or choose **Open notebook** to open or create a notebook that contains all of the available algorithms. If you choose the notebook option, you can then find the Braket algorithm library in the root folder of your notebook.

Hello AHS: Run your first Analog Hamiltonian Simulation

AHS

Analog Hamiltonian simulation (AHS) is a paradigm of quantum computing different from quantum circuits: instead of a sequence of gates, each acting only on a couple of qubits at a time, an AHS program is defined by the time- and space-dependent parameters of the Hamiltonian in question. The Hamiltonian of a system encodes its energy levels and the effects of external forces, which together govern the time evolution of its states. For an N-qubit systems, the Hamiltonian can be represented by a $2^N X 2^N$ square matrix of complex numbers.

Quantum devices capable of performing AHS will tune their parameters (for example amplitude and detuning of a coherent driving field) to closely approximate the time evolution of the quantum system under the custom Hamiltonian. The AHS paradigm is suitable for simulating static and dynamic properties of quantum systems of many interacting particles. Purpose-built QPUs, such as the Aquila device from QuEra can simulate the time evolution of systems with sizes that are otherwise infeasible on classical hardware.

Interacting spin chain

For a canonical example of a system of many interacting particles, let us consider a ring of eight spins (each of which can be in “up” |↑# and “down” |↓# states). Albeit small, this model system already exhibits a handful of interesting phenomena of naturally occurring magnetic materials. In this example, we will show how to prepare a so-called anti-ferromagnetic order, where consecutive spins point in opposite directions.

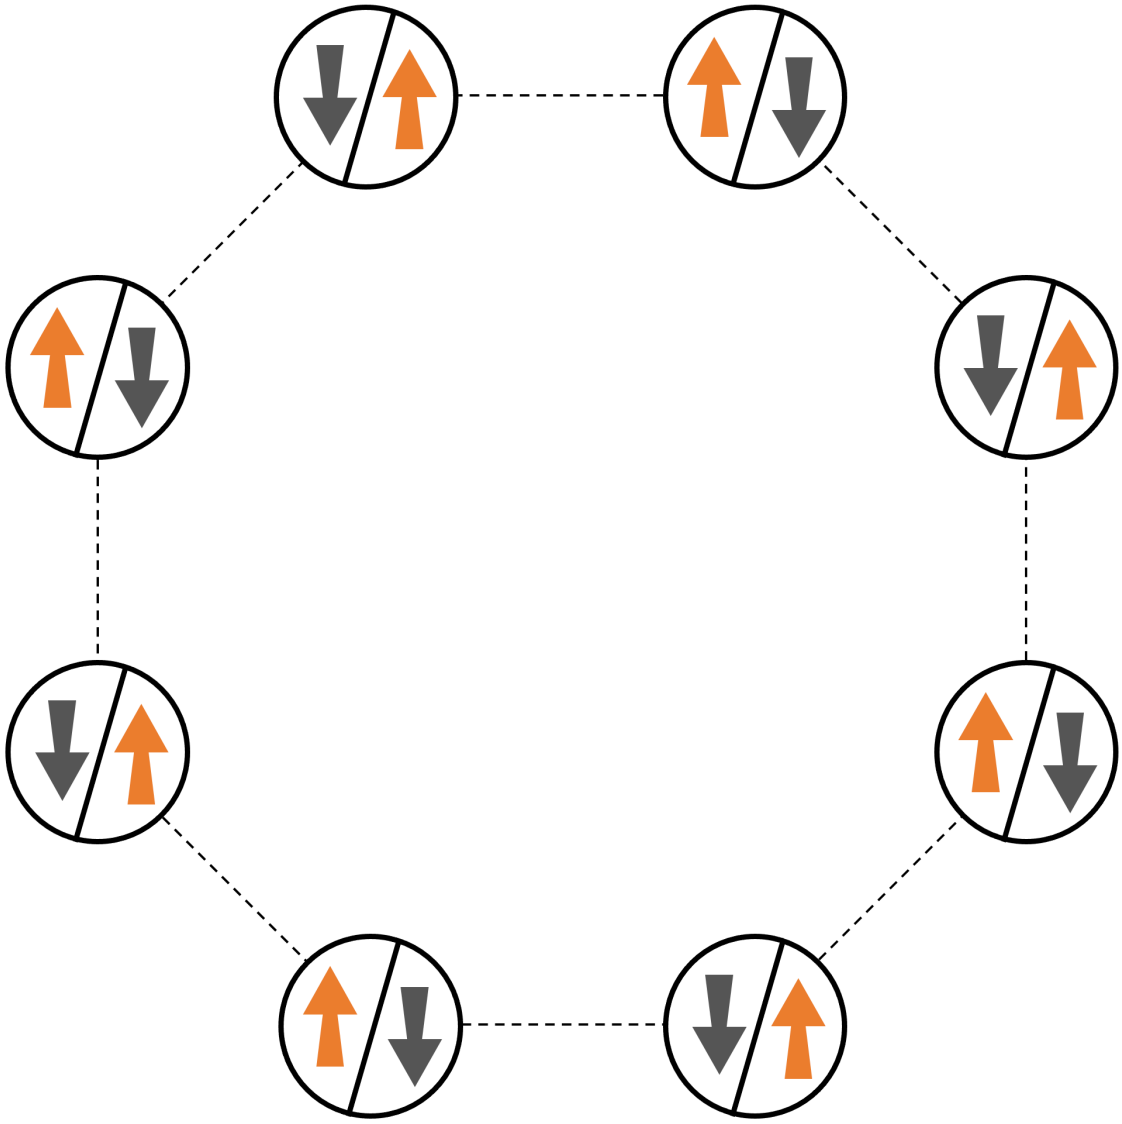

Arrangement

We will use one neutral atom to stand for each spin, and the “up” and “down” spin states will be encoded in excited Rydberg state and ground state of the atoms, respectively. First, we create the 2-d arrangement. We can program the above ring of spins with the following code.

Prerequisites: You need to pip install the Braket SDK. (If you are using a Braket hosted notebook instance, this SDK comes pre-installed with the notebooks.) To reproduce the plots, you also need to separately install matplotlib with the shell command `pip install matplotlib`.

```
import numpy as np
import matplotlib.pyplot as plt  # required for plotting

from braket.ahs.atom_arrangement import AtomArrangement

a = 5.7e-6  # nearest-neighbor separation (in meters)
```

```
register = AtomArrangement()
register.add(np.array([0.5, 0.5 + 1/np.sqrt(2)]) * a)
register.add(np.array([0.5 + 1/np.sqrt(2), 0.5]) * a)
register.add(np.array([0.5 + 1/np.sqrt(2), - 0.5]) * a)
register.add(np.array([0.5, - 0.5 - 1/np.sqrt(2)]) * a)
register.add(np.array([-0.5, - 0.5 - 1/np.sqrt(2)]) * a)
register.add(np.array([-0.5 - 1/np.sqrt(2), - 0.5]) * a)
register.add(np.array([-0.5 - 1/np.sqrt(2), 0.5]) * a)
register.add(np.array([-0.5, 0.5 + 1/np.sqrt(2)]) * a)
```

which we can also plot with

```
fig, ax = plt.subplots(1, 1, figsize=(7,7))
xs, ys = [register.coordinate_list(dim) for dim in (0, 1)]
ax.plot(xs, ys, 'r.', ms=15)
for idx, (x, y) in enumerate(zip(xs, ys)):
    ax.text(x, y, f" {idx}", fontsize=12)
plt.show()  # this will show the plot below in an ipython or jupyter session
```

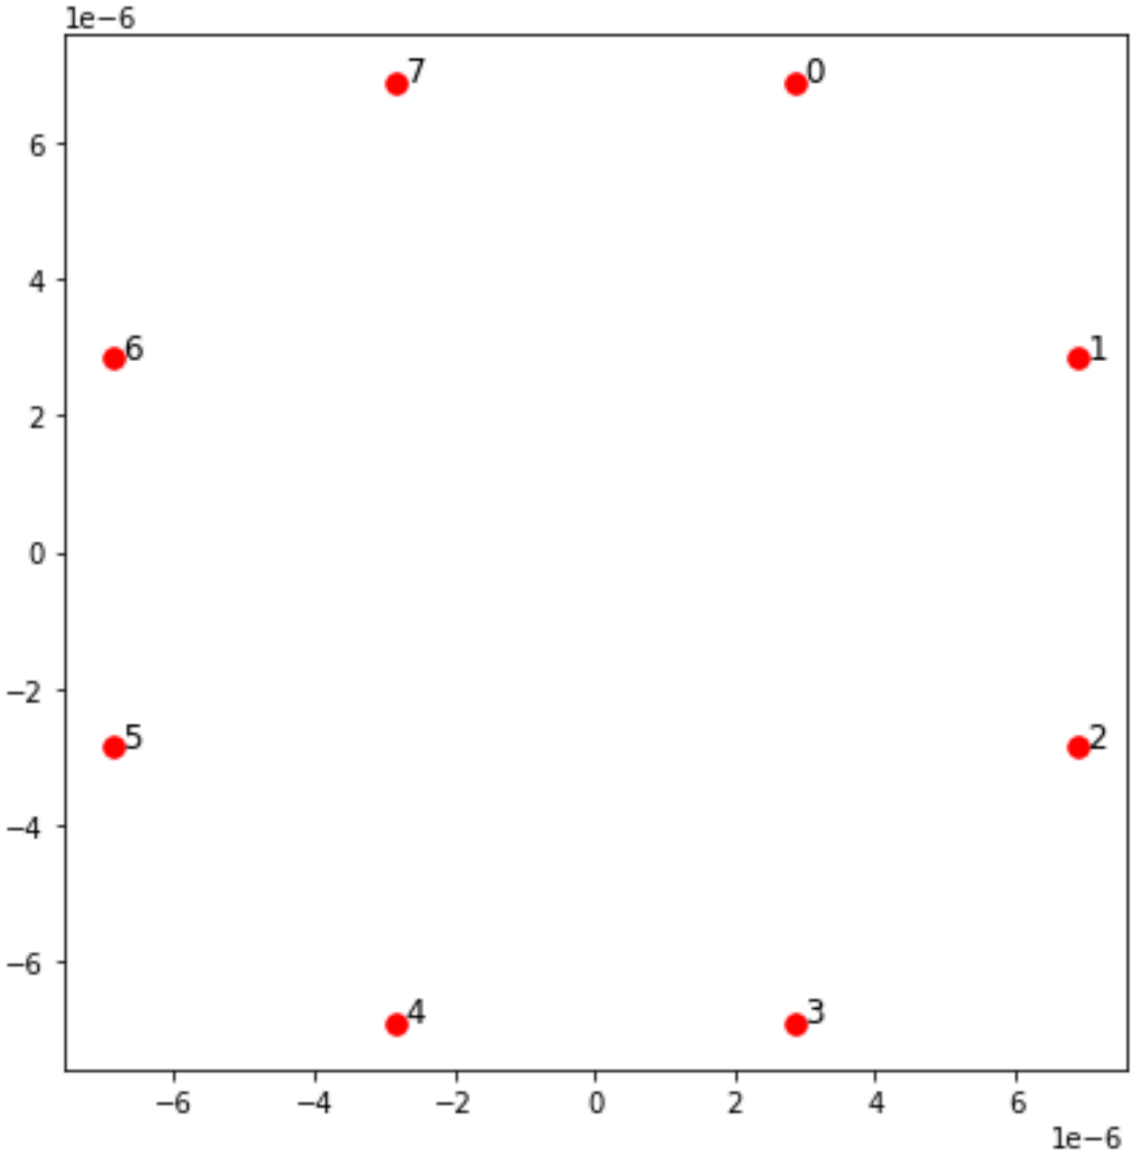

Interaction

To prepare the anti-ferromagnetic phase, we need to induce interactions between neighboring spins. We use the van der Waals interaction for this, which is natively implemented by neutral atom devices (such as the Aquila device from QuEra). Using the spin-representation, the Hamiltonian term for this interaction can be expressed as a sum over all spin pairs (j,k).

$$H_{\text{interaction}} = \sum_{j=1}^{N-1} \sum_{k=j+1}^{N} V_{j,k} n_j n_k$$

Here, nj=$|\uparrow_j\rangle\langle\#_j|$ is an operator that takes the value of 1 only if spin j is in the “up” state, and 0 otherwise. The strength is $V_{j,k}=C_6/(d_{j,k})^6$, where C_6 is the fixed coefficient, and $d_{j,k}$ is the Euclidean distance between spins j and k. The immediate effect of this interaction term is that any state where both spin j and spin k are “up” have elevated energy (by the amount $V_{j,k}$). By carefully designing the rest of the AHS program, this interaction will prevent neighboring spins from both being in the “up” state, an effect commonly known as "Rydberg blockade."

Driving field

At the beginning of the AHS program, all spins (by default) start in their "down" state, they are in a so-called ferromagnetic phase. Keeping an eye on our goal to prepare the anti-ferromagnetic phase, we specify a time-dependent coherent driving field that smoothly transitions the spins from this state to a many-body state where the "up" states are preferred. The corresponding Hamiltonian can be written as

$$H_{\text{drive}}(t) = \sum_{k=1}^{N} \frac{1}{2}\Omega(t)\left[e^{i\phi(t)}S_{-,k} + e^{-i\phi(t)}S_{+,k}\right] - \sum_{k=1}^{N} \Delta(t) n_k$$

where Ω(t),ϕ(t),Δ(t) are the time-dependent, global amplitude (aka Rabi frequency), phase, and detuning of the driving field affecting all spins uniformly. Here $S_{-,k}=|\downarrow_k\rangle\langle\#_k|$and $S_{+,k}=(S_{-,k})^{\dagger}=|\uparrow_k\rangle\langle\#_k|$ are the lowering and raising operators of spin k, respectively, and $n_k=|\uparrow_k\rangle\langle\#_k|$ is the same operator as before. The Ω part of the driving field coherently couples the "down" and the "up" states of all spins simultaneously, while the Δ part controls the energy reward for "up" states.

To program a smooth transition from the ferromagnetic phase to the anti-ferromagnetic phase, we specify the driving field with the following code.

```
from braket.timings.time_series import TimeSeries
from braket.ahs.driving_field import DrivingField

# smooth transition from "down" to "up" state
time_max = 4e-6  # seconds
time_ramp = 1e-7  # seconds
omega_max = 6300000.0  # rad / sec
delta_start = -5 * omega_max
delta_end = 5 * omega_max

omega = TimeSeries()
omega.put(0.0, 0.0)
omega.put(time_ramp, omega_max)
omega.put(time_max - time_ramp, omega_max)
omega.put(time_max, 0.0)

delta = TimeSeries()
delta.put(0.0, delta_start)
delta.put(time_ramp, delta_start)
delta.put(time_max - time_ramp, delta_end)
delta.put(time_max, delta_end)

phi = TimeSeries().put(0.0, 0.0).put(time_max, 0.0)

drive = DrivingField(
    amplitude=omega,
    phase=phi,
    detuning=delta
)
```

We can visualize the time series of the driving field with the following script.

```
fig, axes = plt.subplots(3, 1, figsize=(12, 7), sharex=True)

ax = axes[0]
time_series = drive.amplitude.time_series
ax.plot(time_series.times(), time_series.values(), '.-');
ax.grid()
ax.set_ylabel('Omega [rad/s]')

ax = axes[1]
time_series = drive.detuning.time_series
ax.plot(time_series.times(), time_series.values(), '.-');
```

```
ax.grid()
ax.set_ylabel('Delta [rad/s]')

ax = axes[2]
time_series = drive.phase.time_series
# Note: time series of phase is understood as a piecewise constant function
ax.step(time_series.times(), time_series.values(), '.-', where='post');
ax.set_ylabel('phi [rad]')
ax.grid()
ax.set_xlabel('time [s]')

plt.show()  # this will show the plot below in an ipython or jupyter session
```

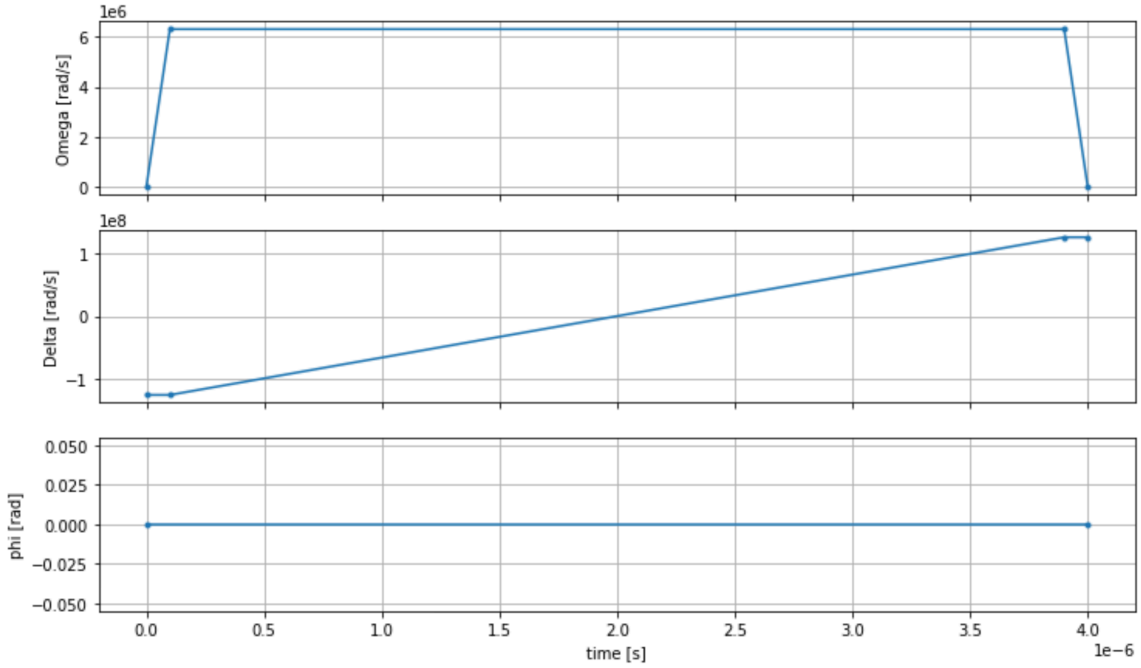

AHS program

The register, the driving field, (and the implicit van der Waals interactions) make up the analog Hamiltonian simulation program `ahs_program`.

```
from braket.ahs.analog_hamiltonian_simulation import AnalogHamiltonianSimulation

ahs_program = AnalogHamiltonianSimulation(
    register=register,
    hamiltonian=drive
)
```

Running on local simulator

Since this example is small (less than 15 spins), before running it on an AHS-compatible QPU, we can run it on the local AHS simulator which comes with the Braket SDK. Since the local simulator is available for free with the Braket SDK, this is best practice to ensure that our code can correctly execute.

Here, we can set the number of shots to a high value (say, 1 million) because the local simulator tracks the time evolution of the quantum state and draws samples from the final state; hence, increasing the number of shots, while increasing the total runtime only marginally.

```
from braket.devices import LocalSimulator
device = LocalSimulator("braket_ahs")

result_simulator = device.run(
    ahs_program,
    shots=1_000_000
).result()  # takes about 5 seconds
```

Analyzing simulator results

We can aggregate the shot results with the following function that infers the state of each spin (which may be “d” for “down”, “u” for “up”, or “e” for empty site), and counts how many times each configuration occurred across the shots.

```
from collections import Counter

def get_counts(result):
    """Aggregate state counts from AHS shot results

    A count of strings (of length = # of spins) are returned, where
    each character denotes the state of a spin (site):
       e: empty site
       u: up state spin
       d: down state spin

    Args:
       result
 (braket.tasks.analog_hamiltonian_simulation_quantum_task_result.AnalogHamiltonianSimulationQuantumTask

    Returns
        dict: number of times each state configuration is measured

    """
    state_counts = Counter()
    states = ['e', 'u', 'd']
    for shot in result.measurements:
        pre = shot.pre_sequence
        post = shot.post_sequence
        state_idx = np.array(pre) * (1 + np.array(post))
        state = "".join(map(lambda s_idx: states[s_idx], state_idx))
        state_counts.update((state,))
    return dict(state_counts)

counts_simulator = get_counts(result_simulator)  # takes about 5 seconds
print(counts_simulator)
```

```
{'udududud': 330944, 'dudududu': 329576, 'dududdud': 38033, ...}
```

Here `counts` is a dictionary that counts the number of times each state configuration is observed across the shots. We can also visualize them with the following code.

```
from collections import Counter

def has_neighboring_up_states(state):
    if 'uu' in state:
        return True
    if state[0] == 'u' and state[-1] == 'u':
        return True
    return False
```

```
def number_of_up_states(state):
    return Counter(state)['u']

def plot_counts(counts):
    non_blockaded = []
    blockaded = []
    for state, count in counts.items():
        if not has_neighboring_up_states(state):
            collection = non_blockaded
        else:
            collection = blockaded
        collection.append((state, count, number_of_up_states(state)))

    blockaded.sort(key=lambda _: _[1], reverse=True)
    non_blockaded.sort(key=lambda _: _[1], reverse=True)

    for configurations, name in zip((non_blockaded,
                                     blockaded),
                                     ('no neighboring "up" states',
                                     'some neighboring "up" states')):
        plt.figure(figsize=(14, 3))
        plt.bar(range(len(configurations)), [item[1] for item in configurations])
        plt.xticks(range(len(configurations)))
        plt.gca().set_xticklabels([item[0] for item in configurations], rotation=90)
        plt.ylabel('shots')
        plt.grid(axis='y')
        plt.title(f'{name} configurations')
        plt.show()

plot_counts(counts_simulator)
```

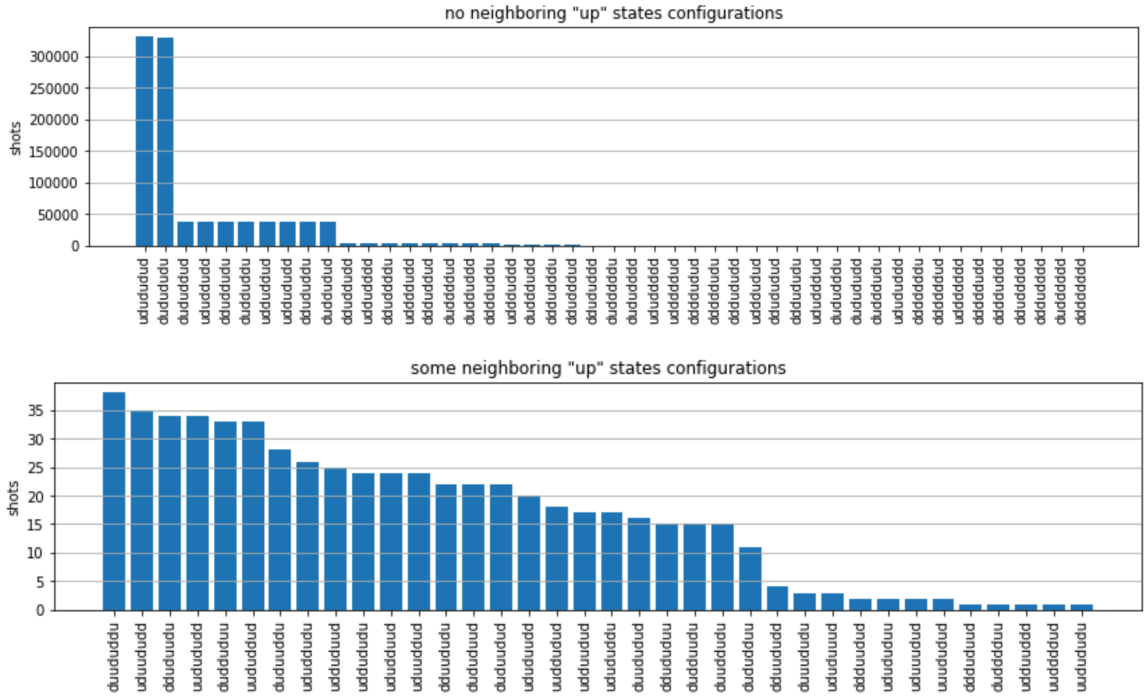

From the plots, we can read the following observations the verify that we successfully prepared the anti-ferromagnetic phase.

1. Generally, non-blockaded states (where no two neighboring spins are in the "up" state) are more common than states where at least one pair of neighboring spins are both in "up" states.
2. Generally, states with more "up" excitations are favored, unless the configuration is blockaded.
3. The most common states are indeed the perfect anti-ferromagnetic states `"dudududu"` and `"udududud"`.
4. The second most common states are the ones where there is only 3 "up" excitations with consecutive separations of 1, 2, 2. This shows that the van der Waals interaction has an affect (albeit much smaller) on next-nearest neighbors too.

Running on QuEra's Aquila QPU

Prerequisites: Apart from pip installing the Braket SDK, if you are new to Amazon Braket, please make sure that you have completed the necessary Get Started steps.

> **Note**
> If you are using a Braket hosted notebook instance, the Braket SDK comes pre-installed with the instance.

With all dependencies installed, we can connect to the Aquila QPU.

```
from braket.aws import AwsDevice

aquila_qpu = AwsDevice("arn:aws:braket:us-east-1::device/qpu/quera/Aquila")
```

To make our AHS program suitable for the QuEra machine, we need to round all values to comply with the levels of precision allowed by the Aquila QPU. (These requirements are governed by the device parameters with "Resolution" in their name. We can see them by executing `aquila_qpu.properties.dict()` in a notebook. For more details of capabilities and requirements of Aquila, see the Introduction to Aquila notebook.) We can do this by calling the `discretize` method.

```
discretized_ahs_program = ahs_program.discretize(aquila_qpu)
```

Now we can run the program (running only 100 shots for now) on the Aquila QPU.

> **Note**
> Running this program on the Aquila processor will incur a cost. The Amazon Braket SDK includes a Cost Tracker that enables customers to set cost limits as well as track their costs in near real-time.

```
task = aquila_qpu.run(discretized_ahs_program, shots=100)

metadata = task.metadata()
task_arn = metadata['quantumTaskArn']
task_status = metadata['status']

print(f"ARN: {task_arn}")
print(f"status: {task_status}")
```

```
task ARN: arn:aws:braket:us-east-1:123456789012:quantum-task/12345678-90ab-
cdef-1234-567890abcdef
task status: CREATED
```

Due to the large variance of how long a task may take to run (depending on availability windows and QPU utilization), it is a good idea to note down the task ARN, so we can check its status at a later time with the following code snippet.

```
# Optionally, in a new python session

from braket.aws import AwsQuantumTask

SAVED_TASK_ARN = "arn:aws:braket:us-east-1:123456789012:quantum-task/12345678-90ab-
cdef-1234-567890abcdef"

task = AwsQuantumTask(arn=SAVED_TASK_ARN)
metadata = task.metadata()
task_arn = metadata['quantumTaskArn']
task_status = metadata['status']

print(f"ARN: {task_arn}")
print(f"status: {task_status}")
```

```
*[Output]*
task ARN: arn:aws:braket:us-east-1:123456789012:quantum-task/12345678-90ab-
cdef-1234-567890abcdef
task status: COMPLETED
```

Once the status is COMPLETED (which can also be checked from the tasks page of the Amazon Braket console), we can query the results with:

```
result_aquila = task.result()
```

Analyzing QPU results

Using the same `get_counts` functions as before, we can compute the counts:

```
counts_aquila = get_counts(result_aquila)
print(counts_aquila)
```

```
*[Output]*
{'udududud': 24, 'dudududu': 17, 'dududdud': 3, ...}
```

and plot them with `plot_counts`:

```
plot_counts(counts_aquila)
```

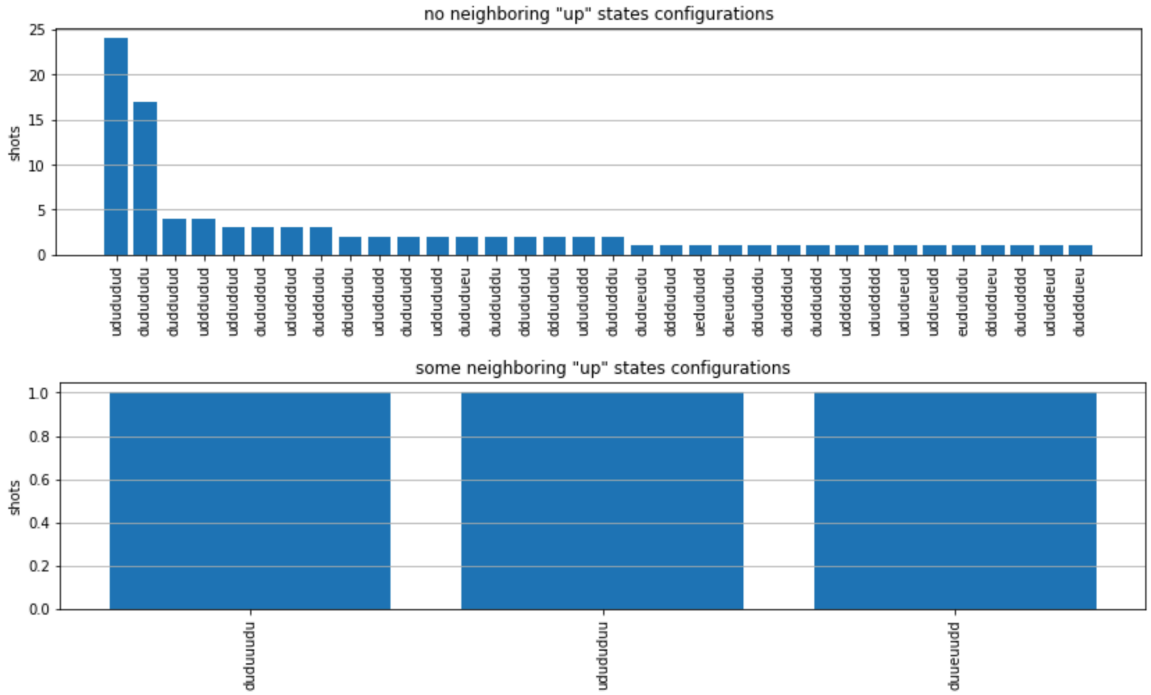

Note that a small fraction of shots have empty sites (marked with “e”). This is due to a 1—2% per atom preparation imperfections of the Aquila QPU. Apart from this, the results match with the simulation within the expected statistical fluctuation due to small number of shots.

Next

Congratulations, you have now run your first AHS workload on Amazon Braket using the local AHS simulator and the Aquila QPU.

To learn more about Rydberg physics, analog Hamiltonian simulation and the Aquila device, refer to our example notebooks.

Work with Amazon Braket

This section shows you how to design quantum circuits, submit these problems as tasks to devices, and monitor the tasks with the Amazon Braket SDK.

The following are the main means of interacting with resources on Amazon Braket.

- The Amazon Braket Console provides device information and status to help you create, manage, and monitor your resources and tasks.
- Submit and run quantum tasks through the Amazon Braket Python SDK, as well as through the console. The SDK is accessible through preconfigured Amazon Braket notebooks.
- The Amazon Braket API is accessible through the Amazon Braket Python SDK and notebooks. You can make calls directly to the API if you're building applications that work with quantum computing programmatically.

The examples throughout this section demonstrate how you can work with the Amazon Braket API directly using the Amazon Braket Python SDK along with the AWS Python SDK for Braket (Boto3).

More about the Amazon Braket Python SDK

To work with the Amazon Braket Python SDK, first install the AWS Python SDK for Braket (Boto3) so that you can communicate with the AWS API. You can think of the Amazon Braket Python SDK as a convenient wrapper around Boto3 for quantum customers.

- Boto3 contains interfaces you need to tap into the AWS API. (Note that Boto3 is a large Python SDK that talks to the AWS API. Most AWS services support a Boto3 interface.)
- The Amazon Braket Python SDK contains software modules for circuits, gates, devices, result types, and other parts of a quantum task. Each time you create a program, you import the modules you need for that task.
- The Amazon Braket Python SDK is accessible through notebooks, which are pre-loaded with all of the modules and dependencies you need for running quantum tasks.
- You can import modules from the Amazon Braket Python SDK into any Python script if you do not wish to work with notebooks.

After you've installed Boto3, an overview of steps for creating a task through the Amazon Braket Python SDK resembles the following:

1. (Optionally) Open your notebook.
2. Import the SDK modules you need for your circuits.
3. Specify a QPU or simulator.
4. Instantiate the circuit.
5. Run the circuit.
6. Collect the results.

The examples in this section show details of each step.

For more examples, see the Amazon Braket Examples repository on GitHub.

In this section:

Construct circuits in the SDK

This section provides examples of defining a circuit, viewing available gates, extending a circuit, and viewing gates that each device supports. It also contains instructions on how to manually allocate qubits, instruct the compiler to run your circuits exactly as defined, and build noisy circuits with a noise simulator.

Amazon Braket also supports working directly at the pulse level defined gates with certain QPUs. To learn more, refer to the Pulse Control on Amazon Braket (p. 109) section.

In this section:

Gates and circuits

Quantum gates and circuits are defined in the `braket.circuits` class of the Amazon Braket Python SDK. From the SDK, you can instantiate a new circuit object by calling `Circuit()`.

Example: Define a circuit

The example starts by defining a sample circuit of four qubits (labelled `q0`, `q1`, `q2`, and `q3`) consisting of standard, single-qubit Hadamard gates and two-qubit CNOT gates. You can visualize this circuit by calling the `print` function as the following example shows.

```
# import the circuit module
from braket.circuits import Circuit

# define circuit with 4 qubits
my_circuit = Circuit().h(range(4)).cnot(control=0, target=2).cnot(control=1, target=3)
print(my_circuit)
```

```
T  : |0| 1 |

q0 : -H-C---
        |
q1 : -H-|-C-
        | |
q2 : -H-X-|-
          |
```

```
q3 : -H---X-

T  : |0| 1 |
```

Example: See all available gates

The following example shows how to look at all the available gates in Amazon Braket.

```
import string
from braket.circuits import Gate
# print all available gates in Amazon Braket
gate_set = [attr for attr in dir(Gate) if attr[0] in string.ascii_uppercase]
print(gate_set)
```

The output from this code lists all of the gates.

```
['CCNot', 'CNot', 'CPhaseShift', 'CPhaseShift00', 'CPhaseShift01', 'CPhaseShift10',
 'CSwap', 'CV', 'CY', 'CZ', 'ECR', 'H', 'I', 'ISwap', 'PSwap', 'PhaseShift', 'Rx', 'Ry',
 'Rz', 'S', 'Si', 'Swap', 'T', 'Ti', 'Unitary', 'V', 'Vi', 'X', 'XX', 'XY', 'Y', 'YY', 'Z',
 'ZZ']
```

Any of these gates can be appended to a circuit by calling the method for that type of circuit. For example, you'd call `circ.h(0)`, to add a Hadamard gate to the first qubit.

> **Note**
> Gates are appended in place, and the example that follows adds all of the gates listed in the previous example to the same circuit.

```
circ = Circuit()
# toffoli gate with q0, q1 the control qubits and q2 the target.
circ.ccnot(0, 1, 2)
# cnot gate
circ.cnot(0, 1)
# controlled-phase gate that phases the |11> state, cphaseshift(phi) =
 diag((1,1,1,exp(1j*phi))), where phi=0.15 in the examples below
circ.cphaseshift(0, 1, 0.15)
# controlled-phase gate that phases the |00> state, cphaseshift00(phi) =
 diag([exp(1j*phi),1,1,1])
circ.cphaseshift00(0, 1, 0.15)
# controlled-phase gate that phases the |01> state, cphaseshift01(phi) =
 diag([1,exp(1j*phi),1,1])
circ.cphaseshift01(0, 1, 0.15)
# controlled-phase gate that phases the |10> state, cphaseshift10(phi) =
 diag([1,1,exp(1j*phi),1])
circ.cphaseshift10(0, 1, 0.15)
# controlled swap gate
circ.cswap(0, 1, 2)
# swap gate
circ.swap(0,1)
# phaseshift(phi)= diag([1,exp(1j*phi)])
circ.phaseshift(0,0.15)
# controlled Y gate
circ.cy(0, 1)
# controlled phase gate
circ.cz(0, 1)
# Echoed cross-resonance gate applied to q0, q1
circ = Circuit().ecr(0,1)
# X rotation with angle 0.15
circ.rx(0, 0.15)
# Y rotation with angle 0.15
circ.ry(0, 0.15)
```

```
# Z rotation with angle 0.15
circ.rz(0, 0.15)
# Hadamard gates applied to q0, q1, q2
circ.h(range(3))
# identity gates applied to q0, q1, q2
circ.i([0, 1, 2])
# iswap gate, iswap = [[1,0,0,0],[0,0,1j,0],[0,1j,0,0],[0,0,0,1]]
circ.iswap(0, 1)
# pswap gate, PSWAP(phi) = [[1,0,0,0],[0,0,exp(1j*phi),0],[0,exp(1j*phi),0,0],[0,0,0,1]]
circ.pswap(0, 1, 0.15)
# X gate applied to q1, q2
circ.x([1, 2])
# Y gate applied to q1, q2
circ.y([1, 2])
# Z gate applied to q1, q2
circ.z([1, 2])
# S gate applied to q0, q1, q2
circ.s([0, 1, 2])
# conjugate transpose of S gate applied to q0, q1
circ.si([0, 1])
# T gate applied to q0, q1
circ.t([0, 1])
# conjugate transpose of T gate applied to q0, q1
circ.ti([0, 1])
# square root of not gate applied to q0, q1, q2
circ.v([0, 1, 2])
# conjugate transpose of square root of not gate applied to q0, q1, q2
circ.vi([0, 1, 2])
# exp(i(XX+YY) theta/4), where theta=0.15 in the examples below
circ.xx(0, 1, 0.15)
# exp(-iXY theta/2)
circ.xy(0, 1, 0.15)
# exp(-iYY theta/2)
circ.yy(0, 1, 0.15)
# exp(-iZZ theta/2)
circ.zz(0, 1, 0.15)
```

Apart from the pre-defined gate set, you also can apply self-defined unitary gates to the circuit. These can be single-qubit gates (as shown in the following source code) or multi-qubit gates applied to the qubits defined by the `targets` parameter.

```
import numpy as np
# apply a general unitary
my_unitary = np.array([[0, 1],[1, 0]])
circ.unitary(matrix=my_unitary, targets=[0])
```

Example: Extend existing circuits

You can extend existing circuits by adding instructions. An `Instruction` is a quantum directive that describes the task to perform on a quantum device. `Instruction` operators include objects of type `Gate` only.

```
# import the Gate and Instruction modules
from braket.circuits import Gate, Instruction

# add instructions directly.
circ = Circuit([Instruction(Gate.H(), 4), Instruction(Gate.CNot(), [4, 5])])

# or with add_instruction/add functions
instr = Instruction(Gate.CNot(), [0, 1])
circ.add_instruction(instr)
circ.add(instr)
```

```
# specify where the circuit is appended
circ.add_instruction(instr, target=[3, 4])
circ.add_instruction(instr, target_mapping={0: 3, 1: 4})

# print the instructions
print(circ.instructions)
# if there are multiple instructions, you can print them in a for loop
for instr in circ.instructions:
     print(instr)

# instructions can be copied
new_instr = instr.copy()
# appoint the instruction to target
new_instr = instr.copy(target=[5])
new_instr = instr.copy(target_mapping={0: 5})
```

Example: View the gates that each device supports

Simulators support all gates in the Braket SDK, but QPU devices support a smaller subset. You can find the supported gates of a device in the device properties.

```
# import the device module
from braket.aws import AwsDevice

device = AwsDevice("arn:aws:braket:::device/qpu/ionq/ionQdevice")

# get device name
device_name = device.name
# show supportedQuantumOperations (supported gates for a device)
device_operations = device.properties.dict()['action']['braket.ir.jaqcd.program']
['supportedOperations']
print('Quantum Gates supported by {}:\n {}'.format(device_name, device_operations))
```

```
Quantum Gates supported by IonQ Device:
 ['x', 'y', 'z', 'rx', 'ry', 'rz', 'h', 'cnot', 's', 'si', 't', 'ti', 'v', 'vi', 'xx',
 'yy', 'zz', 'swap', 'i']
```

```
device = AwsDevice("arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2")

# get device name
device_name = device.name
# show supportedQuantumOperations (supported gates for a device)
device_operations = device.properties.dict()['action']['braket.ir.jaqcd.program']
['supportedOperations']
print('Quantum Gates supported by {}:\n {}'.format(device.name, device_operations))
```

```
Quantum Gates supported by Aspen-M-2:
 ['cz', 'xy', 'ccnot', 'cnot', 'cphaseshift', 'cphaseshift00', 'cphaseshift01',
 'cphaseshift10', 'cswap', 'h', 'i', 'iswap', 'phaseshift', 'pswap', 'rx', 'ry', 'rz', 's',
 'si', 'swap', 't', 'ti', 'x', 'y', 'z']
```

Supported gates may need to be compiled into native gates before they can run on quantum hardware. When you submit a circuit, Amazon Braket performs this compilation automatically.

Manual qubit allocation

When you run a quantum circuit on quantum computers from Rigetti, you can optionally use manual qubit allocation to control which qubits are used for your algorithm. The Amazon Braket Console and

the Amazon Braket SDK help you to inspect the most recent calibration data of your selected quantum processing unit (QPU) device, so you can select the best qubits for your experiment.

Manual qubit allocation enables you to run circuits with greater accuracy and to investigate individual qubit properties. Researchers and advanced users optimize their circuit design based on the latest device calibration data and can obtain more accurate results.

The following example demonstrates how to allocate qubits explicitly.

```
circ = Circuit().h(0).cnot(0, 7)  # Indices of actual qubits in the QPU
my_task = device.run(circ, s3_location, shots=100, disable_qubit_rewiring=True)
```

For more information, see the Amazon Braket examples on GitHub, or more specifically, this notebook: Allocating Qubits on QPU Devices.

> **Note**
> The OQC compiler does not support setting `disable_qubit_rewiring=True`. Setting this flag to `True` yields the following error: `An error occurred (ValidationException) when calling the CreateQuantumTask operation: Device arn:aws:braket:eu-west-2::device/qpu/oqc/Lucy does not support disabled qubit rewiring`.

Verbatim compilation

When you run a quantum circuit on quantum computers from Rigetti, IonQ, or Oxford Quantum Circuits (OQC), you can direct the compiler to run your circuits exactly as defined without any modifications. Using verbatim compilation, you can specify either that an entire circuit be preserved precisely (supported by Rigetti, IonQ, and OQC) as specified or that only specific parts of it be preserved (supported by Rigetti only). When developing algorithms for hardware benchmarking or error mitigation protocols, you need have the option to exactly specify the gates and circuit layouts that you're running on the hardware. Verbatim compilation gives you direct control over the compilation process by turning off certain optimization steps, thereby ensuring that your circuits run exactly as designed.

Verbatim compilation is currently supported on Rigetti, IonQ, and Oxford Quantum Circuits (OQC) devices and requires the use of native gates. When using verbatim compilation, it is advisable to check the topology of the device to ensure that gates are called on connected qubits and that the circuit uses the native gates supported on the hardware. The following example shows how to programmatically access the list of native gates supported by a device.

```
device.properties.paradigm.nativeGateSet
```

For Rigetti, qubit rewiring must be turned off by setting `disableQubitRewiring=True` for use with verbatim compilation. If `disableQubitRewiring=False` is set when using verbatim boxes in a compilation, the quantum circuit fails validation and does not run.

If verbatim compilation is enabled for a circuit and run on a QPU that does not support it, an error is generated indicating that an unsupported operation has caused the task to fail. As more quantum hardware natively support compiler functions, this feature will be expanded to include these devices. Devices that support verbatim compilation include it as a supported operation when queried with the following code.

```
from braket.aws import AwsDevice
from braket.device_schema.device_action_properties import DeviceActionType
device = AwsDevice("arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2")
device.properties.action[DeviceActionType.OPENQASM].supportedPragmas
```

There is no additional cost associated with using verbatim compilation. You continue to be charged for tasks executed on Braket QPU devices, notebook instances, and managed simulators based on

current rates as specified on the Amazon Braket Pricing page. For more information, see the Verbatim compilation example notebook.

> **Note**
> If you are using OpenQASM to write your circuits for the OQC and IonQ devices, and you wish to map your circuit directly to the physical qubits, you need to use the `#pragma braket verbatim` as the `disableQubitRewiring` flag is completely ignored by OpenQASM.

Noise simulation

To instantiate the local noise simulator you can change the backend as follows.

```
device = LocalSimulator(backend="braket_dm")
```

You can build noisy circuits in two ways:

1. Build the noisy circuit from the bottom up.
2. Take an existing, noise-free circuit and inject noise throughout.

The following example shows the approaches using a simple circuit with depolarizing noise and a custom Kraus channel.

```
# Bottom up approach
# apply depolarizing noise to qubit 0 with probability of 0.1
circ = Circuit().x(0).x(1).depolarizing(0, probability=0.1)

# create an arbitrary 2-qubit Kraus channel
E0 = scipy.stats.unitary_group.rvs(4) * np.sqrt(0.8)
E1 = scipy.stats.unitary_group.rvs(4) * np.sqrt(0.2)
K = [E0, E1]

# apply a two-qubit Kraus channel to qubits 0 and 2
circ = circ.kraus([0,2], K)
```

```
# Inject noise approach
# define phase damping noise
noise = Noise.PhaseDamping(gamma=0.1)
# the noise channel is applied to all the X gates in the circuit
circ = Circuit().x(0).y(1).cnot(0,2).x(1).z(2)
circ_noise = circ.copy()
circ_noise.apply_gate_noise(noise, target_gates = Gate.X)
```

Running a circuit is the same user experience as before, as shown in the following two examples.

Example 1

```
task = device.run(circ, s3_location)
```

Or

Example 2

```
task = device.run(circ_noise, s3_location)
```

For more examples, see the Braket introductory noise simulator example

Inspecting the circuit

Quantum circuits in Amazon Braket have a pseudo-time concept called `Moments`. Each qubit can experience a single gate per `Moment`. The purpose of `Moments` is to make circuits and their gates easier to address and to provide a temporal structure.

Note
Moments generally do not correspond to the real time at which gates are executed on a QPU.

The depth of a circuit is given by the total number of Moments in that circuit. You can view the circuit depth calling the method `circuit.depth` as shown in the following example.

```
# define a circuit with parametrized gates
circ = Circuit().rx(0, 0.15).ry(1, 0.2).cnot(0,2).zz(1, 3, 0.15).x(0)
print(circ)
print('Total circuit depth:', circ.depth)
```

```
T  : |   0    |    1     |2|

q0 : -Rx(0.15)-C----------X-
                |
q1 : -Ry(0.2)--|-ZZ(0.15)---
               | |
q2 : ----------X-|----------
                 |
q3 : ------------ZZ(0.15)---

T  : |   0    |    1     |2|
Total circuit depth: 3
```

The total circuit depth of the circuit above is 3 (shown as moments `0`, 1, and 2). You can check the gate operation for each moment.

`Moments` functions as a dictionary of *key-value* pairs.

- The key is `MomentsKey()`, which contains pseudo-time and qubit information.
- The value is assigned in the type of `Instructions()`.

```
moments = circ.moments
for key, value in moments.items():
    print(key)
    print(value, "\n")
```

```
MomentsKey(time=0, qubits=QubitSet([Qubit(0)]))
Instruction('operator': Rx('angle': 0.15, 'qubit_count': 1), 'target':
 QubitSet([Qubit(0)]))

MomentsKey(time=0, qubits=QubitSet([Qubit(1)]))
Instruction('operator': Ry('angle': 0.2, 'qubit_count': 1), 'target': QubitSet([Qubit(1)]))

MomentsKey(time=1, qubits=QubitSet([Qubit(0), Qubit(2)]))
Instruction('operator': CNot('qubit_count': 2), 'target': QubitSet([Qubit(0), Qubit(2)]))

MomentsKey(time=1, qubits=QubitSet([Qubit(1), Qubit(3)]))
Instruction('operator': ZZ('angle': 0.15, 'qubit_count': 2), 'target': QubitSet([Qubit(1),
 Qubit(3)]))
```

```
MomentsKey(time=2, qubits=QubitSet([Qubit(0)]))
Instruction('operator': X('qubit_count': 1), 'target': QubitSet([Qubit(0)]))
```

You can also add gates to a circuit through `Moments`.

```
new_circ = Circuit()
instructions = [Instruction(Gate.S(), 0),
                Instruction(Gate.CZ(), [1,0]),
                Instruction(Gate.H(), 1)
]
new_circ.moments.add(instructions)
print(new_circ)
```

```
T  : |0|1|2|

q0 : -S-Z---
       |
q1 : ---C-H-

T  : |0|1|2|
```

Result types

Amazon Braket can return different types of results when a circuit is measured using `ResultType`. A circuit can return the following types of results.

- `Amplitude` returns the amplitude of specified quantum states in the output wave function. It is available on the SV1 and local simulators only.
- `Expectation` returns the expectation value of a given observable, which can be specified with the `Observable` class introduced later in this chapter. The target qubits used to measure the observable must be specified, and the number of specified targets must equal the number of qubits on which the observable acts. If no targets are specified, the observable must operate only on 1 qubit and it is applied to all qubits in parallel.
- `Probability` returns the probabilities of measuring computational basis states. If no targets are specified, `Probability` returns the probability of measuring all basis states. If targets are specified, only the marginal probabilities of the basis vectors on the specified qubits are returned.
- `Reduced density matrix` returns a density matrix for a subsystem of specified target qubits from a system of qubits. To limit the size of this result type, Braket limits the number of target qubits to a maximum of 8.
- `StateVector` returns the full state vector. It is available on the local simulator.
- `Sample` returns the measurement counts of a specified target qubit set and observable. If no targets are specified, the observable must operate only on 1 qubit and it is applied to all qubits in parallel. If targets are specified, the number of specified targets must equal the number of qubits on which the observable acts.
- `Variance` returns the variance (`mean([x-mean(x)]`^{2}`)`) of the specified target qubit set and observable as the requested result type. If no targets are specified, the observable must operate only on 1 qubit and it is applied to all qubits in parallel. Otherwise, the number of targets specified must equal the number of qubits to which the observable can be applied.

The supported result types for different devices:

	Local sim	SV1	DM1	TN1	Rigetti	IonQ	OQC

Amplitude	Y	Y	N	N	N	N	N
Expectation	Y	Y	Y	Y	Y	Y	Y
Probability	Y	Y	Y	N	Y*	Y	Y
Reduced density matrix	Y	N	Y	N	N	N	N
State vector	Y	N	N	N	N	N	N
Sample	Y	Y	Y	Y	Y	Y	Y
Variance	Y	Y	Y	Y	Y	Y	Y

Note
* Rigetti only supports probability result types of up to 40 qubits.

You can check the supported result types by examining the device properties, as shown in the following example.

```
device = AwsDevice("arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2")

# print the result types supported by this device
for iter in device.properties.action['braket.ir.jaqcd.program'].supportedResultTypes:
    print(iter)
```

```
name='Sample' observables=['x', 'y', 'z', 'h', 'i'] minShots=10 maxShots=100000
name='Expectation' observables=['x', 'y', 'z', 'h', 'i'] minShots=10 maxShots=100000
name='Variance' observables=['x', 'y', 'z', 'h', 'i'] minShots=10 maxShots=100000
name='Probability' observables=None minShots=10 maxShots=100000
```

To call a `ResultType`, append it to a circuit, as shown in the following example.

```
from braket.circuits import Observable

circ = Circuit().h(0).cnot(0, 1).amplitude(state=["01", "10"])
circ.probability(target=[0, 1])
circ.probability(target=0)
circ.expectation(observable=Observable.Z(), target=0)
circ.sample(observable=Observable.X(), target=0)
circ.state_vector()
circ.variance(observable=Observable.Z(), target=0)

# print one of the result types assigned to the circuit
print(circ.result_types[0])
```

Observables

Amazon Braket includes an `Observable` class, which can be used to specify an observable to be measured.

You can apply at most one unique non-identity observable to each qubit. If you specify two or more different non-identity observables to the same qubit, you see an error. For this purpose, each factor of a tensor product counts as an individual observable, so it is permissible to have multiple tensor products acting on the same qubit, provided that the factor acting on that qubit is the same.

The `Observable` class includes the following observables.

```
Observable.I()
Observable.H()
Observable.X()
Observable.Y()
Observable.Z()

# get the eigenvalues of the observable
print("Eigenvalue:", Observable.H().eigenvalues)
# or whether to rotate the basis to be computational basis
print("Basis rotation gates:",Observable.H().basis_rotation_gates)

# get the tensor product of observable for the multi-qubit case
tensor_product = Observable.Y() @ Observable.Z()
# view the matrix form of an observable by using
print("The matrix form of the observable:\n",Observable.Z().to_matrix())
print("The matrix form of the tensor product:\n",tensor_product.to_matrix())

# also factorize an observable in the tensor form
print("Factorize an observable:",tensor_product.factors)

# self-define observables given it is a Hermitian
print("Self-defined Hermitian:",Observable.Hermitian(matrix=np.array([[0, 1],[1, 0]])))
```

```
Eigenvalue: [ 1 -1]
Basis rotation gates: (Ry('angle': -0.7853981633974483, 'qubit_count': 1),)
The matrix form of the observable:
 [[ 1.+0.j  0.+0.j]
 [ 0.+0.j -1.+0.j]]
The matrix form of the tensor product:
 [[ 0.+0.j  0.+0.j  0.-1.j  0.-0.j]
 [ 0.+0.j -0.+0.j  0.-0.j  0.+1.j]
 [ 0.+1.j  0.+0.j  0.+0.j  0.+0.j]
 [ 0.+0.j -0.-1.j  0.+0.j -0.+0.j]]
Factorize an observable: (Y('qubit_count': 1), Z('qubit_count': 1))
Self-defined Hermitian: Hermitian('qubit_count': 1, 'matrix': [[0.+0.j 1.+0.j], [1.+0.j 0.
+0.j]])
```

Run your circuits with OpenQASM 3.0

Amazon Braket now supports OpenQASM 3.0 for gate-based quantum devices and simulators. This user guide provides information about the subset of OpenQASM 3.0 supported by Braket. Braket customers now have the choice of submitting Braket circuits with the SDK (p. 53) or by directly providing OpenQASM 3.0 strings to all gate-based devices with the Amazon Braket API and the Amazon Braket Python SDK.

The topics in this guide walk you through various examples of how to complete the following tasks.

- Create and submit OpenQASM tasks on different Braket devices (p. 67)
- Access the supported operations and result types (p. 75)
- Simulate noise with OpenQASM (p. 77)
- Use verbatim compilation with OpenQASM (p. 79)
- Troubleshoot OpenQASM issues

This guide also provides an introduction to certain hardware-specific features that can be implemented with OpenQASM 3.0 on Braket and links to further resources.

In this section:

What is OpenQASM 3.0?

The Open Quantum Assembly Language (OpenQASM) is an intermediate representation for quantum instructions. OpenQASM is an open-source framework and is widely used for the specification of quantum programs for gate-based devices. With OpenQASM, users can program the quantum gates and measurement operations that form the building blocks of quantum computation. The previous version of OpenQASM (2.0) was used by a number of quantum programming libraries to describe simple programs.

The new version of OpenQASM (3.0) extends the previous version to include more features, such as pulse-level control, gate timing, and classical control flow to bridge the gap between end-user interface and hardware description language. Details and specification on the current version 3.0 are available on the GitHub OpenQASM 3.x Live Specification. OpenQASM's future development is governed by the OpenQASM 3.0 Technical Steering Committee, of which AWS is a member alongside IBM, Microsoft, and the University of Innsbruck.

When to use OpenQASM 3.0

OpenQASM provides an expressive framework to specify quantum programs through low-level controls that are not architecture specific, making it well suited as a representation across multiple gate-based devices. The Braket support for OpenQASM furthers its adoption as a consistent approach to developing gate-based quantum algorithms, reducing the need for users to learn and maintain libraries in multiple frameworks.

If you have existing libraries of programs in OpenQASM 3.0, you can adapt them for use with Braket rather than completely rewriting these circuits. Researchers and developers should also benefit from an increasing number of available third-party libraries with support for algorithm development in OpenQASM.

How OpenQASM 3.0 works

Support for OpenQASM 3.0 from Braket provides feature parity with the current Intermediate Representation. This means that anything you can do today on hardware devices and managed simulators with Braket, you can do with OpenQASM using the Braket API. You can run OpenQASM 3.0 programs by directly supplying OpenQASM strings to all gate-based devices in a manner that is similar to how circuits are currently supplied to devices on Braket. Braket users can also integrate third-party libraries that support OpenQASM 3.0. The rest of this guide details how to develop OpenQASM representations for use with Braket.

Prerequisites

To use OpenQASM 3.0 on Amazon Braket, you must have version v1.8.0 of the Amazon Braket Python Schemas and v1.17.0 or higher of the Amazon Braket Python SDK.

If you are a first time user of Amazon Braket, you need to enable Amazon Braket. For instructions, see Enable Amazon Braket.

What OpenQASM features does Braket support?

The following section lists the OpenQASM 3.0 data types, statements, and pragma instructions supported by Braket.

In this section:

Supported OpenQASM data types

The following OpenQASM data types are supported by Amazon Braket.

- Non-negative integers are used for (virtual and physical) qubit indices:
  - `cnot q[0], q[1];`
  - `h $0;`
- Floating-point numbers or constants may be used for gate rotation angles:
  - `rx(-0.314) $0;`
  - `rx(pi/4) $0;`
- Arrays of complex numbers (with the OpenQASM `im` notation for imaginary part) are allowed in result type pragmas for defining general hermitian observables and in unitary pragmas:
  - `#pragma braket unitary [[0, -1im], [1im, 0]] q[0]`
  - `#pragma braket result expectation hermitian([[0, -1im], [1im, 0]]) q[0]`

Supported OpenQASM statements

The following OpenQASM statements are supported by Amazon Braket.

- `Header: OPENQASM 3;`
- Classic bit declarations:
  - `bit b1;` (equivalently, `creg b1;`)
  - `bit[10] b2;` (equivalently, `creg b2[10];`)
- Qubit declarations:
  - `qubit b1;` (equivalently, `qreg b1;`)
  - `qubit[10] b2;` (equivalently, `qreg b2[10];`)
- Indexing within arrays: `q[0]`
- specification of physical qubits: `$0`

- Supported gates and operations on a device:
  - `h $0;`
  - `iswap q[0], q[1];`

  Note
  A device's supported gates can be found in the device properties for OpenQASM actions; no gate definitions are needed to use these gates.

- Verbatim box statements. Currently, we do not support box duration notation. Native gates and physical qubits are required in verbatim boxes.

```
#pragma braket verbatim
box{
    rx(0.314) $0;
}
```

- Measurement and measurement assignment on qubits or a whole qubit register.
  - `measure $0;`
  - `measure q;`
  - `measure q[0];`
  - `b = measure q;`
  - `measure q # b;`

Braket OpenQASM pragmas

The following OpenQASM pragma instructions are supported by Amazon Braket.

- Noise pragmas
  - `#pragma braket noise bit_flip(0.2) q[0]`
  - `#pragma braket noise phase_flip(0.1) q[0]`
  - `#pragma braket noise pauli_channel`
- Verbatim pragmas
  - `#pragma braket verbatim`
- Result type pragmas
  - Basis invariant result types:
    - State vector: `#pragma braket result state_vector`
    - Density matrix: `#pragma braket result density_matrix`
  - Z basis result types:
    - Amplitude: `#pragma braket result amplitude "01"`
    - Probability: `#pragma braket result probability q[0], q[1]`
  - Basis rotated result types
    - Expectation: `#pragma braket result expectation x(q[0]) @ y([q1])`
    - Variance: `#pragma braket result variance hermitian([[0, -1im], [1im, 0]]) $0`
    - Sample: `#pragma braket result sample h($1)`

  Note
  OpenQASM 3.0 is backwards compatible with OpenQASM 2.0, so programs written using 2.0 can run on Braket. However the features of OpenQASM 3.0 supported by Braket do have some

minor syntax differences, such as `qreg` vs `creg` and `qubit` vs `bit`. There are also differences in measurement syntax, and these need to be supported with their correct syntax.

Advanced feature support for OpenQASM on the Local Simulator

The `LocalSimulator` supports advanced OpenQASM features which are not offered as part of Braket's QPU's or managed simulators. The following list of features are only supported in the `LocalSimulator`:

- Gate modifiers
- OpenQASM built-in gates
- Classical variables
- Classical operations
- Custom gates
- Classical control
- Input
- QASM files
- Subroutines

For examples of each advanced feature, see this sample notebook. For the full OpenQASM specification, see the OpenQASM website.

Supported Operations and Grammar with OpenPulse

Supported OpenPulse Data Types

Cal blocks:

```
cal {
    ...
}
```

Defcal blocks:

```
// 1 qubit
defcal x $0 {
...
}

// 1 qubit w. input parameters as constants
defcal my_rx(pi) $0 {
...
}

// 1 qubit w. input parameters as free parameters
defcal my_rz(angle theta) $0 {
...
}

// 2 qubit (above gate args are also valid)
defcal cz $1, $0 {
...
}
```

Frames:

```
frame my_frame = newframe(port_0, 4.5e9, 0.0);
```

Waveforms:

```
// prebuilt
waveform my_waveform_1 = constant(1e-6, 1.0);

//arbitrary
waveform my_waveform_2 = {0.1 + 0.1im, 0.1 + 0.1im, 0.1, 0.1};
```

Custom Gate Calibration Example:

```
cal {
    waveform wf1 = constant(1e-6, 0.25);
}

defcal my_x $0 {
   play(wf1, q0_rf_frame);
}

defcal my_cz $1, $0 {
    barrier q0_q1_cz_frame, q0_rf_frame;
    play(q0_q1_cz_frame, wf1);
    delay[300ns] q0_rf_frame
    shift_phase(q0_rf_frame, 4.366186381749424);
    delay[300ns] q0_rf_frame;
    shift_phase(q0_rf_frame.phase, 5.916747563126659);
    barrier q0_q1_cz_frame, q0_rf_frame;
    shift_phase(q0_q1_cz_frame, 2.183093190874712);
}

bit[2] ro;
my_x $0;
my_cz $1,$0;
c[0] = measure $0;
```

Arbitrary pulse example:

```
bit[2] ro;
cal {
    waveform wf1 = {0.1 + 0.1im, 0.1 + 0.1im, 0.1, 0.1};
    barrier q0_drive, q0_q1_cross_resonance;
    play(q0_q1_cross_resonance, wf1);
    delay[300ns] q0_drive;
    shift_phase(q0_drive, 4.366186381749424);
    delay[300dt] q0_drive;
   barrier q0_drive, q0_q1_cross_resonance;
   play(q0_q1_cross_resonance, wf1);
    ro[0] = capture_v0(r0_measure);
    ro[1] = capture_v0(r1_measure);
}
```

Create and submit an example OpenQASM 3.0 task

You can use the Amazon Braket Python SDK, Boto3, or the AWS CLI to submit OpenQASM 3.0 tasks to an Amazon Braket device.

In this section:

An example OpenQASM 3.0 program

To create a OpenQASM 3.0 task, you can start with a simple OpenQASM 3.0 program (ghz.qasm) that prepares a GHZ state as shown in the following example.

```
// ghz.qasm
// Prepare a GHZ state
OPENQASM 3;

qubit[3] q;
bit[3] c;

h q[0];
cnot q[0], q[1];
cnot q[1], q[2];

c = measure q;
```

Use the Python SDK to create OpenQASM 3.0 tasks

You can use the Amazon Braket Python SDK to submit this program to an Amazon Braket device with the following code.

```
with open("ghz.qasm", "r") as ghz:
    ghz_qasm_string = ghz.read()

# import the device module
from braket.aws import AwsDevice
# choose the Rigetti device
device = AwsDevice("arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2")
from braket.ir.openqasm import Program

program = Program(source=ghz_qasm_string)
my_task = device.run(program)

# You can also specify an optional s3 bucket location and number of shots,
 # if you so choose, when running the program
s3_location = ("amazon-braket-my-bucket", "openqasm-tasks")
my_task = device.run(
    program,
    s3_location,
    shots=100,
)
```

Use Boto3 to create OpenQASM 3.0 tasks

You can also use AWS Python SDK for Braket (Boto3) to create the quantum tasks using OpenQASM 3.0 strings, as shown in the following example. The following code snippet references ghz.qasm that prepares a GHZ state as shown above.

```
import boto3
import json

my_bucket = "amazon-braket-my-bucket"
```

```
s3_prefix = "openqasm-tasks"

with open("ghz.qasm") as f:
    source = f.read()

action = {
    "braketSchemaHeader": {
        "name": "braket.ir.openqasm.program",
        "version": "1"
    },
    "source": source
}
device_parameters = {}
device_arn = "arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2"
shots = 100

braket_client = boto3.client('braket', region_name='us-west-1')
rsp = braket_client.create_quantum_task(
    action=json.dumps(
        action
    ),
    deviceParameters=json.dumps(
        device_parameters
    ),
    deviceArn=device_arn,
    shots=shots,
    outputS3Bucket=my_bucket,
    outputS3KeyPrefix=s3_prefix,
)
```

Use the AWS CLI to create OpenQASM 3.0 tasks

The AWS Command Line Interface (CLI) can also be used to submit OpenQASM 3.0 programs, as shown in the following example.

```
aws braket create-quantum-task \
    --region "us-west-1" \
    --device-arn "arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2" \
    --shots 100 \
    --output-s3-bucket "amazon-braket-my-bucket" \
    --output-s3-key-prefix "openqasm-tasks" \
    --action '{
        "braketSchemaHeader": {
            "name": "braket.ir.openqasm.program",
            "version": "1"
        },
        "source": $(cat ghz.qasm)
    }'
```

Support for OpenQASM on different Braket Devices

For devices supporting OpenQASM 3.0, the `action` field supports a new action through the `GetDevice` response, as shown in the following example for the Rigetti and IonQ devices.

```
//OpenQASM as available with the Rigetti device capabilities
{
    "braketSchemaHeader": {
        "name": "braket.device_schema.rigetti.rigetti_device_capabilities",
        "version": "1"
    },
    "service": {...},
```

```
    "action": {
        "braket.ir.jaqcd.program": {...},
        "braket.ir.openqasm.program": {
            "actionType": "braket.ir.openqasm.program",
            "version": [
                "1"
            ],
            ….
        }
    }
}

//OpenQASM as available with the IonQ device capabilities
{
    "braketSchemaHeader": {
        "name": "braket.device_schema.ionq.ionq_device_capabilities",
        "version": "1"
    },
    "service": {...},
    "action": {
        "braket.ir.jaqcd.program": {...},
        "braket.ir.openqasm.program": {
            "actionType": "braket.ir.openqasm.program",
            "version": [
                "1"
            ],
            ….
        }
    }
}
```

For devices supporting pulse control, the `pulse` field has been added as a new field in the `GetDevice` response, as shown in the following example for the Rigetti and OQC devices.

```
// Rigetti
{
  "pulse": {
    "braketSchemaHeader": {
      "name": "braket.device_schema.pulse.pulse_device_action_properties",
      "version": "1"
    },
    "supportedQhpTemplateWaveforms": {
      "constant": {
        "functionName": "constant",
        "arguments": [
          {
            "name": "length",
            "type": "float",
            "optional": false
          },
          {
            "name": "iq",
            "type": "complex",
            "optional": false
          }
        ]
      },
      ...
    },
    "ports": {
      "q0_ff": {
        "portId": "q0_ff",
        "direction": "tx",
        "portType": "ff",
```

```
        "dt": 1e-9,
        "centerFrequencies": [
          375000000
        ]
      },
      ...
    },
    "supportedFunctions": {
      "shift_phase": {
        "functionName": "shift_phase",
        "arguments": [
          {
            "name": "frame",
            "type": "frame",
            "optional": false
          },
          {
            "name": "phase",
            "type": "float",
            "optional": false
          }
        ]
      },
     ...
    },
    "frames": {
      "q0_q1_cphase_frame": {
        "frameId": "q0_q1_cphase_frame",
        "portId": "q0_ff",
        "frequency": 462475694.24460185,
        "centerFrequency": 375000000,
        "phase": 0,
        "associatedGate": "cphase",
        "qubitMappings": [
          0,
          1
        ]
      },
      ...
    },
    "supportsLocalPulseElements": false,
    "supportsDynamicFrames": false,
    "supportsNonNativeGatesWithPulses": false,
    "validationParameters": {
      "MAX_SCALE": 4,
      "MAX_AMPLITUDE": 1,
      "PERMITTED_FREQUENCY_DIFFERENCE": 400000000
    }
  }
}

// OQC

{
  "pulse": {
    "braketSchemaHeader": {
      "name": "braket.device_schema.pulse.pulse_device_action_properties",
      "version": "1"
    },
    "supportedQhpTemplateWaveforms": {
      "gaussian": {
        "functionName": "gaussian",
        "arguments": [
          {
            "name": "length",
            "type": "float",
```

```
          "optional": false
        },
        {
          "name": "sigma",
          "type": "float",
          "optional": false
        },
        {
          "name": "amplitude",
          "type": "float",
          "optional": true
        },
        {
          "name": "zero_at_edges",
          "type": "bool",
          "optional": true
        }
      ]
    },
    ...
  },
  "ports": {
    "channel_1": {
      "portId": "channel_1",
      "direction": "tx",
      "portType": "port_type_1",
      "dt": 5e-10,
      "qubitMappings": [
        0
      ]
    },
    ...
  },
  "supportedFunctions": {
    "new_frame": {
      "functionName": "new_frame",
      "arguments": [
        {
          "name": "port",
          "type": "port",
          "optional": false
        },
        {
          "name": "frequency",
          "type": "float",
          "optional": false
        },
        {
          "name": "phase",
          "type": "float",
          "optional": true
        }
      ]
    },
    ...
  },
  "frames": {
    "q0_drive": {
      "frameId": "q0_drive",
      "portId": "channel_1",
      "frequency": 5500000000,
      "centerFrequency": 5500000000,
      "phase": 0,
      "qubitMappings": [
        0
      ]
```

```
      },
      ...
    },
    "supportsLocalPulseElements": false,
    "supportsDynamicFrames": true,
    "supportsNonNativeGatesWithPulses": true,
    "validationParameters": {
      "MAX_SCALE": 1,
      "MAX_AMPLITUDE": 1,
      "PERMITTED_FREQUENCY_DIFFERENCE": 1,
      "MIN_PULSE_LENGTH": 8e-9,
      "MAX_PULSE_LENGTH": 0.00012
    }
  }
}
```

The above fields detail the following:

Ports:

Describes pre-made external (`extern`) device ports declared on the QPU as well as the associated properties of the given port. All ports listed in this structure are pre-declared as valid identifiers within the `OpenQASM 3.0` program submitted by the user. The additional properties for a port include:

- Port id (portId)
  - The port name declared as an identifier in OpenQASM 3.0
- Direction (direction)
  - The direction of the port. Drive ports will transmit pulses (direction “tx”), while measurement ports will receive pulses (direction “rx”)
- Port type (portType)
  - The type of action this port is responsible for (ex: drive, capture, ff - fast-flux)
- Dt (dt)
  - The time in seconds that represents a single sample time step on the given port
- Qubit mappings (qubitMappings)
  - The associated qubits with the given port
- Center frequencies (centerFrequencies)
  - A list of the associated center frequencies for all pre-declared or user-defined frames on the port. For more information, please refer to Frames
- QHP Specific Properties (qhpSpecificProperties)
  - An optional map detailing existing properties about the port specific to the QHP

Frames:

Describes pre-made external frames declared on the QPU as well as associated properties about the frames. All frames listed in this structure are pre-declared as valid identifiers within the `OpenQASM 3.0` program submitted by the user. The additional properties for a frame include:

- Frame Id (frameId)
  - The frame name declared as an identifier in OpenQASM 3.0
- Port Id (portId)
  - The associated hardware port for the frame
- Frequency (frequency)
  - The default initial frequency of the frame
- Center Frequency (centerFrequency)

  - The center of the frequency bandwidth for the frame. Typically frames may only be adjusted to a certain bandwidth around the center frequency. As a result, frequency adjustments should ensure they stick with a given delta of the center frequency. The bandwidth value can be found in the validation parameters
- Phase (phase)
  - The default initial phase of the frame
- Associated Gate (associatedGate)
  - The associated gates with the given frame
- Qubit Mappings (qubitMappings)
  - The associated qubits with the given frame
- QHP Specific Properties (qhpSpecificProperties)
  - An optional map detailing existing properties about the frame specific to the QHP

SupportsDynamicFrames:

Describes whether or not a frame can be declared in `cal` or `defcal` blocks via the OpenPulse `newframe` function. If this is false, only frames listed in the frame structure may be used within the program.

SupportedFunctions:

Describes the OpenPulse functions that are supported for the device, as well as the associated arguments, argument types and return types for the given functions. Example usage of the OpenPulse functions can be found on the OpenPulse specification. At this time Braket supports:

- shift_phase
  - Shifts the phase of a frame by a specified value
- set_phase
  - Sets the phase of frame to the specified value
- shift_frequency
  - Shifts the frequency of a frame by a specified value
- set_frequency
  - Sets the frequency of frame to the specified value
- play
  - Schedules a waveform
- capture_v0
  - Returns the value on a capture frame to a bit register

SupportedQhpTemplateWaveforms:

Describes the pre-built waveform functions available on the device and the associated arguments and types. By default Braket Pulse offers pre-built waveform routines on all devices, which are:

Constant

$$Constant(t, \tau, iq) = iq$$

τ is the length of the waveform, iq is a complex number

```
def constant(length, iq)
```

Gaussian

$$Gaussian(t, \tau, \sigma, A = 1, ZaE = 0) = \frac{A}{1 - ZaE * \exp\left(-\frac{1}{2}\left(\frac{\tau}{2\sigma}\right)^2\right)} \left[\exp\left(-\frac{1}{2}\left(\frac{t - \frac{\tau}{2}}{\sigma}\right)^2\right) - ZaE * \exp\left(-\frac{1}{2}\left(\frac{\tau}{2\sigma}\right)^2\right)\right]$$

τ is the length of the waveform, σ the width of the Gaussian, A the amplitude. If setting ZaE to True, the Gaussian is offset and rescaled such that it is equals to zero at the start and end of the waveform, and reaches A at maximum.

```
def gaussian(length, sigma, amplitude=1, zero_at_edges=False)
```

DRAG Gaussian

$$DRAG\_Gaussian(t, \tau, \sigma, \beta, A = 1, ZaE = 0) = \frac{A}{1 - ZaE * \exp\left(-\frac{1}{2}\left(\frac{\tau}{2\sigma}\right)^2\right)} \left(1 - i\beta\frac{t - \frac{\tau}{2}}{\sigma^2}\right) \left[\exp\left(-\frac{1}{2}\left(\frac{t - \frac{\tau}{2}}{\sigma}\right)^2\right) - ZaE * \exp\left(-\frac{1}{2}\left(\frac{\tau}{2\sigma}\right)^2\right)\right]$$

τ is the length of the waveform, σ the width of the gaussian, β a free parameter and A the amplitude. If setting ZaE to True,The Derivative Removal by Adiabatic Gate (DRAG) Gaussian is offset and rescaled such that it is equals to zero at the start and end of the waveform, and the real part reaches A at maximum. The DRAG waveform is described in this paper.

```
def drag_gaussian(length, sigma, beta, amplitude=1, zero_at_edges=False)
```

SupportsLocalPulseElements:

Describes whether or not pulse elements, such as ports, frames, and waveforms may be defined locally in `defcal` blocks. If the value is `false` elements must be defined in `cal` blocks.

SupportsNonNativeGatesWithPulses:

This details whether or not non-native gates can be used in combination with pulse programs. For example assuming this was false, a program would not be able to use an H gate (non-native), without first defining the gate via `defcal`, for the used qubit, when using a pulse program. The list of native gates can be found in the `nativeGateSet` key in the device capabilities.

ValidationParameters:

Describes pulse element validation boundaries, including:

- Maximum Scale / Maximum Amplitude values for waveforms (arbitrary and pre-built)
- Maximum frequency bandwidth from supplied center frequency in Hz
- Minimum pulse length/duration in seconds
- Maximum pulse length/duration in seconds

Supported Operations, Results and Result Types with OpenQASM

To find out which OpenQASM 3.0 features each device supports, you can refer to the `braket.ir.openqasm.program` key in the `action` field on the device capabilities output. For example, the following are the supported operations and result types available for the Braket Managed State Vector simulator SV1.

```
...
```

```
 "action": {
   "braket.ir.jaqcd.program": {
     ...
   },
"braket.ir.openqasm.program": {
     "version": [
       "1.0"
     ],
     "actionType": "braket.ir.openqasm.program",
     "supportedOperations": [
       "ccnot",
       "cnot",
       "cphaseshift",
       "cphaseshift00",
       "cphaseshift01",
       "cphaseshift10",
       "cswap",
       "cy",
       "cz",
       "h",
       "i",
       "iswap",
       "pswap",
       "phaseshift",
       "rx",
       "ry",
       "rz",
       "s",
       "si",
       "swap",
       "t",
       "ti",
       "v",
       "vi",
       "x",
       "xx",
       "xy",
       "y",
       "yy",
       "z",
       "zz"
     ],
     "supportedPragmas": [
       "braket_unitary_matrix"
     ],
     "forbiddenPragmas": [],
     "maximumQubitArrays": 1,
     "maximumClassicalArrays": 1,
     "forbiddenArrayOperations": [
       "concatenation",
       "negativeIndex",
       "range",
       "rangeWithStep",
       "slicing",
       "selection"
     ],
     "requiresAllQubitsMeasurement": true,
     "supportsPhysicalQubits": false,
     "requiresContiguousQubitIndices": true,
     "disabledQubitRewiringSupported": false,
     "supportedResultTypes": [
       {
         "name": "Sample",
         "observables": [
           "x",
           "y",
```

```
            "z",
            "h",
            "i",
            "hermitian"
          ],
          "minShots": 1,
          "maxShots": 100000
        },
        {
          "name": "Expectation",
          "observables": [
            "x",
            "y",
            "z",
            "h",
            "i",
            "hermitian"
          ],
          "minShots": 0,
          "maxShots": 100000
        },
        {
          "name": "Variance",
          "observables": [
            "x",
            "y",
            "z",
            "h",
            "i",
            "hermitian"
          ],
          "minShots": 0,
          "maxShots": 100000
        },
        {
          "name": "Probability",
          "minShots": 1,
          "maxShots": 100000
        },
        {
          "name": "Amplitude",
          "minShots": 0,
          "maxShots": 0
        }
      ]
    }
  },
...
```

Simulate noise with OpenQASM3

To simulate noise with OpenQASM3, you use *pragma* instructions to add noise operators. For example, to simulate the noisy version of the GHZ program (p. 68) provided previously, you can submit the following OpenQASM program.

```
// ghz.qasm
// Prepare a GHZ state
OPENQASM 3;

qubit[3] q;
bit[3] c;

h q[0];
```

```
#pragma braket noise depolarizing(0.75) q[0] cnot q[0], q[1];
#pragma braket noise depolarizing(0.75) q[0]
#pragma braket noise depolarizing(0.75) q[1] cnot q[1], q[2];
#pragma braket noise depolarizing(0.75) q[0]
#pragma braket noise depolarizing(0.75) q[1]

c = measure q;
```

Specifications for all supported pragma noise operators are provided in the following list.

```
#pragma braket noise bit_flip(<float>) <qubit>
#pragma braket noise phase_flip(<float>) <qubit>
#pragma braket noise pauli_channel(<float>, <float>, <float>)  <qubit>
#pragma braket noise depolarizing(<float in [0,3/4]>) <qubit>
#pragma braket noise two_qubit_depolarizing(<float in [0,15/16]>) <qubit>, <qubit>
#pragma braket noise two_qubit_dephasing(<float in [0,3/4]>) <qubit>, <qubit>
#pragma braket noise amplitude_damping(<float in [0,1]>) <qubit>
#pragma braket noise generalized_amplitude_damping(<float in [0,1]> <float in [0,1]>)
 <qubit>
#pragma braket noise phase_damping(<float in [0,1]>) <qubit>
#pragma braket noise kraus([[<complex m0_00>, ], ...], [[<complex m1_00>, ], ...], ...)
 <qubit>[, <qubit>]     // maximum of 2 qubits and maximum of 4 matrices for 1 qubit, 16
 for 2
) <qubit>[, <qubit>]     // maximum of 2 qubits
```

Kraus Operator

In order to generate a Kraus operator, you can iterate through a list of matrices, printing each element of the matrix as a complex expression.

When using Kraus operators, remember the following:

- The number of qubits must not exceed 2. The current definition in the schemas sets this limit.
- The length of the argument list must be a multiple of 8. This means it must be composed only of 2x2 matrices.
- The total length does not exceed 2^{2*num_qubits} matrices. This means 4 matrices for 1 qubit and 16 for 2 qubits.
- All supplied matrices are completely positive trace preserving (CPTP).
- The product of the Kraus operators with their transpose conjugates need to add up to an identity matrix.

Qubit rewiring with OpenQASM3

Amazon Braket supports the physical qubit notation within OpenQASM on Rigetti devices (to learn more see this page). When using physical qubits with the naive rewiring strategy, ensure that the qubits are connected on the selected device. Alternatively, if qubit registers are used instead, the PARTIAL rewiring strategy is enabled by default on Rigetti devices.

```
// ghz.qasm
// Prepare a GHZ state
OPENQASM 3;

h $0;
cnot $0, $1;
cnot $1, $2;

measure $0;
measure $1;
```

```
measure $2;
```

Verbatim Compilation with OpenQASM3

When you run a quantum circuit on quantum computers from Rigetti, OQC, and IonQ, you can direct the compiler to run your circuits exactly as defined, without any modifications. This feature is known as *verbatim compilation*. With Rigetti devices, you can specify precisely what gets preserved—either an entire circuit or only specific parts of it. To preserve only specific parts of a circuit, you'll need to use native gates within the preserved regions. Currently, IonQ and OQC only support verbatim compilation for the entire circuit, so every instruction in the circuit needs to be enclosed in a verbatim box.

With OpenQASM, you can specify a verbatim pragma around a box of code that is untouched and not optimized by the low-level compilation routine of the hardware. The following code example shows how to use the `#pragma braket verbatim`.

```
OPENQASM 3;

bit[2] c;

#pragma braket verbatim
box{
    rx(0.314159) $0;
    rz(0.628318) $0, $1;
    cz $0, $1;
}

c[0] = measure $0;
c[1] = measure $1;
```

For more information on verbatim compilation, see the Verbatim compilation sample notebook.

The Braket console

OpenQASM 3.0 tasks are available and can be managed within the Amazon Braket console. On the console, you have the same experience submitting tasks in OpenQASM 3.0 as you had submitting existing tasks.

More resources

OpenQASM is available in all Amazon Braket Regions.

For an example notebook for getting started with OpenQASM on Amazon Braket, see Braket Tutorials GitHub.

Submitting tasks to QPUs and simulators

Amazon Braket provides access to several devices that can run quantum tasks. You can submit tasks individually or you can set up task batching.

QPUs

You can submit tasks to QPUs at any time, but the task runs within certain availability windows that are displayed on the **Devices** page of the Amazon Braket console. You can retrieve the results of the task with the task ID, which is introduced in the next section.

- **IonQ** : `arn:aws:braket:::device/qpu/ionq/ionQdevice`

- **OQC Lucy** : `arn:aws:braket:eu-west-2::device/qpu/oqc/Lucy`
- **QuEra Aquila** : `arn:aws:braket:us-east-1::device/qpu/quera/Aquila`
- **Rigetti Aspen-11** : `arn:aws:braket:::device/qpu/rigetti/Aspen-11`
- **Rigetti Aspen-M-2** : `arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2`
- **Xanadu Borealis** : `arn:aws:braket:us-east-1::device/qpu/xanadu/Borealis`

Simulators

- **Managed density matrix simulator, DM1** : `arn:aws:braket:::device/quantum-simulator/amazon/dm1`
- **Managed state vector simulator, SV1** : `arn:aws:braket:::device/quantum-simulator/amazon/sv1`
- **Managed tensor network simulator, TN1** : `arn:aws:braket:::device/quantum-simulator/amazon/tn1`
- **The local simulator** : `LocalSimulator()`

> **Note**
> You can cancel tasks in the CREATED state for QPUs and managed simulators. You can cancel tasks in the QUEUED state on a best-effort basis for managed simulators and QPUs. Note that QPU QUEUED tasks are unlikely to be cancelled successfully during QPU availability windows.

In this section:

Example tasks on Amazon Braket

This section walks through the stages of running an example task, from selecting the device to viewing the result. As a best practice for Amazon Braket, we recommend that you begin by running the circuit on a simulator, such as SV1.

In this section:

Specify the device

First, select and specify the device for your task. This example shows how to choose the simulator, SV1.

```
# choose the managed simulator to run the circuit
from braket.aws import AwsDevice
device = AwsDevice("arn:aws:braket:::device/quantum-simulator/amazon/sv1")
```

You can view some of the properties of this device as follows:

```
print (device.name)
for iter in device.properties.action['braket.ir.jaqcd.program']:
    print(iter)
```

```
SV1
('version', ['1.0', '1.1'])
('actionType', <DeviceActionType.JAQCD: 'braket.ir.jaqcd.program'>)
('supportedOperations', ['ccnot', 'cnot', 'cphaseshift', 'cphaseshift00', 'cphaseshift01',
 'cphaseshift10', 'cswap', 'cy', 'cz', 'h', 'i', 'iswap', 'pswap', 'phaseshift', 'rx',
 'ry', 'rz', 's', 'si', 'swap', 't', 'ti', 'unitary', 'v', 'vi', 'x', 'xx', 'xy', 'y',
 'yy', 'z', 'zz'])
('supportedResultTypes', [ResultType(name='Sample', observables=['x', 'y', 'z', 'h',
 'i', 'hermitian'], minShots=1, maxShots=100000), ResultType(name='Expectation',
 observables=['x', 'y', 'z', 'h', 'i', 'hermitian'], minShots=0, maxShots=100000),
 ResultType(name='Variance', observables=['x', 'y', 'z', 'h', 'i', 'hermitian'],
 minShots=0, maxShots=100000), ResultType(name='Probability', observables=None, minShots=1,
 maxShots=100000), ResultType(name='Amplitude', observables=None, minShots=0, maxShots=0)])
```

Submit an example task

Submit an example task to run on the managed simulator.

```
# create a circuit with a result type
circ = Circuit().rx(0, 1).ry(1, 0.2).cnot(0,2).variance(observable=Observable.Z(),
 target=0)
# add another result type
circ.probability(target=[0, 2])

# set up S3 bucket (where results are stored)
my_bucket = "amazon-braket-your-s3-bucket-name" # the name of the bucket
my_prefix = "your-folder-name" # the name of the folder in the bucket
s3_location = (my_bucket, my_prefix)

# submit the task to run
my_task = device.run(circ, s3_location, shots=1000, poll_timeout_seconds = 100,
 poll_interval_seconds = 10)
# The positional argument for the S3 bucket is optional if you want to specify a bucket
 other than the default

# get results of the task
result = my_task.result()
```

The `device.run()` command creates a task through the `CreateQuantumTask` API. After a short initialization time, the task is queued until capacity exists to run the task on a device. In this case, the device is the managed simulator SV1. After the device completes the computation, Amazon Braket writes the results to the Amazon S3 location specified in the call. The positional argument `s3_location` is required for all devices except the local simulator.

Note
The Braket quantum task action is limited to 3MB in size.

Specify shots

The shots argument refers to the number of desired measurement shots. Simulators such as SV1 support two simulation modes.

- For shots = 0, the simulator performs an exact simulation, returning the true values for all result types. (Not available on TN1.)

- For non-zero values of shots, the simulator samples from the output distribution to emulate the shot noise of real QPUs. QPU devices only allow shots > 0.

For information about the maximum number of shots per task, please refer to Braket Quotas (p. 23).

Poll for results

When executing `my_task.result()`, the SDK begins polling for a result with the parameters you define upon task creation:

- `poll_timeout_seconds` is the number of seconds to poll the task before it times out when running the task on the managed simulator and or QPU devices. The default value is 432,000 seconds, which is 5 days.
- **Note:** For QPU devices such as Rigetti and IonQ, we recommend that you allow a few days. If your polling timeout is too short, results may not be returned within the polling time. For example, when a QPU is unavailable, a local timeout error is returned.
- `poll_interval_seconds` is the frequency with which the task is polled. It specifies how often you call the Braket API to get the status when the task is run on the managed simulator and on QPU devices. The default value is 1 second.

This asynchronous execution facilitates the interaction with QPU devices that are not always available. For example, a device could be unavailable during a regular maintenance window.

The returned result contains a range of metadata associated with the task. You can check the measurement result with the following commands:

```
print('Measurement results:\n',result.measurements)
print('Counts for collapsed states:\n',result.measurement_counts)
print('Probabilities for collapsed states:\n',result.measurement_probabilities)
```

```
Measurement results:
 [[1 0 1]
 [0 0 0]
 [1 0 1]
 ...
 [0 0 0]
 [0 0 0]
 [0 0 0]]
Counts for collapsed states:
 Counter({'000': 761, '101': 226, '010': 10, '111': 3})
Probabilities for collapsed states:
 {'101': 0.226, '000': 0.761, '111': 0.003, '010': 0.01}
```

View the example results

Because you've also specified the `ResultType`, you can view the returned results. The result types appear in the order in which they were added to the circuit.

```
print('Result types include:\n', result.result_types)
print('Variance=',result.values[0])
print('Probability=',result.values[1])

# you can plot the result and do some analysis
import matplotlib.pyplot as plt
plt.bar(result.measurement_counts.keys(), result.measurement_counts.values());
plt.xlabel('bitstrings');
```

```
plt.ylabel('counts');
```

```
Result types include:
 [ResultTypeValue(type={'observable': ['z'], 'targets': [0], 'type': 'variance'},
 value=0.7062359999999999), ResultTypeValue(type={'targets': [0, 2], 'type':
 'probability'}, value=array([0.771, 0.   , 0.   , 0.229]))]
Variance= 0.7062359999999999
Probability= [0.771 0.    0.    0.229]
```

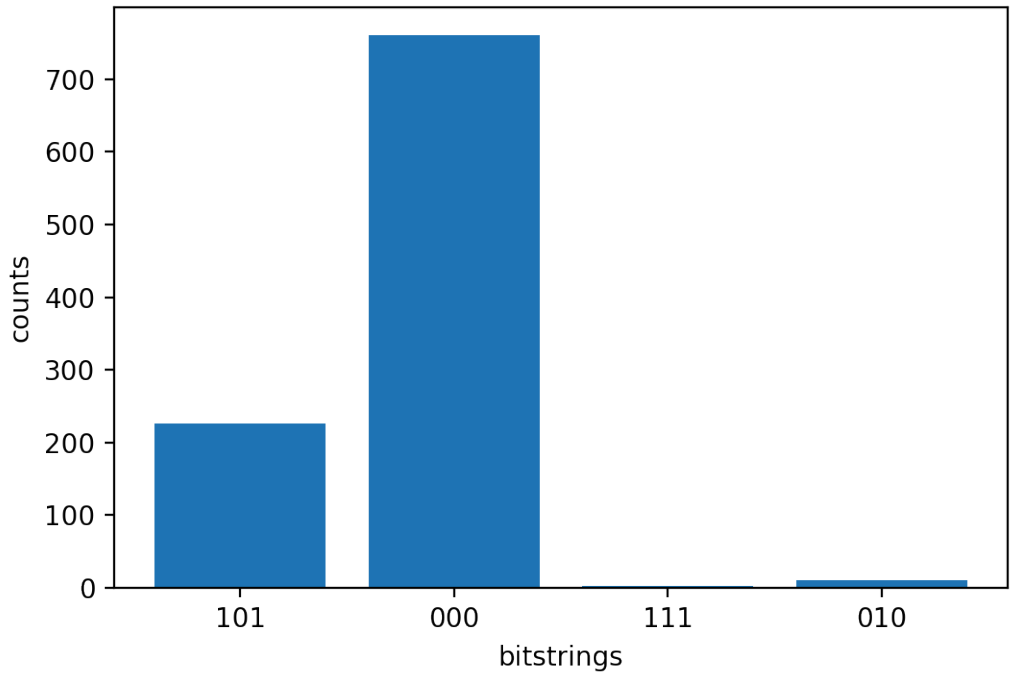

Submitting tasks to a QPU

Amazon Braket allows you to run a quantum circuit on a QPU device. The following example shows how to submit a task to the Rigetti or the IonQ device.

Choose the Rigetti device, then look at the associated connectivity graph

```
# import the QPU module
from braket.aws import AwsDevice
# choose the Rigetti device
device = AwsDevice("arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2")

# take a look at the device connectivity graph
device.properties.dict()['paradigm']['connectivity']
```

```
{'fullyConnected': False,
 'connectivityGraph': {'0': ['1', '7'],
  '1': ['0', '16'],
  '2': ['3', '15'],
  '3': ['2', '4'],
  '4': ['3', '5'],
```

```
  '5': ['4', '6'],
  '6': ['5', '7'],
  '7': ['0', '6'],
  '11': ['12', '26'],
  '12': ['13', '11'],
  '13': ['12', '14'],
  '14': ['13', '15'],
  '15': ['2', '14', '16'],
  '16': ['1', '15', '17'],
  '17': ['16'],
  '20': ['21', '27'],
  '21': ['20', '36'],
  '22': ['23', '35'],
  '23': ['22', '24'],
  '24': ['23', '25'],
  '25': ['24', '26'],
  '26': ['11', '25', '27'],
  '27': ['20', '26'],
  '30': ['31', '37'],
  '31': ['30', '32'],
  '32': ['31', '33'],
  '33': ['32', '34'],
  '34': ['33', '35'],
  '35': ['22', '34', '36'],
  '36': ['21', '35', '37'],
  '37': ['30', '36']}}
```

The preceding dictionary `connectivityGraph` contains information about the connectivity of the current Rigetti device.

Choose the IonQ device

For the IonQ device, the `connectivityGraph` is empty, as shown in the following example, because the device offers *all-to-all* connectivity. Therefore, a detailed `connectivityGraph` is not needed.

```
# or choose the IonQ device
device = AwsDevice("arn:aws:braket:::device/qpu/ionq/ionQdevice")

# take a look at the device connectivity graph
device.properties.dict()['paradigm']['connectivity']
```

```
{'fullyConnected': True, 'connectivityGraph': {}}
```

For any QPU device, when you launch a task, remember to specify the S3 bucket in which to store your results. You have the option to adjust the shots (default=1000), the `poll_timeout_seconds` (default = 432000 = 5 days), and the `poll_interval_seconds` (default = 1) when you submit the task, as shown in the following example.

```
my_task = device.run(circ, s3_location, shots=100, poll_timeout_seconds = 100,
 poll_interval_seconds = 10)
```

The IonQ and Rigetti devices compile the provided circuit into their respective native gate sets automatically, and they map the abstract qubit indices to physical qubits on the respective QPU.

> **Note**
> QPU devices have limited capacity. You can expect longer wait times when capacity is reached.

Amazon Braket can run QPU tasks within certain availability windows, but you can still submit tasks any time (24/7) because all corresponding data and metadata are stored reliably in your S3 bucket. As shown in the next section, you can recover your task using `AwsQuantumTask` and your unique task ID.

Running a task with the local simulator

You can send tasks directly to a local simulator for rapid prototyping and testing. This simulator runs in your local environment, so you do not need to specify an Amazon S3 location. The results are computed directly in your session. To run a task on the local simulator, you must only specify the shots parameter.

> **Note**
> The execution speed and maximum number of qubits the local simulator can process depends on the Amazon Braket notebook instance type, or on your local hardware specifications.

The following commands are all identical and instantiate the state vector (noise free) local simulator.

```
# import the LocalSimulator module
from braket.devices import LocalSimulator
# the following are identical commands
device = LocalSimulator()
device = LocalSimulator("default")
device = LocalSimulator(backend="default")
device = LocalSimulator(backend="braket_sv")
```

Then run a task with the following.

```
my_task = device.run(circ, shots=1000)
```

To instantiate the local density matrix (noise) simulator customers change the backend as follows.

```
# import the LocalSimulator module
from braket.devices import LocalSimulator
device = LocalSimulator(backend="braket_dm")
```

Task batching

Task batching is available on every Amazon Braket device, except the local simulator. Batching is especially useful for tasks you run on the managed simulators (TN1 or SV1) because they can process multiple tasks in parallel. To help you set up various tasks, Amazon Braket provides example notebooks.

Batching allows you to launch tasks in parallel. For example, if you wish to make a calculation that requires 10 tasks and the circuits in those tasks are independent of each other, it is a good idea to use batching. That way, you don't have to wait for one task to be complete before another task begins.

The following example shows how to run a batch of tasks:

```
circuits = [bell for _ in range(5)]
batch = device.run_batch(circuits, s3_folder, shots=100)
print(batch.results()[0].measurement_counts)  # The result of the first task in the batch
```

For more information, see the Amazon Braket examples on GitHub or Quantum task batching, which has more specific information about batching.

About task batching and costs

A few caveats to keep in mind regarding task batching and billing costs:

- By default, task batching retries all time out or fail tasks 3 times.

- A batch of long running tasks, such as 34 qubits for SV1, can incur large costs. Be sure to double check the `run_batch` assignment values carefully before you start a batch of tasks. We do not recommend using TN1 with `run_batch`.
- TN1 can incur costs for failed rehearsal phase tasks (see the TN1 description for more information). Automatic retries can add to the cost and so we recommend setting the number of 'max_retries' on batching to 0 when using TN1 (see Quantum Task Batching, Line 186).

Task batching and PennyLane

Take advantage of batching when you're using PennyLane on Amazon Braket by setting `parallel = True` when you instantiate an Amazon Braket device, as shown in the following example.

```
device = qml.device("braket.aws.qubit",device_arn="arn:aws:braket:::device/quantum-
simulator/amazon/sv1",wires=wires,s3_destination_folder=s3_folder,parallel=True,)
```

For more information about batching with PennyLane, see Parallelized optimization of quantum circuits.

Set up SNS notifications (optional)

You can set up notifications through the Amazon Simple Notification Service (SNS) so that you receive an alert when your Amazon Braket task is complete. Active notifications are useful if you expect a long wait time; for example, when you submit a large task or when you submit a task outside of a device's availability window. If you do not want to wait for the task to complete, you can set up an SNS notification.

An Amazon Braket notebook walks you through the setup steps. For more information, see the Amazon Braket examples on GitHub and, specifically, the example notebook for setting up notifications.

Submit an analog program using QuEra's Aquila

This page provides a comprehensive documentation about the capabilities of the Aquila machine from QuEra. Details covered here are the following: 1) The parameterized Hamiltonian simulated by Aquila, 2) AHS program parameters, 3) AHS result content, 4) Aquila capabilities parameter. We suggest using Ctrl +F text search to find parameters relevant to your questions.

Hamiltonian

The Aquila machine from QuEra simulates the following (time-dependent) Hamiltonian natively:

$$H(t) = \sum_{k=1}^{N} H_{\text{drive},k}(t) + \sum_{k=1}^{N-1} \sum_{l=k+1}^{N} V_{\text{vdw},k,l}$$

where

- $H_{\text{drive},k}(t) = (\,{}^{1}/_{2}\,\Omega(t)e^{i\phi(t)}S_{-,k} + {}^{1}/_{2}\,\Omega(t)e^{-i\phi(t)}\,S_{+,k}) + (-\Delta_{\text{global}}(t)n_k)$,
  - $\Omega(t)$ is the time-dependent, global driving amplitude (aka Rabi frequency), in units of (rad / us)
  - $\phi(t)$ is the time-dependent, global phase, measured in radians
  - $S_{-,k}$ and $S_{+,k}$ are the spin lowering and raising operators of atom k (in the basis $|\#\rangle=|g\rangle$, $|\#\rangle=|r\rangle$, they are $S_{-}=|g\rangle\langle r|$, $S_{+}=(S_{-})^{\dagger}=|r\rangle\langle g|$)
  - $\Delta_{\text{global}}(t)$ is the time-dependent, global detuning
  - n_k is the projection operator on the Rydberg state of atom k (i.e. $n=|r\rangle\langle r|$)

- $V_{vdw,k,l}=C_6/(d_{k,l})^6 n_k n_l$,
  - C_6 is the van der Waals coefficient, in units of (rad / s) * (m)^6
  - $d_{k,l}$ is the Euclidean distance between atom k and l, measured in um.

The users have control over the following parameters via the Braket AHS program schema.

- 2-d atom arrangement (x_k and y_k coordinates of each atom k, in units of um), which controls the pairwise atomic distances $d_{k,l}$ with k,l=1,2,...N
- $\Omega(t)$, the time-dependent, global Rabi frequency, in units of (rad / s)
- $\phi(t)$, the time-dependent, global phase, in units of (rad / s)
- $\Delta_{global}(t)$, the time-dependent, global detuning, in units of (rad / s)

Note
The user cannot control which levels are involved (i.e. $S_-, S_{+,n}$ operators are fixed) nor the strength of the Rydberg-Rydberg interaction coefficient (C_6).

Braket AHS program schema

braket.ir.ahs.program_v1.Program object (example)

```
Program(
    braketSchemaHeader=BraketSchemaHeader(
        name='braket.ir.ahs.program',
        version='1'
    ),
    setup=Setup(
        ahs_register=AtomArrangement(
            sites=[[Decimal('0'), Decimal('0')]],
            filling=[1]
        )
    ),
    hamiltonian=Hamiltonian(
        drivingFields=[
            DrivingField(
                amplitude=PhysicalField(
                    time_series=TimeSeries(
                        values=[Decimal('0'), Decimal('15700000.0')],
                        times=[Decimal('0'), Decimal('0.000001')]
                    ),
                    pattern='uniform'
                ),
                phase=PhysicalField(
                    time_series=TimeSeries(
                        values=[Decimal('0'), Decimal('0')],
                        times=[Decimal('0'), Decimal('0.000001')]
                    ),
                    pattern='uniform'
                ),
                detuning=PhysicalField(
                    time_series=TimeSeries(
                        values=[Decimal('-54000000.0'), Decimal('54000000.0')],
                        times=[Decimal('0'), Decimal('0.000001')]
                    ),
                    pattern='uniform'
                )
            )
        ],
        shiftingFields=[]
    )
```

```
)
```

JSON (example)

```
{
    "braketSchemaHeader": {
        "name": "braket.ir.ahs.program",
        "version": "1"
    },
    "setup": {
        "ahs_register": {
            "sites": [[0E-7, 0E-7]],
            "filling": [1]
        }
    },
    "hamiltonian": {
        "drivingFields": [
            {
                "amplitude": {
                    "time_series": {
                        "values": [0.0, 15700000.0],
                        "times": [0E-9, 0.000001000]
                    },
                    "pattern": "uniform"
                },
                "phase": {
                    "time_series": {
                        "values": [0E-7, 0E-7],
                        "times": [0E-9, 0.000001000]
                    },
                    "pattern": "uniform"
                },
                "detuning": {
                    "time_series": {
                        "values": [-54000000.0, 54000000.0],
                        "times": [0E-9, 0.000001000]
                    },
                    "pattern": "uniform"
                }
            }
        ],
        "shiftingFields": []
    }
}
```

Main fields

Program field	type	description
setup.ahs_register.sites	List[List[Decimal]]	List of 2-d coordinates where the tweezers trap atoms
setup.ahs_register.filling	List[int]	Marks atoms that occupy the trap sites with 1, and empty sites with 0
hamiltonian.drivingFields[].amplitude.time_series.times	List[Decimal]	time points of driving amplitude, Omega(t)

Program field	type	description
hamiltonian.drivingFields[].amplitude.time_series.values	List[Decimal]	values of driving amplitude, Omega(t)
hamiltonian.drivingFields[].amplitude.pattern	str	spatial pattern of driving amplitude, Omega(t); must be 'uniform'
hamiltonian.drivingFields[].phase.time_series.times	List[Decimal]	time points of driving phase, phi(t)
hamiltonian.drivingFields[].phase.time_series.values	List[Decimal]	values of driving phase, phi(t)
hamiltonian.drivingFields[].phase.pattern	str	spatial pattern of driving phase, phi(t); must be 'uniform'
hamiltonian.drivingFields[].detuning.time_series.times	List[Decimal]	time points of driving detuning, Delta_global(t)
hamiltonian.drivingFields[].detuning.time_series.values	List[Decimal]	values of driving detuning, Delta_global(t)
hamiltonian.drivingFields[].detuning.pattern	str	spatial pattern of driving detuning, Delta_global(t); must be 'uniform'
hamiltonian.shiftingFields	List	must be empty

Metadata fields

Program field	type	description
braketSchemaHeader.name	str	name of the schema; must be 'braket.ir.ahs.program'
braketSchemaHeader.version	str	version of the schema

Braket AHS task result schema

braket.tasks.analog_hamiltonian_simulation_quantum_task_result.AnalogHamiltonianSimulationQuantumTaskR
(example)

```
AnalogHamiltonianSimulationQuantumTaskResult(
    task_metadata=TaskMetadata(
        braketSchemaHeader=BraketSchemaHeader(
            name='braket.task_result.task_metadata',
            version='1'
        ),
        id='arn:aws:braket:us-east-1:123456789012:quantum-task/12345678-90ab-
cdef-1234-567890abcdef',
```

```
        shots=2,
        deviceId='arn:aws:braket:us-east-1::device/qpu/quera/Aquila',
        deviceParameters=None,
        createdAt='2022-10-25T20:59:10.788Z',
        endedAt='2022-10-25T21:00:58.218Z',
        status='COMPLETED',
        failureReason=None
    ),
    measurements=[
        ShotResult(
            status=<AnalogHamiltonianSimulationShotStatus.SUCCESS: 'Success'>,

            pre_sequence=array([1, 1, 1, 1]),
            post_sequence=array([0, 1, 1, 1])
        ),

        ShotResult(
            status=<AnalogHamiltonianSimulationShotStatus.SUCCESS: 'Success'>,

            pre_sequence=array([1, 1, 0, 1]),
            post_sequence=array([1, 0, 0, 0])
        )
    ]
)
```

JSON (example)

```
{
    "braketSchemaHeader": {
        "name": "braket.task_result.analog_hamiltonian_simulation_task_result",
        "version": "1"
    },
    "taskMetadata": {
        "braketSchemaHeader": {
            "name": "braket.task_result.task_metadata",
            "version": "1"
        },
        "id": "arn:aws:braket:us-east-1:123456789012:quantum-task/12345678-90ab-
cdef-1234-567890abcdef",
        "shots": 2,
        "deviceId": "arn:aws:braket:us-east-1::device/qpu/quera/Aquila",

        "createdAt": "2022-10-25T20:59:10.788Z",
        "endedAt": "2022-10-25T21:00:58.218Z",
        "status": "COMPLETED"

    },
    "measurements": [
        {
            "shotMetadata": {"shotStatus": "Success"},
            "shotResult": {
                "preSequence": [1, 1, 1, 1],
                "postSequence": [0, 1, 1, 1]
            }
        },
        {
            "shotMetadata": {"shotStatus": "Success"},
            "shotResult": {
                "preSequence": [1, 1, 0, 1],
                "postSequence": [1, 0, 0, 0]
            }
        }
    ],
    "additionalMetadata": {
        "action": {...}
```

```
        "queraMetadata": {
            "braketSchemaHeader": {
                "name": "braket.task_result.quera_metadata",
                "version": "1"
            },
            "numSuccessfulShots": 100
        }
    }
}
```

Main fields

Task result field	type	description
measurements[].shotResult.preSequence	List[int]	Pre-sequence measurement bits (one for each atomic site) for each shot: 0 if site is empty, 1 if site is filled, measured before the sequences of pulses that run the quantum evolution
measurements[].shotResult.postSequence	List[int]	Post-sequence measurement bits for each shot: 0 if atom is in Rydberg state or site is empty, 1 if atom is in ground state, measured at the end of the sequences of pulses that run the quantum evolution

Metadata fields

Task result field	type	description
braketSchemaHeader.name	str	name of the schema; must be 'braket.task_result.analog_ham
braketSchemaHeader.version	str	version of the schema
taskMetadata.braketSchemaHeader.name	str	name of the schema; must be 'braket.task_result.task_metada
taskMetadata.braketSchemaHeader.version	str	version of the schema
taskMetadata.id	str	The ID of the task. For AWS tasks, this is the task ARN.
taskMetadata.shots	int	The number of shots for the task
taskMetadata.shots.deviceId	str	The ID of the device on which the task ran. For

Task result field	type	description
		AWS devices, this is the device ARN.
taskMetadata.shots.createdAt	str	The timestamp of creation; the format must be in ISO-8601/ RFC3339 string format YYYY-MM-DDTHH:mm:ss.sssZ. Default is None.
taskMetadata.shots.endedAt	str	The timestamp of when the task ended; the format must be in ISO-8601/ RFC3339 string format YYYY-MM-DDTHH:mm:ss.sssZ. Default is None.
taskMetadata.shots.status	str	The status of the task (CREATED, QUEUED, RUNNING, COMPLETED, FAILED). Default is None.
taskMetadata.shots.failureReason	str	The failure reason of the task. Default is None.
additionalMetadata.action	braket.ir.ahs.program_v1.Program	(See the Braket AHS program schema (p. 87) section)
additionalMetadata.action.braketSchemaHeader.name	str	queraMetadata.name	name of the schema; must be 'braket.task_result.quera_meta
additionalMetadata.action.braketSchemaHeader.version	str	queraMetadata.version	version of the schema
additionalMetadata.action.numSuccessfulShots	int	number of completely successful shots; must be equal to the requested number of shots

Task result field	type	description
measurements[].shotMetadata.shotStatus	int	The status of the shot, (Success, Partial success, Failure); must be "Success"

QuEra device properties schema

braket.device_schema.quera.quera_device_capabilities_v1.QueraDeviceCapabilities (example)

```
QueraDeviceCapabilities(
    service=DeviceServiceProperties(
        braketSchemaHeader=BraketSchemaHeader(
            name='braket.device_schema.device_service_properties',
            version='1'
        ),
        executionWindows=[
            DeviceExecutionWindow(
                executionDay=<ExecutionDay.TUESDAY: 'Tuesday'>,
                windowStartHour=datetime.time(16, 0),
                windowEndHour=datetime.time(20, 0)
            )
        ],
        shotsRange=(1, 1000),
        deviceCost=DeviceCost(
            price=0.01,
            unit='shot'
        ),
        deviceDocumentation=
            DeviceDocumentation(
                imageUrl='https://
a.b.cdn.console.awsstatic.com/59534b58c709fc239521ef866db9ea3f1aba73ad3ebcf60c23914ad8c5c5c878/
a6cfc6fca26cf1c2e1c6.png',
                summary='Analog quantum processor based on neutral atom arrays',
                externalDocumentationUrl='https://www.quera.com/aquila'
            ),
            deviceLocation='Boston, USA',
            updatedAt=datetime.datetime(2022, 10, 24, 15, 45, tzinfo=datetime.timezone.utc)
    ),
    action={
        <DeviceActionType.AHS: 'braket.ir.ahs.program'>: DeviceActionProperties(
                version=['1'],
                actionType=<DeviceActionType.AHS: 'braket.ir.ahs.program'>
            )
    },
    deviceParameters={},
    braketSchemaHeader=BraketSchemaHeader(
        name='braket.device_schema.quera.quera_device_capabilities',
        version='1'
    ),
    paradigm=QueraAhsParadigmProperties(
        braketSchemaHeader=BraketSchemaHeader(
            name='braket.device_schema.quera.quera_ahs_paradigm_properties',
            version='1'
        ),
        qubitCount=256,
        lattice=Lattice(
            area=Area(
                width=Decimal('0.000075'),
```

```
                height=Decimal('0.0001')
            ),
            geometry=Geometry(
                spacingRadialMin=Decimal('0.000004'),
                spacingVerticalMin=Decimal('0.000004'),
                positionResolution=Decimal('1E-7'),
                numberSitesMax=256
            )
        ),
        rydberg=Rydberg(
            c6Coefficient=Decimal('5.42E-24'),
            rydbergGlobal=RydbergGlobal(
                rabiFrequencyRange=(Decimal('0.0'), Decimal('15800000.0')),
                rabiFrequencyResolution=Decimal('400.0'),
                rabiFrequencySlewRateMax=Decimal('250000000000000.0'),
                detuningRange=(Decimal('-125000000.0'), Decimal('125000000.0')),
                detuningResolution=Decimal('0.2'),
                detuningSlewRateMax=Decimal('2500000000000000.0'),
                phaseRange=(Decimal('-99.0'), Decimal('99.0')),
                phaseResolution=Decimal('5E-7'),
                timeResolution=Decimal('1E-9'),
                timeDeltaMin=Decimal('1E-8'),
                timeMin=Decimal('0.0'),
                timeMax=Decimal('0.000004')
            )
        ),
        performance=Performance(
            lattice=PerformanceLattice(
                positionErrorAbs=Decimal('1E-7')
            ),
            rydberg=PerformanceRydberg(
                rydbergGlobal=PerformanceRydbergGlobal(
                    rabiFrequencyErrorRel=Decimal('0.02')
                )
            )
        )
    )
)
```

JSON (example)

```
{
    "service": {
        "braketSchemaHeader": {
            "name": "braket.device_schema.device_service_properties",
            "version": "1"
        },
        "executionWindows": [
            {
                "executionDay": "Tuesday",
                "windowStartHour": "16:00:00",
                "windowEndHour": "20:00:00"
            }
        ],
        "shotsRange": [1, 1000],
        "deviceCost": {
            "price": 0.01,
            "unit": "shot"
        },

        "deviceDocumentation": {
            "imageUrl": "https://
a.b.cdn.console.awsstatic.com/59534b58c709fc239521ef866db9ea3f1aba73ad3ebcf60c23914ad8c5c5c878/
a6cfc6fca26cf1c2e1c6.png",
            "summary": "Analog quantum processor based on neutral atom arrays",
```

```
            "externalDocumentationUrl": "https://www.quera.com/aquila"
        },
        "deviceLocation": "Boston, USA",
        "updatedAt": "2022-10-24T15:45:00+00:00"
    },
    "action": {
        "braket.ir.ahs.program": {
            "version": ["1"],
            "actionType": "braket.ir.ahs.program"
        }
    },
    "deviceParameters": {},
    "braketSchemaHeader": {
        "name": "braket.device_schema.quera.quera_device_capabilities",
        "version": "1"
    },
    "paradigm": {
        "braketSchemaHeader": {
            "name": "braket.device_schema.quera.quera_ahs_paradigm_properties",
            "version": "1"
        },
        "qubitCount": 256,
        "lattice": {
            "area": {
                "width": 7.5e-05,
                "height": 0.0001
            },
            "geometry": {
                "spacingRadialMin": 4e-06,
                "spacingVerticalMin": 4e-06,
                "positionResolution": 1e-07,
                "numberSitesMax": 256
            }
        },
        "rydberg": {
            "c6Coefficient": 5.42e-24,
            "rydbergGlobal": {
                "rabiFrequencyRange": [0.0,15800000.0],
                "rabiFrequencyResolution": 400.0,
                "rabiFrequencySlewRateMax": 250000000000000.0,
                "detuningRange": [-125000000.0, 125000000.0],
                "detuningResolution": 0.2,
                "detuningSlewRateMax": 2500000000000000.0,
                "phaseRange": [-99.0, 99.0],
                "phaseResolution": 5e-07,
                "timeResolution": 1e-09,
                "timeDeltaMin": 1e-08,
                "timeMin": 0.0,
                "timeMax": 4e-06
            }
        },
        "performance": {
            "lattice": {
                "positionErrorAbs": 1e-07
            },
            "rydberg": {
                "rydbergGlobal": {
                    "rabiFrequencyErrorRel": 0.02
                }
            }
        }
    }
}
```

Capabilities fields

Task result field	type	description
paradigm.qubitCount	int	the number of qubits
paradigm.lattice.area.width	Decimal	Largest allowed difference between x coordinates of any two sites (measured in meters)
paradigm.lattice.area.height	Decimal	Largest allowed difference between y coordinates of any two sites (measured in meters)
paradigm.lattice.geometry.spacingRadialMin	Decimal	Minimum radial spacing between any two sites in the lattice (measured in meters)
paradigm.lattice.geometry.spacingVerticalMin	Decimal	Minimum spacing between any two rows in the lattice (measured in meters)
paradigm.lattice.geometry.positionResolution	Decimal	Resolution with which site positions can be specified (measured in meters)
paradigm.lattice.geometry.numberSitesMax	int	Maximum number of sites that can be placed in the lattice
paradigm.rydberg.c6Coefficient	Decimal	Rydberg-Rydberg C6 interaction coefficient (measured in (rad/s)*m^6)
paradigm.rydberg.rydbergGlobal.rabiFrequencyRange	Tuple[Decimal, Decimal]	Achievable Rabi frequency range for the global Rydberg drive time

Task result field	type	description
		series (measured in rad/s)
paradigm.rydberg.rydbergGlobal.rabiFrequencyResolution	Decimal	Resolution with which global Rabi frequency amplitude can be specified (measured in rad/s)
paradigm.rydberg.rydbergGlobal.rabiFrequencySlewRateMax	Decimal	Maximum slew rate for changing the global Rabi frequency (measured in (rad/s)/s)
paradigm.rydberg.rydbergGlobal.detuningRange	Tuple[Decimal, Decimal]	Achievable detuning range for the global Rydberg pulse (measured in rad/s)
paradigm.rydberg.rydbergGlobal.detuningResolution	Decimal	Resolution with which global detuning can be specified (measured in rad/s)
paradigm.rydberg.rydbergGlobal.detuningSlewRateMax	Decimal	Maximum slew rate for detuning (measured in (rad/s)/s)
paradigm.rydberg.rydbergGlobal.phaseRange	Tuple[Decimal, Decimal]	Achievable phase range for the global Rydberg pulse (measured in rad)
paradigm.rydberg.rydbergGlobal.phaseResolution	Decimal	Resolution with which global Rabi frequency phase can be specified (measured in rad)
paradigm.rydberg.rydbergGlobal.timeResolution	Decimal	Resolution with which times for global Rydberg drive parameters can be specified (measured in s)

Task result field	type	description
paradigm.rydberg.rydbergGlobal.timeDeltaMin	Decimal	Minimum time step with which times for global Rydberg drive parameters can be specified (measured in s)
paradigm.rydberg.rydbergGlobal.timeMin	Decimal	Minimum duration of Rydberg drive (measured in s)
paradigm.rydberg.rydbergGlobal.timeMax	Decimal	Maximum duration of Rydberg drive (measured in s)

Performance fields

Task result field	type	description
paradigm.peformance.lattice.positionErrorAbs	Decimal	Error between target and actual site position (measured in meters)
paradigm.peformance.rydberg.rydbergGlobal.rabiFrequencyErrorRel	Decimal	random error in the Rabi frequency, relative (unitless)

Service properties fields

Task result field	type	description
service.executionWindows[].executionDay	ExecutionDay	Days of the execution window; must be 'Everyday', 'Weekdays', 'Weekend', 'Monday', 'Tuesday', 'Wednesday', Thursday', 'Friday', 'Saturday' or 'Sunday'
service.executionWindows[].windowStartHour	datetime.time	UTC 24-hour format of the time when the execution window starts
service.executionWindows[].windowEndHour	datetime.time	UTC 24-hour format of the time when the execution window ends
service.qpu_capabilities.service.shotsRange	Tuple[int, int]	Minimum and maximum number of shots for the device

Task result field	type	description
service.qpu_capabilities.service.deviceCost.price	float	Price of the device in terms of US dollars
service.qpu_capabilities.service.deviceCost.unit	str	unit for charging the price, e.g: 'minute', 'hour', 'shot', 'task'

Metadata fields

Task result field	type	description
paradigm.braketSchemaHeader.name	str	name of the schema; must be 'braket.device_schema.quera.quera_ah
paradigm.braketSchemaHeader.version	str	version of the schema
action[].version	str	version of the AHS program schema
action[].actionType	ActionType	AHS program schema name; must be 'braket.ir.ahs.program'
service.braketSchemaHeader.name	str	name of the schema; must be 'braket.device_schema.device_service_
service.braketSchemaHeader.version	str	version of the schema
service.deviceDocumentation.imageUrl	str	URL for the image of the device
service.deviceDocumentation.summary	str	brief description on the device
service.deviceDocumentation.externalDocumentationUrl	str	external documentation URL
service.deviceLocation	str	geographic location fo the device
service.updatedAt	datetime	time when the device properties were last updated

Monitoring and tracking tasks

After you submit a task, you can keep track of its status through the Amazon Braket SDK and console. When the task completes, Braket saves the results in your specified Amazon S3 location. Completion may take some time, especially for QPU devices, depending on the length of the queue. Status types include:

- `CREATED`
- `RUNNING`

- COMPLETED
- FAILED
- CANCELLED

In this section:

Tracking tasks from the Amazon Braket SDK

The command `device.run(…)` defines a task with a unique task ID. You can query and track the status with `task.state()` as shown in the following example.

Note: `task = device.run()` is an asynchronous operation, which means that you can keep working while the system processes your task in the background.

Retrieve a result

When you call `task.result()`, the SDK begins polling Amazon Braket to see whether the task is complete. The SDK uses the polling parameters you defined in `.run()`. After the task is complete, the SDK retrieves the result from the S3 bucket and returns it as a `QuantumTaskResult` object.

```
# create a circuit, specify the device and run the circuit
circ = Circuit().rx(0, 0.15).ry(1, 0.2).cnot(0,2)
device = AwsDevice("arn:aws:braket:::device/quantum-simulator/amazon/sv1")
task = device.run(circ, s3_location, shots=1000)

# get ID and status of submitted task
task_id = task.id
status = task.state()
print('ID of task:', task_id)
print('Status of task:', status)
# wait for job to complete
while status != 'COMPLETED':
    status = task.state()
    print('Status:', status)
```

```
ID of task:
arn:aws:braket:us-west-2:123412341234:quantum-task/b68ae94b-1547-4d1d-aa92-1500b82c300d
Status of task: QUEUED
Status: QUEUED
Status: QUEUED
Status: QUEUED
Status: QUEUED
Status: QUEUED
Status: QUEUED
Status: QUEUED
Status: RUNNING
Status: RUNNING
Status: COMPLETED
```

Cancel a task

To cancel a task, call the `cancel()` method, as shown in the following example.

```
# cancel task
```

```
task.cancel()
status = task.state()
print('Status of task:', status)
```

```
Status of task: CANCELLING
```

Check the metadata

You can check the metadata of the finished task, as shown in the following example.

```
# get the metadata of the task
metadata = task.metadata()
# example of metadata
shots = metadata['shots']
date = metadata['ResponseMetadata']['HTTPHeaders']['date']
# print example metadata
print("{} shots taken on {}.".format(shots, date))

# print name of the s3 bucket where the result is saved
results_bucket = metadata['outputS3Bucket']
print('Bucket where results are stored:', results_bucket)
# print the s3 object key (folder name)
results_object_key = metadata['outputS3Directory']
print('S3 object key:', results_object_key)

# the entire look-up string of the saved result data
look_up = 's3://'+results_bucket+'/'+results_object_key
print('S3 URI:', look_up)
```

```
1000 shots taken on Wed, 05 Aug 2020 14:44:22 GMT.
Bucket where results are stored: amazon-braket-123412341234
S3 object key: simulation-output/b68ae94b-1547-4d1d-aa92-1500b82c300d
S3 URI: s3://amazon-braket-123412341234/simulation-output/b68ae94b-1547-4d1d-
aa92-1500b82c300d
```

Retrieve a task or result

If your kernel dies after you submit the task or if you close your notebook or computer, you can reconstruct the `task` object with its unique ARN (task ID). Then you can call `task.result()` to get the result from the S3 bucket where it is stored.

```
from braket.aws import AwsSession, AwsQuantumTask

# restore task with unique arn
task_load = AwsQuantumTask(arn=task_id)
# retrieve the result of the task
result = task_load.result()
```

Advanced logging

You can record the whole task-processing process using a logger. These advanced logging techniques allow you to see the background polling and create a record for later debugging.

To use the logger, we recommend changing the `poll_timeout_seconds` and `poll_interval_seconds` parameters, so that a task can be long-running and the task status is logged continuously, with results saved to a file. You can transfer this code to a Python script instead of a Jupyter notebook, so that the script can run as a process in the background.

Configure the logger

First, configure the logger so that all logs are written into a text file automatically, as shown in the following example lines.

```
# import the module
import logging
from datetime import datetime

# set filename for logs
log_file = 'device_logs-'+datetime.strftime(datetime.now(), '%Y%m%d%H%M%S')+'.txt'
print('Task info will be logged in:', log_file)

# create new logger object
logger = logging.getLogger("newLogger")

# configure to log to file device_logs.txt in the appending mode
logger.addHandler(logging.FileHandler(filename=log_file, mode='a'))

# add to file all log messages with level DEBUG or above
logger.setLevel(logging.DEBUG)
```

```
Task info will be logged in: device_logs-20200803203309.txt
```

Create and run the circuit

Now you can create a circuit, submit it to a device to run, and see what happens as shown in this example.

```
# define circuit
circ_log = Circuit().rx(0, 0.15).ry(1, 0.2).rz(2, 0.25).h(3).cnot(control=0,
 target=2).zz(1, 3, 0.15).x(4)
print(circ_log)
# define backend
device = AwsDevice("arn:aws:braket:::device/quantum-simulator/amazon/sv1")
# define what info to log
logger.info(
    device.run(circ_log, s3_location,
               poll_timeout_seconds=1200, poll_interval_seconds=0.25, logger=logger,
 shots=1000)
    .result().measurement_counts
)
```

Check the log file

You can check what is written into the file by entering the following command.

```
# print logs
! cat {log_file}
```

```
Task arn:aws:braket:us-west-2:123412341234:quantum-
task/5088ec6c-89cf-4338-9750-9f5bb12a0dc4: start polling for completion
Task arn:aws:braket:us-west-2:123412341234:quantum-
task/5088ec6c-89cf-4338-9750-9f5bb12a0dc4: task status CREATED
Task arn:aws:braket:us-west-2:123412341234:quantum-
task/5088ec6c-89cf-4338-9750-9f5bb12a0dc4: task status CREATED
Task arn:aws:braket:us-west-2:123412341234:quantum-
task/5088ec6c-89cf-4338-9750-9f5bb12a0dc4: task status QUEUED
Task arn:aws:braket:us-west-2:123412341234:quantum-
task/5088ec6c-89cf-4338-9750-9f5bb12a0dc4: task status RUNNING
```

```
Task arn:aws:braket:us-west-2:123412341234:quantum-
task/5088ec6c-89cf-4338-9750-9f5bb12a0dc4: task status RUNNING
Task arn:aws:braket:us-west-2:123412341234:quantum-
task/5088ec6c-89cf-4338-9750-9f5bb12a0dc4: task status COMPLETED
Counter({'00001': 493, '00011': 493, '01001': 5, '10111': 4, '01011': 3, '10101': 2})
```

Get the ARN from the log file

From the log file output that's returned, as shown in the previous example, you can obtain the ARN information. With the ARN ID, you can retrieve the result of the completed task.

```
# parse log file for arn
with open(log_file) as openfile:
    for line in openfile:
        for part in line.split():
            if "arn:" in part:
                arn = part
                break
# remove final semicolon in logs
arn = arn[:-1]

# with this arn you can restore again task from unique arn
task_load = AwsQuantumTask(arn=arn, aws_session=AwsSession())

# get results of task
result = task_load.result()
```

Monitoring tasks through the Amazon Braket console

Amazon Braket offers a convenient way of monitoring the task through the Amazon Braket console. All submitted tasks are listed in the **Tasks** field as shown in the following figure. This service is *Region-specific*, which means that you can only view those tasks created in the specific AWS Region.

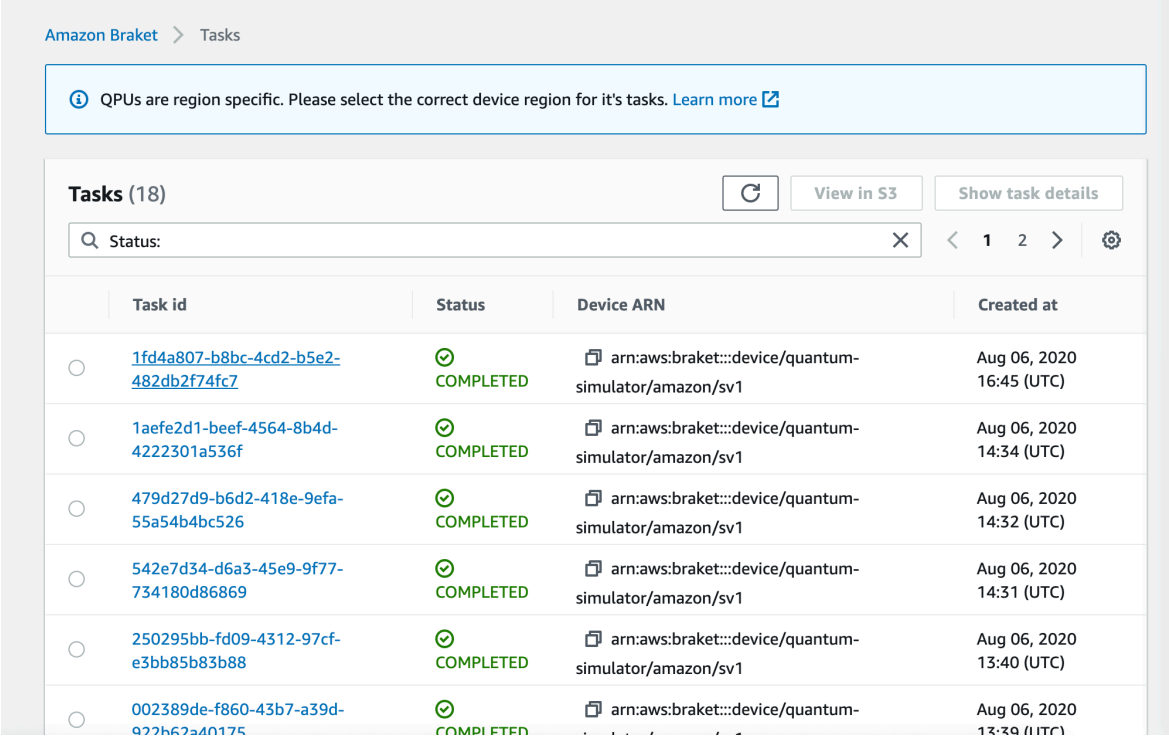

You can search for particular tasks through the navigation bar. The search can be based on Task ARN (ID), status, device, and creation time. The options appear automatically when you select the navigation bar, as shown in the following example.

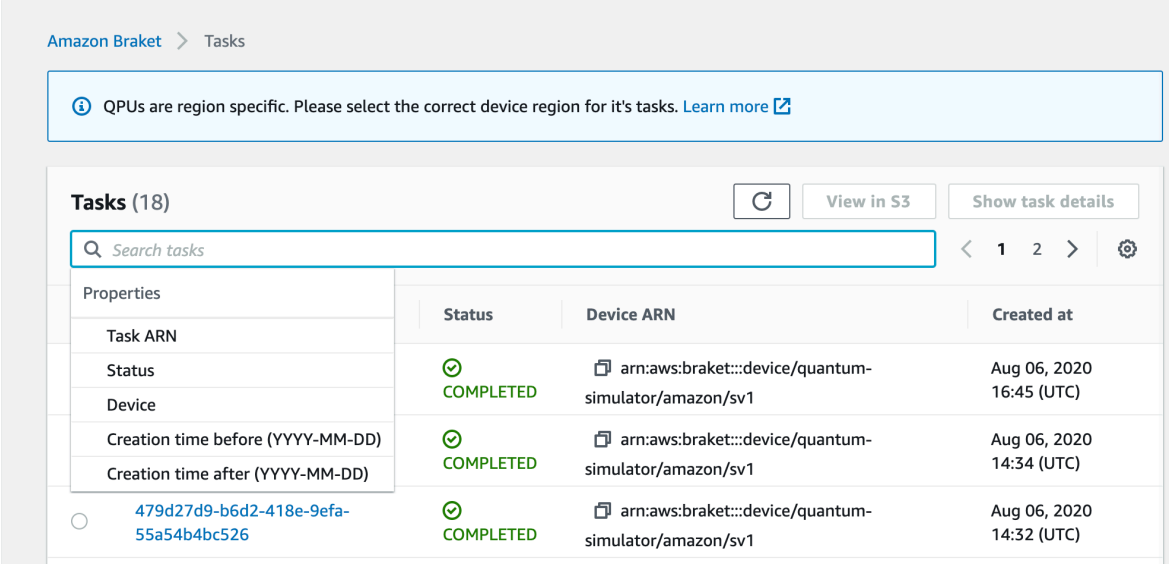

The following image shows an example of searching for a task based on its unique task ID, which can be obtained by calling `task.id`.

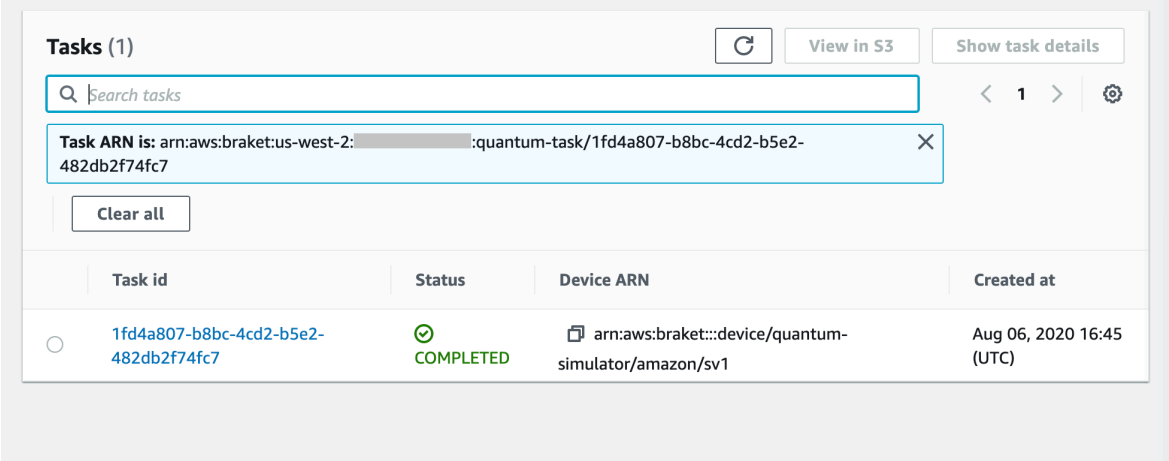

Working with Boto3

Boto3 is the AWS SDK for Python. With Boto3, Python developers can create, configure, and manage AWS services, such as Amazon Braket. Boto3 provides an object-oriented API, as well as low-level access to Amazon Braket.

Follow the instructions in the Boto3 Quickstart guide to learn how to install and configure Boto3.

Boto3 provides the core functionality that works along with the Amazon Braket Python SDK to help you configure and run your quantum tasks. Python customers always need to install Boto3, because that is the core implementation. If you want to make use of additional helper methods, you also need to install the Amazon Braket SDK.

For example, when you call `CreateQuantumTask`, the Amazon Braket SDK submits the request to Boto3, which then calls the AWS API.

In this section:

Turn on the Amazon Braket Boto3 client

To use Boto3 with Amazon Braket, you must import Boto3 and then define a client that you use to connect to the Amazon Braket API. In the following example, the Boto3 client is named `braket`.

> **Note**
> For backwards compatibility with older versions of BraketSchemas, OpenQASM information is omitted from `GetDevice` API calls. To get this information, the user-agent needs to present a recent version of the BraketSchemas (1.8.0 or later). The Braket SDK automatically reports this for you. If you do not see OpenQASM results in the `GetDevice` response when using a Braket SDK, you may need to set the `AWS_EXECUTION_ENV` environment variable to configure the user-agent. See the code examples provided in the GetDevice does not return OpenQASM results error (p. 173) topic for how to do this for the AWS CLI, Boto3, and the Go, Java, and JavaScript/TypeScript SDKs.

```
import boto3
import botocore

client = boto3.client("braket",
 config=botocore.client.Config(user_agent_extra="BraketSchemas/1.8.0"))
```

Now that you have a `braket` client established, you can make requests and process responses from the Amazon Braket service. You can get more detail on request and response data in the API Reference.

The following examples show how to work with devices and quantum tasks.

Search for devices

- `search_devices(**kwargs)`

Search for devices using the specified filters.

```
# Pass search filters and optional parameters when sending the
#  request and capture the response
response = braket.search_devices(filters=[{
    'name': 'deviceArn',
    'values': ['arn:aws:braket:::device/quantum-simulator/amazon/sv1']
}], maxResults=10)

print(f"Found {len(response['devices'])} devices")
```

```
for i in range(len(response['devices'])):
    device = response['devices'][i]
    print(device['deviceArn'])
```

Retrieve a device

- `get_device(deviceArn)`

Retrieve the devices available in Amazon Braket.

```
# Pass the device ARN when sending the request and capture the repsonse
response = braket.get_device(deviceArn='arn:aws:braket:::device/quantum-simulator/amazon/
sv1')

print(f"Device {response['deviceName']} is {response['deviceStatus']}")
```

Create a quantum task

- `create_quantum_task(**kwargs)`

Create a quantum task.

```
# Create parameters to pass into create_quantum_task()
kwargs = {
    # Create a Bell pair
    'action': '{"braketSchemaHeader": {"name": "braket.ir.jaqcd.program", "version": "1"},
 "results": [], "basis_rotation_instructions": [], "instructions": [{"type": "h", "target":
 0}, {"type": "cnot", "control": 0, "target": 1}]}',
    # Specify the SV1 Device ARN
    'deviceArn': 'arn:aws:braket:::device/quantum-simulator/amazon/sv1',
    # Specify 2 qubits for the Bell pair
    'deviceParameters': '{"braketSchemaHeader": {"name":
 "braket.device_schema.simulators.gate_model_simulator_device_parameters",
 "version": "1"}, "paradigmParameters": {"braketSchemaHeader": {"name":
 "braket.device_schema.gate_model_parameters", "version": "1"}, "qubitCount": 2}}',
    # Specify where results should be placed when the quantum task completes.
    #  You must ensure the S3 Bucket exists before calling create_quantum_task()
    'outputS3Bucket': 'amazon-braket-examples',
    'outputS3KeyPrefix': 'boto-examples',
    # Specify number of shots for the quantum task
    'shots': 100
}

# Send the request and capture the response
response = braket.create_quantum_task(**kwargs)

print(f"Quantum task {response['quantumTaskArn']} created")
```

Retrieve a quantum task

- `get_quantum_task(quantumTaskArn)`

Retrieve the specified quantum task.

```
# Pass the quantum task ARN when sending the request and capture the response
```

```
response = braket.get_quantum_task(quantumTaskArn='arn:aws:braket:us-
west-1:123456789012:quantum-task/ce78c429-cef5-45f2-88da-123456789012')

print(response['status'])
```

Search for quantum tasks

- `search_quantum_tasks(**kwargs)`

Search for tasks that match the specified filter values.

```
# Pass search filters and optional parameters when sending the
#  request and capture the response
response = braket.search_quantum_tasks(filters=[{
    'name': 'deviceArn',
    'operator': 'EQUAL',
    'values': ['arn:aws:braket:::device/quantum-simulator/amazon/sv1']
}], maxResults=25)

print(f"Found {len(response['quantumTasks'])} quantum tasks")

for n in range(len(response['quantumTasks'])):
    task = response['quantumTasks'][n]
    print(f"Quantum task {task['quantumTaskArn']} for {task['deviceArn']} is
 {task['status']}")
```

Cancel quantum task

- `cancel_quantum_task(quantumTaskArn)`

Cancel the specified task.

```
# Pass the quantum task ARN when sending the request and capture the response
response = braket.cancel_quantum_task(quantumTaskArn='arn:aws:braket:us-
west-1:123456789012:quantum-task/ce78c429-cef5-45f2-88da-123456789012')

print(f"Quantum task {response['quantumTaskArn']} is {response['cancellationStatus']}")
```

Configure AWS CLI profiles for Boto3 and the Amazon Braket SDK

The Amazon Braket SDK relies upon the default AWS CLI credentials, unless you explicitly specify otherwise. We recommend that you keep the default when you run on a managed Amazon Braket notebook because you must provide an IAM role that has permissions to launch the notebook instance.

Optionally, if you run your code locally (on an Amazon EC2 instance, for example), you can establish named AWS CLI profiles. You can give each profile a different permission set, rather than regularly overwriting the default profile.

This section provides a brief explanation of how to configure such a CLI `profile` and how to incorporate that profile into Amazon Braket so that API calls are made with the permissions from that profile.

In this section:

Step 1: Configure a local AWS CLI `profile`

It is beyond the scope of this document to explain how to create IAM users and how to configure a non-default profile. For information on these topics, see:

- Configure an IAM User
- Establish a CLI `profile`

To use Amazon Braket, you must provide this IAM user — and the associated CLI `profile` — with the necessary Braket permissions. For instance, you can attach the **AmazonBraketFullAccess** policy.

Step 2: Establish a Boto3 session object

```
from boto3 import Session

# Insert CLI profile name here
boto_sess = Session(profile_name=`profile`)
```

Note
If the expected API calls have Region-based restrictions that are not aligned with your `profile` default Region, you can specify a Region for the Boto3 session as shown in the following example.

```
# Insert CLI profile name _and_ region
boto_sess = Session(profile_name=`profile`, region_name=`region`)
```

For the argument designated as `region`, substitute a value that corresponds to one of the AWS Regions in which Amazon Braket is available such as `us-east-1`, `us-west-1`, and so forth.

Step 3: Incorporate the Boto3 session into the Braket AwsSession

The following example shows how to initialize a Boto3 Braket session and instantiate a device in that session.

```
from braket.aws import AwsSession, AwsDevice

# Initialize Braket session with Boto3 Session credentials
aws_session = AwsSession(boto_session=boto_sess)

# Instantiate any Braket QPU device with the previously initiated AwsSession
sim_arn = 'arn:aws:braket:::device/quantum-simulator/amazon/sv1'
device = AwsDevice(sim_arn, aws_session=aws_session)
```

After this setup is complete, you can submit quantum tasks to that instantiated `AwsDevice` object (by calling the `device.run(...)` command for example). All API calls made by that device can leverage the IAM credentials associated with the CLI profile that you previously designated as `profile`.

Pulse Control on Amazon Braket

This section instructs the user on how to utilize pulse control on various QPUs.

In this section:

Braket Pulse

Some of the devices on Amazon Braket allow users to submit circuits using pulses directly, which are the analog signals that control the qubits on a quantum computer. Braket Pulse is an Amazon Braket feature that enables you to work with pulses. You can access Braket Pulse through the Braket SDK, using OpenQASM 3.0, or directly through the Braket APIs. First, let's introduce some key concepts for pulse-control on Amazon Braket.

Port

A port is a software abstraction representing any input/output hardware component controlling qubits. It allows a hardware vendor to provide an interface with which users can interact to manipulate and observe qubits. Ports are characterized by a single string which represents the name of the connector. It also exposes the minimum time increment that specifies how finely the waveforms can be designed.

```
from braket.pulse import Port
Port0 = Port(“channel_0”)
```

Frame

A frame is a software abstraction that acts as both a clock within the quantum program with its time being incremented on each usage and a stateful carrier signal defined by a frequency and a phase. When transmitting signals to the qubit, a frame determines its carrier frequency, its phase offset and the time at which the waveform envelope is emitted. In Braket Pulse, constructing frames depends on the device, which can either expose a list of predefined frames or let customers instantiate new frames providing a desired port, frequency, and phase.

```
from braket.pulse import Frame
device = AwsDevice("arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2")
drive_frame = device.frame[“q0_rf_frame”]

device = AwsDevice("arn:aws:braket:eu-west-2::device/qpu/oqc/Lucy")
readout_frame = Frame(name="r0_measure", port=port0, frequency=5e9, phase=0)
```

Waveform

A waveform is a time-dependent envelope that can be used to emit signals on an output port or capture signals via an input port. Customers have the possibility to specify their waveforms directly via a list

of complex numbers or via a waveform template that will be used to generate a list by the hardware provider.

```
from braket.pulse import ArbitraryWaveform, ConstantWaveform
cst_wfm = ConstantWaveform(length=1e-7, iq=0.1)
arb_wf = ArbitraryWaveform(amplitudes=np.linspace(0, 100))
```

Braket Pulse provides a standard library of waveforms including a constant waveform, a gaussian waveform and a Derivative Removal by Adiabatic Gate, or DRAG waveform. You can retrieve the waveform data via the `sample` function to draw the shape of the waveform for instance:

```
gaussian_waveform = GaussianWaveform(1e-7, 25e-9, 0.1)
x = np.arange(0, gaussian_waveform.length, drive_frame.port.dt)
plt.plot(x, gaussian_waveform.sample(drive_frame.port.dt))
```

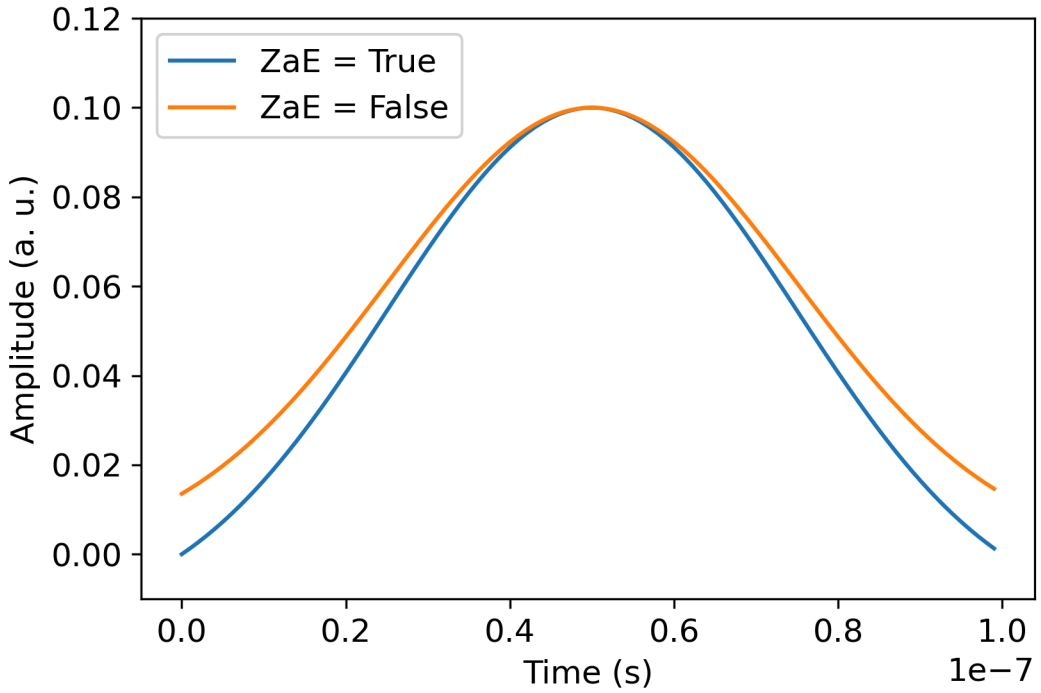

The image above depicts the Gaussian waveforms created from `GaussianWaveform`. We choose a pulse length of 100ns, a width of 25ns and an amplitude of 0.1 (arbitrary units). The waveforms are centered in the pulse window. `GaussianWaveform` accepts a boolean argument `zero_at_edges` (ZaE in the legend) which, when set to True, offsets the Gaussian waveform such that the points at t=0 and `t=length` are at zero and rescales its amplitude such that the maximum value corresponds to the `amplitude` argument.

Now that we have covered the basic concepts for pulse-level access, next we will see how to construct a circuit using gates and pulses.

Roles of frames and ports

Here we will give a description of every predefined frame and port for each device. We will also briefly discuss the mechanisms involved when pulses are played on certain frames.

Rigetti

Frames

Rigetti devices support predefined frames that have their frequency and phase calibrated to be on resonance with the associated qubit. The naming convention is `q{i}[_q{j}]_{role}_frame` where `{i}` refers to the first qubit number, `{j}` to the second qubit number in case the frame serves to activate a two-qubit interaction, and `{role}` is referring to the role of the frame as described below.

- `rf` is the frame to drive the 0-1 transition of the qubit. Pulses are transmitted as microwave transient signals of frequency and phase previously provided via the `set` and `shift` functions. The time-dependent amplitude of the signal is given by the waveform played on the frame. The frame plugs a single-qubit, off-diagonal interaction. See Krantz et al. and Rahamim et al. for more details.
- `rf_f12` is similar to `rf` and its parameters are targeting the 1-2 transition.
- `ro_rx` is used to achieve dispersive readout of the qubit via a coupled coplanar waveguide. The frequency, phase and the full set of parameter for the readout waveform are precalibrated. It is currently used via the `capture_v0` that does not require any argument besides the frame identifier.
- `ro_tx` is used to transmit signals from the resonator. It is currently unused.
- `cz` is a frame calibrated to enable the two-qubit cz gate. As all the frame associated to a `ff` port, it turns on a entangling interaction via the flux line by modulating the tunable qubit of the pair on resonance with its neighbor. See Reagor et al., Caldwell et al. and Didier et al. for more information about the entangling mechanism.
- `cphase` is a frame calibrated to enable the two-qubit `cphaseshift` gate and is linked to a `ff` port. See paragraph about the cz frame for more information about the entangling mechanism.
- `xy` is a frame calibrated to enable the two-qubit XY(θ) gates and is linked to a ff port. See paragraph about the `cz` frame and Abrams et al. for more information about the entangling mechanism and how to achieve XY gates.

As frames based on the `ff` port shift the frequency of the tunable qubit, all the other driving frames related to the qubit will be dephased by an amount that is related to the amplitude and the duration of the frequency shift. Consequently, you must compensate this effect by adding a corresponding phase shift to the frames of the neighboring qubits.

Ports

The Rigetti devices provide a list of ports that can be inspected via the device capabilities. Port names follow the convention `q{i}_{type}` with `{i}` the qubit number and `{type}` the type of the port. Please notice that not all the qubits have a complete set of ports. Next is the list of types with a description.

- `rf` represents the main interface to drive the single-qubit transition. It is associated with the `rf` and `rf_f12` frames. It is capacitively coupled to the qubit, allowing microwave driving in the gigahertz range.
- `ro_tx` serves to transmit signals to the readout resonator capacitively coupled to the qubit. Readout signal delivery is multiplexed eight-fold by octagon.
- `ro_rx` serves to receive signals from the readout resonator coupled to the qubit.
- `ff` is the port representing the fast-flux line inductively coupled to the qubit which allows to tune the frequency of the transmon. Only qubits designed to be highly tunable have an `ff` port. The port serves to activate qubit-qubit interaction as there is a static capacitive coupling between each pair of neighboring transmons.

For more information of the architecture, please see Valery et al..

OQC

OQC devices support predefined frames that have their frequency and phase calibrated to be on resonance with the associated qubit. The naming convention is:

Frames

- driving frame: `q{i}[_q{j}]_{role}` where `{i}` refers to the first qubit number, `{j}` to the second qubit number in case the frame serves to activate a two-qubit interaction, and `{role}` is refers to the role of the frame as described below.
- qubit readout frame: `r{i}_{role}` where `{i}` refers to the qubit number and `{role}` is referring to the frame role as described below.

It is recommended to use each frame for its designed role:

- `drive` is used as the main frame to drive the 0-1 transition of the qubit. Pulses are transmitted as microwave transient signals of frequency and phase previously provided via the `set` and `shift` functions. The time-dependent amplitude of the signal is given by the waveform played on the frame. The frame plugs a single-qubit, off-diagonal interaction. See Krantz et al. and Rahamim et al. for more details.
- `second_state` is equivalent to the `drive` frame but its frequency is tuned on resonance with the 1-2 transition.
- `measure` is for readout. The frequency, phase, and the full set of parameters for the readout waveform are precalibrated. It is currently used via the `capture_v0` that does not require any argument besides the frame identifier.
- `acquire` is used to capture signals from the resonator. It is currently unused.
- `cross_resonance` activates the cross resonance interaction between the qubits `i` and `j` by driving the control qubit `i` at the transition frequency of the target qubit `j`. Consequently, the frame frequency is set of the frequency of the target qubit. The interaction occurs with a rate proportional to the amplitude of this cross-resonant drive. Different types of crosstalks induces unwanted effects which requires corrections. See Patterson et al. for more information about the cross-resonance interaction with coaxially-shaped transmon qubits ('coaxmons').
- `cross_resonance_cancellation` allows you to add corrections to suppress deleterious effects induced by crosstalks when the cross-resonance interaction is activated. The initial frame frequency is set to the transition frequency of the control qubit `i`. See Patterson et al. for more information about the cancellation method.

Ports

The OQC devices provide a list of ports that can be inspected via the device capabilities. The previously described frames are associated with ports that are identified by their id `channel_{N}` where `{N}` is an integer. Ports are the interface to control lines (direction `tx`) and readout resonators (direction `rx`) connected to coaxmons. Each qubit is associated to one control line and one readout resonator. The transmission port is the interface for single-qubit and two-qubit manipulation. The reception port serves for qubit readout.

Hello Pulse

Here, you will learn how to construct a simple Bell pair directly with pulses, and execute this pulse program on the Rigetti device. A Bell pair is a two-qubit circuit consisting of a Hadamard gate on the first qubit, followed by a CNOT gate between the first and second qubits. Creating entangled states with pulses requires specific mechanisms that are dependent on the hardware type and device architecture.

Specifically here, we will not use a native mechanism to create the CNOT gate. We will therefore use specific waveforms and frames that enable the CZ gate natively. In this example, we will create a Hadamard gate using the single-qubit native gates Rx and Rz and express the CZ gate using pulses.

First, let's import the necessary libraries. In addition to the `Circuit` class, you will now also need to import the `PulseSequence` class.

```
from braket.aws import AwsDevice
from braket.pulse import PulseSequence, ArbitraryWaveform, GaussianWaveform

from braket.circuits import Circuit
import braket.circuits.circuit as circuit
```

Next, instantiate a new braket device using the ARN of the Rigetti Aspen M-2 device. Refer to the Devices page on the Amazon Braket console to view the layout of the Rigetti Aspen-M-2 device.

```
a=10 #specifies the control qubit
b=113 #specifies the target qubit
device = AwsDevice("arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2")
```

As the Hadamard gate is not a native gate of the Rigetti device, it cannot be used in combination with pulses. You need to therefore decompose it into a sequence of the Rx and Rz native gates.

```
import numpy as np
import matplotlib.pyplot as plt
@circuit.subroutine(register=True)
def rigetti_native_h(q0):
    return (
        Circuit()
        .rz(q0, np.pi)
        .rx(q0, np.pi/2)
        .rz(q0, np.pi/2)
        .rx(q0, -np.pi/2)
    )
```

For the CZ gate, we will use an arbitrary waveform with parameters (amplitude, rise/fall time, duration) that have been predetermined by the hardware provider during a calibration stage. This waveform will be applied on the `q10_q113_cz_frame`. Customers who want a more recent version of the arbitrary waveform used here can retrieve one from QCS, the cloud computing platform for Rigetti. This requires the user to create a QCS account.

```
a_b_cz_wfm = ArbitraryWaveform([0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
 0.00017888439538396808, 0.00046751103636033026, 0.0011372942989106456,
 0.002577059611929697, 0.005443941944632366, 0.010731922770068104, 0.01976701723583167,
 0.03406712171899736, 0.05503285980691202, 0.08350670755829034, 0.11932853352131022,
 0.16107456696238298, 0.20614055551722368, 0.2512065440720643, 0.292952577513137,
 0.328774403476157, 0.3572482512275353, 0.3782139893154499, 0.3925140937986156,
 0.40154918826437913, 0.4068371690898149, 0.4097040514225177, 0.41114381673553674,
 0.411813599998087, 0.4121022266390633, 0.4122174383870584, 0.41226003881132406,
 0.4122746298554775, 0.4122792591252675, 0.4122806196003006, 0.41228098995582513,
 0.41228108334474756, 0.4122811051578895, 0.4122811098772742, 0.4122811108230642,
 0.4122811109986316, 0.41228111102881937, 0.41228111103362725, 0.4122811110343365,
 0.41228111103443343, 0.4122811110344457, 0.4122811110344471, 0.41228111103444737,
 0.41228111103444737, 0.41228111103444737, 0.41228111103444737, 0.41228111103444737,
 0.41228111103444737, 0.41228111103444737, 0.41228111103444737, 0.41228111103444737,
 0.41228111103444737, 0.41228111103444737, 0.41228111103444737, 0.41228111103444737,
 0.41228111103444737, 0.41228111103444737, 0.41228111103444737, 0.41228111103444737,
 0.41228111103444737, 0.41228111103444737, 0.41228111103444737, 0.41228111103444737,
 0.4122811110344471, 0.4122811110344457, 0.41228111103443343, 0.4122811110343365,
```

```
 0.41228111103362725, 0.41228111102881937, 0.4122811109986316, 0.4122811108230642,
 0.4122811098772742, 0.4122811051578895, 0.41228108334474756, 0.41228098995582513,
 0.4122806196003006, 0.4122792591252675, 0.4122746298554775, 0.41226003881132406,
 0.4122174383870584, 0.4121022266390633, 0.411813599998087, 0.41114381673553674,
 0.4097040514225176, 0.4068371690898149, 0.40154918826437913, 0.3925140937986155,
 0.37821398931544986, 0.3572482512275351, 0.32877440347615655, 0.2929525775131368,
 0.2512065440720641, 0.20614055551722307, 0.16107456696238268, 0.11932853352131002,
 0.08350670755829034, 0.05503285980691184, 0.03406712171899729, 0.01976701723583167,
 0.010731922770068058, 0.005443941944632366, 0.002577059611929697, 0.001137294298910622 9,
 0.0004675110363603302 6, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0])
a_b_cz_frame = device.frames[f'q{a}_q{b}_cz_frame']

dt = a_b_cz_frame.port.dt
a_b_cz_wfm_duration = len(a_b_cz_wfm.amplitudes)*dt
print('CZ pulse duration:', a_b_cz_wfm_duration*1e9, 'ns')
```

This should return:

```
CZ pulse duration: 124 ns
```

Now we can construct the CZ gate using the waveform defined above. Recall that the CZ gate consists of a phase flip of the target qubit if the control qubit is in the |1> state.

```
phase_shift_a=1.1733407221086924
phase_shift_b=6.269846678712192

a_rf_frame = device.frames[f'q{a}_rf_frame']
b_rf_frame = device.frames[f'q{b}_rf_frame']

frames = [a_rf_frame, b_rf_frame, a_b_cz_frame]

cz_pulse_sequence = (
    PulseSequence()
    .barrier(frames)
    .play(a_b_cz_frame, a_b_cz_wfm)
    .delay(a_rf_frame, a_b_cz_wfm_duration)
    .shift_phase(a_rf_frame, phase_shift_a)
    .delay(b_rf_frame, a_b_cz_wfm_duration)
    .shift_phase(b_rf_frame, phase_shift_b)
    .barrier(frames)
)
```

The `a_b_cz_wfm` waveform is played on a frame that is associated to a fast-flux port. Its role is to shift the qubit frequency to activate a qubit-qubit interaction, see Roles of frames and ports (p. 110). As the frequency varies, the qubit frames rotate at different rates than the single-qubit `rf` frames that are kept untouched: the latter ones are getting dephased. These phase shifts were calibrated via Ramsey sequences beforehand and are provided here as hardcoded information via `phase_shift_a` and `phase_shift_b` (Full period). We correct this dephasing by using `shift_phase` instructions on `rf` frames. Note that this sequence will only work in programs where no XY frame related to qubit `a` and `b` is used as we do not compensate for the phase shift that occurs on these frames. This is the case for this single Bell pair program which uses only `rf` and `cz` frames (see Caldwell et al. for more information).

Now you are ready to create a Bell pair with pulses:

```
bell_circuit_pulse = (
    Circuit()
    .rigetti_native_h(a)
    .rigetti_native_h(b)
    .pulse_gate([a, b], cz_pulse_sequence)
    .rigetti_native_h(b)
)
```

```
print(bell_circuit_pulse)
```

```
T  : |   0    |   1    |   2    |    3    |4 |   5    |   6    |   7    |    8    |

q5 : -Rz(3.14)-Rx(1.57)-Rz(1.57)-Rx(-1.57)-PG-------------------------------------
                                            |
q6 : -Rz(3.14)-Rx(1.57)-Rz(1.57)-Rx(-1.57)-PG-Rz(3.14)-Rx(1.57)-Rz(1.57)-Rx(-1.57)-

T  : |   0    |   1    |   2    |    3    |4 |   5    |   6    |   7    |    8    |
```

Let's run this Bell pair on the Rigetti device. Note that running this code block will incur a charge. For more information about pricing, refer to this page. We recommend that you test your circuits using a small amount of shots to ensure that it can run on the device, before increasing the shot count.

```
task = device.run(bell_pair_pulses, shots=100)

counts = task.result().measurement_counts

plt.bar(sorted(counts), [counts[k] for k in sorted(counts)])
```

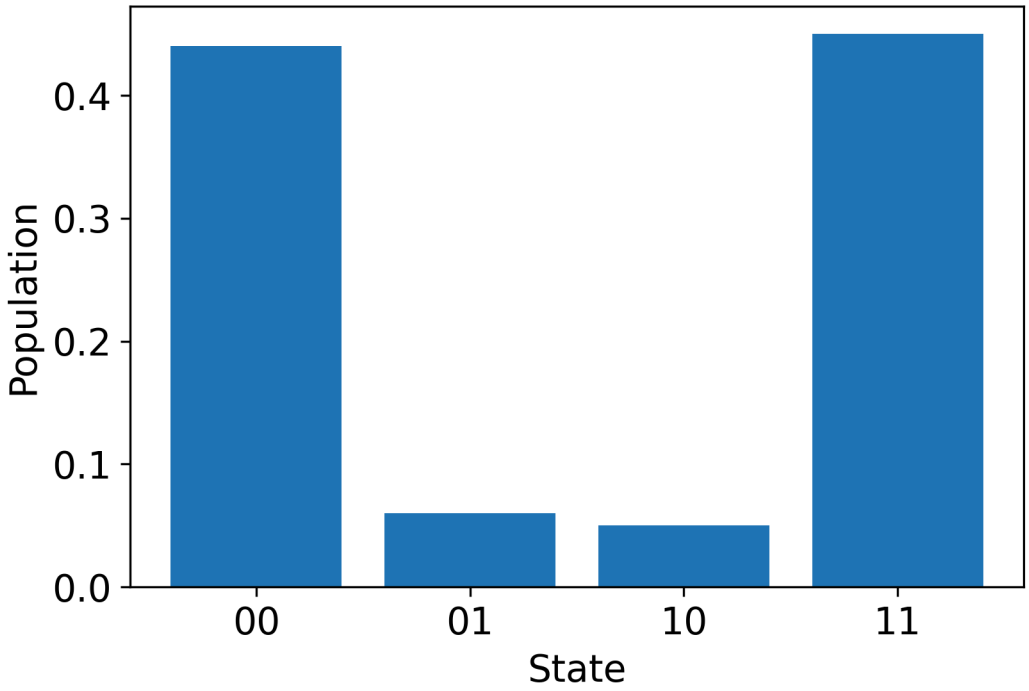

Hello Pulse using OpenPulse

OpenPulse is a language for specifying pulse-level control of a general quantum device and is part of the OpenQASM 3.0 specification. Amazon Braket supports OpenPulse for directly programming pulses using the OpenQASM 3.0 representation.

Amazon Braket uses OpenPulse as the underlying intermediate representation for expressing pulses in native instructions. OpenPulse supports the addition of instruction calibrations in the form of `defcal`

(short for “define calibration”) declarations which allow the programmer to specify an implementation of a gate instruction within a lower-level control grammar.

In this example, we will construct a Bell circuit using OpenQASM 3.0 and OpenPulse on a device using frequency-tunable transmons. Recall that a Bell circuit is a two-qubit circuit that consists of a Hadamard gate on the first qubit followed by a CNOT gate between the two qubits. As CNOT gates differ from CZ gates only via a basis transform, here we will define a Bell pair using Hadamard and CZ gates instead as the device provides a simpler way to create CZ gates for this demonstration.

Begin with defining the Hadamard gate using native gates of the device:

```
client = boto3.client('braket', region_name='us-west-1')
defcal h $10 {
    rz(pi) $10;
    rx(pi/2) $10;
    rz(pi/2) $10;
    rx(-pi/2) $10;
}
defcal h $113 {
    rz(pi) $113;
    rx(pi/2) $113;
    rz(pi/2) $113;
    rx(-pi/2) $113;
}
```

For the CZ gate, we will use an arbitrary waveform with parameters (amplitude, rise/fall time, duration) that have been predetermined beforehand. This waveform will be applied on the `q10_q113_cz_frame`.

```
cal {
    waveform q10_q113_cz_wfm = {0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
 0.00021019328936380065, 0.0008644760431374357, 0.003028104466884364, 0.009058671036471382,
 0.02322670104785881, 0.05128438687551476, 0.09812230691191462, 0.16403303942241076,
 0.24221990600377236, 0.32040677258513395, 0.38631750509563006, 0.43315542513203,
 0.4612131109596859, 0.4753811409710734, 0.4814117075406603, 0.48357533596440727,
 0.48422961871818093, 0.4843963766398558, 0.4844321964728096, 0.4844386806183817,
 0.4844396697373718, 0.48443979687791755, 0.48443981064783953, 0.48443981190433844,
 0.4844398120009317, 0.48443981200718716, 0.4844398120075284, 0.48443981200754405,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
```

```
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.48443981200754405, 0.4844398120075284, 0.48443981200718716,
 0.4844398120009317, 0.48443981190433844, 0.48443981064783953, 0.48443979687791755,
 0.4844396697373718, 0.4844386806183817, 0.4844321964728096, 0.4843963766398558,
 0.48422961871818093, 0.48357533596440727, 0.4814117075406603, 0.4753811409710734,
 0.46121311095968553, 0.4331554251320285, 0.38631750509562957, 0.32040677258513167,
 0.24221990600377236, 0.16403303942240913, 0.0981223069119151, 0.0512843868755143,
 0.023226701047858084, 0.009058671036471328, 0.0030281044668842563, 0.0008644760431374626,
 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0};
}
defcal cz $10, $113 {
    barrier q10_rf_frame, q113_rf_frame, q10_q113_cz_frame;
    play(q10_q113_cz_frame, q10_q113_cz_wfm);
    delay[124ns] q10_rf_frame;
    shift_phase(q10_rf_frame, 1.1733407221086924);
    delay[124ns] q113_rf_frame;
    shift_phase(q113_rf_frame, 6.269846678712192);
    barrier q10_rf_frame, q113_rf_frame, q10_q113_cz_frame;
}
```

The duration of the `q10_q113_cz_wfm` waveform is 124 samples, which corresponds to 124ns as the minimum time increment `dt` is 1ns.

The `q10_q113_cz_wfm` waveform is played on a frame that is bound to a fast-flux port. Its role is to shift the qubit frequency to activate a qubit-qubit interaction, see Roles of frames and ports (p. 110). As the frequency varies, the qubit frames rotate at different rates compared to the single-qubit `rf` frames that are kept untouched: the latter ones are getting dephased. This dephasing can be measured with Ramsey sequences during a calibration stage and compensated with `shift_phase` instructions on `rf` and `xy` frames (see Caldwell et al. for more information).

Customers can then simply execute the Bell pair circuit where we decomposed CNOT with a couple of Hadamard and CZ gates.

```
bit[2] c;
```

```
h $10;
h $113;
cz $10, $113;
h $113;
c[0] = measure $10;
c[1] = measure $113;
```

The full OpenQASM 3.0 representation for the Bell circuit constructed using a combination of native gates and pulses is as follows:

```
// bell_pair_with_pulse.qasm
OPENQASM 3.0;
cal {
    waveform q10_q113_cz_wfm = {0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0,
 0.00021019328936380065, 0.0008644760431374357, 0.003028104466884364, 0.009058671036471382,
 0.0232267010478581, 0.05128438687551476, 0.09812230691191462, 0.1640330394224107б,
 0.24221990600377236, 0.32040677258513395, 0.38631750509563006, 0.43315542513203,
 0.4612131109596859, 0.4753811409710734, 0.4814117075406603, 0.48357533596440727,
 0.48422961871818093, 0.4843963766398558, 0.4844321964728096, 0.4844386806183817,
 0.4844396697373718, 0.48443979687791755, 0.48443981064783953, 0.48443981190433844,
 0.4844398120009317, 0.48443981200718716, 0.4844398120075284, 0.48443981200754405,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
```

```
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.4844398120075447, 0.4844398120075447, 0.4844398120075447,
 0.4844398120075447, 0.48443981200754405, 0.4844398120075284, 0.48443981200718716,
 0.4844398120009317, 0.48443981190433844, 0.48443981064783953, 0.48443979687791755,
 0.4844396697373718, 0.4844386806183817, 0.4844321964728096, 0.4843963766398558,
 0.48422961871818093, 0.48357533596440727, 0.4814117075406603, 0.4753811409710734,
 0.46121311095968553, 0.4331554251320285, 0.38631750509562957, 0.32040677258513167,
 0.24221990600377236, 0.16403303942240913, 0.0981223069119151, 0.0512843868755143,
 0.023226701047858084, 0.009058671036471328, 0.0030281044668842563, 0.0008644760431374626,
 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0};
}
defcal h $10 {
    rz(pi) $10;
    rx(pi/2) $10;
    rz(pi/2) $10;
    rx(-pi/2) $10;
}
defcal h $113 {
    rz(pi) $113;
    rx(pi/2) $113;
    rz(pi/2) $113;
    rx(-pi/2) $113;
}
defcal cz $10, $113 {
    barrier q10_rf_frame, q113_rf_frame, q10_q113_cz_frame;
    play(q10_q113_cz_frame, q10_q113_cz_wfm);
    delay[124ns] q10_rf_frame;
    shift_phase(q10_rf_frame, 1.1733407221086924);
    delay[124ns] q113_rf_frame;
    shift_phase(q113_rf_frame, 6.269846678712192);
    barrier q10_rf_frame, q113_rf_frame, q10_q113_cz_frame;
}
bit[2] c;
h $10;
h $113;
cz $10, $113;
h $113;
c[0] = measure $10;
c[1] = measure $113;
```

You can now use the Braket SDK to execute this OpenQASM 3.0 program on the Rigetti device using the code below:

```
# import the device module
from braket.aws import AwsDevice
from braket.ir.openqasm import Program

client = boto3.client('braket', region_name='us-west-1')

with open("pulse.qasm", "r") as pulse:
    pulse_qasm_string = pulse.read()

# choose the Rigetti device
```

```
device = AwsDevice("arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2")

program = Program(source=pulse_qasm_string)
my_task = device.run(program)

# You can also specify an optional s3 bucket location and number of shots,
 # if you so choose, when running the program
s3_location = ("amazon-braket-my-bucket", "openqasm-tasks")
my_task = device.run(
    program,
    s3_location,
    shots=100,
)
```

Amazon Braket Hybrid Jobs User Guide

This section provides instructions about how to set up and manage hybrid jobs in Amazon Braket.

You can access hybrid jobs in Braket using:

- The Amazon Braket Python SDK.
- The Amazon Braket console.
- The Amazon Braket API.

In this section:

What is a Hybrid Job

With Hybrid Jobs, you can run algorithms that use both classical and quantum compute resources. These hybrid quantum-classical algorithms can help you optimize the performance of the quantum devices currently available. These algorithms are designed to maximize the benefits and reduce the drawbacks of both types of computational systems and are generally used to find the ground state or global minimum of a particular system.

You can submit Hybrid Jobs to any quantum processing unit (QPU) or quantum simulator. You'll need to define how the algorithm is implemented by submitting a Python script and perhaps also configuring input data or hyperparameters. This script, along with any input data and hyperparameters, is part of the environment that defines how the Hybrid Job is run. This environment is then placed inside a container, which is accessed by the Hybrid Job once it's been submitted. To learn more about Hybrid Jobs and how to use them, visit the Getting started with Amazon Braket Hybrid Jobs tutorial.

When to use Amazon Braket Hybrid Jobs

Amazon Braket Hybrid Jobs enables you to easily run hybrid quantum-classical algorithms, such as the Variational Quantum Eigensolver (VQE) and the Quantum Approximate Optimization Algorithm (QAOA), that combine classical compute resources with quantum computing devices to optimize the performance of today's quantum systems. Amazon Braket Hybrid Jobs provides three main benefits:

1. **Convenience**: Amazon Braket Hybrid Jobs simplifies setting up and managing your compute environment and keeping it running while your hybrid algorithm runs. You just need to provide your algorithm script and select a quantum device (either a quantum processing unit or a simulator) on which to run. Amazon Braket waits for the target device to become available, spins up the classical resources, runs the workload in pre-built container environments, returns the results to Amazon Simple Storage Service (Amazon S3), and releases the compute resources.
2. **Metrics**: Amazon Braket Hybrid Jobs provides on-the-fly insights into running algorithms and delivers customizable algorithm metrics in near real-time to Amazon CloudWatch and the Amazon Braket console so you can track the progress of your algorithms.
3. **Performance**: Amazon Braket Hybrid Jobs provides better performance than running hybrid algorithms from your own environment. While your job is running, it has priority access to the selected target QPU: tasks from your job run ahead of other tasks queued on the device. This results in shorter and more predictable runtimes for hybrid algorithms.

Run a job with Amazon Braket Hybrid Jobs

To run a job with Amazon Braket Hybrid Jobs, you first need to define your algorithm. You can define it by writing the *algorithm script* and, optionally, other dependency files using the Amazon Braket Python SDK or PennyLane. If you want to use other (open source or proprietary) libraries, you can define your own custom container image using Docker, which includes these libraries. For more information, see Bring your own container (BYOC) (p. 146).

In either case, next you create a job using the Amazon Braket API, where you provide your algorithm script or container, select the target quantum device the job is to use, and then choose from a variety of optional settings. The default values provided for these optional settings work for the majority of use cases. For the target device to run your Hybrid Job, you have a choice between a QPU, an on-demand simulator (such as SV1, DM1 or TN1), or the classical job instance itself. With an on-demand simulator or QPU, your hybrid jobs container makes API calls to a remote device. With the embedded simulators, the simulator is embedded in the same container as your algorithm script. The lightning simulators from PennyLane are embedded with the default pre-built jobs container for you to use. If you run your code using an embedded PennyLane simulator or a custom simulator, you can specify an instance type as well as how many instances you wish to use. Refer to the Amazon Braket Pricing page for the costs associated with each choice.

If your target device is a managed or embedded simulator, Amazon Braket starts running the job right away. It spins up the job instance (you can customize the instance type in the API call), runs your algorithm, writes the results to Amazon S3, and releases your resources. This release of resources ensures that you only pay for what you use.

The total number of concurrent jobs per quantum processing unit (QPU) is restricted. Queues are used to control the number of jobs allowed to run so as not to exceed the limit allowed. If your target device is a QPU, your job first enters the *job queue* of the selected QPU. Once your job has moved up to first position and the device is ready to start a new job, Amazon Braket spins up the job instance needed and runs your job on the device. For the duration of your algorithm, your job has priority access, meaning that tasks from your job run ahead of other tasks queued up on the device. You are only billed when your job starts, which means you're not billed for any wait time in the job queue.

> **Note**
> Devices are regional and your job runs in the same AWS Region as your primary device.

In both the simulator and QPU target scenarios, you have the option to define custom algorithm metrics, such as the energy of your Hamiltonian, as part of your algorithm. These metrics are automatically reported to Amazon CloudWatch and from there, they display in near real-time in the Amazon Braket console.

> **Note**
> If you wish to use a GPU based instance, be sure to use one of the GPU-based simulators available with the embedded simulators on Braket (for example, `lightning.gpu`). If you choose one of the CPU-based embedded simulators (for example, `lightning.qubit`, or `braket:default-simulator`), the GPU will not be used and you may incur unnecessary costs.

Create your first job

This section shows you how create a basic first job.

Before you run your first job, you must ensure that you have sufficient permissions to proceed with this task. To determine that you have the correct permissions, select **Permissions** from the menu on left side of the Braket Console. The **Permissions management for Amazon Braket** page helps you verifiy whether one of your existing roles has permissions that are sufficient to run your job or guides you through the creation of a default role that can be used to run your job if you do not already have such a role.

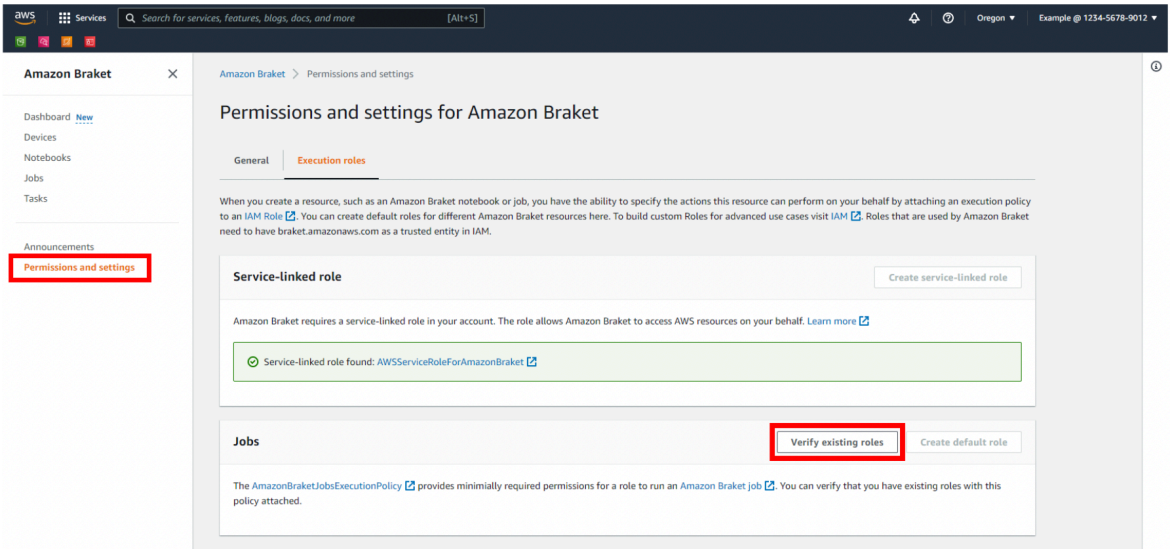

To verify that you have roles with suficient permissions to execute a job, select the **Verify existing role** button. If you do, you get a message that the roles were found. To see the names of the roles and their role ARNs, select the **Show roles** button.

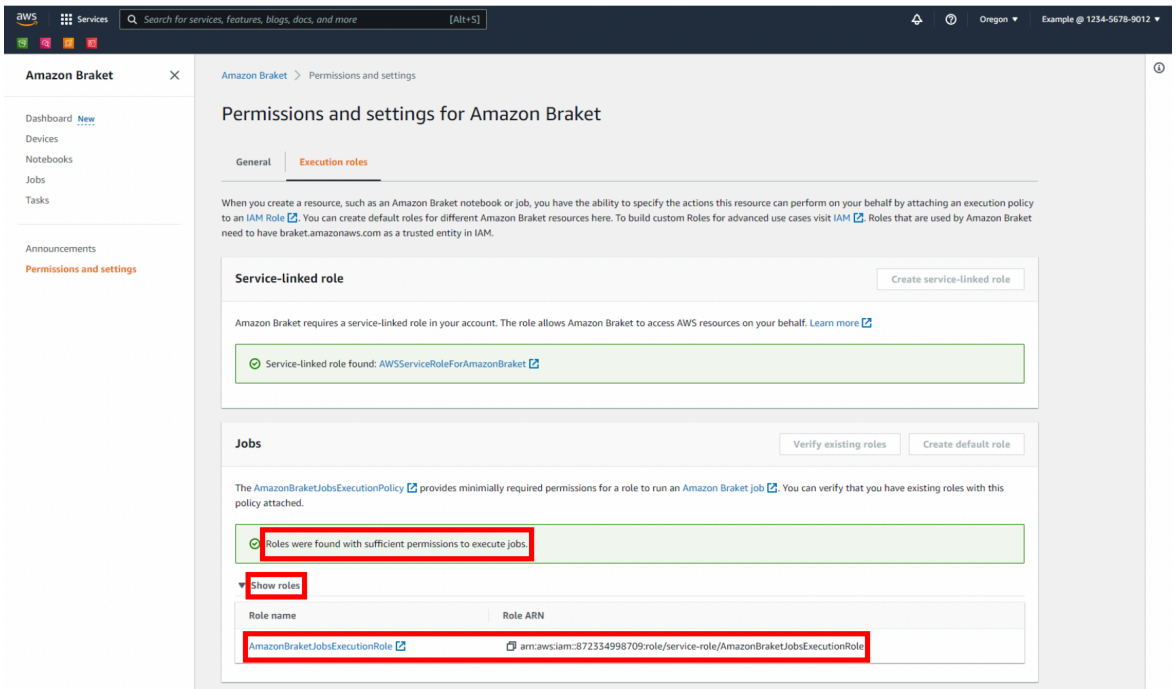

If you do not have a role with sufficient permissions to execute a job, you get a message that no such role was found. Select the **Create default role** button to obtain a role with sufficient permissions.

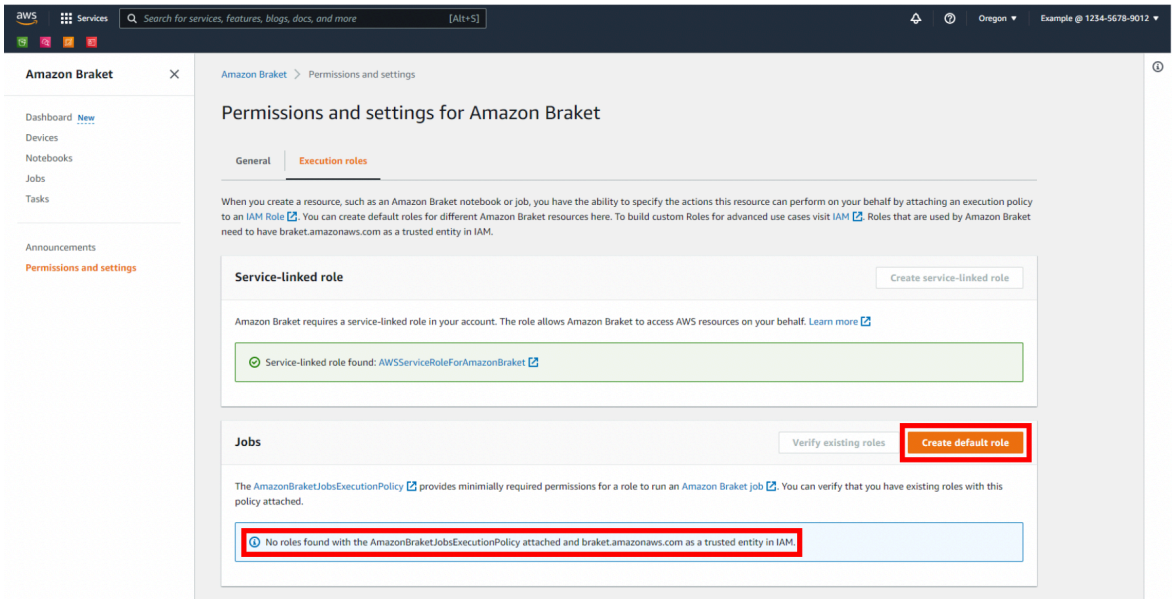

If the role was created successfully, you get a message confirming this.

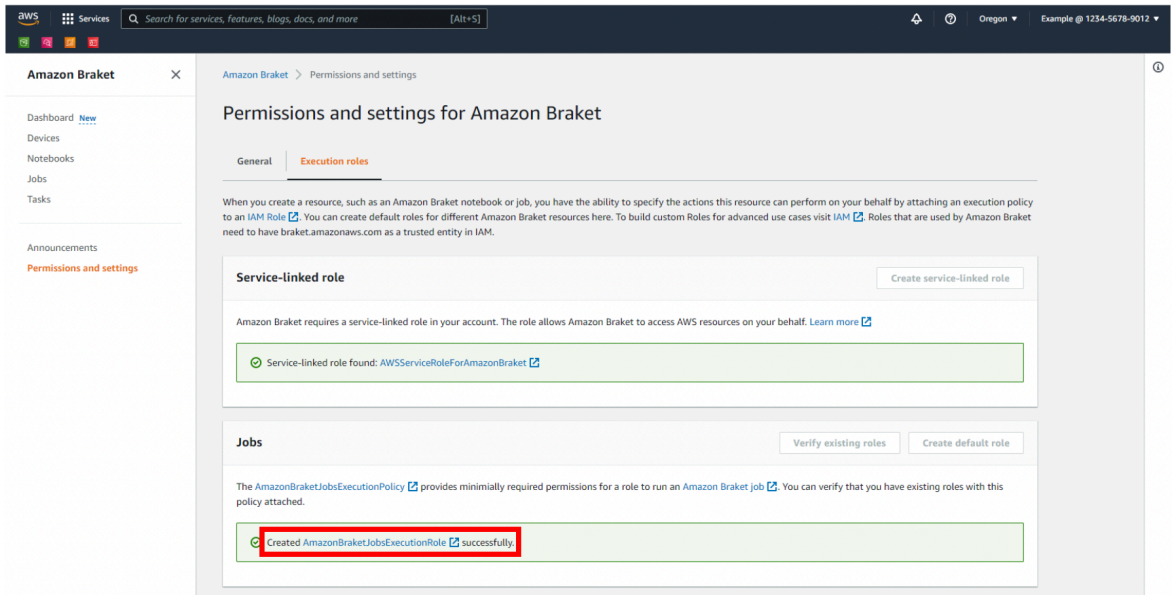

If you do not have permissions to make this inquiry, you will be denied access. In this case, contact your internal AWS administrator.

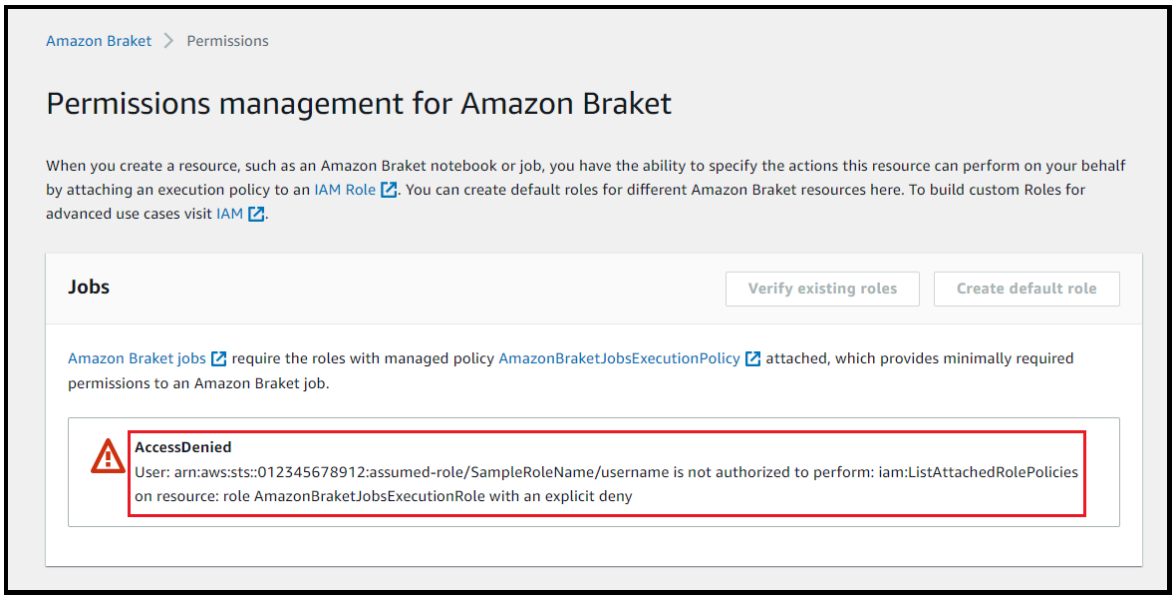

Once you have a role with permissions to run a job, you are ready to proceed. The key piece of your first Braket job is the *algorithm script*. It defines the algorithm you want to run and contains the classical logic and quantum tasks that are part of your algorithm. In addition to your algorithm script, you can provide other dependency files. The algorithm script together with its dependencies is called the *source module*. The *entry point* defines the first file or function to run in your source module when the job starts.

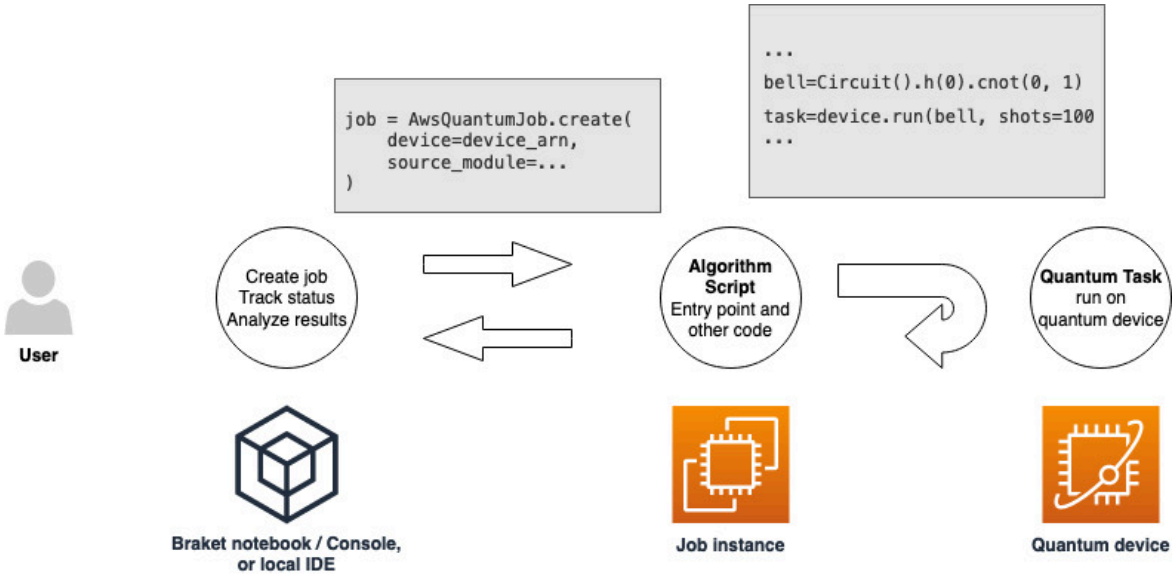

First, consider the following basic example of an algorithm script that creates five bell states and prints the corresponding measurement results.

```
import os

from braket.aws import AwsDevice
from braket.circuits import Circuit

def start_here():

    print("Test job started!!!!!")
```

```
    # Use the device declared in the job script
    device = AwsDevice(os.environ["AMZN_BRAKET_DEVICE_ARN"])

    bell = Circuit().h(0).cnot(0, 1)
    for count in range(5):
        task = device.run(bell, shots=100)
        print(task.result().measurement_counts)

    print("Test job completed!!!!!")
```

Save this file with the name *algorithm_script.py* in your current working directory on your Braket notebook or local environment. The algorithm_script.py file has `start_here()` as the planned entry point.

Next, create a Python file or Python notebook in the same directory as the algorithm_script.py file. This script kicks off the job and handles any asynchronous processing, such as printing the status or key outcomes that we are interested in. At a minimum, this script needs to specify your job script and your primary device.

> **Note**
> For more information about how to create a Braket notebook or upload a file, such as the *algorithm_script.py* file, in the same directory as the notebooks, see Run your first circuit using the Amazon Braket Python SDK (p. 38)

For this basic first case, you target a simulator. Whichever type of quantum device you target, a simulator or an actual quantum processing unit (QPU), the device you specify with `device` in the following script is used to schedule the job and is available to the algorithm scripts as the environment variable `AMZN_BRAKET_DEVICE_ARN`.

> **Note**
> You can only use devices that are available in the AWS Region of your job. The Amazon Braket SDK autoselects this AWS Region. For example, a job in us-east-1 can use IonQ, SV1, DM1, and TN1 devices, but not Rigetti devices.

If you choose a quantum computer instead of a simulator, Braket schedules your jobs to run all of their tasks with priority access.

```
from braket.aws import AwsQuantumJob

job = AwsQuantumJob.create(
    "arn:aws:braket:::device/quantum-simulator/amazon/sv1",
    source_module="algorithm_script.py",
    entry_point="algorithm_script:start_here",
    wait_until_complete=True
)
```

The parameter `wait_until_complete=True` sets a verbose mode so that your job prints output from the actual job as it's running. You should see an output similar to the following example.

```
job = AwsQuantumJob.create(
    "arn:aws:braket:::device/quantum-simulator/amazon/sv1",
    source_module="algorithm_script.py",
    entry_point="algorithm_script:start_here",
    wait_until_complete=True,
)
Initializing Braket Job: arn:aws:braket:us-west-2:<accountid>:job/braket-job-
default-1631915042705
.....................................
.
.
.
```

```
Completed 36.1 KiB/36.1 KiB (692.1 KiB/s) with 1 file(s) remaining#015download:
 s3://braket-external-assets-preview-us-west-2/HybridJobsAccess/models/
braket-2019-09-01.normal.json to ../../braket/additional_lib/original/
braket-2019-09-01.normal.json
Running Code As Process
Test job started!!!!!
Counter({'00': 55, '11': 45})
Counter({'11': 59, '00': 41})
Counter({'00': 55, '11': 45})
Counter({'00': 58, '11': 42})
Counter({'00': 55, '11': 45})
Test job completed!!!!!
Code Run Finished
2021-09-17 21:48:05,544 sagemaker-training-toolkit INFO     Reporting training SUCCESS
```

Alternatively, you can access the log output from Amazon CloudWatch. To do this, go to the **Log groups** tab on the left menu of the job detail page, select the log group `aws/braket/jobs`, and then choose the log stream that contains the last part of your `job-arn` in the name. In the example above, this is braket-job-default-1631915042705/algo-1-1631915190.

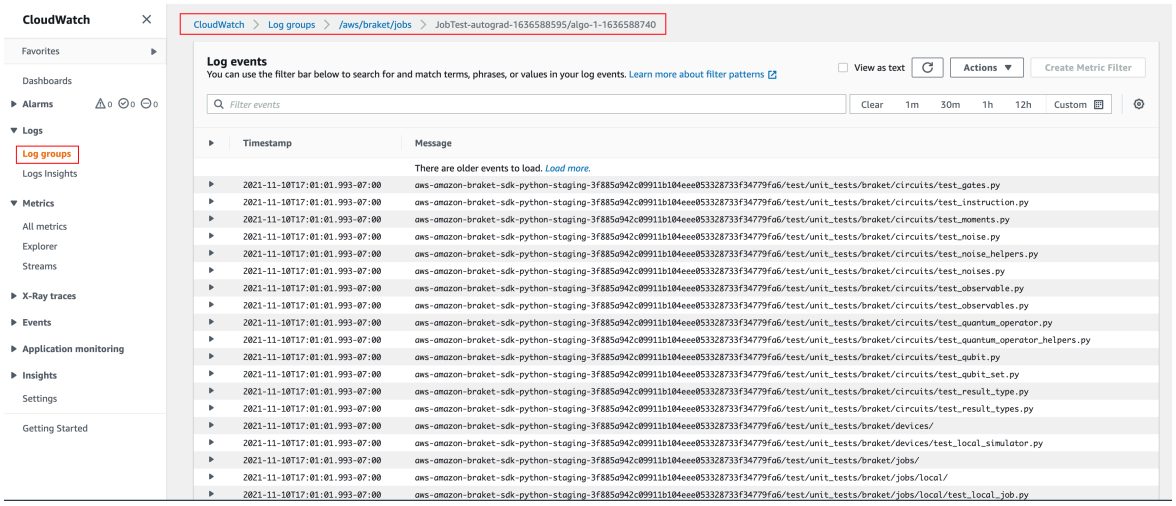

You can also view the status of the job in the console by selecting the **Jobs** page and then choose **Settings**.

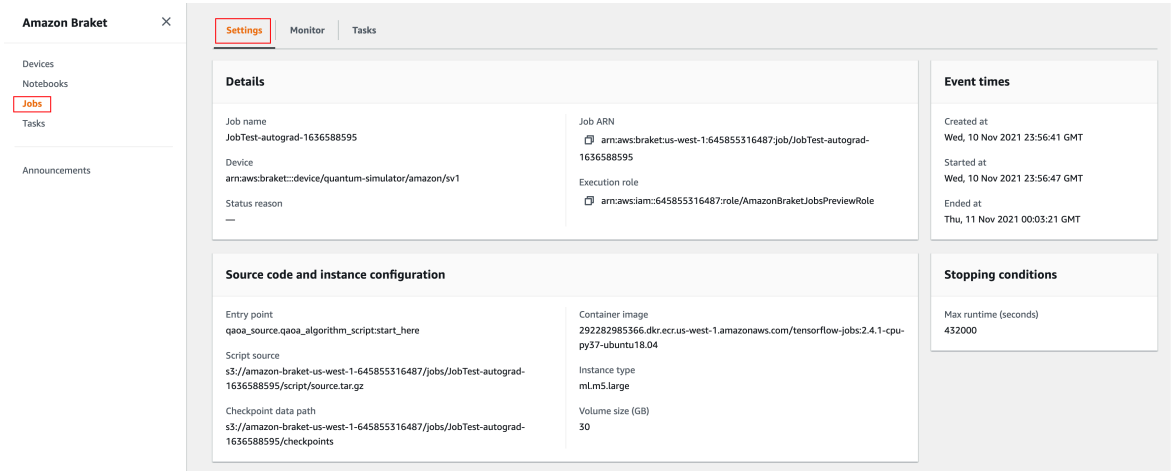

Your job produces some artifacts in Amazon S3 while it runs. The default S3 bucket name is `amazon-braket-<region>-<accountid>` and the content is in the `jobs/<jobname>` directory. You can

configure the S3 locations where these artifacts are stored by specifying a different `code_location` when the job is created with the Braket Python SDK.

> **Note**
> This S3 bucket must be located in the same AWS Region as your job script.

The `jobs/<jobname>` directory contains a subfolder with the output from the entry point script in a model.tar.gz file. There is also a directory called `script` that contains your algorithm script artifacts in a source.tar.gz file. The results from your actual quantum tasks are in the directory named `jobs/<jobname>/tasks`.

Inputs, outputs, environmental variables, and helper functions

In addition to the file or files that makes up your complete algorithm script, your job can have additional inputs and outputs. When your job starts, Amazon Braket copies inputs provided as part of the job creation into the container that runs the algorithm script. When the job completes, all outputs defined during the algorithm are copied to the Amazon S3 location specified.

> **Note**
> *Algorithm metrics* are reported in real time and do not follow this output procedure.

Amazon Braket also provides several environment variables and helper functions to simplify the interactions with container inputs and outputs.

This section explains the key concepts of the `AwsQuantumJob.create` function provided by the Amazon Braket Python SDK and their mapping to the container file structure.

In this section:

Inputs

Input data: You can provide input data using `input_data`. Specify that `input_data` is a keyword in the `AwsQuantumJob.create` function in the SDK. This data is copied to the container filesystem at the location given by the environment variable `"AMZN_BRAKET_INPUT_DIR"`. For an example, see the QAOA with Amazon Braket Hybrid Jobs and PennyLane Jupyter notebook.

> **Note**
> When the input data is large (>1GB), there will be a long wait time before the job is submitted. This is due to the fact that the local input data will first be uploaded to an S3 bucket, then the S3 path will be added to the job request, and, finally, the job request is submitted to Braket service.

Hyperparameters: If you pass in `hyperparameters`, they are available under the environment variable `"AMZN_BRAKET_HP_FILE"`. Hyperparameters are passed directly to the API (not via S3) when creating a job.

Checkpoints: To specify a `job-arn` whose checkpoint you want to use in a new job, use the `copy_checkpoints_from_job` command. This command copies over the checkpoint data to the

`checkpoint_configs3Uri` of the new job, making it available at the path given by the environment variable `AMZN_BRAKET_CHECKPOINT_DIR` while the job runs. The default is `None`, meaning checkpoint data from another job will not be used in the new job.

Outputs

Tasks: Task results are stored in the S3 location specified in `output_data_config`. If you don't specify this value, it defaults to `s3://amazon-braket-<region>-<accountID>/jobs/<job-name>/tasks`.

Job results: Everything that your algorithm script saves to the directory given by the environment variable `"AMZN_BRAKET_JOB_RESULTS_DIR"` is copied to the S3 location specified in `output_data_config`. If you don't specify this value, it defaults to `s3://amazon-braket-<region>-<accountID>/jobs/<job-name>/data`. We provide the SDK helper function **`save_job_result`** , which you can use to store results conveniently and in the correct format from your algorithm script.

Checkpoints: If you want to use checkpoints, you can save them in the directory given by the environment variable `AMZN_BRAKET_CHECKPOINT_DIR"`. You can also use the SDK helper function `save_job_checkpoint` instead.

Algorithm metrics: You can define algorithm metrics as part of your algorithm script that are emitted to Amazon CloudWatch and displayed in real time in the Amazon Braket console while your job is running. For an example of how to use algorithm metrics, see Use Amazon Braket Hybrid Jobs to run a QAOA algorithm_ (p. 138).

Environmental variables

Amazon Braket provides several environment variables to simplify the interactions with container inputs and outputs. The folllowing code lists the environmental variables that Braket uses.

```
# the input data directory opt/braket/input/data
os.environ["AMZN_BRAKET_INPUT_DIR"]
# the output directory opt/braket/model to write ob results to
os.environ["AMZN_BRAKET_JOB_RESULTS_DIR"]
# the name of the job
os.environ["AMZN_BRAKET_JOB_NAME"]
# the checkpoint directory
os.environ["AMZN_BRAKET_CHECKPOINT_DIR"]
# the hyperparameter
os.environ["AMZN_BRAKET_HP_FILE"]
# the device ARN (AWS Resource Number)
os.environ["AMZN_BRAKET_DEVICE_ARN"]
# the output S3 bucket, as specified in the CreateJob request's OutputDataConfig
os.environ["AMZN_BRAKET_OUT_S3_BUCKET"]
# the entry point as specified in the CreateJob request's ScriptModeConfig
os.environ["AMZN_BRAKET_SCRIPT_ENTRY_POINT"]
# the compression type as specified in the CreateJob request's ScriptModeConfig
os.environ["AMZN_BRAKET_SCRIPT_COMPRESSION_TYPE"]
# the S3 location of the user's script as specified in the CreateJob request's
 ScriptModeConfig
os.environ["AMZN_BRAKET_SCRIPT_S3_URI"]
# the S3 location where the SDK would store the task results by default for the job
os.environ["AMZN_BRAKET_TASK_RESULTS_S3_URI"]
# the S3 location where the job results would be stored, as specified in CreateJob
 request's OutputDataConfig
os.environ["AMZN_BRAKET_JOB_RESULTS_S3_PATH"]
# the string that should be passed to CreateQuantumTask's jobToken parameter for quantum
 tasks created in the job container
```

```
os.environ["AMZN_BRAKET_JOB_TOKEN"]
```

Helper functions

Amazon Braket provides several helper functions to simplify the interactions with container inputs and outputs. The following example demonstrates how to use them.

```
save_job_result() # helper function to save your job results
save_job_checkpoint() # helper function to save checkpoints
load_job_checkpoint() # helper function to load a previously saved job checkpoints
```

Save job results

You can save the results generated by the algorithm script so that they are available from the job object in the job script as well as from the output folder in Amazon S3 (in a tar-zipped file named model.tar.gz). The output must be saved in a file using a JavaScript Object Notation (JSON) format. To save the results of the jobs, you add the following lines commented with #ADD to the algorithm script.

```
from braket.aws import AwsDevice
from braket.circuits import Circuit
from braket.jobs import save_job_result #ADD

def start_here():

    print("Test job started!!!!!")

    device = AwsDevice(os.environ['AMZN_BRAKET_DEVICE_ARN'])

    results = []  #ADD

    bell = Circuit().h(0).cnot(0, 1)
    for count in range(5):
        task = device.run(bell, shots=100)
        print(task.result().measurement_counts)
        results.append(task.result().measurement_counts)  #ADD

        save_job_result({ "measurement_counts": results })  #ADD

    print("Test job completed!!!!!")
```

You can then display the results of the job from your job script by appending the line **`print(job.result())`** commented with #ADD.

```
import time
from braket.aws import AwsQuantumJob

job = AwsQuantumJob.create(
    source_module="algorithm_script.py",
    entry_point="algorithm_script:start_here",
    device_arn="arn:aws:braket:::device/quantum-simulator/amazon/sv1",
)

print(job.arn)
while job.state() not in AwsQuantumJob.TERMINAL_STATES:
    print(job.state())
    time.sleep(10)

print(job.state())
```

```
print(job.result())   #ADD
```

In this example, we have removed `wait_until_complete=True` to suppress verbose output. You can add it back in for debugging. When you run this job, it outputs the identifier and the `job-arn`, followed by the state of the job every 10 seconds until the job is `COMPLETED`, after which it shows you the results of the bell circuit. See the following example.

```
arn:aws:braket:us-west-2:111122223333:job/braket-job-default-1234567890123
INITIALIZED
RUNNING
RUNNING
RUNNING
RUNNING
RUNNING
RUNNING
RUNNING
RUNNING
RUNNING
RUNNING
...
RUNNING
RUNNING
COMPLETED
{'measurement_counts': [{'11': 53, '00': 47},..., {'00': 51, '11': 49}]}
```

Save and restart jobs using checkpoints

You can save intermediate iterations of your jobs using checkpoints. In the algorithm script example from the previous section, you would add the following lines commented with #ADD to create checkpoint files.

```
from braket.aws import AwsDevice
from braket.circuits import Circuit
from braket.jobs import save_job_checkpoint  #ADD
import os

def start_here():

    print("Test job starts!!!!!")

    device = AwsDevice(os.environ["AMZN_BRAKET_DEVICE_ARN"])

    #ADD the following code
    job_name = os.environ["AMZN_BRAKET_JOB_NAME"]
     save_job_checkpoint(
     checkpoint_data={"data": f"data for checkpoint from {job_name}"},
     checkpoint_file_suffix="checkpoint-1",
     ) #End of ADD

    bell = Circuit().h(0).cnot(0, 1)
    for count in range(5):
        task = device.run(bell, shots=100)
        print(task.result().measurement_counts)

    print("Test job completed!!!!!")
```

When you run the job, it creates the file *<jobname>-checkpoint-1.json* in your job artifacts in the checkpoints directory with a default `/opt/jobs/checkpoints` path. The job script remains unchanged unless you want to change this default path.

If you want to load a job from a checkpoint generated by a previous job, the algorithm script uses `from braket.jobs import load_job_checkpoint`. The logic to load in your algorithm script is as follows.

```
checkpoint_1 = load_job_checkpoint(
    "previous_job_name",
    checkpoint_file_suffix="checkpoint-1",
)
```

After loading this checkpoint, you can continue your logic based on the content loaded to `checkpoint-1`.

Note
The *checkpoint_file_suffix* must match the suffix previously specified when creating the checkpoint.

Your orchestration script needs to specify the `job-arn` from the previous job with the line commented with #ADD.

```
job = AwsQuantumJob.create(
    source_module="source_dir",
    entry_point="source_dir.algorithm_script:start_here",
    device_arn="arn:aws:braket:::device/quantum-simulator/amazon/sv1",
    copy_checkpoints_from_job="arn:aws:braket:<region>:<account-id>:job/
<previous_job_name>", #ADD
    )
```

Define the environment for your algorithm script

Amazon Braket supports three environments defined by containers for your algorithm script:

- A base container (the default, if no `image_uri` is specified)
- A container with Tensorflow and PennyLane
- A container with PyTorch and PennyLane

The following table provides details about the containers and the libraries they include.

Amazon Braket containers

Type	PennyLane-TensorFlow:2.4.1-gpu-py37-ubuntu18.04	PennyLane-PyTorch1.9.1-gpu-py38-ubuntu20.04	Braket-Base:1.0.0-cpu-py37-ubuntu18.04
Base	111122223333.dkr.ecr.us-east-1.amazonaws.com/amazon-braket-tensorflow-jobs:2.4.1-gpu-py37-cu110-ubuntu18.04	111122223333.dkr.ecr.us-west-2.amazonaws.com/amazon-braket-pytorch-jobs:1.9.1-gpu-py38-cu111-ubuntu20.04	ubuntu:18.04
Inherited Libraries	• awscli • numpy • pandas • scipy	• awscli • numpy • pandas • scipy	
Additional Libraries	• amazon-braket-default-simulator	• amazon-braket-default-simulator	• amazon-braket-default-simulator

Type	PennyLane-TensorFlow:2.4.1-gpu-py37-ubuntu18.04	PennyLane-PyTorch1.9.1-gpu-py38-ubuntu20.04	Braket-Base:1.0.0-cpu-py37-ubuntu18.04
	• amazon-braket-pennylane-plugin • amazon-braket-schemas • amazon-braket-sdk • ipykernel • keras • matplotlib • networkx • openbabel • PennyLane • protobuf • psi4 • rsa • PennyLane-Lightning-gpu • cuQuantum	• amazon-braket-pennylane-plugin • amazon-braket-schemas • amazon-braket-sdk • ipykernel • keras • matplotlib • networkx • openbabel • PennyLane • protobuf • psi4 • rsa • PennyLane-Lightning-gpu • cuQuantum	• amazon-braket-pennylane-plugin • amazon-braket-schemas • amazon-braket-sdk • awscli • boto3 • ipykernel • matplotlib • networkx • numpy • openbabel • pandas • PennyLane • protobuf • psi4 • rsa • scipy

You can view and access the open source container definitions at aws/amazon-braket-containers. Choose the container that best matches your use case. The container must be in the AWS Region from which you invoke your job. You specify the container image when you create a job by adding one of the following three arguments to your `create(...)` call in the job script. You can install additional dependencies into the container you choose at runtime (at the cost of startup or runtime) because the Amazon Braket containers have internet connectivity. The following example is for the us-west-2 Region.

- **Base image** image_uri="292282985366.dkr.ecr.us-west-2.amazonaws.com/amazon-braket-base-jobs:1.0-cpu-py37-ubuntu18.04"
- **Tensorflow image** image_uri="292282985366.dkr.ecr.us-east-1.amazonaws.com/amazon-braket-tensorflow-jobs:2.4.1-gpu-py37-cu110-ubuntu18.04"
- **PyTorch image** image_uri="292282985366.dkr.ecr.us-west-2.amazonaws.com/amazon-braket-pytorch-jobs:1.9.1-gpu-py38-cu111-ubuntu20.04"

The `image-uris` can also be retrieved using the `retrieve_image()` function in the Amazon Braket SDK. The following example shows how to retrieve them from the us-west-2 AWS Region.

```
from braket.jobs.image_uris import retrieve_image, Framework

image_uri_base = retrieve_image(Framework.BASE, "us-west-2")
image_uri_tf = retrieve_image(Framework.PL_TENSORFLOW, "us-west-2")
image_uri_pytorch = retrieve_image(Framework.PL_PYTORCH, "us-west-2")
```

Use hyperparameters

You can define hyperparameters needed by your algorithm, such as the learning rate or the depth of your Ansatz, when you create a job. Hyperparameter values are used to tune an algorithm for optimal performance. This is currently a manual process in Amazon Braket. You specify the hyperparameter

values that you want to test when searching for the optimal set of values. You load the hyperparameters in your algorithm script with the following code.

```
import json

with open(hp_file, "r") as f:
    hyperparams = json.load(f)
```

You pass in the hyperparameter during job creation by adding the following code to your job script.

```
job = AwsQuantumJob.create(
    ...
    hyperparameters={
        "param-1": "first parameter",
        "param-2": "second parameter",
        ...
    },
)
```

Configure the job instance to run your algorithm script

Depending on your algorithm, you may have different requirements. By default, Amazon Braket runs your algorithm script on an `ml.m5.large` instance. However, you can customize this instance type when you create a job using the following import and configuration argument.

```
from braket.jobs.config import InstanceConfig

job = AwsQuantumJob.create(
    ...
    instance_config=InstanceConfig(instanceType="ml.p3.8xlarge"), # Use NVIDIA Tesla V100
 instance with 4 GPUs.
    ...
    ),
```

If you are running an embedded simulation and have specified a local device in the device configuration, you will be able to additionally request more than one instance in the InstanceConfig by specifying the instanceCount and setting it to be greater than one. The upper limit is 5. For instance, you can choose 3 instances as follows.

```
from braket.jobs.config import InstanceConfig
job = AwsQuantumJob.create(
    ...
    instance_config=InstanceConfig(instanceType="ml.p3.8xlarge", instanceCount=3), # Use 3
 NVIDIA Tesla V100
    ...
    ),
```

When you use multiple instances, consider distributing your job using the data parallel feature. See the following example notebook for more details on how-to see this Braket example.

The following three tables list the available instance types and specs for standard, compute optimized, and accelerated computing instances.

Note
To view the default classical compute instance quotas for Hybrid Jobs, see this page (p. 23).

Standard Instances	vCPU	Memory
ml.m5.large (default)	2	8 GiB
ml.m5.xlarge	4	16 GiB
ml.m5.2xlarge	8	32 GiB
ml.m5.4xlarge	16	64 GiB
ml.m5.12xlarge	48	192 GiB
ml.m5.24xlarge	96	384 GiB
ml.m4.xlarge	4	16 GiB
ml.m4.2xlarge	8	32 GiB
ml.m4.4xlarge	16	64 GiB
ml.m4.10xlarge	40	256 GiB

Compute Optimized Instances	vCPU	Memory
ml.c4.xlarge	4	7.5 GiB
ml.c4.2xlarge	8	15 GiB
ml.c4.4xlarge	16	30 GiB
ml.c4.8xlarge	36	192 GiB
ml.c5.xlarge	4	8 GiB
ml.c5.2xlarge	8	16 GiB
ml.c5.4xlarge	16	32 GiB
ml.c5.9xlarge	36	72 GiB
ml.c5.18xlarge	72	144 GiB
ml.c5n.xlarge	4	10.5 GiB
ml.c5n.2xlarge	8	21 GiB
ml.c5n.4xlarge	16	42 GiB
ml.c5n.9xlarge	36	96 GiB
ml.c5n.18xlarge	72	192 GiB

Accelerated Compting Instances	vCPU	Memory
ml.p2.xlarge	4	61 GiB
ml.p2.8xlarge	32	488 GiB

Accelerated Compting Instances	vCPU	Memory
ml.p2.16xlarge	64	732 GiB
ml.p3.2xlarge	8	61 GiB
ml.p3.8xlarge	32	244 GiB
ml.p3.16xlarge	64	488 GiB
ml.g4dn.xlarge	4	16 GiB
ml.g4dn.2xlarge	8	32 GiB
ml.g4dn.4xlarge	16	64 GiB
ml.g4dn.8xlarge	32	128 GiB
ml.g4dn.12xlarge	48	192 GiB
ml.g4dn.16xlarge	64	256 GiB

Note
p3 instances are not available in us-west-1. If your job is unable to provision requested ML compute capacity, use another Region.

Each instance uses a default configuration of data storage (SSD) of 30 GB. But you can adjust the storage in the same way that you configure the `instanceType`. The following example shows how to increase the total storage to 50 GB.

```
from braket.jobs.config import InstanceConfig

job = AwsQuantumJob.create(
    ...
    instance_config=InstanceConfig(
        instanceType="ml.p3.8xlarge",
        volumeSizeInGb=50,
    ),
    ...
    ),
```

Cancel a job

You may need to cancel a job in a non-terminal state. This can be done either in the console or with code.

To cancel your job in the console, select the job to cancel from the **Jobs** page and then select **Cancel job** from the **Actions** dropdown menu.

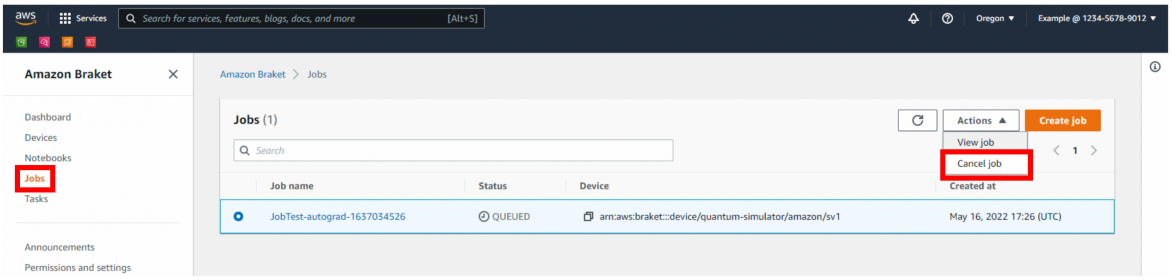

To confirm the cancellation, enter *cancel* into the input field when prompted and then select **OK**.

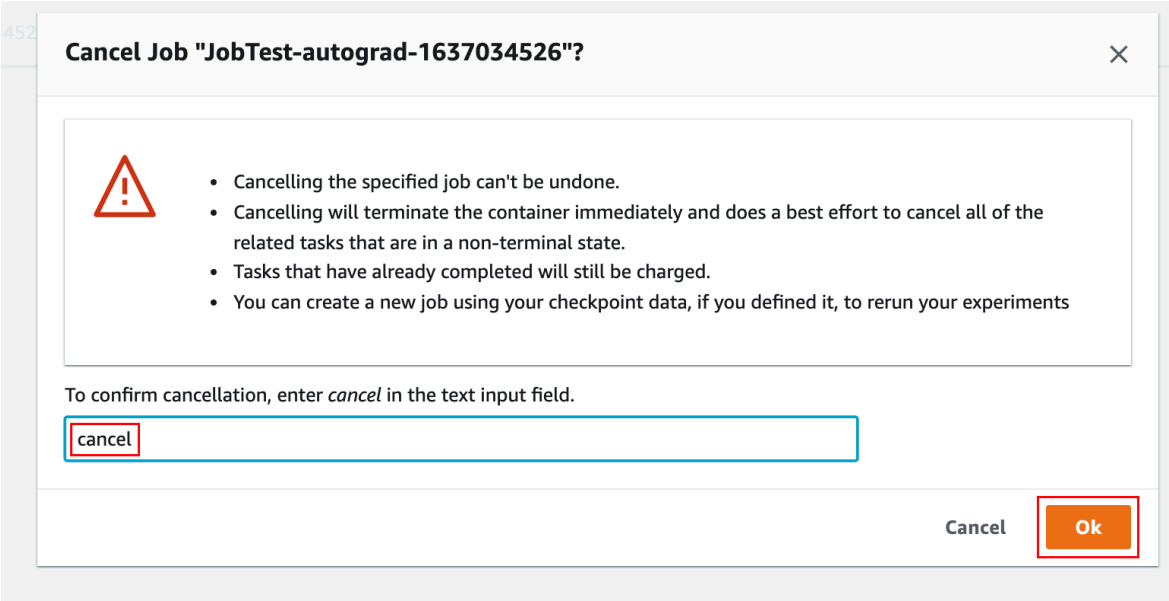

To cancel your job using code from the Braket Python SDK, use the `job_arn` to identify the job and then call the `cancel` command on it as shown in following code.

```
job = AwsQuantumJob(arn=job_arn)
job.cancel()
```

The `cancel` command terminates the classical job container immediately and does a best effort to cancel all of the related tasks that are still in a non-terminal state.

Use Amazon Braket Hybrid Jobs to run a QAOA algorithm

In this section, you'll use what you've learned to write an actual hybrid program using PennyLane. You use the algorithm script to address a Quantum Approximate Optimization Algorithm (QAOA) problem. It creates a cost function corresponding to a classical Max Cut optimization problem, specifies a parametrized quantum circuit, and uses a simple gradient descent method to optimize the parameters so that the cost function is minimized. In this example, we generate the problem graph in the algorithm script for simplicity, but for more typical use cases it is considered a best practice to provide the problem specification through a dedicated channel in the input data configuration.

```
import os
import json
import time

from braket.jobs import save_job_result
from braket.jobs.metrics import log_metric

import networkx as nx
import pennylane as qml
from pennylane import numpy as np
from matplotlib import pyplot as plt
```

```
def init_pl_device(device_arn, num_nodes, shots, max_parallel):
    return qml.device(
        "braket.aws.qubit",
        device_arn=device_arn,
        wires=num_nodes,
        shots=shots,
        # Set s3_destination_folder=None to output task results to a default folder
        s3_destination_folder=None,
        parallel=True,
        max_parallel=max_parallel,
    )

def start_here():
    input_dir = os.environ["AMZN_BRAKET_INPUT_DIR"]
    output_dir = os.environ["AMZN_BRAKET_JOB_RESULTS_DIR"]
    job_name = os.environ["AMZN_BRAKET_JOB_NAME"]
    checkpoint_dir = os.environ["AMZN_BRAKET_CHECKPOINT_DIR"]
    hp_file = os.environ["AMZN_BRAKET_HP_FILE"]
    device_arn = os.environ["AMZN_BRAKET_DEVICE_ARN"]

    # Read the hyperparameters
    with open(hp_file, "r") as f:
        hyperparams = json.load(f)

    p = int(hyperparams["p"])
    seed = int(hyperparams["seed"])
    max_parallel = int(hyperparams["max_parallel"])
    num_iterations = int(hyperparams["num_iterations"])
    stepsize = float(hyperparams["stepsize"])
    shots = int(hyperparams["shots"])

    # Generate random graph
    num_nodes = 6
    num_edges = 8
    graph_seed = 1967
    g = nx.gnm_random_graph(num_nodes, num_edges, seed=graph_seed)

    # Output figure to file
    positions = nx.spring_layout(g, seed=seed)
    nx.draw(g, with_labels=True, pos=positions, node_size=600)
    plt.savefig(f"{output_dir}/graph.png")

    # Set up the QAOA problem
    cost_h, mixer_h = qml.qaoa.maxcut(g)

    def qaoa_layer(gamma, alpha):
        qml.qaoa.cost_layer(gamma, cost_h)
        qml.qaoa.mixer_layer(alpha, mixer_h)

    def circuit(params, **kwargs):
        for i in range(num_nodes):
            qml.Hadamard(wires=i)
        qml.layer(qaoa_layer, p, params[0], params[1])

    dev = init_pl_device(device_arn, num_nodes, shots, max_parallel)

    np.random.seed(seed)
    cost_function = qml.ExpvalCost(circuit, cost_h, dev, optimize=True)
    params = 0.01 * np.random.uniform(size=[2, p])

    optimizer = qml.GradientDescentOptimizer(stepsize=stepsize)
    print("Optimization start")

    for iteration in range(num_iterations):
        t0 = time.time()
```

```
        # Evaluates the cost, then does a gradient step to new params
        params, cost_before = optimizer.step_and_cost(cost_function, params)
        # Convert cost_before to a float so it's easier to handle
        cost_before = float(cost_before)

        t1 = time.time()

        if iteration == 0:
            print("Initial cost:", cost_before)
        else:
            print(f"Cost at step {iteration}:", cost_before)

        # Log the current loss as a metric
        log_metric(
            metric_name="Cost",
            value=cost_before,
            iteration_number=iteration,
        )

        print(f"Completed iteration {iteration + 1}")
        print(f"Time to complete iteration: {t1 - t0} seconds")

    final_cost = float(cost_function(params))
    log_metric(
        metric_name="Cost",
        value=final_cost,
        iteration_number=num_iterations,
    )

    # We're done with the job, so save the result.
    # This will be returned in job.result()
    save_job_result({"params": params.numpy().tolist(), "cost": final_cost})
```

The job script is fairly similar to the previous scripts, except that it also tracks and prints metrics and logs produced from the algorithm script and downloads the results to your local directory.

```
import boto3
import time

from braket.aws import AwsQuantumJob, AwsSession
from braket.jobs.image_uris import Framework, retrieve_image
from braket.jobs.metrics_data.definitions import MetricType

device_arn = "arn:aws:braket:::device/quantum-simulator/amazon/sv1"

hyperparameters = {
    # Number of tasks per iteration = 2 * (num_nodes + num_edges) * p + 1
    "p": "2",
    "seed": "1967",
    # Maximum number of simultaneous tasks allowed
    "max_parallel": "10",
    # Number of total optimization iterations, including those from previous checkpoint (if
 any)
    "num_iterations": "5",
    # Step size / learning rate for gradient descent
    "stepsize": "0.1",
    # Shots for each circuit execution
    "shots": "1000",
}

# Use either the TensorFlow or PyTorch container for PennyLane
region = AwsSession().region
image_uri = retrieve_image(Framework.PL_TENSORFLOW, region)
# image_uri = retrieve_image(Framework.PL_PYTORCH, region)
```

```
start_time = time.time()

job = AwsQuantumJob.create(
    image_uri=image_uri,
    entry_point="qaoa_source.algorithm_script:start_here",
    device=device_arn,
    source_module="qaoa_source",
    hyperparameters=hyperparameters,
)

print(job.arn)
while job.state() not in AwsQuantumJob.TERMINAL_STATES:
    print(job.state())
    time.sleep(10)

end_time = time.time()
print(job.state())

print(end_time - start_time)

print(job.metadata())

print(job.result())

# Metrics may not show up immediately, so wait for 120 seconds
time.sleep(120)
print(job.metrics())
# Print out logs from CloudWatch
print(job.logs())
# Download outputs to local directory
job.download_result()
```

Accelerate your hybrid workloads with embedded simulators from PennyLane

Let's look at how you can use embedded simulators from PennyLane on Amazon Braket Hybrid Jobs to run hybrid workloads. Pennylane's GPU-based simulator, `lightning.gpu`, uses the Nvidia cuQuantum library to accelerate circuit simulations. The GPU simulator is pre-configured in all of the Braket job containers that users can use out of the box. In this page, we show you how to use `lightning.gpu` to speed up your hybrid workloads.

Using `lightning.gpu` for Quantum Approximate Optimization Algorithm workloads

Consider the Quantum Approximate Optimization Algorithm (QAOA) examples from this notebook. To select an embedded simulator, you specify the `device` argument to be a string of the form: `"local:<provider>/<simulator_name>"`. For example, you would set `"local:pennylane/lightning.gpu"` for `lightning.gpu`. The device string you give to the Hybrid Job when you launch is passed to the job as the environment variable `"AMZN_BRAKET_DEVICE_ARN"`.

```
device_string = os.environ["AMZN_BRAKET_DEVICE_ARN"]
prefix, device_name = device_string.split("/")
device = qml.device(simulator_name, wires=n_wires)
```

In this page, let's compare the two PennyLane state vector simulators `lightning.qubit` (which is CPU-based) and `lightning.gpu` (which is GPU-based). You'll need to provide the simulators with some custom gate decompositions in order to compute various gradients.

Now you're ready to prepare the job launching script. You'll run the QAOA algorithm using two instance types: `m5.2xlarge` and `p3.2xlarge`. The `m5.2xlarge` instance type is comparable to a standard developer laptop. The `p3.2xlarge` is an accelerated computing instance that has a single NVIDIA Volta GPU with 16GB of memory.

The `hyperparameters` for all your jobs will be the same. All you need to do to try out different instances and simulators is change two lines as follows.

```
# Specify device that the job will primarily be targeting
device = "local:pennylane/lightning.qubit"
# Run on a CPU based instance with about as much power as a laptop
instance_config = InstanceConfig(instanceType='ml.m5.2xlarge')
```

or:

```
# Specify device that the job will primarily be targeting
device = "local:pennylane/lightning.gpu"
# Run on an inexpensive GPU based instance
instance_config = InstanceConfig(instanceType='ml.p3.2xlarge')
```

Note
If you specify the `instance_config` as using a GPU-based instance, but choose the `device` to be the CPU-based simulator (`lightning.qubit`), the GPU **will not** be used. Make sure to use the GPU-based simulator if you wish to target the GPU!

First, you can create two jobs and solve Max-Cut with QAOA on a graph with 18 vertices. This translates to an 18-qubit circuit—relatively small and feasible to run quickly on your laptop or the `m5.2xlarge` instance.

```
num_nodes = 18
num_edges = 24
seed = 1967

graph = nx.gnm_random_graph(num_nodes, num_edges, seed=seed)

# And similarly for the p3 job
m5_job = AwsQuantumJob.create(
    device=device,
    source_module="qaoa_source",
    job_name="qaoa-m5-" + str(int(time.time())),
    image_uri=image_uri,
    # Relative to the source_module
    entry_point="qaoa_source.qaoa_algorithm_script",
    copy_checkpoints_from_job=None,
    instance_config=instance_config,
    # general parameters
    hyperparameters=hyperparameters,
    input_data={"input-graph": input_file_path},
    wait_until_complete=True,
)
```

The mean iteration time for the `m5.2xlarge` instance is about 25 seconds, while for the `p3.2xlarge` instance it's about 12 seconds. For this 18-qubit workflow, the GPU instance gives us a 2x speedup. If you look at the Amazon Braket Hybrid Jobs pricing page, you can see that the cost per minute for an `m5.2xlarge` instance is $0.00768, while for the `p3.2xlarge` instance it's $0.06375. To run for 5

total iterations, as you did here, would cost $0.016 using the CPU instance or $0.06375 using the GPU instance — both pretty inexpensive!

Now let's make the problem harder, and try solving a Max-Cut problem on a 24-vertex graph, which will translate to 24 qubits. Run the jobs again on the same two instances and compare the cost.

Note
You'll see that the time to run this job on the CPU instance may be about five hours!

```
num_nodes = 24
num_edges = 36
seed = 1967

graph = nx.gnm_random_graph(num_nodes, num_edges, seed=seed)

# And similarly for the p3 job
m5_big_job = AwsQuantumJob.create(
    device=device,
    source_module="qaoa_source",
    job_name="qaoa-m5-big-" + str(int(time.time())),
    image_uri=image_uri,
    # Relative to the source_module
    entry_point="qaoa_source.qaoa_algorithm_script",
    copy_checkpoints_from_job=None,
    instance_config=instance_config,
    # general parameters
    hyperparameters=hyperparameters,
    input_data={"input-graph": input_file_path},
    wait_until_complete=True,
)
```

The mean iteration time for the `m5.2xlarge` instance is roughly an hour, while for the `p3.2xlarge` instance it's roughly two minutes. For this larger problem, the GPU instance is an order of magnitude faster! All you had to do to benefit from this speedup was to change two lines of code, swapping out the instance type and the local simulator used. To run for 5 total iterations, as was done here, would cost about $2.27072 using the CPU instance or about $0.775625 using the GPU instance. The CPU usage is not only more expensive, but also takes more time to run. Accelerating this workflow with a GPU instance available on AWS, using PennyLane's simulator backed by NVIDIA CuQuantum, allows you to run workflows with intermediate qubit counts (between 20 and 30) for less total cost and in less time. This means you can experiment with quantum computing even for problems that are too big to run quickly on your laptop or a similarly-sized instance.

Quantum machine learning and data parallelism

If your workload type is quantum machine learning (QML) that trains on datasets, you can further accelerate your workload using data parallelism. In QML, the model contains one or more quantum circuits. The model may or may not also contain classical neural nets. When training the model with the dataset, the parameters in the model are updated to minimize the loss function. A loss function is usually defined for a single data point, and the total loss for the average loss over the whole dataset. In QML, the losses are usually computed in serial before averaging to total loss for gradient computations. This procedure is time consuming, especially when there are hundreds of data points.

Because the loss from one data point does not depend on other data points, the losses can be evaluated in parallel! Losses and gradients associated with different data points can be evaluated at the same time. This is known as data parallelism. With SageMaker's distributed data parallel library, Amazon Braket Hybrid Jobs make it easier for you to leverage data parallelism to accelerate your training.

Consider the following QML workload for data parallelism which uses the Sonar dataset dataset from the well-known UCI repository as an example for binary classification. The Sonar dataset have 208 data points each with 60 features that are collected from sonar signals bouncing off materials. Each

data points is either labeled as "M" for mines or "R" for rocks. Our QML model consists of an input layer, a quantum circuit as a hidden layer, and an output layer. The input and output layers are classical neural nets implemented in PyTorch. The quantum circuit is integrated with the PyTorch neural nets using PennyLane's qml.qnn module. See our example notebooks for more detail about the workload. Like the QAOA example above, you can harness the power of GPU by using GPU-based simulators like PennyLane's `lightning.gpu` to improve the performance over CPU-based simulators.

To create a job, you can call `AwsQuantumJob.create` and specify the algorithm script, device, and other configurations through its keyword arguments.

```
instance_config = InstanceConfig(instanceType='ml.p3.2xlarge')

hyperparameters={"nwires": "10",
                 "ndata": "32",
                 ...
}

job = AwsQuantumJob.create(
    device="local:pennylane/lightning.gpu",
    source_module="qml_source",
    entry_point="qml_source.train_single",
    hyperparameters=hyperparameters,
    instance_config=instance_config,
    ...
)
```

In order to use data parallelism, you need to modify few lines of code in the algorithm script for the SageMaker distributed library to correctly parallelize the training. First, you import the `smdistributed` package which does most of the heavy-lifting for distributing your workloads across multiple GPUs and multiple instances. This package is preconfigured in the Braket PyTorch and TensorFlow containers. The `dist` module tells our algorithm script what the total number of GPUs for the training (`world_size`) is as well as the `rank` and `local_rank` of a GPU core. `rank` is the absolute index of a GPU across all instances, while `local_rank` is the index of a GPU within an instance. For example, if there are four instances each with eight GPUs allocated for the training, the `rank` ranges from 0 to 31 and the `local_rank` ranges from 0 to 7.

```
import smdistributed.dataparallel.torch.distributed as dist

dp_info = {
    "world_size": dist.get_world_size(),
    "rank": dist.get_rank(),
    "local_rank": dist.get_local_rank(),
}
batch_size //= dp_info["world_size"] // 8
batch_size = max(batch_size, 1)
```

Next, you define a `DistributedSampler` according to the `world_size` and `rank` and then pass it into the data loader. This sampler avoids GPUs accessing the same slice of a dataset.

```
train_sampler = torch.utils.data.distributed.DistributedSampler(
    train_dataset,
    num_replicas=dp_info["world_size"],
    rank=dp_info["rank"]
)
train_loader = torch.utils.data.DataLoader(
    train_dataset,
    batch_size=batch_size,
    shuffle=False,
    num_workers=0,
    pin_memory=True,
```

```
    sampler=train_sampler,
)
```

Next, you use the `DistributedDataParallel` class to enable data parallelism.

```
from smdistributed.dataparallel.torch.parallel.distributed import DistributedDataParallel
 as DDP

model = DressedQNN(qc_dev).to(device)
model = DDP(model)
torch.cuda.set_device(dp_info["local_rank"])
model.cuda(dp_info["local_rank"])
```

The above are the changes you need to use data parallelism. In QML, you often want to save results and print training progress. If each GPU executes the saving and printing command, the log will be flooded with the repeated information and the results will overwrite each other. To avoid this, you can only save and print from the GPU that has `rank` 0.

```
if dp_info["rank"]==0:
    print('elapsed time: ', elapsed)
    torch.save(model.state_dict(), f"{output_dir}/test_local.pt")
    save_job_result({"last loss": loss_before})
```

Amazon Braket Hybrid Jobs supports `ml.p3.16xlarge` instance types for the SageMaker distributed data parallel library. You configure the instance type through the `InstanceConfig` argument in Hybrid Jobs. For the SageMaker distributed data parallel library to know that data parallelism is enabled, you need to add two additional hyperparameters, `"sagemaker_distributed_dataparallel_enabled"` setting to `"true"` and `"sagemaker_instance_type"` setting to the instance type you are using. These two hyperparameters are used by `smdistributed` package. Your algorithm script does not need to explicitly use them. In Amazon Braket SDK, it provides a convenient keyword argument `distribution`. With `distribution="data_parallel"` in job creation, the Amazon Braket SDK automatically inserts the two hyperparameters for you. If you use the Amazon Braket API, you need to include these two hyperparameters.

With the instance and data parallelism configured, you can now submit your job. There are 8 GPUs in a `ml.p3.16xlarge` instance. When you set `instanceCount=1` , the workload is distributed across the 8 GPUs in the instance. When you set `instanceCount` greater than one, the workload is distributed across GPUs available in all instances. When using multiple instances, each instance incurs a charge based on how much time you use it. For example, when you use four instances, the billable time is four times the run time per instance because there are four instances running your workloads at the same time.

```
instance_config = InstanceConfig(instanceType='ml.p3.16xlarge',
                                 instanceCount=1,
)

hyperparameters={"nwires": "10",
                 "ndata": "32",
                 ...,
}

job = AwsQuantumJob.create(
    device="local:pennylane/lightning.gpu",
    source_module="qml_source",
    entry_point="qml_source.train_dp",
    hyperparameters=hyperparameters,
    instance_config=instance_config,
    distribution="data_parallel",
    ...
)
```

Note
In the above job creation, `train_dp.py` is the modified algorithm script for using data parallelism. Keep in mind that data parallelism only works correctly when you modify your algorithm script according to the above section. If the data parallelism option is enabled without a correctly modified algorithm script, the job may throw errors, or each GPU may repeatedly process the same data slice, which is inefficient.

Let's compare the run time and cost in an example where when train a model with a 26-qubit quantum circuit for the binary classification problem mentioned above. The `ml.p3.16xlarge` instance used in this example costs $0.4692 per minute. Without data parallelism, it takes the simulator about 45 minutes to train the model for 1 epoch (i.e., over 208 data points) and it costs about $20. With data parallelism across 1 instance and 4 instances, it only takes 6 minutes and 1.5 minutes respectively, which translates to roughly $2.8 for both. By using data parallelism across 4 instances, you not only improve the run time by 30x, but also reduce costs by an order of magnitude!

Build and debug a job with local mode

If you are building a new hybrid algorithm, local mode helps you to debug and test your algorithm script. Local mode is a feature that allows you to run code you plan to use in Amazon Braket Hybrid Jobs, but without needing Braket to manage the infrastructure for running the job. Instead, you run jobs locally on your Braket Notebook instance or on some other preferred client such as a laptop or desktop computer. In local mode, you can still send quantum tasks to actual devices, but you do not get the performance benefits when running against an actual QPU while in local mode.

To use local mode, modify `AwsQuantumJob` to `LocalQuantumJob` wherever it occurs. For instance, to run the example from Create your first job (p. 123), edit the job script as follows.

```
from braket.jobs.local import LocalQuantumJob

job = LocalQuantumJob.create(
    device="arn:aws:braket:::device/quantum-simulator/amazon/sv1",
    source_module="algorithm_script.py",
    entry_point="algorithm_script:start_here",
)
```

Note
Docker, which is already pre-installed in the Amazon Braket notebooks, needs to be installed in your local environment to use this feature. Instructions for installing Docker can be found here. In addition, not all parameters are supported in local mode.

Bring your own container (BYOC)

Amazon Braket Hybrid Jobs provides three pre-built containers for running code in different environments that are described in the Define the environment for your algorithm script (p. 133) topic. If one of these containers supports your use case, you only have to provide your algorithm script when you create a job. Minor missing dependencies can be added from your algorithm script using pip.

Note
The Strawberry Fields plugin is not natively included in the Amazon Braket hybrid jobs containers. If you wish to run a hybrid workload using the Xanadu device, you have to BYOC with the Strawberry Fields dependencies included.

If none of these containers support your use case or if you wish to expand on them, Amazon Braket Hybrid Jobs supports running jobs with your own custom Docker container image. This BYOC capability enables you to run code in an environment that has software installed that you specify. The container

starts running a specified entrypoint script, which you may then use to access custom user code and data. For examples, see the Amazon Braket Hybrid Jobs tutorials.

To run Amazon Braket Hybrid Jobs in your own container:

1. **Set up prerequisites**: To build and upload your custom container, you must have Docker installed. Amazon Braket notebooks come pre-installed with Docker.
2. **Create your container**: Create a new directory with a Dockerfile to install and set up the software and the environment your script needs to run. You can download an example file by following the instructions in the Define the environment for your algorithm script (p. 133) topic. You may need to add additional packages to support your container. This file may have more software than you need, in which case you may speed up the build process by removing the unnecessary components. Note that the `sagemaker-training` component is necessary and should not be removed. You must specify the initial script to run when your container starts with the environment variable `ENV SAGEMAKER_PROGRAM your_file.py`, where `your_file.py` is a file that exists in the directory `/opt/ml/code` inside of your container. When creating a job, the user is able to specify an Amazon S3 location with code and specify an entry point to run. Braket Jobs creates environment variables for your script to find user input code to run. The Amazon S3 location and entry point to the user code are exposed to the initial script through the environment variables `AMZN_BRAKET_SCRIPT_S3_URI` and `AMZN_BRAKET_SCRIPT_ENTRY_POINT` respectively. The user input may be compressed and a compression type can be accessed through the variable `AMZN_BRAKET_SCRIPT_COMPRESSION_TYPE`. It is the responsibility of your script, `your_file.py`, to use this information to run the user code. An example script that runs user code on user data is `braket_container.py` in the examples you can download as described in the Define the environment for your algorithm script (p. 133) topic.
3. **Create an Amazon ECR repository**: The repository you create must be private. If you specify a public image, the Amazon Braket Hybrid Jobs client throws an error when it attempts to validate the image location. Also, the AWS Region of the container image must match the AWS Region where you run the job. To set up your repository, access the Amazon ECR console. Then, select **Create repository**, choose **private** for visibility setting, and give your new repository a name. For a full user guide on setting up a repository, see Getting started with Amazon ECR using the AWS Management Console. Note that the standard role created for jobs only has permissions to download container images beginning with the prefix "amazon-braket-". You will need to add inline permissions to your job role if you wish to use container images with a different name (see current policy restrictions here). If you plan to run jobs from an account that is different than the account you used to create the repository, you must allow the account used to run the jobs to pull the images you have created from the account used to create them. The actions that you must allow on your repository are `ecr:DescribeImages` and `ecr:BatchGetImage`. To add these permissions to the image from the Amazon ECR console, select your image and choose **permissions** in the menu. You can add a rule to allow these two actions for the account you wish to allow to run jobs with your container.
4. **Build the image, and push it to the repository**: You must authenticate Docker to upload to Amazon ECR. You can retrieve an authentication token and authenticate your Docker client to your registry with the AWS CLI using the following code.

```
aws ecr get-login-password --region ${your_region} | docker login --username AWS --
password-stdin ${aws_account_id}.dkr.ecr.${your_region}.amazonaws.com
```

Now you should be able to build your image and push it to your repository.

```
cd ${your_docker_directory}
docker build -t ${your_job_container} .
docker tag ${your_job_container}:latest ${aws_account_id}.dkr.ecr.
${your_region}.amazonaws.com/${your_job_container}:latest
docker push ${aws_account_id}.dkr.ecr.${your_region}.amazonaws.com/
${your_job_container}:latest
```

To create a job with your own container, call AwsQuantumJob.create with the argument image_uri specified. You can use a QPU, an on-demand simulator, or run your code locally on the classical processor available with Amazon Braket Hybrid jobs. To run your code on the classical processor, specify the instanceType and the instanceCount you wish to use by updating the InstanceConfig. Note that if you specify an instance_count > 1, you will need to make sure that your code has the ability to run across multiple hosts. The upper limit for the number of instances you can choose is 5.

Let's look at an example.

```
job = AwsQuantumJob.create(
    source_module="source_dir",
    entry_point="source_dir.algorithm_script:start_here",
    image_uri="111122223333.dkr.ecr.us-west-2.amazonaws.com/my-byoc-container:latest",
    instance_config=InstanceConfig(instanceType="ml.p3.8xlarge", instanceCount=3),
    device="local:braket/braket.local.qubit",
    # ...
)
```

Note
The device arn allows you to track the simulator you used as job metadata. Acceptable values must follow the format device = "local:<provider>/<simulator_name>". Remember that <provider> and <simulator_name> must consist only of letters, numbers, _, - and .. The string is limited to 256 characters.

If you plan to use BYOC and you are not using the Amazon Braket SDK to create tasks, you should pass the value of the environmental variable AMZN_BRAKET_JOB_TOKEN to the jobToken parameter in the CreateQuantumTask request. Failing to do so causes the tasks to not get priority and to be billed as regular standalone tasks.

Configure the default bucket in AwsSession

Providing your own AwsSession gives you greater flexibility, for example, in the location of your default bucket. By default, an AwsSession has a default bucket location of f"amazon-braket-{id}-{region}". But you can override that default when creating an AwsSession. Users can optionally pass in an AwsSession object into AwsQuantumJob.create with the parameter name aws_session as shown in the following code example.

```
aws_session = AwsSession(default_bucket="other-default-bucket")

# then you can use that AwsSession when creating a job
job = AwsQuantumJob.create(
    ...
    aws_session=aws_session
)
```

Interact with jobs directly using the API

You can access and interact with Amazon Braket Hybrid Jobs directly using the API. However, defaults and convenience methods are not available when using the API directly.

Note
We strongly recommend that you interact with Amazon Braket Hybrid Jobs using the Amazon Braket Python SDK. It offers convenient defaults and protections that help your jobs run successfully.

This topic covers the basics of using the API. If you choose to use the API, keep in mind that this approach can be more complex and be prepared for several iterations to get your job to run.

To use the API, your account should have a role with the `AmazonBraketFullAccess` managed policy.

> **Note**
> For more information on how to obtain a role with the `AmazonBraketFullAccess` managed policy, see the Enable Amazon Braket (p. 33) page.

Additionally, you need an **execution role**. This role will be passed to the service. You can create the role using the **Amazon Braket console**. Use the **Execution roles** tab on the **Permissions and settings** page to create a default role for jobs.

The `CreateJob` API requires that you specify all the required parameters for the job. To use Python, compress your algorithm script files to a tar bundle, such as an input.tar.gz file, and run the following script. Update the parts of the code within angled brackets (<>) to match your account information and entry point that specify the path, file, and method where your job starts.

```
from braket.aws import AwsDevice, AwsSession
import boto3
from datetime import datetime

s3_client = boto3.client("s3")
client = boto3.client("braket")

project_name = "job-test"
job_name = project_name + "-" + datetime.strftime(datetime.now(), "%Y%m%d%H%M%S")
bucket = "amazon-braket-<your_bucket>"
s3_prefix = job_name

job_script = "input.tar.gz"
job_object = f"{s3_prefix}/script/{job_script}"
s3_client.upload_file(job_script, bucket, job_object)

input_data = "inputdata.csv"
input_object = f"{s3_prefix}/input/{input_data}"
s3_client.upload_file(input_data, bucket, input_object)

job = client.create_job(
    jobName=job_name,
    roleArn="arn:aws:iam::<your_account>:role/service-role/AmazonBraketJobsExecutionRole",
  # https://docs.aws.amazon.com/braket/latest/developerguide/braket-manage-
access.html#about-amazonbraketjobsexecution
    algorithmSpecification={
        "scriptModeConfig": {
            "entryPoint": "<your_execution_module>:<your_execution_method>",
            "containerImage": {"uri": "292282985366.dkr.ecr.us-west-1.amazonaws.com/amazon-
braket-base-jobs:1.0-cpu-py37-ubuntu18.04"}   # Change to the specific region you are using
            "s3Uri": f"s3://{bucket}/{job_object}",
            "compressionType": "GZIP"
        }
    },
    inputDataConfig=[
        {
            "channelName": "hellothere",
            "compressionType": "NONE",
            "dataSource": {
                "s3DataSource": {
                    "s3Uri": f"s3://{bucket}/{s3_prefix}/input",
                    "s3DataType": "S3_PREFIX"
                }
            }
        }
    ],
```

```
    outputDataConfig={
        "s3Path": f"s3://{bucket}/{s3_prefix}/output"
    },
    instanceConfig={
        "instanceType": "ml.m5.large",
        "instanceCount": 1,
        "volumeSizeInGb": 1
    },
    checkpointConfig={
        "s3Uri":  f"s3://{bucket}/{s3_prefix}/checkpoints",
        "localPath": "/opt/omega/checkpoints"
    },
    deviceConfig={
        "priorityAccess": {
            "devices": [
                "arn:aws:braket:us-west-1::device/qpu/rigetti/Aspen-M-2"
            ]
        }
    },
    hyperParameters={
        "hyperparameter key you wish to pass": "<hyperparameter value you wish to pass>",
    },
    stoppingCondition={
        "maxRuntimeInSeconds": 1200,
        "maximumTaskLimit": 10
    },
)
```

Once you create your job, you can access the job details through the `GetJob` API or the console. To get the job details from the Python session in which you ran the `createJob` code as in the previous example, use the following Python command.

```
getJob = client.get_job(jobArn=job["jobArn"])
```

To cancel a job, call the `CancelJob` API with the Amazon Resource Name of the job ('JobArn').

```
cancelJob = client.cancel_job(jobArn=job["jobArn"])
```

You can specify checkpoints as part of the `createJob` API using the `checkpointConfig` parameter.

```
    checkpointConfig = {
        "localPath" : "/opt/omega/checkpoints",
        "s3Uri": f"s3://{bucket}/{s3_prefix}/checkpoints"
    },
```

Note
The localPath of `checkpointConfig` cannot start with any of the following reserved paths: `/opt/ml`, `/opt/braket`, `/tmp`, or `/usr/local/nvidia`.

Use PennyLane with Amazon Braket

Hybrid algorithms are algorithms that contain classical and quantum instructions. The classical instructions are executed on classical hardware (an EC2 instance or your laptop), and the quantum instructions are executed either on a simulator or on a quantum computer. We recommend that you run hybrid algorithms using the hybrid jobs feature. For more guidance, see When to use Amazon Braket Jobs (p. 122).

Amazon Braket enables you to set up and run hybrid quantum algorithms with the assistance of the **Amazon Braket PennyLane plugin**, or with the **Amazon Braket Python SDK** and example notebook repositories. Amazon Braket example notebooks, based on the SDK, enable you to set up and run certain hybrid algorithms without the PennyLane plugin. However, we recommend PennyLane because it provides a much easier and richer experience.

About hybrid quantum algorithms

Hybrid quantum algorithms are important to the industry today because contemporary quantum computing devices generally produce noise, and therefore, errors. Every quantum gate added to a computation increases the chance of adding noise; therefore, long-running algorithms can be overwhelmed by noise, which results in faulty computation.

Pure quantum algorithms such as Shor's (QPE example) or Grover's (Grover's example) require thousands, or millions, of operations. For this reason, they can be impractical for existing quantum devices, which are generally referred to as *noisy intermediate-scale quantum* (NISQ) devices.

In hybrid quantum algorithms, quantum processing units (QPUs) work as co-processors for classic CPUs, specifically to speed up certain calculations in a classical algorithm. Circuit executions become much shorter, within reach of the capabilities of today's devices.

Amazon Braket with PennyLane

Amazon Braket provides support for PennyLane, an open-source software framework built around the concept of *quantum differentiable programming*. This framework allows you to train quantum circuits in the same way that you would train a neural network to find solutions for computational problems in quantum chemistry, quantum machine learning, and optimization.

The PennyLane library provides interfaces to familiar machine learning tools, including PyTorch and TensorFlow, to make training quantum circuits fast, easy and intuitive.

- **The PennyLane Library** -– PennyLane is pre-installed in Amazon Braket notebooks. For access to Amazon Braket devices from PennyLane, open a notebook and import the PennyLane library with this command:

```
import pennylane as qml
```

Tutorial notebooks help you get started quickly. Alternatively, you can use PennyLane on Amazon Braket from any IDE of your choice.

- **The Amazon Braket PennyLane plugin** — To use your own IDE, you can install the Amazon Braket PennyLane plugin manually. The plugin connects PennyLane with the Amazon Braket Python SDK, so

you can run circuits in PennyLane on Amazon Braket devices. To install the the PennyLane plugin, use this command:

```
pip install amazon-braket-pennylane-plugin
```

The following example demonstrates how to set up access to Amazon Braket devices in PennyLane:

```
# to use SV1
import pennylane as qml
s3 = ("my-bucket", "my-prefix")
sv1 = qml.device("braket.aws.qubit", device_arn="arn:aws:braket:::device/quantum-simulator/
amazon/sv1", s3_destination_folder=s3, wires=2)

# to run a circuit:
@qml.qnode(sv1)
def circuit(x):
qml.RZ(x, wires=0)
qml.CNOT(wires=[0,1])
qml.RY(x, wires=1)
return qml.expval(qml.PauliZ(1))

result = circuit(0.543)


#To use the local sim:

local = qml.device("braket.local.qubit", wires=2)
```

You can find information and PennyLane tutorial examples in the Amazon Braket examples repository.

The Amazon Braket PennyLane plugin enables you to switch between Amazon Braket QPU and simulator devices in PennyLane with a single line of code. It offers two Amazon Braket quantum devices to work with PennyLane:

- `braket.aws.qubit` for running with the Amazon Braket service's quantum devices, including QPUs and simulators
- `braket.local.qubit` for running with the Amazon Braket SDK's local simulator

The Amazon Braket PennyLane plugin is open source. It can be installed from the PennyLane Plugin GitHub repository.

To find out more, visit the PennyLane Documentation.

Hybrid algorithms in Amazon Braket example notebooks

Amazon Braket does provide a variety of example notebooks that do not rely on the PennyLane plugin for running hybrid algorithms. You can get started with any of these Amazon Braket hybrid example notebooks that illustrate *variational methods*, such as the Quantum Approximate Optimization Algorithm (QAOA) or Variational Quantum Eigensolver (VQE).

The Amazon Braket example notebooks rely on the Amazon Braket Python SDK. The SDK provides a framework to interact with quantum computing hardware devices through Amazon Braket. It is an open source library that is designed to assist you with the quantum portion of your hybrid workflow.

You can explore Amazon Braket further, with any of these other example notebooks.

Hybrid algorithms with embedded PennyLane simulators

Amazon Braket Hybrid Jobs now comes with five high performance CPU- and GPU-based simulators from PennyLane. This family of simulators can be embedded directly within your jobs container and includes the fast state-vector `lightning.qubit` simulator, the `lightning.gpu` simulator accelerated using NVIDIA's cuQuantum library, and others. These simulators are ideally suited for variational algorithms such as quantum machine learning that can benefit from advanced methods such as the adjoint differentiation method. You can run these embedded simulators on one or multiple CPU or GPU instances.

With hybrid jobs, you can now run your variational algorithm code using a combination of a classical co-processor and a QPU, an Amazon Braket managed simulator such as SV1, or directly using the embedded simulator from PennyLane.

The embedded simulator is already available with the Hybrid jobs container, you simply need to specify the `device` (e.g `device="local:pennylane/lightning.gpu"`) in the Braket SDK as shown below. To use the `lightning.gpu` simulator, you also need to specify a GPU instance in the `InstanceConfig` as shown in the code snippet below:

```
job = AwsQuantumJob.create(
      device="local:pennylane/lightning.gpu",
      source_module="algorithm_script.py",
      instance_config=InstanceConfig(instanceType="ml.p3.8xlarge")
      )
```

and refer to the device in your algorithm code using the environment variable `AMZN_BRAKET_DEVICE_ARN`:

```
import os
import pennylane as qml

device_string = os.environ["AMZN_BRAKET_DEVICE_ARN"]
prefix, simulator_name = device_string.split("/")
device = qml.device(simulator_name, wires=n_wires)
```

Refer to the example notebook to get started with using a PennyLane embedded simulator with hybrid jobs.

Security in Amazon Braket

This chapter helps you understand how to apply the shared responsibility model when using Amazon Braket. It shows you how to configure Amazon Braket to meet your security and compliance objectives. You also learn how to use other AWS services that help you to monitor and secure your Amazon Braket resources.

Cloud security at AWS is the highest priority. As an AWS customer, you benefit from a data center and network architecture that is built to meet the requirements of the most security-sensitive organizations. You are responsible for other factors, including the sensitivity of your data, your company's requirements, and applicable laws and regulations.

Shared responsibility for security

Security is a shared responsibility between AWS and you. The shared responsibility model describes this as security *of* the cloud and security *in* the cloud:

- **Security of the cloud** – AWS is responsible for protecting the infrastructure that runs AWS services in the AWS Cloud. AWS also provides you with services that you can use securely. Third-party auditors regularly test and verify the effectiveness of our security as part of the AWS Compliance Programs. To learn about the compliance programs that apply to Amazon Braket, see AWS Services in Scope by Compliance Program.
- **Security in the cloud** – You are responsible for maintaining control over your content that is hosted on this AWS infrastructure. This content includes the security configuration and management tasks for the AWS services that you use.

Data protection

For more information about data privacy, see the Data Privacy FAQ.

For data protection purposes, we recommend that you protect AWS account credentials and set up individual user accounts with AWS Identity and Access Management (IAM). That way, each user is given only the permissions necessary to fulfill their job duties. We also recommend that you secure your data in the following ways:

- Use multi-factor authentication (MFA) with each account.
- Use SSL/TLS to communicate with AWS resources. We recommend TLS 1.2 or later.
- Set up API and user activity logging with AWS CloudTrail.
- Use AWS encryption solutions, along with all default security controls within AWS services.
- Use advanced managed security services such as Amazon Macie, which assists in discovering and securing personal data that is stored in Amazon S3.
- If you require FIPS 140-2 validated cryptographic modules when accessing AWS through a command line interface or an API, use a FIPS endpoint. For more information about the available FIPS endpoints, see Federal Information Processing Standard (FIPS) 140-2.

We strongly recommend that you never put sensitive identifying information, such as your customer account numbers, into free-form fields such as a **Name** field. This includes when you work with Braket or other AWS services using the console, API, CLI, or AWS SDKs. Any data that you enter into Braket or other services might get picked up for inclusion in diagnostic logs. When you provide a URL to an external server, do not include credentials information in the URL to validate your request to that server.

Data retention

After 90 days, Amazon Braket automatically removes all task IDs and other metadata associated with your tasks. As a result of this data retention policy, these tasks and results are no longer retrievable by search from the Amazon Braket console, although they remain stored in your S3 bucket.

If you need access to historical tasks and results that are stored in your S3 bucket for longer than 90 days, you must keep a separate record of your task ID and other metadata associated with that data. Be sure to save the information prior to 90 days. You can use that saved information to retrieve the historical data.

Managing access to Amazon Braket

This chapter describes the permissions that are required to run Amazon Braket, or to restrict the access of specific IAM users and roles. You can grant (or deny) the required permissions to any IAM user or role in your account. To do so, attach the appropriate Amazon Braket policy to that user or role in the account, as given in this chapter.

As a prerequisite, you must enable Amazon Braket. To enable Amazon Braket, be sure to sign in as a user or role that has (1) administrator permissions or (2) is assigned the **AmazonBraketFullAccess** policy and has permissions to create S3 buckets.

In this section:

Amazon Braket resources

Amazon Braket creates one type of resource, which is the *quantum-task* resource. Here is the form of the ARN for that resource type:

- **Resource Name:** *AWS::Service::Braket*
- **ARN Regex:** *arn:${Partition}:braket:${Region}:${Account}:quantum-task/${RandomId}*

Notebooks and roles

Noteboooks are another type of resource that Amazon Braket utilizes on your behalf. A notebook is an Amazon SageMaker resource, which Braket is able to share. The notebooks require a specific IAM role to function: a role with a name that begins with `AmazonBraketServiceSageMakerNotebook`.

To create a notebook, you must use a role with Admin permissions or that has the following inline policy attached to it.

```
{
```

```
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "iam:CreateRole",
            "Resource": "arn:aws:iam::*:role/service-role/
AmazonBraketServiceSageMakerNotebookRole-*"
        },
        {
            "Effect": "Allow",
            "Action": "iam:CreatePolicy",
            "Resource": "arn:aws:iam::*:policy/service-role/
AmazonBraketServiceSageMakerNotebookAccess-*"
        },
        {
            "Effect": "Allow",
            "Action": "iam:AttachRolePolicy",
            "Resource": "arn:aws:iam::*:role/service-role/
AmazonBraketServiceSageMakerNotebookRole-*",
            "Condition": {
                "StringLike": {
                    "iam:PolicyARN": [
                        "arn:aws:iam::aws:policy/AmazonBraketFullAccess",
                        "arn:aws:iam::*:policy/service-role/
AmazonBraketServiceSageMakerNotebookAccess-*"
                    ]
                }
            }
        }
    ]
}
```

To create the role, follow the steps given in the Create a notebook page or have your administrator create it for you. Ensure that the **AmazonBraketFullAccess** policy is attached.

After you've created the role, you can reuse that role for all notebooks you launch in the future.

About the AmazonBraketFullAccess policy

The **AmazonBraketFullAccess** policy grants permissions for Amazon Braket operations, including permissions for these tasks:

- **Amazon Elastic Container Registry** – to read and download container images to be used for Amazon Braket Hybrid Jobs feature. The containers must conform to the format "arn:aws:ecr:::repository/amazon-braket"
- **Keep AWS CloudTrail logs** – for all *describe*, *get*, and *list* actions as well as starting and stopping queries, testing metrics filters, and filtering log events. The AWS CloudTrail log file contains a record of all Amazon Braket API activity that occurs in your account.
- **Utilize roles to control resources** – to create a service-linked role in your account. The service-linked role has access to AWS resources on your behalf. It can be used only by the Amazon Braket service. Also to pass in IAM roles to the Amazon Braket `CreateJob` API and to create a role and attach a policy scoped to AmazonBraketFullAccess to the role.
- **Create log groups, log events, and query log groups. Maintain usage log files for your account** – to create, store, and view logging information about Amazon Braket usage in your account. Query metrics on jobs log groups. Encompass the proper Braket path and allowing putting log data. Put metric data in CloudWatch.
- **Create and Store data in Amazon S3 buckets, and list all buckets** – to create S3 buckets, list the S3 buckets in your account, and to put objects into and get objects from any bucket in your account whose name begins with *amazon-braket-*. These permissions are required for Amazon Braket to put files containing results from processed tasks into the bucket and to retrieve them from the bucket.

- **Pass IAM roles** – to pass in IAM roles to the `CreateJob` API.
- **Amazon SageMaker Notebook** – to create and manage SageMaker Notebook instances scoped to the resource from "arn:aws:sagemaker:::notebook-instance/amazon-braket-"

Policy contents

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "s3:GetObject",
                "s3:PutObject",
                "s3:ListBucket",
                "s3:CreateBucket",
                "s3:PutBucketPublicAccessBlock",
                "s3:PutBucketPolicy"
            ],
            "Resource": "arn:aws:s3:::amazon-braket-*"
        },
        {
            "Effect": "Allow",
            "Action": [
                "s3:ListAllMyBuckets"
            ],
            "Resource": "*"
        },
        {
            "Effect": "Allow",
            "Action": [
                "ecr:GetDownloadUrlForLayer",
                "ecr:BatchGetImage",
                "ecr:BatchCheckLayerAvailability"
            ],
            "Resource": "arn:aws:ecr:*:*:repository/amazon-braket*"
        },
        {
            "Effect": "Allow",
            "Action": [
                "ecr:GetAuthorizationToken"
            ],
            "Resource": "*"
        },
        {
            "Effect": "Allow",
            "Action": [
                "logs:Describe*",
                "logs:Get*",
                "logs:List*",
                "logs:StartQuery",
                "logs:StopQuery",
                "logs:TestMetricFilter",
                "logs:FilterLogEvents"
            ],
            "Resource": "arn:aws:logs:*:*:log-group:/aws/braket*"
        },
        {
            "Effect": "Allow",
            "Action": [
                "iam:ListRoles",
                "iam:ListRolePolicies",
                "iam:GetRole",
```

```
                "iam:GetRolePolicy",
                "iam:ListAttachedRolePolicies"
            ],
            "Resource": "*"
        },
        {
            "Effect": "Allow",
            "Action": [
                "sagemaker:ListNotebookInstances"
            ],
            "Resource": "*"
        },
        {
            "Effect": "Allow",
            "Action": [
                "sagemaker:CreatePresignedNotebookInstanceUrl",
                "sagemaker:CreateNotebookInstance",
                "sagemaker:DeleteNotebookInstance",
                "sagemaker:DescribeNotebookInstance",
                "sagemaker:StartNotebookInstance",
                "sagemaker:StopNotebookInstance",
                "sagemaker:UpdateNotebookInstance",
                "sagemaker:ListTags",
                "sagemaker:AddTags",
                "sagemaker:DeleteTags"
            ],
            "Resource": "arn:aws:sagemaker:*:*:notebook-instance/amazon-braket-*"
        },
        {
            "Effect": "Allow",
            "Action": [
                "sagemaker:DescribeNotebookInstanceLifecycleConfig",
                "sagemaker:CreateNotebookInstanceLifecycleConfig",
                "sagemaker:DeleteNotebookInstanceLifecycleConfig",
                "sagemaker:ListNotebookInstanceLifecycleConfigs",
                "sagemaker:UpdateNotebookInstanceLifecycleConfig"
            ],
            "Resource": "arn:aws:sagemaker:*:*:notebook-instance-lifecycle-config/amazon-
braket-*"
        },
        {
            "Effect": "Allow",
            "Action": "braket:*",
            "Resource": "*"
        },
        {
            "Effect": "Allow",
            "Action": "iam:CreateServiceLinkedRole",
            "Resource": "arn:aws:iam::*:role/aws-service-role/braket.amazonaws.com/
AWSServiceRoleForAmazonBraket*",
            "Condition": {
                "StringEquals": {
                    "iam:AWSServiceName": "braket.amazonaws.com"
                }
            }
        },
        {
            "Effect": "Allow",
            "Action": [
                "iam:PassRole"
            ],
            "Resource": "arn:aws:iam::*:role/service-role/
AmazonBraketServiceSageMakerNotebookRole*",
            "Condition": {
                "StringLike": {
                    "iam:PassedToService": [
```

```
                        "sagemaker.amazonaws.com"
                    ]
                }
            }
        },
        {
            "Effect": "Allow",
            "Action": [
                "iam:PassRole"
            ],
            "Resource": "arn:aws:iam::*:role/service-role/AmazonBraketJobsExecutionRole*",
            "Condition": {
                "StringLike": {
                    "iam:PassedToService": [
                        "braket.amazonaws.com"
                    ]
                }
            }
        },
        {
            "Effect": "Allow",
            "Action": [
                "logs:GetQueryResults"
            ],
            "Resource": [
                "arn:aws:logs:*:*:log-group:*"
            ]
        },
        {
            "Effect": "Allow",
            "Action": [
                "logs:PutLogEvents",
                "logs:CreateLogStream",
                "logs:CreateLogGroup"
            ],
            "Resource": "arn:aws:logs:*:*:log-group:/aws/braket*"
        },
        {
            "Effect": "Allow",
            "Action": "cloudwatch:PutMetricData",
            "Resource": "*",
            "Condition": {
                "StringEquals": {
                    "cloudwatch:namespace": "/aws/braket"
                }
            }
        }
    ]
}
```

About the AmazonBraketJobsExecutionPolicy policy

The **AmazonBraketJobsExecutionPolicy** policy grants permissions for execution roles used in Amazon Braket Hybrid Jobs

- **Amazon Elastic Container Registry** - permissions to read and download container images to be used for Amazon Braket Hybrid Jobs feature. Containers must conform to the format "arn:aws:ecr:*:*:repository/amazon-braket*"
- **Create log groups and log events and query log groups. Maintain usage log files for your account** – Create, store, and view logging information about Amazon Braket usage in your account. Query metrics on jobs log groups. Encompass the proper Braket path and allowing putting log data. Put metric data in CloudWatch.

- **Store data in Amazon S3 buckets** – list the S3 buckets in your account, put objects into and get objects from any bucket in your account that starts with *amazon-braket-* in its name. These permissions are required for Amazon Braket to put files containing results from processed tasks into the bucket, and to retrieve them from the bucket.
- **Pass IAM roles** passing in IAM roles to the CreateJob API. Roles must conform to the format arn:aws:iam::*:role/service-role/AmazonBraketJobsExecutionRole*.

```
"Version": "2012-10-17",
"Statement": [
 {
  "Effect": "Allow",
  "Action": [
   "s3:GetObject",
   "s3:PutObject",
   "s3:ListBucket",
   "s3:CreateBucket",
   "s3:PutBucketPublicAccessBlock",
   "s3:PutBucketPolicy"
  ],
  "Resource": "arn:aws:s3:::amazon-braket-*"
 },
 {
  "Effect": "Allow",
  "Action": [
   "ecr:GetDownloadUrlForLayer",
   "ecr:BatchGetImage",
   "ecr:BatchCheckLayerAvailability"
  ],
  "Resource": "arn:aws:ecr:*:*:repository/amazon-braket*"
 },
 {
  "Effect": "Allow",
  "Action": [
   "ecr:GetAuthorizationToken"
  ],
  "Resource": "*"
 },
 {
  "Effect": "Allow",
  "Action": [
   "braket:CancelJob",
   "braket:CancelQuantumTask",
   "braket:CreateJob",
   "braket:CreateQuantumTask",
   "braket:GetDevice",
   "braket:GetJob",
   "braket:GetQuantumTask",
   "braket:SearchDevices",
   "braket:SearchJobs",
   "braket:SearchQuantumTasks",
   "braket:ListTagsForResource",
   "braket:TagResource",
   "braket:UntagResource"
  ],
  "Resource": "*"
 },
 {
  "Effect": "Allow",
  "Action": [
   "iam:PassRole"
  ],
  "Resource": "arn:aws:iam::*:role/service-role/AmazonBraketJobsExecutionRole*",
  "Condition": {
```

```
    "StringLike": {
     "iam:PassedToService": [
      "braket.amazonaws.com"
     ]
    }
   }
  },
  {
   "Effect": "Allow",
   "Action": [
    "iam:ListRoles"
   ],
   "Resource": "arn:aws:iam::*:role/*"
  },
  {
   "Effect": "Allow",
   "Action": [
    "logs:GetQueryResults"
   ],
   "Resource": [
    "arn:aws:logs:*:*:log-group:*"
   ]
  },
  {
   "Effect": "Allow",
   "Action": [
    "logs:PutLogEvents",
    "logs:CreateLogStream",
    "logs:CreateLogGroup",
    "logs:GetLogEvents",
    "logs:DescribeLogStreams",
    "logs:StartQuery",
    "logs:StopQuery"
   ],
   "Resource": "arn:aws:logs:*:*:log-group:/aws/braket*"
  },
  {
   "Effect": "Allow",
   "Action": "cloudwatch:PutMetricData",
   "Resource": "*",
   "Condition": {
    "StringEquals": {
     "cloudwatch:namespace": "/aws/braket"
    }
   }
  }
 ]
}
```

Restrict user access to certain devices

To restrict access for certain users to certain Amazon Braket devices, you can add a *deny permissions* policy to a specific IAM role.

The following example restricts access to all QPUs for the AWS account `012345678901`.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
    "Effect": "Deny",
    "Action": [
      "braket:CreateQuantumTask",
```

```
      "braket:CancelQuantumTask",
      "braket:GetQuantumTask",
      "braket:SearchQuantumTasks",
      "braket:GetDevice",
      "braket:SearchDevices"
    ],
    "Resource": [
    "arn:aws:braket:*:*:device/qpu/*"
    ]
   }
  ]
}
```

To adapt this code, substitute the Amazon Resource Number (ARN) of the restricted device for the string shown in the previous example. This string provides the **Resource** value. In Amazon Braket, a device represents a QPU or simulator that you can call to run quantum tasks. The devices available are listed on the Devices page. There are two schemas used to specify access to these devices:

- `arn:aws:braket:<region>:<account id>:device/qpu/<provider>/<device_id>`
- `arn:aws:braket:<region>:<account id>:device/quantum-simulator/<provider>/<device_id>`

Here are examples for various types of device access

- To select all QPUs across all regions: `arn:aws:braket:*:*:device/qpu/*`
- To select all QPUs in the us-west-2 region ONLY: `arn:aws:braket:us-west-2:012345678901:device/qpu/*`
- Equivalently, to select all QPUs in the us-west-2 region ONLY (since devices are a service resource, not a customer resource): `arn:aws:braket:us-west-2:* :device/qpu/*`
- To restrict access to all managed simulator devices: `arn:aws:braket:* :012345678901:device/quantum-simulator/*`
- To restric access to the IonQ device in us-east-1 region: `arn:aws:braket:us-east-1:012345678901:device/ionq/ionQdevice`
- To restrict access to devices from a certain provider (for example, to Rigetti QPU devices): `arn:aws:braket:* :012345678901:device/qpu/rigetti/*`
- To restrict access to TN1 device: `arn:aws:braket:* :012345678901:device/quantum-simulator/amazon/tn1`

Amazon Braket updates to AWS managed policies

The following table provides details about updates to AWS managed policies for Amazon Braket since this service began tracking these changes.

Change	Description	Date
AmazonBraketFullAccess (p. 156) - Full access policy for Amazon Braket	Added s3:ListAllMyBuckets permissions to allow users to view and inspect the buckets created and used for Amazon Braket.	March 31, 2022
AmazonBraketFullAccess (p. 156) - Full access policy for Amazon Braket	Amazon Braket adjusted iam:PassRole permissions for AmazonBraketFullAccess to include the `service-role/` path.	November 29, 2021

Change	Description	Date
AmazonBraketJobsExecutionPolicy (p. 159) - Jobs execution policy for Amazon Braket Hybrid Jobs	Amazon Braket updated the jobs execution role ARN to include the `service-role/` path.	November 29, 2021
Amazon Braket started tracking changes	Amazon Braket started tracking changes for its AWS managed policies.	November 29, 2021

Restrict user access to certain notebook instances

To restrict access for certain users to specific Amazon Braket notebook instances, you can add a *deny permissions* policy to a specific IAM role, user, or group.

The following example uses policy variables to efficiently restrict permissions to start, stop, and access specific notebook instances in the AWS account `012345678901`, which is named according to the IAM user who should have access (e.g. user `Alice` would have access to a notebook instance named `amazon-braket-Alice`):

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Deny",
      "Action": [
        "sagemaker:CreateNotebookInstance",
        "sagemaker:DeleteNotebookInstance",
        "sagemaker:UpdateNotebookInstance",
        "sagemaker:CreateNotebookInstanceLifecycleConfig",
        "sagemaker:DeleteNotebookInstanceLifecycleConfig",
        "sagemaker:UpdateNotebookInstanceLifecycleConfig"
      ],
      "Resource": "*"
    },
    {
      "Effect": "Deny",
      "Action": [
        "sagemaker:DescribeNotebookInstance",
        "sagemaker:StartNotebookInstance",
        "sagemaker:StopNotebookInstance",
      ],
      "NotResource": [
        "arn:aws:sagemaker:*:012345678901:notebook-instance/amazon-braket-${aws:username}"
      ]
    },
    {
      "Effect": "Deny",
      "Action": [
        "sagemaker:CreatePresignedNotebookInstanceUrl"
      ],
      "NotResource": [
        "arn:aws:sagemaker:*:012345678901:notebook-instance/amazon-braket-${aws:username}*"
      ]
    }
  ]
}
```

Restrict user access to certain S3 buckets

To restrict access for certain users to specific Amazon S3 buckets, you can add a deny policy to a specific IAM role, user, or group.

The following example restricts permissions to retrieve and place objects into a specific S3 bucket (`arn:aws:s3:::amazon-braket-us-east-1-012345678901-Alice`) and also restricts the listing of those objects.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Deny",
      "Action": [
        "s3:ListBucket"
      ],
      "NotResource": [
        "arn:aws:s3:::amazon-braket-us-east-1-012345678901-Alice"
      ]
    },
    {
      "Effect": "Deny",
      "Action": [
        "s3:GetObject"
      ],
      "NotResource": [
        "arn:aws:s3:::amazon-braket-us-east-1-012345678901-Alice/*"
      ]
    }
  ]
}
```

To restrict access to the bucket for a certain notebook instance, you can add the above policy to the notebook execution role.

Amazon Braket service-linked role

When you enable Amazon Braket, a *service-linked role* is created in your account.

A service-linked role is a unique type of IAM role that, in this case, is linked directly to Amazon Braket. The Amazon Braket service-linked role is predefined to include all the permissions that Braket requires when calling other AWS services on your behalf.

A service-linked role makes setting up Amazon Braket easier because you don't have to add the necessary permissions manually. Amazon Braket defines the permissions of its service-linked roles. Unless you change these definitions, only Amazon Braket can assume its roles. The defined permissions include the *trust policy* and the *permissions policy*. The permissions policy cannot be attached to any other IAM entity.

The service-linked role that Amazon Braket sets up is part of the AWS Identity and Access Management (IAM) service-linked roles capability. For information about other AWS services that support service-linked roles, see AWS Services That Work with IAM and look for the services that have **Yes** in the **Service-Linked Role** column. Choose a **Yes** with a link to view the service-linked role documentation for that service.

Service-linked role permissions for Amazon Braket

Amazon Braket uses the `AWSServiceRoleForAmazonBraket` service-linked role that trusts the braket.amazonaws.com entity to assume the role.

You must configure permissions to allow an IAM entity (such as a user, group, or role) to create, edit, or delete a service-linked role. For more information, see Service-Linked Role Permissions in the *IAM User Guide*.

The service-linked role in Amazon Braket is granted the following permissions by default:

- **Amazon S3** – permissions to list the buckets in your account, and put objects into and get objects from any bucket in your account with a name that starts with *amazon-braket-*.
- **Amazon CloudWatch Logs** – permissions to list and create log groups, create the associated log streams, and put events into the log group created for Amazon Braket.

The following policy is attached to the **`AWSServiceRoleForAmazonBraket`** service-linked role:

```
{"Version": "2012-10-17",
    "Statement": [
        {"Effect": "Allow",
            "Action": [
                "s3:GetObject",
                "s3:PutObject",
                "s3:ListBucket"
            ],
            "Resource": "arn:aws:s3:::amazon-braket*"
        },
        {"Effect": "Allow",
            "Action": [
                "logs:Describe*",
                "logs:Get*",
                "logs:List*",
                "logs:StartQuery",
                "logs:StopQuery",
                "logs:TestMetricFilter",
                "logs:FilterLogEvents"
            ],
            "Resource": "arn:aws:logs:*:*:log-group:/aws/braket/*"
        },
        {"Effect": "Allow",
            "Action": "braket:*",
            "Resource": "*"
        },
        {"Effect": "Allow",
            "Action": "iam:CreateServiceLinkedRole",
            "Resource": "arn:aws:iam::*:role/aws-service-role/braket.amazonaws.com/
AWSServiceRoleForAmazonBraket*",
            "Condition": {"StringEquals": {"iam:AWSServiceName": "braket.amazonaws.com"
                }
            }
        }
    ]
}
```

Resilience in Amazon Braket

The AWS global infrastructure is built around AWS Regions and Availability Zones.

Each Region provides multiple *availability zones* that are physically separated and isolated. These availability zones (AZs) are connected through low-latency, high-throughput, and highly redundant networking. As a result, availability zones are more highly available, fault tolerant, and scalable than traditional single- or multiple-datacenter infrastructures.

You can design and operate applications and databases that fail over between AZs automatically, without interruption.

For more information about AWS Regions and availability zones, see AWS Global Infrastructure.

Compliance validation for Amazon Braket

Your compliance responsibility when using Amazon Braket is determined by the sensitivity of your data, your company's compliance objectives, and applicable laws and regulations. AWS provides the following resources to help with compliance:

- Security and Compliance Quick Start Guides – These deployment guides discuss architectural considerations and provide steps for deploying security- and compliance-focused baseline environments on AWS.
- AWS Compliance Resources – This collection of workbooks and guides might apply to your industry and location.

For a list of AWS services in scope of specific compliance programs, see AWS Services in Scope by Compliance Program. For general information, see AWS Compliance Programs.

Infrastructure Security in Amazon Braket

As a managed service, Amazon Braket is protected by the AWS global network security procedures that are described in the Amazon Web Services: Overview of Security Processes whitepaper.

For access to Amazon Braket through the network, you make calls to published AWS APIs. Clients must support Transport Layer Security (TLS) 1.2 or later. Clients also must support cipher suites with perfect forward secrecy (PFS) such as Ephemeral Diffie-Hellman (DHE) or Elliptic Curve Ephemeral Diffie-Hellman (ECDHE). Most modern systems such as Java 7 and later support these modes.

Additionally, requests must be signed using an access key ID and a secret access key that is associated with an IAM principal. Or you can use the AWS Security Token Service (AWS STS) to generate temporary security credentials to sign requests.

Security of Amazon Braket Hardware Providers

QPUs on Amazon Braket are hosted by third-party hardware providers. When you run your task on a QPU, Amazon Braket uses the DeviceARN as an identifier when sending the circuit to the specified QPU for processing.

If you use Amazon Braket for access to quantum computing hardware operated by one of the third-party hardware providers, your circuit and its associated data are processed by hardware providers outside of facilities operated by AWS. Information about the physical location and AWS Region where each QPU is available can be found in the **Device Details** section of the Amazon Braket console.

Your content is anonymized. Only the content necessary to process the circuit is sent to third parties. AWS account information is not transmitted to third parties.

All data is encrypted at rest and in transit. Data is decrypted for processing only. Amazon Braket third-party providers are not permitted to store or use your content for purposes other than processing your circuit. Once the circuit completes, the results are returned to Amazon Braket and stored in your S3 bucket.

The security of Amazon Braket third-party quantum hardware providers is audited periodically, to ensure that standards of network security, access control, data protection, and physical security are met.

Troubleshoot

Solve common problems you might find when working with Amazon Braket.

Topics

In this section:

Access denied exception

If you receive an **AccessDeniedException**, as shown in the following image, when enabling or using Braket, then you are most likely attempting to enable or use Braket in a region where your restricted role does not have access.

In such cases, you should contact your internal AWS administrator to understand which of the following conditions apply:

- if there are role restrictions preventing access to a region
- if the role you are attempting to use is permitted to use Braket

If your role does not have access to a given region when using Braket, then you will be unable to use devices in that particular region.

A task is failing creation

If you receive an error along the lines of "An error occurred (ValidationException) when calling the CreateQuantumTask operation: Caller doesn't have access to amazon-braket-…" make sure you are referring to an existing s3_folder since we do not auto create new Amazon S3 buckets and prefixes for you.

If you are accessing the API directly and getting an error like "Failed to create quantum task: Caller doesn't have access to s3://MY_BUCKET" make sure you are not including s3:// in the Amazon S3 bucket path.

An SDK feature does not work

Make sure your SDK (and schemas) are up-to-date. From the notebook or your python editor run

```
!pip install --upgrade amazon-braket-sdk
```

Make sure your SDK and schemas are up-to-date. To update the SDK from the notebook or your python editor run the following:

```
pip install --upgrade amazon-braket-schemas
```

If you are accessing Amazon Braket from your own client make sure your AWS region is set to one supported by Amazon Braket.

A job fails due to an exceeded quota

My job fails with `ServiceQuotaExceededException`

A job running tasks against the Amazon Braket simulators can fail to be created if you exceed the concurrent task limit for the simulator device you are targeting. The service limits are explained in the Quotas (p. 23) topic. If you are running multiple jobs from your account, which run concurrent tasks against a simulator device, you could encounter this error.

To see the number of concurrent tasks against a specific simulator device, use the `search-quantum-tasks` API, as shown in the following code example.

```
DEVICE_ARN=arn:aws:braket:::device/quantum-simulator/amazon/sv1
task_list=""
for status_value in "CREATED" "QUEUED" "RUNNING" "CANCELLING"; do
    tasks=$(aws braket search-quantum-tasks --filters name=status,operator=EQUAL,values=
${status_value} name=deviceArn,operator=EQUAL,values=$DEVICE_ARN --max-results 100 --query
 'quantumTasks[*].quantumTaskArn' --output text)
    task_list="$task_list $tasks"
done;
echo "$task_list" | tr -s ' \t' '[\n*]' | sort | uniq
```

You can also view the created tasks against a device using Amazon CloudWatch metrics: **Braket** > **By Device**.

To avoid running into these errors, you can either:

1. Request a service quota increase for the number of concurrent tasks for the simulator device. This is only applicable to the SV1 device.
2. Handle `ServiceQuotaExceeded` exceptions in your code and retry.

Something stopped working in your notebook instance

If some components of your notebook stop working, try the following:

1. Download any notebooks you created or modified to a local drive.
2. Stop your notebook instance.
3. Delete your notebook instance.
4. Create new notebook instance with a different name.

5. Upload the notebooks to the new instance.

Troubleshoot OpenQASM

This section provides troubleshooting pointers that might be useful when encountering errors using OpenQASM 3.0.

In this section:

Include statement error

Braket currently doesn't have a standard gate library file to be included in OpenQASM programs. For example, the following example raises a parser error.

```
OPENQASM 3;
include "standardlib.inc";
```

This code generates the error message: `No terminal matches '"' in the current parser context, at line 2 col 17.`

Non-contiguous qubits error

Using non-contiguous qubits on devices that `requiresContiguousQubitIndices` be set to `true` in the device capability result in an error.

When running tasks on simulators and IonQ, the following program triggers the error.

```
OPENQASM 3;

qubit[4] q;

h q[0];
cnot q[0], q[2];
cnot q[0], q[3];
```

This code generates the error message: `Device requires contiguous qubits. Qubit register q has unused qubits q[1], q[4].`

Mixing physical qubits with virtual qubits error

Mixing physical qubits with virtual qubits in the same program is not allowed and results in an error. The following code generates the error.

```
OPENQASM 3;

qubit[2] q;
cnot q[0], $1;
```

This code generates the error message: `[line 4] mixes physical qubits and qubits registers.`

Requesting result types and measuring qubits in the same program error

Requesting result types and that qubits are explicitly measured in the same program results in an error. The following code generates the error.

```
OPENQASM 3;

qubit[2] q;

h q[0];
cnot q[0], q[1];
measure q;

#pragma braket result expectation x(q[0]) @ z(q[1])
```

This code generates the error message: `Qubits should not be explicitly measured when result types are requested.`

Classical and qubit register limits exceeded error

Only one classical register and one qubit register are allowed. The following code generates the error.

```
OPENQASM 3;

qubit[2] q0;
qubit[2] q1;
```

This code generates the error message: `[line 4] cannot declare a qubit register. Only 1 qubit register is supported.`

Box not preceded by a verbatim pragma error

All boxes must be preceded by a verbatim pragma. The following code generates the error.

```
box{
rx(0.5) $0;
}
```

This code generates the error message: `In verbatim boxes, native gates are required. x is not a device native gate.`

Verbatim boxes missing native gates error

Verbatim boxes should have native gates and physical qubits. The following code generates the native gates error.

```
#pragma braket verbatim
box{
x $0;
}
```

This code generates the error message: `In verbatim boxes, native gates are required. x is not a device native gate.`

Verbatim boxes missing physical qubits error

Verbatim boxes must have physical qubits. The following code generates the missing physical qubits error.

```
qubit[2] q;

#pragma braket verbatim
box{
rx(0.1) q[0];
}
```

This code generates the error message: `Physical qubits are required in verbatim box.`

The verbatim pragma is missing "braket" error

You must include “braket” in the verbatim pragma. The following code generates the error.

```
#pragma braket verbatim          // Correct
#pragma verbatim                 // wrong
```

This code generates the error message: `You must include “braket” in the verbatim pragma`

Single qubits cannot be indexed error

Single qubits cannot be indexed. The following code generates the error.

```
OPENQASM 3;

qubit q;
h q[0];
```

This code generates the error: `[line 4] single qubit cannot be indexed.`

However, single qubit arrays can be indexed as follows:

```
OPENQASM 3;

qubit[1] q;
h q[0];   // This is valid
```

The physical qubits in a two qubit gate are not connected error

To use physical qubits, first confirm that the device uses physical qubits by checking `device.properties.action[DeviceActionType.OPENQASM].supportPhysicalQubits` and then verify the connectivity graph by checking `device.properties.paradigm.connectivity.connectivityGraph` or `device.properties.paradigm.connectivity.fullyConnected`.

```
OPENQASM 3;

cnot $0, $14;
```

This code generates the error message: `[line 3] has disconnected qubits 0 and 14`

GetDevice does not return OpenQASM results error

If you do not see OpenQASM results in the GetDevice response when using a Braket SDK, you may need to set AWS_EXECUTION_ENV environment variable to configure user-agent. See the code examples provided below for how to do this for the Go and Java SDKs.

To set AWS_EXECUTION_ENV environment variable to configure user-agent when using the AWS CLI:

```
% export AWS_EXECUTION_ENV="aws-cli BraketSchemas/1.8.0"
# Or for single execution
% AWS_EXECUTION_ENV="aws-cli BraketSchemas/1.8.0" aws braket <cmd> [options]
```

To set AWS_EXECUTION_ENV environment variable to configure user-agent when using Boto3:

```
import boto3
import botocore

client = boto3.client("braket",
 config=botocore.client.Config(user_agent_extra="BraketSchemas/1.8.0"))
```

To set AWS_EXECUTION_ENV environment variable to configure user-agent when using the JavaScript/TypeScript (SDK v2):

```
import Braket from 'aws-sdk/clients/braket';
const client = new Braket({ region: 'us-west-2', credentials: AWS_CREDENTIALS,
 customUserAgent: 'BraketSchemas/1.8.0'  });
```

To set AWS_EXECUTION_ENV environment variable to configure user-agent when using the JavaScript/TypeScript (SDK v3):

```
import { Braket } from '@aws-sdk/client-braket';
const client = new Braket({ region: 'us-west-2', credentials: AWS_CREDENTIALS,
 customUserAgent: 'BraketSchemas/1.8.0' });
```

To set AWS_EXECUTION_ENV environment variable to configure user-agent when using the Go SDK:

```
os.Setenv("AWS_EXECUTION_ENV", "BraketGo BraketSchemas/1.8.0")
mySession := session.Must(session.NewSession())
```

```
svc := braket.New(mySession)
```

To set AWS_EXECUTION_ENV environment variable to configure user-agent when using the Java SDK:

```
ClientConfiguration config = new ClientConfiguration();
config.setUserAgentSuffix("BraketSchemas/1.8.0");
BraketClient braketClient =
 BraketClientBuilder.standard().withClientConfiguration(config).build();
```

Local simulator support warning

The `LocalSimulator` supports advanced features in OpenQASM that may not be available on QPUs or managed simulators. If your program contains language features specific only to the `LocalSimulator`, as seen in the following example, you will receive a warning.

```
qasm_string = """
qubit[2] q;

h q[0];
ctrl @ x q[0], q[1];
"""
qasm_program = Program(source=qasm_string)
```

This code generates the warning: `This program uses OpenQASM language features only supported in the LocalSimulator. Some of these features may not be supported on QPUs or managed simulators.`

For more information on supported OpenQASM features, click here (p. 66).

Amazon VPC endpoints for Amazon Braket

You can establish a private connection between your VPC and Amazon Braket by creating an interface VPC endpoint. Interface endpoints are powered by AWS PrivateLink, a technology that enables access to Braket APIs without an internet gateway, NAT device, VPN connection, or AWS Direct Connect connection. Instances in your VPC don't need public IP addresses to communicate with Braket APIs.

Each interface endpoint is represented by one or more Elastic Network Interfaces in your subnets.

With PrivateLink, traffic between your VPC and Braket does not leave the Amazon network, which increases the security of data that you share with cloud-based applications, because it reduces your data's exposure to the public internet. For more information, see Interface VPC endpoints (AWS PrivateLink) in the Amazon VPC User Guide.

Considerations for Amazon Braket VPC endpoints

Before you set up an interface VPC endpoint for Braket, ensure that you review Interface endpoint properties and limitations in the *Amazon VPC User Guide*.

Braket supports making calls to all of its API actions from your VPC.

By default, full access to Braket is allowed through the VPC endpoint. You can control access if you specify VPC endpoint policies. For more information, see Controlling access to services with VPC endpoints in the *Amazon VPC User Guide*.

Set up Braket and PrivateLink

To use AWS PrivateLink with Amazon Braket, you must create an Amazon Virtual Private Cloud (Amazon VPC) endpoint as an interface, and then connect to the endpoint through the Amazon Braket API service.

Here are the general steps of this process, which are explained in detail in later sections.

- Configure and launch an Amazon VPC to host your AWS resources. If you already have a VPC, you can skip this step.
- Create an Amazon VPC endpoint for Braket
- Connect and run Braket tasks through your endpoint

Step 1: Launch an Amazon VPC if needed

Remember that you can skip this step if your account already has a VPC in operation.

A VPC controls your network settings, such as the IP address range, subnets, route tables, and network gateways. Essentially, you are launching your AWS resources in a custom virtual network. For more information about VPCs, see the Amazon VPC User Guide.

Open the Amazon VPC console and create a new VPC with subnets, security groups, and network gateways.

Step 2: Create an interface VPC endpoint for Braket

You can create a VPC endpoint for the Braket service using either the Amazon VPC console or the AWS Command Line Interface (AWS CLI). For more information, see Creating an interface endpoint in the *Amazon VPC User Guide*.

To create a VPC endpoint in the console, open the Amazon VPC console, open the **Endpoints** page, and proceed to create the new endpoint. Make note of the endpoint ID for later reference. It is required as part of the –`endpoint-url` flag when you are making certain calls to the Braket API.

Create the VPC endpoint for Braket using the following service name:

- `com.amazonaws.substitute_your_region.braket`

Note: If you enable private DNS for the endpoint, you can make API requests to Braket using its default DNS name for the Region, for example, `braket.us-east-1.amazonaws.com`.

For more information, see Accessing a service through an interface endpoint in the *Amazon VPC User Guide*.

Step 3: Connect and run Braket tasks through your endpoint

After you have created a VPC endpoint, you can run CLI commands that include the `endpoint-url` parameter to specify interface endpoints to the API or runtime, such as the following example:

```
aws braket search-quantum-tasks --endpoint-url
 VPC_Endpoint_ID.braket.substituteYourRegionHere.vpce.amazonaws.com
```

If you enable private DNS hostnames for your VPC endpoint, you don't need to specify the endpoint as a URL in your CLI commands. Instead, the Amazon Braket API DNS hostname, which the CLI and Braket SDK use by default, resolves to your VPC endpoint. It has the form shown in the following example:

```
https://braket.substituteYourRegionHere.amazonaws.com
```

The blog post called Direct access to Amazon SageMaker notebooks from Amazon VPC by using an AWS PrivateLink endpoint provides an example of how to set up an endpoint to make secure connections to SageMaker notebooks, which are similar to Amazon Braket notebooks.

If you're following the steps in the blog post, remember to substitute the name **Amazon Braket** for **Amazon SageMaker**. For **Service Name** enter `com.amazonaws.us-east-1.braket` or substitute your correct AWS Region name into that string, if your Region is not *us-east-1*.

More about creating an endpoint

- For information about how to create a VPC with private subnets, see Create a VPC with private subnets
- For information about creating and configuring an endpoint using the Amazon VPC console or the AWS CLI, see Creating an Interface Endpoint in the *Amazon VPC User Guide*.

- For information about creating and configuring an endpoint using AWS CloudFormation, see the AWS::EC2::VPCEndpoint resource in the *AWS CloudFormation User Guide*.

Control access with Amazon VPC endpoint policies

To control connectivity access to Amazon Braket, you can attach an AWS Identity and Access Management (IAM) endpoint policy to your Amazon VPC endpoint. The policy specifies the following information:

- The principal (user or role) that can perform actions.
- The actions that can be performed.
- The resources on which actions can be performed.

For more information, see Controlling access to services with VPC endpoints in the *Amazon VPC User Guide*.

Example: VPC endpoint policy for Braket actions

The following example shows an endpoint policy for Braket. When attached to an endpoint, this policy grants access to the listed Braket actions for all principals on all resources.

```
{
 "Statement":[
 {
   "Principal":"*",
   "Effect":"Allow",
   "Action":[
     “braket:action-1",
     “braket:action-2",
     “braket:action-3”
     ],
   "Resource":"*"
   }
  ]
}
```

You can create complex IAM rules by attaching multiple endpoint policies. For more information and examples, see:

- Amazon Virtual Private Cloud Endpoint Policies for Step Functions
- Creating Granular IAM Permissions for Non-Admin Users
- Controlling Access to Services with VPC Endpoints

Tagging Amazon Braket resources

A *tag* is a custom attribute label that you assign or that AWS assigns to an AWS resource. A tag is *metadata* that tells more about your resource. Each tag consists of a *key* and a *value*. Together these are known as *key-value pairs*. For tags that you assign, you define the key and value.

In the Amazon Braket console, you can navigate to a task or a notebook and view the list of tags associated with it. You can add a tag, remove a tag, or modify a tag. You can tag a task or notebook upon creation, and then manage associated tags through the console, AWS CLI, or API.

Using tags

Tags can organize your resources into categories that are useful to you. For example, you can assign a "Department" tag to specify the department that owns this resource.

Each tag has two parts:

- A tag key (for example, *CostCenter*, *Environment*, or *Project*). Tag keys are case sensitive.
- An optional field known as a tag value (for example, *111122223333* or *Production*). Omitting the tag value is the same as using an empty string. Like tag keys, tag values are case sensitive.

Tags help you do the following things:

- **Identify and organize your AWS resources.** Many AWS services support tagging, so you can assign the same tag to resources from different services to indicate that the resources are related.
- **Track your AWS costs.** You activate these tags on the AWS Billing and Cost Management dashboard. AWS uses the tags to categorize your costs and deliver a monthly cost allocation report to you. For more information, see Use cost allocation tags in the *AWS Billing and Cost Management User Guide*.
- **Control access to your AWS resources.** For more information, see Controlling access using tags in the *IAM User Guide*.

More about AWS and tags

- For general information on tagging, including naming and usage conventions, see Tagging AWS Resources in the *AWS General Reference*.
- For information about restrictions on tagging, see Tag naming limits and requirements in the *AWS General Reference*.
- For best practices and tagging strategies, see Tagging best practices and AWS Tagging Strategies.
- For a list of services that support using tags, see the *Resource Groups Tagging API Reference*.

The following sections provide more specific information about tags for Amazon Braket.

Supported resources in Amazon Braket

The following resource type in Amazon Braket supports tagging:

- `quantum-task` resource
- **Resource Name:** `AWS::Service::Braket`
- **ARN Regex:** `arn:${Partition}:braket:${Region}:${Account}:quantum-task/${RandomId}`

Note: You can apply and manage tags for your Amazon Braket notebooks in the Amazon Braket console, by using the console to navigate to the notebook resource, although the notebooks actually are Amazon SageMaker resources. For more information, see Notebook Instance Metadata in the SageMaker documentation.

Tag restrictions

The following basic restrictions apply to tags on Amazon Braket resources:

- Maximum number of tags that you can assign to a resource: 50
- Maximum key length: 128 Unicode characters
- Maximum value length: 256 Unicode characters
- Valid characters for key and value: `a-z, A-Z, 0-9, space`, and these characters: `_ . : / = + -` and `@`
- Keys and values are case sensitive.
- Don't use `aws` as a prefix for keys; it's reserved for AWS use.

Managing tags in Amazon Braket

You set tags as *properties* on a *resource*. You can view, add, modify, list, and delete tags through the Amazon Braket console, the Amazon Braket API, or the AWS CLI. For more information, see the Amazon Braket API reference.

Add tags

You can add tags to taggable resources at the following times:

- **When you create the resource:** Use the console, or include the `Tags` parameter with the `Create` operation in the AWS API.
- **After you create the resource:** Use the console to navigate to the task or notebook resource, or call the `TagResource` operation in the AWS API.

To add tags to a resource when you create it, you also need permission to create a resource of the specified type.

View tags

You can view the tags on any of the taggable resources in Amazon Braket by using the console to navigate to the task or notebook resource, or by calling the AWS `ListTagsForResource` API operation.

You can use the following AWS API command to view tags on a resource:

- **AWS API:** `ListTagsForResource`

Edit tags

You can edit tags by using the console to navigate to the task or notebook resource or you can use the following command to modify the value for a tag attached to a taggable resource. When you specify a tag key that already exists, the value for that key is overwritten:

- **AWS API:** `TagResource`

Remove tags

You can remove tags from a resource by specifying the keys to remove, by using the console to navigate to the task or notebook resource, or when calling the `UntagResource` operation.

- **AWS API:** `UntagResource`

Example of CLI tagging in Amazon Braket

If you're working with the AWS CLI, here is an example command showing how to create a tag that applies to a quantum task you create for the SV1 simulator with parameter settings of the Rigetti QPU. Notice that the tag is specified at the end of the example command. In this case, **Key** is given the value `state` and **Value** is given the value `Washington`.

```
aws braket create-quantum-task --action /
"{\"braketSchemaHeader\": {\"name\": \"braket.ir.jaqcd.program\", /
    \"version\": \"1\"}, /
    \"instructions\": [{\"angle\": 0.15, \"target\": 0, \"type\": \"rz\"}], /
    \"results\": null, /
    \"basis_rotation_instructions\": null}" /
  --device-arn "arn:aws:braket:::device/quantum-simulator/amazon/sv1" /
  --output-s3-bucket  "my-example-braket-bucket-name" /
  --output-s3-key-prefix "my-example-username"  /
  --shots 100   /
  --device-parameters /
  "{\"braketSchemaHeader\": /
     {\"name\": \"braket.device_schema.rigetti.rigetti_device_parameters\", /
      \"version\": \"1\"}, \"paradigmParameters\": /
       {\"braketSchemaHeader\": /
         {\"name\": \"braket.device_schema.gate_model_parameters\", /
          \"version\": \"1\"}, /
          \"qubitCount\": 2}}" /
          --tags {\"state\":\”Washington\"}
```

Tagging with the Amazon Braket API

- If you're using the Amazon Braket API to set up tags on a resource, call the `TagResourceAPI`.

```
aws braket tag-resource --resource-arn $YOUR_TASK_ARN --tags {\"city\":
\"Seattle\"}
```

- To remove tags from a resource, call the `UntagResourceAPI`.

```
aws braket list-tags-for-resource --resource-arn $YOUR_TASK_ARN
```

- To list all tags that are attached to a particular resource, call the ListTagsForResourceAPI.

```
aws braket tag-resource --resource-arn $YOUR_TASK_ARN --tag-keys "[\"city\",
\"state\"]"
```

Monitoring Amazon Braket with Amazon CloudWatch

You can monitor Amazon Braket using Amazon CloudWatch, which collects raw data and processes it into readable, near real-time metrics. You view historical information generated up to 15 months ago or search metrics that have been updated in the last 2 weeks in the Amazon CloudWatch console to gain a better perspective on how Amazon Braket is performing. To learn more, see Using CloudWatch metrics.

Amazon Braket Metrics and Dimensions

Metrics are the fundamental concept in CloudWatch. A metric represents a time-ordered set of data points that are published to CloudWatch. Every metric is characterized by a set of dimensions. To learn more about metrics dimensions in CloudWatch, see CloudWatch dimensions.

Amazon Braket sends the following metric data, specific to Amazon Braket, into the Amazon CloudWatch metrics:

Task Metrics

Metrics are available if tasks exist. They are displayed under **AWS/Braket/By Device** in the CloudWatch console.

Metric	Description
Count	Number of tasks.
Latency	This metric is emitted when a task has completed. It represents the total time from task initialization to completion.

Dimensions for Task Metrics

The task metrics are published with a dimension based on the `deviceArn` parameter, which has the form **arn:aws:braket:::device/xxx**.

Supported Devices

For a list of supported devices and device ARNs, see Braket devices.

> **Note**
> You can view the CloudWatch log streams for Amazon Braket notebooks by navigating to the **Notebook detail** page on the Amazon SageMaker console. Additional Amazon Braket notebook settings are available through the SageMaker console.

Events and automated actions for Amazon Braket with Amazon EventBridge

Amazon EventBridge monitors status change events in Amazon Braket tasks. Events from Amazon Braket are delivered to EventBridge, almost in real time. You can write simple rules that indicate which events interest you, including automated actions to take when an event matches a rule. Automatic actions that can be triggered include these:

- Invoking an AWS Lambda function
- Activating an AWS Step Functions state machine
- Notifying an Amazon SNS topic

EventBridge monitors these Amazon Braket status change events:

- The state of task changes

Amazon Braket guarantees delivery of task status change events. These events are delivered at least once, but possibly out of order.

For more information, see the Events and Event Patterns in EventBridge.

In this section:

Monitor task status with EventBridge

With EventBridge, you can create rules that define actions to take when Amazon Braket sends notification of a status change regarding a Braket task. For example, you can create a rule that sends you an email message each time the status of a task changes.

1. Log in to AWS using an account that has permissions to use EventBridge and Amazon Braket.
2. Open the Amazon EventBridge console at https://console.aws.amazon.com/events/.
3. Using the following values, create an EventBridge rule:
   - For **Rule type**, choose **Rule with an event pattern**.
   - For **Event source**, choose **Other**.
   - In the **Event pattern** section, choose **Custom patterns (JSON editor)**, and then paste the following event pattern into the text area:

```
{
  "source": [
    "aws.braket"
  ],
  "detail-type": [
```

```
    "Braket Task State Change"
  ]
}
```

To capture all events from Amazon Braket, exclude the `detail-type` section as shown in the following code:

```
{
  "source": [
    "aws.braket"
  ]
}
```

- For **Target types**, choose **AWS service**, and for **Select a target**, choose a target such as an Amazon SNS topic or AWS Lambda function. The target is triggered when a task state change event is received from Amazon Braket.

  For example, use an Amazon Simple Notification Service (SNS) topic to send an email or text message when an event occurs. To do that, first create an Amazon SNS topic using the Amazon SNS console. To learn more, see Using Amazon SNS for user notifications.

For details about creating rules, see Creating Amazon EventBridge rules that react to events.

Example Amazon Braket EventBridge event

For information on the fields for an Amazon Braket Task Status Change event, see Events and Event Patterns in EventBridge.

The following attributes appear in the JSON "detail" field.

- **`quantumTaskArn`** (str): The task for which this event was generated.
- **`status`** (Optional[str]): The status to which the task transitioned.
- **`deviceArn`** (str): The device specified by the user for which this task was created.
- **`shots`** (int): The number of shots requested by the user.
- **`outputS3Bucket`** (str): The output bucket specified by the user.
- **`outputS3Directory`** (str): The output key prefix specified by the user.
- **`createdAt`** (str): The task creation time as an ISO-8601 string.
- **`endedAt`** (Optional[str]): The time at which the task reached a terminal state. This field is present only when the task has transitioned to a terminal state.

The following JSON code shows an example of an Amazon Braket Task Status Change event.

```
{
    "version":"0",
    "id":"6101452d-8caf-062b-6dbc-ceb5421334c5",
    "detail-type":"Braket Task State Change",
    "source":"aws.braket",
    "account":"012345678901",
    "time":"2021-10-28T01:17:45Z",
    "region":"us-east-1",
    "resources":[
        "arn:aws:braket:us-east-1:012345678901:quantum-task/834b21ed-77a7-4b36-a90c-
c776afc9a71e"
    ],
```

```
    "detail":{
        "quantumTaskArn":"arn:aws:braket:us-east-1:012345678901:quantum-
task/834b21ed-77a7-4b36-a90c-c776afc9a71e",
        "status":"COMPLETED",
        "deviceArn":"arn:aws:braket:::device/quantum-simulator/amazon/sv1",
        "shots":"100",
        "outputS3Bucket":"amazon-braket-0260a8bc871e",
        "outputS3Directory":"sns-testing/834b21ed-77a7-4b36-a90c-c776afc9a71e",
        "createdAt":"2021-10-28T01:17:42.898Z",
        "eventName":"MODIFY",
        "endedAt":"2021-10-28T01:17:44.735Z"
      }
}
```

Amazon Braket API logging with CloudTrail

Amazon Braket is integrated with AWS CloudTrail, a service that provides a record of actions taken by a user, role, or an AWS service in Amazon Braket. CloudTrail captures all API calls for Amazon Braket as events. The calls captured include calls from the Amazon Braket console and code calls to the Amazon Braket API operations. If you create a trail, you can enable continuous delivery of CloudTrail events to an Amazon S3 bucket, including events for Amazon Braket. If you do not configure a trail, you can still view the most recent events in the CloudTrail console in **Event history**. Using the information collected by CloudTrail, you can determine the request that was made to Amazon Braket, the IP address from which the request was made, who made the request, when it was made, and additional details.

To learn more about CloudTrail, see the AWS CloudTrail User Guide.

Amazon Braket Information in CloudTrail

CloudTrail is enabled on your AWS account when you create the account. When activity occurs in Amazon Braket, that activity is recorded in a CloudTrail event along with other AWS service events in **Event history**. You can view, search, and download recent events in your AWS account. For more information, see Viewing Events with CloudTrail Event History.

For an ongoing record of events in your AWS account, including events for Amazon Braket, create a trail. A *trail* enables CloudTrail to deliver log files to an Amazon S3 bucket. By default, when you create a trail in the console, the trail applies to all AWS Regions. The trail logs events from all Regions in the AWS partition and delivers the log files to the Amazon S3 bucket that you specify. Additionally, you can configure other AWS services to further analyze and act upon the event data collected in CloudTrail logs. For more information, see the following:

- Overview for Creating a Trail
- CloudTrail Supported Services and Integrations
- Configuring Amazon SNS Notifications for CloudTrail
- Receiving CloudTrail Log Files from Multiple Regions and Receiving CloudTrail Log Files from Multiple Accounts

All Amazon Braket actions are logged by CloudTrail. For example, calls to the `GetQuantumTask` or `GetDevice` actions generate entries in the CloudTrail log files.

Every event or log entry contains information about who generated the request. The identity information helps you determine the following:

- Whether the request was made with root or AWS Identity and Access Management (IAM) user credentials.
- Whether the request was made with temporary security credentials for a role or federated user.
- Whether the request was made by another AWS service.

For more information, see the CloudTrail userIdentity Element.

Understanding Amazon Braket Log File Entries

A trail is a configuration that enables delivery of events as log files to an Amazon S3 bucket that you specify. CloudTrail log files contain one or more log entries. An event represents a single request from any source and includes information about the requested action, the date and time of the action, request parameters, and so on. CloudTrail log files are not an ordered stack trace of the public API calls, so they don't appear in any specific order.

The following example is a log entry for the `GetQuantumTask` action, which gets the details of a quantum task.


```
{
  "eventVersion": "1.05",
  "userIdentity": {
    "type": "AssumedRole",
    "principalId": "foobar",
    "arn": "foobar",
    "accountId": "foobar",
    "accessKeyId": "foobar",
    "sessionContext": {
      "sessionIssuer": {
        "type": "Role",
        "principalId": "foobar",
        "arn": "foobar",
        "accountId": "foobar",
        "userName": "foobar"
      },
      "webIdFederationData": {},
      "attributes": {
        "mfaAuthenticated": "false",
        "creationDate": "2020-08-07T00:56:57Z"
      }
    }
  },
  "eventTime": "2020-08-07T01:00:08Z",
  "eventSource": "braket.amazonaws.com",
  "eventName": "GetQuantumTask",
  "awsRegion": "us-east-1",
  "sourceIPAddress": "foobar",
  "userAgent": "aws-cli/1.18.110 Python/3.6.10
 Linux/4.9.184-0.1.ac.235.83.329.metal1.x86_64 botocore/1.17.33",
  "requestParameters": {
    "quantumTaskArn": "foobar"
  },
  "responseElements": null,
  "requestID": "20e8000c-29b8-4137-9cbc-af77d1dd12f7",
  "eventID": "4a2fdb22-a73d-414a-b30f-c0797c088f7c",
  "readOnly": true,
  "eventType": "AwsApiCall",
  "recipientAccountId": "foobar"
}
```


The following shows a log entry for the `GetDevice` action, which returns the details of a device event.


```
{
  "eventVersion": "1.05",
  "userIdentity": {
    "type": "AssumedRole",
    "principalId": "foobar",
    "arn": "foobar",
    "accountId": "foobar",
    "accessKeyId": "foobar",
```

```
    "sessionContext": {
      "sessionIssuer": {
        "type": "Role",
        "principalId": "foobar",
        "arn": "foobar",
        "accountId": "foobar",
        "userName": "foobar"
      },
      "webIdFederationData": {},
      "attributes": {
        "mfaAuthenticated": "false",
        "creationDate": "2020-08-07T00:46:29Z"
      }
    }
  },
  "eventTime": "2020-08-07T00:46:32Z",
  "eventSource": "braket.amazonaws.com",
  "eventName": "GetDevice",
  "awsRegion": "us-east-1",
  "sourceIPAddress": "foobar",
  "userAgent": "Boto3/1.14.33 Python/3.7.6 Linux/4.14.158-129.185.amzn2.x86_64 exec-env/
AWS_ECS_FARGATE Botocore/1.17.33",
  "errorCode": "404",
  "requestParameters": {
    "deviceArn": "foobar"
  },
  "responseElements": null,
  "requestID": "c614858b-4dcf-43bd-83c9-bcf9f17f522e",
  "eventID": "9642512a-478b-4e7b-9f34-75ba5a3408eb",
  "readOnly": true,
  "eventType": "AwsApiCall",
  "recipientAccountId": "foobar"
}
```

API & SDK Reference Guide for Amazon Braket

Amazon Braket provides APIs, SDKs, and a command line interface that you can use to create and manage notebook instances and train and deploy models.

- Amazon Braket Python SDK (Recommended)
- Amazon Braket API Reference
- AWS Command Line Interface
- AWS SDK for .NET
- AWS SDK for C++
- AWS SDK for Go
- AWS SDK for Java
- AWS SDK for JavaScript
- AWS SDK for PHP
- AWS SDK for Python (Boto)
- AWS SDK for Ruby

You can also get code examples from the Amazon Braket Tutorials GitHub repository.

- Braket Tutorials GitHub

Document history

The following table describes the documentation for this release of Amazon Braket.

- **API version:** April 28, 2022
- **Latest API Reference update:** April 28, 2022
- **Latest documentation update:** November 28, 2022

Change	Description	Date
New algorithm library feature	Added information about the Braket algorithm library, which provides a catalog of pre-built quantum algorithms	November 28, 2022
D-Wave departure	Updated the documentation to accommodate the removal of all D-Wave devices	November 17, 2022
New device QuEra Aquila	Added support for the QuEra Aquila device	October 31, 2022
Support for Braket Pulse	Added support for Braket Pulse, which allows for pulse control to be used on Rigetti and OQC devices	October 20, 2022
Support for IonQ native gates	Added support for the native gate set offered by the IonQ device	September 13, 2022
New instance quotas	Updated the default classical compute instance quotas associated with Hybrid Jobs	August 22, 2022
New service dashboard	Updated console screenshots to include the service dashboard	August 17, 2022
New device Rigetti Aspen-M-2	Added support for the Rigetti Aspen-M-2 device	August 12, 2022
New OpenQASM features	Added OpenQASM features support for the local simulators (braket_sv and braket_dm)	August 4, 2022
New cost tracking procedures	Added how to get near-real time maximum cost estimates for simulators and hardware workloads	July 18, 2022
New Xanadu Borealis device	Added support for the Xanadu Borealis device	June 2, 2022
New onboarding simplification procedures	Added information on how the new and simplified onboarding procedures work	May 16, 2022

New device D-Wave Advantage_system6.1	Added support for the D-Wave Advantage_system6.1 device	May 12, 2022
Support for embedded simulators	Added how to run embedded simulations with hybrid jobs and how to use the PennyLane lightning simulator	May 4, 2022
AmazonBraketFullAccess - Full access policy for Amazon Braket	Added s3:ListAllMyBuckets permissions to allow users to view and inspect the buckets created and used for Amazon Braket	March 31, 2022
Support for OpenQASM	Added OpenQASM 3.0 support for gate-based quantum devices and simulators	March 7, 2022
New Quantum Hardware Provider, Oxford Quantum Circuits and new region, eu-west-2	Added support for OQC and eu-west-2	February 28, 2022
New Rigetti device	Added support for Rigetti Aspen M-1	February 15, 2022
New resource limits	Increased the maximum number of concurrent DM1 and SV1 tasks from 55 to 100	January 5, 2022
New Rigetti device	Added support for Rigetti Aspen-11	December 20, 2021
Retired Rigetti device	Discontinued support for Rigetti Aspen-10 device	December 20, 2021
New result type	Reduced density matrix result type supported by local density matrix simulator and DM1 devices	December 20, 2021
Updated policy description	Amazon Braket updated the role ARN to include the `servicerole/` path. For information on policy updates, see the Amazon Braket updates to AWS managed policies (p. 162) table.	November 29, 2021
Amazon Braket Jobs	User guide for Amazon Braket Hybrid Jobs and API added	November 29, 2021
New Rigetti device	Added support for Rigetti Aspen-10	November 20, 2021
Retired D-Wave device	Discontinued support for D-Wave QPU, Advantage_system1	November 4, 2021

New D-Wave device	Added support for an additional D-Wave QPU, Advantage_system4	October 5, 2021
New noise simulators	Added support for a Density matrix simulator (DM1), which can simulate circuits of up to 17 qubits and a local noise simulator *braket_dm*	May 25, 2021
PennyLane support	Added support for PennyLane on Amazon Braket	December 8, 2020
New simulator	Added support for a Tensor Network Simulator (TN1), which allows larger circuits	December 8, 2020
Task batching	Braket supports customer task batching	November 24, 2020
Manual qubit allocation	Braket supports manual qubit allocation on the Rigetti device	November 24, 2020
Adjustable quotas	Braket supports self-service adjustable quotas for your task resources	October 30, 2020
Support for PrivateLink	You can set up private VPC endpoints for your Braket jobs	October 30, 2020
Support for tags	Braket supports API-based tags for the *quantum-task* resource	October 30, 2020
New D-Wave device	Added support for an additional D-Wave QPU, Advantage_system1	September 29, 2020
Initial release	Initial release of the Amazon Braket documentation	August 12, 2020