

kubectl

This page contains a list of commonly used `kubectl` commands and flags.

Kubectl autocomplete

BASH

```
source <(kubectl completion bash) # setup autocomplete in bash into the current s
echo "source <(kubectl completion bash)" >> ~/.bashrc # add autocomplete permanen
```

You can also use a shorthand alias for `kubectl` that also works with completion:

```
alias k=kubectl
complete -F __start_kubectl k
```

ZSH

```
source <(kubectl completion zsh) # setup autocomplete in zsh into the current sh
echo "[[ $commands[kubectl] ]] && source <(kubectl completion zsh)" >> ~/.zshrc #
```

A Note on --all-namespaces

Appending `--all-namespaces` happens frequently enough where you should be aware of the shorthand for `--all-namespaces` :

```
kubectl -A
```

Kubectl context and configuration

Set which Kubernetes cluster `kubectl` communicates with and modifies configuration information. See [Authenticating Across Clusters with kubeconfig](#) documentation for detailed config file information.

```
kubectl config view # Show Merged kubeconfig settings.

# use multiple kubeconfig files at the same time and view merged config
KUBECONFIG=~/.kube/config:~/.kube/kubconfig2

kubectl config view

# get the password for the e2e user
kubectl config view -o jsonpath='{.users[?(@.name == "e2e")].user.password}'

kubectl config view -o jsonpath='{.users[0].name}' # display the first user
kubectl config view -o jsonpath='{.users[*].name}' # get a list of users
kubectl config get-contexts # display list of contexts
```

```

kubectl config current-context           # display the current-context
kubectl config use-context my-cluster-name # set the default context to my-cluster-name

# add a new user to your kubeconf that supports basic auth
kubectl config set-credentials kubeuser/foo.kubernetes.com --username=kubeuser --password=kubeuser

# permanently save the namespace for all subsequent kubectl commands in that context
kubectl config set-context --current --namespace=ggckad-s2

# set a context utilizing a specific username and namespace.
kubectl config set-context gce --user=cluster-admin --namespace=foo \
  && kubectl config use-context gce

kubectl config unset users.foo          # delete user foo

# short alias to set/show context/namespace (only works for bash and bash-compatible shells)
alias kx='f() { [ "$1" ] && kubectl config use-context $1 || kubectl config current-context; }'
alias kn='f() { [ "$1" ] && kubectl config set-context --current --namespace $1 || kubectl config current-context; }'

```

Kubectl apply

`apply` manages applications through files defining Kubernetes resources. It creates and updates resources in a cluster through running `kubectl apply`. This is the recommended way of managing Kubernetes applications on production. See [Kubectl Book](#).

Creating objects

Kubernetes manifests can be defined in YAML or JSON. The file extension `.yaml`, `.yml`, and `.json` can be used.

```

kubectl apply -f ./my-manifest.yaml           # create resource(s)
kubectl apply -f ./my1.yaml -f ./my2.yaml    # create from multiple files
kubectl apply -f ./dir                       # create resource(s) in all manifests
kubectl apply -f https://git.io/vPieo        # create resource(s) from url
kubectl create deployment nginx --image=nginx # start a single instance of nginx

# create a Job which prints "Hello World"
kubectl create job hello --image=busybox:1.28 -- echo "Hello World"

# create a CronJob that prints "Hello World" every minute
kubectl create cronjob hello --image=busybox:1.28 --schedule="*/1 * * * *" -- echo "Hello World"

kubectl explain pods                        # get the documentation for pod manifests

# Create multiple YAML objects from stdin
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep
spec:
  containers:
  - name: busybox
    image: busybox:1.28
    args:
    - sleep
    - "1000000"
---
apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep-less
spec:
  containers:
  - name: busybox

```

```

    image: busybox:1.28
    args:
    - sleep
    - "1000"
EOF

# Create a secret with several keys
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  password: $(echo -n "s33msi4" | base64 -w0)
  username: $(echo -n "jane" | base64 -w0)
EOF

```

Viewing, finding resources

```

# Get commands with basic output
kubectl get services                # List all services in the namespace
kubectl get pods --all-namespaces   # List all pods in all namespaces
kubectl get pods -o wide            # List all pods in the current namespace
kubectl get deployment my-dep       # List a particular deployment
kubectl get pods                   # List all pods in the namespace
kubectl get pod my-pod -o yaml      # Get a pod's YAML

# Describe commands with verbose output
kubectl describe nodes my-node
kubectl describe pods my-pod

# List Services Sorted by Name
kubectl get services --sort-by=.metadata.name

# List pods Sorted by Restart Count
kubectl get pods --sort-by='.status.containerStatuses[0].restartCount'

# List PersistentVolumes sorted by capacity
kubectl get pv --sort-by=.spec.capacity.storage

# Get the version label of all pods with label app=cassandra
kubectl get pods --selector=app=cassandra -o \
  jsonpath='{.items[*].metadata.labels.version}'

# Retrieve the value of a key with dots, e.g. 'ca.crt'
kubectl get configmap myconfig \
  -o jsonpath='{.data.ca\.crt}'

# Get all worker nodes (use a selector to exclude results that have a label
# named 'node-role.kubernetes.io/control-plane')
kubectl get node --selector='!node-role.kubernetes.io/control-plane'

# Get all running pods in the namespace
kubectl get pods --field-selector=status.phase=Running

# Get ExternalIPs of all nodes
kubectl get nodes -o jsonpath='{.items[*].status.addresses[?(@.type=="ExternalIP"]}'

# List Names of Pods that belong to Particular RC
# "jq" command useful for transformations that are too complex for jsonpath, it can be used like:
sel=${$(kubectl get rc my-rc --output=json | jq -j '.spec.selector | to_entries |
echo ${kubectl get pods --selector=$sel --output=jsonpath='{.items..metadata.name}'

# Show labels for all pods (or any other Kubernetes object that supports labelling)
kubectl get pods --show-labels

```

```

# Check which nodes are ready
JSONPATH='{range .items[*]}{@.metadata.name}:{range @.status.conditions[*]}{@.type}
&& kubectl get nodes -o jsonpath="$JSONPATH" | grep "Ready=True"

# Output decoded secrets without external tools
kubectl get secret my-secret -o go-template='{{range $k,$v := .data}}{{"### "}}{{

# List all Secrets currently in use by a pod
kubectl get pods -o json | jq '.items[].spec.containers[].env[]?.valueFrom.secret

# List all containerIDs of initContainer of all pods
# Helpful when cleaning up stopped containers, while avoiding removal of initCont
kubectl get pods --all-namespaces -o jsonpath='{range .items[*].status.initContai

# List Events sorted by timestamp
kubectl get events --sort-by=.metadata.creationTimestamp

# Compares the current state of the cluster against the state that the cluster wo
kubectl diff -f ./my-manifest.yaml

# Produce a period-delimited tree of all keys returned for nodes
# Helpful when locating a key within a complex nested JSON structure
kubectl get nodes -o json | jq -c 'paths|join(".")'

# Produce a period-delimited tree of all keys returned for pods, etc
kubectl get pods -o json | jq -c 'paths|join(".")'

# Produce ENV for all pods, assuming you have a default container for the pods, d
# Helpful when running any supported command across all pods, not just `env`
for pod in $(kubectl get po --output=jsonpath={.items..metadata.name}); do echo $

# Get a deployment's status subresource
kubectl get deployment nginx-deployment --subresource=status

```

Updating resources

```

kubectl set image deployment/frontend www=image:v2 # Rolling update
kubectl rollout history deployment/frontend # Check the hist
kubectl rollout undo deployment/frontend # Rollback to th
kubectl rollout undo deployment/frontend --to-revision=2 # Rollback to a
kubectl rollout status -w deployment/frontend # Watch rolling
kubectl rollout restart deployment/frontend # Rolling restar

cat pod.json | kubectl replace -f - # Replace a pod

# Force replace, delete and then re-create the resource. Will cause a service out
kubectl replace --force -f ./pod.json

# Create a service for a replicated nginx, which serves on port 80 and connects t
kubectl expose rc nginx --port=80 --target-port=8000

# Update a single-container pod's image version (tag) to v4
kubectl get pod mypod -o yaml | sed 's/\(image: myimage\):.*$/\1:v4/' | kubectl r

kubectl label pods my-pod new-label=awesome # Add a Label
kubectl annotate pods my-pod icon-url=http://goo.gl/XXBTWq # Add an annotat
kubectl autoscale deployment foo --min=2 --max=10 # Auto scale a d

```

Patching resources

```

# Partially update a node

```

```
kubectl patch node k8s-node-1 -p '{"spec":{"unschedulable":true}}'
```

Update a container's image; spec.containers[].name is required because it's a*

```
kubectl patch pod valid-pod -p '{"spec":{"containers":[{"name":"kubernetes-serve-
```

Update a container's image using a json patch with positional arrays

```
kubectl patch pod valid-pod --type='json' -p='[{"op": "replace", "path": "/spec/c
```

Disable a deployment livenessProbe using a json patch with positional arrays

```
kubectl patch deployment valid-deployment --type json -p='[{"op": "remove", "p
```

Add a new element to a positional array

```
kubectl patch sa default --type='json' -p='[{"op": "add", "path": "/secrets/1", "
```

Update a deployment's replica count by patching its scale subresource

```
kubectl patch deployment nginx-deployment --subresource='scale' --type='merge' -p
```

Editing resources

Edit any API resource in your preferred editor.

```
kubectl edit svc/docker-registry # Edit the service named do
KUBE_EDITOR="nano" kubectl edit svc/docker-registry # Use an alternative editor
```

Scaling resources

```
kubectl scale --replicas=3 rs/foo # Scale a repli
kubectl scale --replicas=3 -f foo.yaml # Scale a resou
kubectl scale --current-replicas=2 --replicas=3 deployment/mysql # If the deploy
kubectl scale --replicas=5 rc/foo rc/bar rc/baz # Scale multipl
```

Deleting resources

```
kubectl delete -f ./pod.json # Delete a pod
kubectl delete pod unwanted --now # Delete a pod
kubectl delete pod,service baz foo # Delete pods a
kubectl delete pods,services -l name=myLabel # Delete pods a
kubectl -n my-ns delete pod,svc --all # Delete all po
# Delete all pods matching the awk pattern1 or pattern2
kubectl get pods -n mynamespace --no-headers=true | awk '/pattern1|pattern2/{pri
```

Interacting with running Pods

```
kubectl logs my-pod # dump pod logs (stdout)
kubectl logs -l name=myLabel # dump pod logs, with label n
kubectl logs my-pod --previous # dump pod logs (stdout) for
kubectl logs my-pod -c my-container # dump pod container logs (st
kubectl logs -l name=myLabel -c my-container # dump pod logs, with label n
kubectl logs my-pod -c my-container --previous # dump pod container logs (st
kubectl logs -f my-pod # stream pod logs (stdout)
kubectl logs -f my-pod -c my-container # stream pod container logs (
kubectl logs -f -l name=myLabel --all-containers # stream all pods logs with l
kubectl run -i --tty busybox --image=busybox:1.28 -- sh # Run pod as interactive
kubectl run nginx --image=nginx -n mynamespace # Start a single instance of
```

```
kubectl run nginx --image=nginx                                # Run pod nginx and write its
--dry-run=client -o yaml > pod.yaml

kubectl attach my-pod -i                                       # Attach to Running Container
kubectl port-forward my-pod 5000:6000                         # Listen on port 5000 on the
                                                                # Run command in existing pod
kubectl exec my-pod -- ls /                                    # Interactive shell access to
                                                                # Run command in existing pod
                                                                # Show metrics for a given po
kubectl top pod POD_NAME --containers                         # Show metrics for a given po
```

Copy files and directories to and from containers

```
kubectl cp /tmp/foo_dir my-pod:/tmp/bar_dir                  # Copy /tmp/foo_dir local
kubectl cp /tmp/foo my-pod:/tmp/bar -c my-container          # Copy /tmp/foo local file
kubectl cp /tmp/foo my-namespace/my-pod:/tmp/bar            # Copy /tmp/foo local file
kubectl cp my-namespace/my-pod:/tmp/foo /tmp/bar            # Copy /tmp/foo from a rem
```

Note: `kubectl cp` requires that the 'tar' binary is present in your container image. If 'tar' is not present, `kubectl cp` will fail. For advanced use cases, such as symlinks, wildcard expansion or file mode preservation consider using `kubectl exec`.

```
tar cf - /tmp/foo | kubectl exec -i -n my-namespace my-pod -- tar xf - -C /tmp/ba
kubectl exec -n my-namespace my-pod -- tar cf - /tmp/foo | tar xf - -C /tmp/bar
```

Interacting with Deployments and Services

```
kubectl logs deploy/my-deployment                            # dump Pod logs for a D
kubectl logs deploy/my-deployment -c my-container            # dump Pod logs for a D

kubectl port-forward svc/my-service 5000                    # listen on local port
kubectl port-forward svc/my-service 5000:my-service-port    # listen on local port

kubectl port-forward deploy/my-deployment 5000:6000         # listen on local port
kubectl exec deploy/my-deployment -- ls                      # run command in first
```

Interacting with Nodes and cluster

```
kubectl cordon my-node                                       # Mark my-n
kubectl drain my-node                                         # Drain my-
kubectl uncordon my-node                                       # Mark my-n
kubectl top node my-node                                       # Show metr
kubectl cluster-info                                          # Display a
kubectl cluster-info dump                                     # Dump curr
kubectl cluster-info dump --output-directory=/path/to/cluster-state # Dump curr

# If a taint with that key and effect already exists, its value is replaced as sp
kubectl taint nodes foo dedicated=special-user:NoSchedule
```

Resource types

List all supported resource types along with their shortnames, [API group](#), whether they are [namespaced](#), and [Kind](#):

```
kubectl api-resources
```

Other operations for exploring API resources:

```
kubectl api-resources --namespaced=true      # All namespaced resources
kubectl api-resources --namespaced=false     # All non-namespaced resources
kubectl api-resources -o name                # All resources with simple output (
kubectl api-resources -o wide                # All resources with expanded (aka "
kubectl api-resources --verbs=list,get       # All resources that support the "li
kubectl api-resources --api-group=extensions # All resources in the "extensions"
```

Formatting output

To output details to your terminal window in a specific format, add the `-o` (or `--output`) flag to a supported `kubectl` command.

Output format	Description
<code>-o=custom-columns=<spec></code>	Print a table using a comma separated list of custom columns
<code>-o=custom-columns-file=<filename></code>	Print a table using the custom columns template in the <code><filename></code> file
<code>-o=json</code>	Output a JSON formatted API object
<code>-o=jsonpath=<template></code>	Print the fields defined in a jsonpath expression
<code>-o=jsonpath-file=<filename></code>	Print the fields defined by the jsonpath expression in the <code><filename></code> file
<code>-o=name</code>	Print only the resource name and nothing else
<code>-o=wide</code>	Output in the plain-text format with any additional information, and for pods, the node name is included
<code>-o=yaml</code>	Output a YAML formatted API object

Examples using `-o=custom-columns` :

```
# All images running in a cluster
kubectl get pods -A -o=custom-columns='DATA:spec.containers[*].image'

# All images running in namespace: default, grouped by Pod
kubectl get pods --namespace default --output=custom-columns="NAME:.metadata.name

# All images excluding "k8s.gcr.io/coredns:1.6.2"
kubectl get pods -A -o=custom-columns='DATA:spec.containers[?(@.image!="k8s.gcr.i

# All fields under metadata regardless of name
kubectl get pods -A -o=custom-columns='DATA:metadata.*'
```

More examples in the `kubectl` [reference documentation](#).

Kubectl output verbosity and debugging

Kubectl verbosity is controlled with the `-v` or `--v` flags followed by an integer representing the log level. General Kubernetes logging conventions and the associated log levels are described [here](#).

Verbosity	Description
<code>--v=0</code>	Generally useful for this to <i>always</i> be visible to a cluster operator.
<code>--v=1</code>	A reasonable default log level if you don't want verbosity.
<code>--v=2</code>	Useful steady state information about the service and important log messages that may correlate to significant changes in the system. This is the recommended default log level for most systems.
<code>--v=3</code>	Extended information about changes.
<code>--v=4</code>	Debug level verbosity.
<code>--v=5</code>	Trace level verbosity.
<code>--v=6</code>	Display requested resources.
<code>--v=7</code>	Display HTTP request headers.
<code>--v=8</code>	Display HTTP request contents.
<code>--v=9</code>	Display HTTP request contents without truncation of contents.